

CMPUT 379 Lab

ETLC E1003: Tuesday, 5:00 – 7:50 PM.

Tianyu Zhang, Peiran Yao

CAB 311: Thursday, 2:00 – 4:50 PM.

Max Ellis, Aidan Bush

Assignments

- Don't put your code in a public repo (GitHub has free private repos)
- Make sure to test it on the lab machines
- Remember to name the zip or tar.gz file correctly
 - i.e. Assignment 1: dragonshell.zip or dragonshell.tar.gz
- Remember to name the executable correctly
 - i.e. Assignment 1: dragonshell

Last Week...

- Introduction to networking
- Creating and using sockets
- Getting familiar with TCP and UDP
- A brief look at pthreads
- Race conditions and deadlocks

Today's Lab

- Review on Threads and Thread safe data structures
- Assignment 2: Threadpool and MapReduce

Threads

- Threads allow concurrent processing with shared memory
- In the next assignment you will have to use pthreads

Pthreads

- The POSIX thread library
- Compile and link with **-pthread**

Creating Pthreads

*int pthread_create(pthread_t *, pthread_attr_t *, void * (*)(void*), void *)*

- Starts the given function in a thread
- For default attributes attr = NULL

Defining thread start routines

void *(*start_routine) (void *)

- start_routine is a function pointer
- Return value of type **void *** (pointer to any type)
- Single parameter of type **void ***

```
void* thread_function(void* args_p /* the input */) {  
    int* num = (int *) args_p; /* cast the input to the type you want */  
    /* do something */  
    return (void *) ret; /* return the value as void * pointer */  
}
```


Defining thread start routines

- Pass the function name when calling **pthread_create**

```
pthread_create(&thread_handle, NULL, thread_function, NULL);
```

Passing parameters

*int pthread_create(pthread_t *, pthread_attr_t *, void * (*func)(void*), void *)*

- The last parameter of **pthread_create()** is a pointer to the argument
- For example, to pass an integer, take its address as the last argument

```
int num_to_pass = 42;
pthread_create(&thread_handle, NULL, thread_function, &num_to_pass);
```

- To receive it, cast the last argument back to **int *** in the thread routine


```
void* thread_function(void* args_p /* the input */) {
    int* num = (int *) args_p; /* cast the input to the type you want */
    /* do something */
}
```

Joining Pthreads

```
int pthread_join(pthread_t thread, void **value_ptr);
```

- Wait until the target thread terminates
- Store the return value to where value_ptr points to

Receiving return value

- Values defined in thread routines are stored in the stack, and are destroyed after the thread terminates.
- To persist it, use **malloc** / **new** to store it in the heap.  Remember to free

```
void* thread_function(void* args_p /* the input */) {  
    int* num = malloc(sizeof(int));  *num = 42;  
    return num;
```

- To receive it

```
int* p_ret_value;  
pthread_join(thread, &p_ret_value);  /* now p_ret_value points to 42 */
```

Receiving return value

- For small types (size \leq sizeof(void *)), you can return it directly

```
void* thread_function(void* args_p /* the input */) {  
    int num = 42;  
    return (void*) num;  
}
```

- To receive it

```
int ret;  
pthread_join(thread_handle, (void**)&ret);  
printf("ret: %d\n", ret);
```

Pthread synchronization

- A resource (i.e. variable) shouldn't be accessed by multiple threads at the same time
- Mutex
- Conditional Variable

Too much milk revisited

Consider that Peter and Greg are roommates

Time	Peter	Greg
3:00	Look in fridge, no milk	
3:05	Leave for store	
3:10	Arrive at store	Look in fridge, no milk
3:15	Buy milk	Leave for store
3:20	Arrive at home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive at home, put milk away (oof)

Too much milk revisited

- How to solve this problem talk with your neighbour?
- Share your solutions!

Too much milk revisited

Adding locks!

Time	Peter	Greg
3:00	 Lock the kitchen. Look in fridge, no milk	
3:05	Leave for store	
3:10	Arrive at store	Try to look in fridge, but the kitchen is locked!
3:15	Buy milk	Waiting for Peter 
3:20	Arrive at home, put milk away  Unlock the kitchen.	Look in fridge after Peter unlocks the door
3:25		Enjoy milk  !

Pthread mutex

`pthread_mutex_t`

- Implements a mutually exclusive object and allows locking access to a single thread

- Initialization:

`pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER`

Pthread mutex locking and unlocking

- Locking

*int pthread_mutex_lock(pthread_mutex_t *)*

- Unlocking

*int pthread_mutex_unlock(pthread_mutex_t *)*

- Both return non zero on an error

Exhausted assistant

This time, Peter has an assistant Arthur to help refill milk

Time	Peter	Arthur
3:00		Check the fridge: enough milk
3:05		Check the fridge: enough milk
3:10		Check the fridge: enough milk
3:15	Drink milk 🥛	Check the fridge: no milk
3:20		Buy milk
3:25		Arrive at home, put milk away
3:30		Check the fridge: enough milk

Exhausted assistant

Arthur has a very inefficient way of working. Why?

~~Exhausted~~ Relaxed assistant

Arthur can wait for the fridge to be empty, and Peter can notify Arthur.

Time	Peter	Arthur
3:00		Sleeping 🛏️ 😴
3:05		Sleeping 🛏️ 😴
3:10		Sleeping 🛏️ 😴
3:15	Drink milk 🥛 , notify Arthur 🕒	Buy milk
3:20		Arrive at home, put milk away
3:25		Sleeping 🛏️ 😴
3:30		Sleeping 🛏️ 😴

Pthread condition

`pthread_cond_t`

- Representing a condition to wait for

- Initialization:

`pthread_cond_t cond = PTHREAD_COND_INITIALIZER`

Pthread condition - wait

*int pthread_cond_wait(pthread_cond_t *, pthread_mutex_t *)*

- Allows pthreads to wait on for a condition to be set before continuing
- The mutex must be locked before calling
- When it returns the mutex will be locked and owned by that thread

Pthread condition - waiting

```
pthread_mutex_lock(&kitchen_mutex);  
if (has_milk) {  
    pthread_cond_wait(&no_milk_cond, &kitchen_mutex); // automatically  
release the lock  
} else {  
    buy_milk();  
}  
pthread_mutex_unlock(&kitchen_mutex);
```

Pthread condition - signalling and broadcasting

*int pthread_cond_signal(pthread_cond_t *)*

*int pthread_cond_broadcast(pthread_cond_t *)*

- Signal one or all threads waiting for the condition
- The mutex must be locked before calling

Pthread cheatsheet

	Thread	Process
Starting point	Another function	Current line of code
Creation	pthread_create()	fork()
Get identifier	pthread_self()	getpid()
Waiting / Joining	pthread_wait()	wait() / waitpid()
Terminating	pthread_exit()	exit()

Thread Safe Data Structures

- Avoid writing data at the same time
- Avoid reading data while data is being written
- Lock data when it is being written to
- A straightforward solution: add a global mutex for each data structure
 - Usually, the more you lock, the worse performance you will have

Assignment 2

Thread pool

- Solution for managing multiple threads with multiple jobs
- Thread pool organizes its work queue and threads
 - A mechanism is involved to organize the work queue when adding new works
- Each thread in the threadpool continually checks for work
- Concurrency control for accessing the work queue



Thread pool

You are implementing a thread pool library, make sure it is general



Your implementation will include following

- Data structures for your **thread pool** and your **shared data structure**
- Your thread pool library must include
 - Create / destroy pool
 - In creation take callback functions (function pointers) to process work
 - Have access to the shared data structure


Thread pool - Create



Thread pool - Add works



Thread pool - Add works





Immediately picked by an idle thread
(Not necessarily the first thread, all threads are running parallel, so you need concurrency control)

Thread pool - Add works





Immediately picked by an idle thread
(Not necessarily the first thread, all threads are running parallel, so you need concurrency control)

Thread pool - Add works





After all threads are working on some jobs, new jobs wait in the queue in an order (LJF)
Break tie by FIFO

Thread pool - Add works





After all threads are working on some jobs, new jobs wait in the queue in an order (LJF)
Break tie by FIFO

Thread pool - Add works



After all threads are working on some jobs, new jobs wait in the queue in an order (LJF)
Break tie by FIFO

Thread pool - Add works



After all threads are working on some jobs, new jobs wait in the queue in an order (LJF)
Break tie by FIFO

MapReduce

- Method for solving problems in a distributed setting
- Components
 - Map
 - Reduce
- You will implement a simplified version of MapReduce library for your next assignment
 - It is a library make sure it will work when we test with different Map and Reduce functions

See also: Dean, J., & Ghemawat, S. (2008). *MapReduce: simplified data processing on large clusters*. *Communications of the ACM*, 51(1), 107-113.

MapReduce - Programming model

- Map

- Transform elements of an array, creating any number of K-V pairs for each element.
- Applied on different elements in parallel
- Input type: `list(data)`
- Output type: `list(k, v)`

- Reduce

- Take all values for a given key then generate one or more values
- Applied on different keys in parallel
- Input type: `(k, list(v1))`
- Output type: `list(v2)`

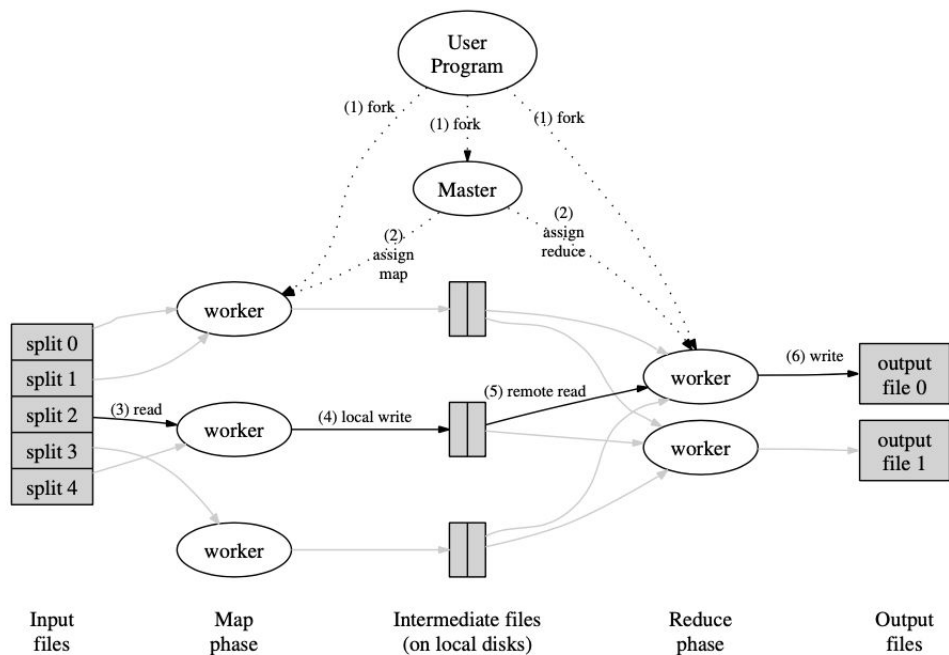
MapReduce - Programming model

- Example: counting words

```
map(document):  
    // document: list of words  
    for each word w in value:  
        EmitIntermediate(w, 1);  
  
reduce(String key, Iterator values):  
    // key: a word  
    // values: a list of counts  
    int result = 0;  
    for each v in values:  
        result += v;  
    Emit(result);
```

MapReduce - Distributed implementation

- You'll need to implement a slightly different version in Assignment 2



Simplified MapReduce Flow

Master thread manages the solution

1. Creates a thread pool contains M threads
2. Wait for all Map threads to complete
3. Creates R Reduce threads
 - Each thread processes a given partition
4. Wait for all Reduce threads to complete

Map Function

- We provide you an implementation of a map function to convert input file into multiple key/value pairs
- Calls ***MR_Emit*** (you implement) whenever a new key/value pair is generated
- ***MR_Emit(key, value)*** inserts the given key/value pair into the correct partition (using ***MR_Partition***) in your intermediate data structure

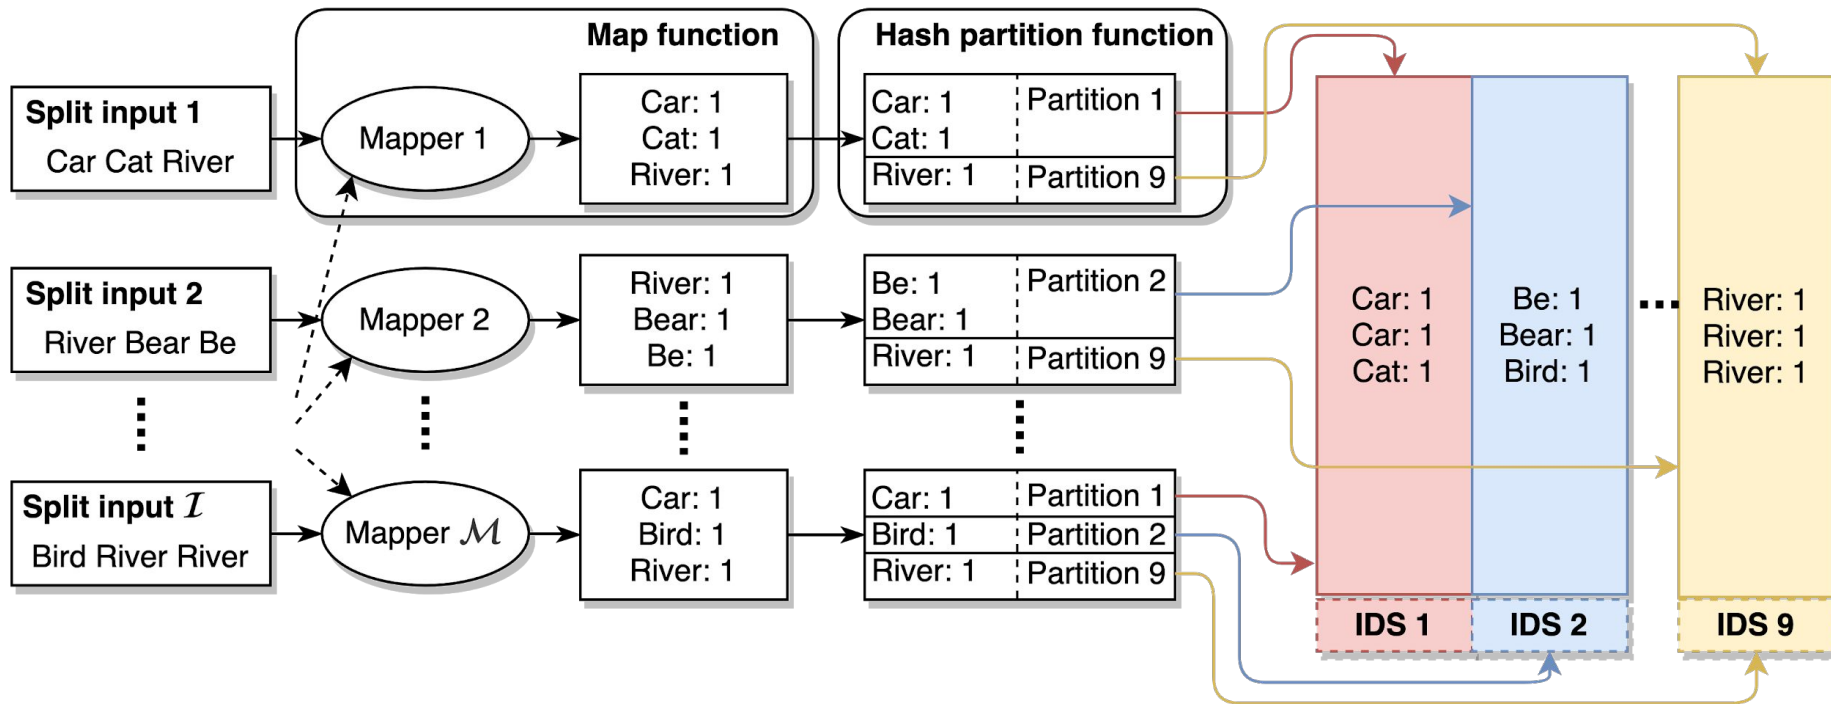
Intermediate Data Structure

- You need to implement your own intermediate data structure, as long as it can store key/value pairs, it must be your own
- Your intermediate data structure must support R partitions (R is the number of reducers)

Map Threads

- Your thread pool library will allow threads to run jobs concurrently
- The Map function will be continually called in each thread until all files have been processed
- You need to manage access to the list of jobs in your work queue

Map



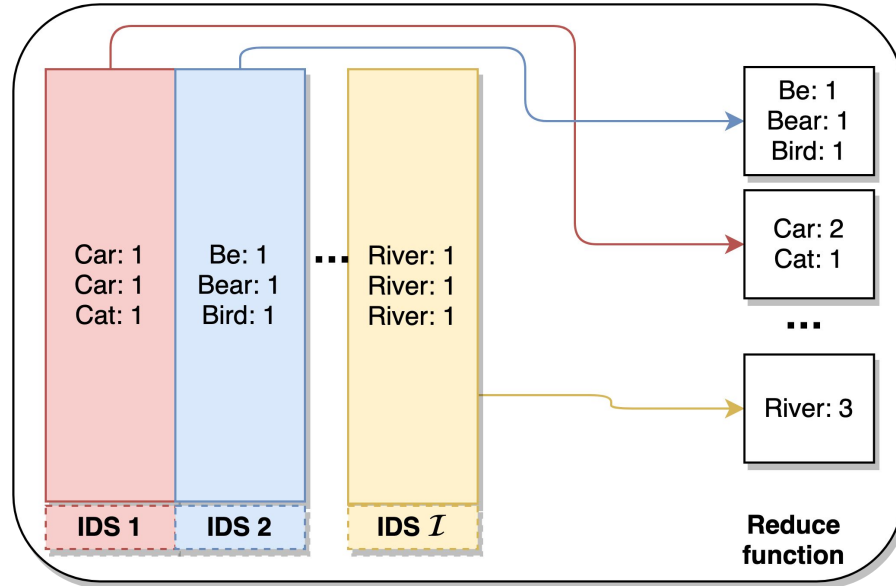
Reduce

- Reducer threads must reduce all keys in their assigned partition
 - You must implement ***MR_ProcessPartition*** (Master Reduce function) to handle the whole partition
- We provide you an implementation of a reduce function that only reduces for a single key
 - It will use ***MR_GetNext*** to retrieve the next value for the given key

Reduce Threads

- R reduce threads are created, one for each partition
 - Hence there must be R partitions
 - The thread function is *MR_ProcessPartition*
- Reducer will call user-defined reduce function iteratively until all key/value pair in the partition is processed
- The partitions is sorted so that all of the same keys are in order

Reduce

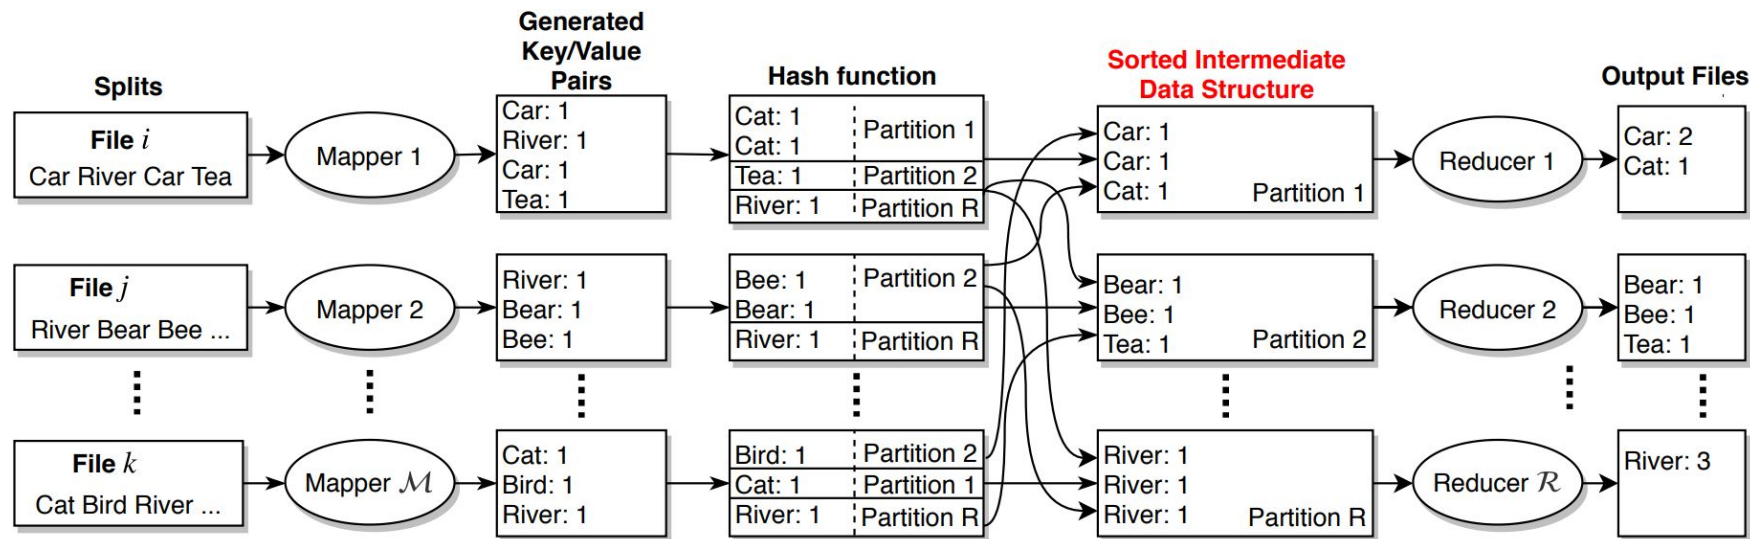


Intermediate Data Structure

- Must use *pthread_mutex* for locking access
- Must be thread safe
- Must be designed and documented individually

MapReduce Summary

- Given list of files
- Call Map on each file using M threads
- Use a hash function to partition keys into R partitions
- Call Reduce for each resulting key using R threads




MapReduce - examples

- Distributed Grep
 - Find all lines matching the given pattern
 - Map: emits a line if it matches a supplied pattern
 - Reduce: return the same line unchanged
- Inverted Index
 - For all words, list all documents in which a word appeared
 - Map: for each word in a document, emit <word, document ID>
 - Reduce: for each word, combine all IDs into a list. Emit <word, list(Document ID)>

MapReduce - more examples

- Spark and Hadoop MapReduce are distributed data-processing libraries inspired by MapReduce.
- Used widely, by Amazon, eBay, Facebook, Hulu, Alibaba Taobao, Yahoo, Yandex



Practice problem 1

- (1) Design a concurrent array data structure that allows multiple threads to access its content without conflicts. Your implementation can
 - (a) Have a global mutex (lock) for the entire array
 - (b) Have one separate mutex for each element of the array

- (2) Create multiple threads that accesses your array concurrently. Measure the performance difference across different designs

Practice problem 2

In Facebook, each user has a list of friends (note that in Facebook, friend is a symmetric relationship). Friends are stored as User -> [List of friends]. For example, user A has friends B C and D, then it will be stored as A -> B C D. Design a pair of map and reduce function that calculates the number of friends in common based on friend lists.

Sample input:

```
A -> B C D
B -> A C D E
C -> A B D E
D -> A B C E
E -> B C D
```

Sample output:

```
(A B) -> (C D)
(A C) -> (B D)
(A D) -> (B C)
(B C) -> (A D E)
.....
```

Practice problem 3

Create a program that using threads accepts concurrent TCP connections which provide integers separated by new lines. From the connections calculate the average of all numbers across all connections.