

# CMPUT 379 Lab

ETLC E1003: Tuesday, 5:00 – 7:50 PM.

Tianyu Zhang, Peiran Yao

CAB 311: Thursday, 2:00 – 4:50 PM.

Max Ellis, Aidan Bush

# Last Week...

- Review on Threads and Thread safe data structures
- Assignment 2: Thread pool and MapReduce

# Today's Lab

- Threading examples
- Fine and coarse grained locks

## Example: joining

- Suppose we'd like to wait until a thread finishes.
- And there is no **pthread\_wait()**

# Joining - Spin Lock

- Simple solution: use global variable and while loop
- But it wastes CPU

```
void *child(void *arg) {  
    printf("child\n");  
    sleep(5);  
    done = 1;  
    return NULL;  
}
```

```
int main(int argc, char *argv[]) {  
    pthread_t p;  
    printf("parent: begin\n");  
    Pthread_create(&p, NULL, child,  
    NULL);  
    while (done == 0)  
        ; // spin  
    printf("parent: end\n");  
    return 0;  
}
```

# Joining - Condition Variable

- Fix: use condition variable to send signals

```
void *child(void *arg) {  
    printf("child\n");  
    sleep(1);  
    Mutex_lock(&m);  
    done = 1;  
    Cond_signal(&c);  
    Mutex_unlock(&m);  
    return NULL;  
}
```

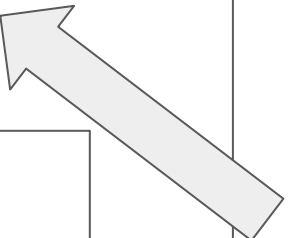
```
int main(int argc, char *argv[]) {  
    pthread_t p;  
    printf("parent: begin\n");  
    Pthread_create(&p, NULL, child,  
    NULL);  
    Mutex_lock(&m);  
    while (done == 0)  
        Cond_wait(&c, &m); // releases  
    lock when going to sleep  
    Mutex_unlock(&m);  
    printf("parent: end\n");  
    return 0;  
}
```

# Joining - Why Use Lock?

- **done** might be changed to 1, and the main thread will miss the signal before it starts waiting

```
void *child(void *arg) {  
    printf("child: begin\n");  
    sleep(1);  
    done = 1; 3  
    printf("child: signal\n");  
    Cond_signal(&c); 4  
    return NULL;  
}
```

```
int main(int argc, char *argv[]) {  
    pthread_t p;  
    printf("parent: begin\n");  
    1 Pthread_create(&p, NULL, child, NULL);  
    Mutex_lock(&m);  
    printf("parent: check condition\n");  
    while (done == 0) { 2  
        sleep(2);  
        printf("parent: wait to be  
signalled...\n");  
        Cond_wait(&c, &m); 5  
    }  
    Mutex_unlock(&m);  
    printf("parent: end\n");  
    return 0;  
}
```



# Joining - Why Use **done**?

- The main thread may also miss the signal before waiting

```
void *child(void *arg) {  
    printf("child: begin\n");  
    Mutex_lock(&m);  
    printf("child: signal\n");  
    Cond_signal(&c);  
    Mutex_unlock(&m);  
    return NULL;  
}
```

```
int main(int argc, char *argv[]) {  
    pthread_t p;  
    printf("parent: begin\n");  
    Pthread_create(&p, NULL, child, NULL);  
    sleep(2);  
    printf("parent: wait to be  
signalled...\n");  
    Mutex_lock(&m);  
    Cond_wait(&c, &m);  
    Mutex_unlock(&m);  
    printf("parent: end\n");  
    return 0;  
}
```



# Example: Producer-Consumer

- Some producers produce items and put them in a queue
- Some consumers take items out of queue and process
- The queue has a limited size

# Producer-Consumer Challenges

- Some producers produce items and put them in a queue
  - Some consumers take items out of queue and process
  - The queue has a limited size
- 
- Concurrent access to a data structure
  - What to do when the queue is full / empty?

# Producer-Consumer Solutions

- Concurrent access to a data structure
  - Use locks to protect it
- What to do when the queue is full?
  - Producers use condition variable to wait for the queue have spaces
  - Consumers notify producers when they removed items from a full queue
- What to do when the queue is empty?
  - Consumers use condition variable to wait for the queue to be filled
  - Producers notify consumers if there is a new item

# Common Bugs - Atomicity Failure

- Check-then-use should be protected by locks.

# Common Bugs - Atomicity Failure

- Check-then-use should be protected by locks.

```
void *thread1(void *arg) {  
    printf("t1: before check\n");  
    if (thd->proc_info) {  
        printf("t1: after check\n");  
        sleep(2);  
        printf("t1: use!\n");  
        printf("%d\n", thd->proc_info->pid);  
    }  
    return NULL;  
}
```

# Common Bugs - Atomicity Failure

- Check-then-use should be protected by locks.


```
void *thread2(void *arg) {  
    printf("                t2: begin\n" );  
    sleep(1); // change to 5 to make the code "work"...  
    printf("                t2: set to NULL\n" );  
    thd->proc_info = NULL;  
    return NULL;  
}
```


# Common Bugs - Atomicity Failure


- Fix: add a lock to **proc\_info**

# Common Bugs - Ordering

- Two threads may have different order of execution

```
int main(int argc, char *argv[]) {  
    printf("ordering: begin\n");  
    mThread =   
    PR_CreateThread(mMain);  
    PR_WaitThread(mThread);  
    printf("ordering: end\n");  
    return 0;  
}
```

```
pr_thread_t *PR_CreateThread(void  
    *(*start_routine)(void *)) {  
    pr_thread_t *p =  
    malloc(sizeof(pr_thread_t));  
    p->State = PR_STATE_INIT;  
    Pthread_create(&p->Tid, NULL,  
    start_routine, NULL);  
    sleep(1);  
    return p;   
}
```

```
void *mMain(void *arg) {  
    printf("mMain: begin\n");  
    int mState = mThread->State;   
    printf("mMain: state is %d\n", mState);  
    return NULL;  
}
```



# Common Bugs - Ordering

- Two threads may have different order of execution
- Fix: use condition variables

```
void *mMain(void *arg) {  
    printf("mMain: begin\n");  
    // wait for thread structure to be  
    initialized  
    Pthread_mutex_lock(&mtLock);  
    while (mtInit == 0)  
        Pthread_cond_wait(&mtCond, &mtLock);  
    Pthread_mutex_unlock(&mtLock);  
  
    int mState = mThread->State;  
    printf("mMain: state is %d\n", mState);  
    return NULL;  
}
```

```
int main(int argc, char *argv[]) {  
    printf("ordering: begin\n");  
    mThread = PR_CreateThread(mMain);  
    // signal: thread has been created,  
    and mThread initialized  
    Pthread_mutex_lock(&mtLock);  
    mtInit = 1;  
    Pthread_cond_signal(&mtCond);  
    Pthread_mutex_unlock(&mtLock)  
    // .....  
}
```

# Common Bugs - Deadlock

- When both threads are waiting for a resource that the other is holding

```
void *thread1(void *arg) {  
    printf("t1: begin\n");  
    printf("t1: try to acquire L1...\n");  
    Pthread_mutex_lock(&L1); 1  
    printf("t1: L1 acquired\n");  
    printf("t1: try to acquire L2...\n");  
    Pthread_mutex_lock(&L2); 3  
    printf("t1: L2 acquired\n");  
    Pthread_mutex_unlock(&L1);  
    Pthread_mutex_unlock(&L2);  
    return NULL;  
}
```

```
void *thread2(void *arg) {  
    printf("t2: begin\n");  
    printf("t2: try to acquire L2...\n");  
    Pthread_mutex_lock(&L2); 2  
    printf("t2: L2 acquired\n");  
    printf("t2: try to acquire L1...\n");  
    Pthread_mutex_lock(&L1); 4  
    printf("t2: L1 acquired\n");  
    Pthread_mutex_unlock(&L1);  
    Pthread_mutex_unlock(&L2);  
    return NULL;  
}
```

# Common Bugs - Deadlock

- Fix: order of resource acquisition matters

# Locking - Fine and Coarse Grained

- Fine grained
  - Locking small sections of a data structure
  - Allows only locking what needs to be
  - Usually faster
- Coarse grained
  - Locking larger sections of a data structure
  - Locking large sections of a data structure that are not modified
  - Usually slower

# Advantages of Fine Grained Locks

- Allows multiple threads to access shared data more freely
- Improved efficiency from allowing more threads to access data in parallel

# Example - Concurrent Array

- Suppose we wish to share an array across threads
- There are possible design choices
  - 1 global mutex
  - 1 local mutex per element
  - 1 mutex per block of element

# Example - Concurrent Array

- Each thread randomly replaces an element

```
void* random_access(void * args) { // thread function
    unsigned seed = (unsigned) time(NULL); // rand() has a global lock. Use
    rand_r() instead.
    for (int i = 0; i < 1024000; i++) {
        replace(rand_r(&seed) % ARRAY_SIZE, rand_r(&seed)); // randomly
        replace an element
    }
    return NULL;
}
```

# Example - Concurrent Array

- Lock the entire array (coarse)

```
int shared_array[ARRAY_SIZE];

pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER; // 1 mutex for the entire
array

int replace(unsigned int index, int new_val) {
    int old_val;
    pthread_mutex_lock(&mutex); // lock the whole array
    old_val = shared_array[index];
    shared_array[index] = new_val;
    pthread_mutex_unlock(&mutex);
    return old_val;
}
```



# Example - Concurrent Array

- Lock single element (most fine-grained)

```
int shared_array[ARRAY_SIZE];

pthread_mutex_t mutex[ARRAY_SIZE]; // one mutex per element

int replace(unsigned int index, int new_val) {
    int old_val;
    pthread_mutex_lock(&mutex[index]); // only lock the element
    old_val = shared_array[index];
    shared_array[index] = new_val;
    pthread_mutex_unlock(&mutex[index]);
    return old_val;
}
```

# Example - Concurrent Array

- Comparing the efficiency
- 32 elements
- 16 threads

```
$ time ./fine
./fine  1.16s user 0.67s system 353% cpu 0.518 total

$ time ./coarse
./coarse  1.46s user 5.45s system 343% cpu 2.008 total
```

# FYI

- When you use a lab computer (using SSH or physically be there), please make sure to clean-up your session before logging out!
- Check any running processes by command “**ps**”, and kill all processes you created.
- We received a warning from the department regarding this. Lots of people put “dragonshell” to the background by “Ctrl+Z” and then log out with no cleanup.