

CMPUT 379 Lab

ETLC E1003: Tuesday, 5:00 – 7:50 PM.

Tianyu Zhang, Peiran Yao

CAB 311: Thursday, 2:00 – 4:50 PM.

Max Ellis, Aidan Bush

Last Week...

- Threading examples
- Fine and coarse grained locks

Today's Lab

- Semaphore
- FAQ

More example on Condition Variable

```
// increment counter a few times
// wake up watch_count thread when reaching COUNT_LIMIT
void *inc_count(void *t) {
    int my_id = *(int*)t;
    for (int i=0; i < NUM_INC; ++i) {
        pthread_mutex_lock(&count_mutex);
        count++;
        // check the value of count and signal waiting thread when
        // condition is reached. This occurs while mutex is locked
        if (count == COUNT_LIMIT) {
            pthread_cond_signal(&count_cond);
            printf("inc_count: thread %d, count = %d Threshold reached.\n", my_id, count);
        }
        printf("inc_count: thread %d, count = %d, unlocking mutex\n", my_id, count);
        pthread_mutex_unlock(&count_mutex);
        sleep(1); // do some "work" so threads can alternate on mutex lock
    }
    return nullptr;
}
```

More example on Condition Variable

```
void *watch_count(void *t) { // wait until signalled, then add 125
    int my_id = *(int*)t;
    printf("Starting watch_count: thread %d\n", my_id);
    // Lock mutex and wait for signal. pthread_cond_wait will unlock
    // mutex while it waits. Also, if COUNT_LIMIT is reached before
    // this function is run by the waiting thread, the loop will be
    // skipped to prevent pthread_cond_wait from never returning
    pthread_mutex_lock(&count_mutex);
    while (count < COUNT_LIMIT) {
        pthread_cond_wait(&count_cond, &count_mutex);
        // check whether we actually received a signal
        if (count >= COUNT_LIMIT) {
            printf("watch_count: thread %d signal received.\n", my_id);
            count += 125;
            printf("watch_count: thread %d count now = %d.\n", my_id, count);
        }
    }
    pthread_mutex_unlock(&count_mutex);
    return 0;
}
```

More example on Condition Variable

```
int main () {
    pthread_t *threads = new pthread_t[NUM_THREADS];
    int *ids = new int[NUM_THREADS];
    // initialize mutex and condition variable objects
    pthread_mutex_init(&count_mutex, 0);
    pthread_cond_init(&count_cond, 0);
    ids[0] = 0;
    pthread_create(&threads[i], nullptr, watch_count, (void *)&ids[0]);
    for (int i=1; i < NUM_THREADS; ++i) {
        ids[i] = i;
        pthread_create(&threads[i], nullptr, inc_count, (void *)&ids[i]);
    }
    // wait for all threads to complete
    for (int i=0; i < NUM_THREADS; ++i)
        pthread_join(threads[i], nullptr);
    printf("Main(): Waited on %d threads. Done.\n", NUM_THREADS);
    // clean up and exit
    pthread_mutex_destroy(&count_mutex);
    pthread_cond_destroy(&count_cond);
    delete [] threads;
    delete [] ids;
    return 0;
}
```

More example on Condition Variable

```
inc_count: thread 1, count = 1, unlocking mutex
Starting watch_count: thread 0
inc_count: thread 2, count = 2, unlocking mutex
inc_count: thread 1, count = 3, unlocking mutex
inc_count: thread 2, count = 4, unlocking mutex
inc_count: thread 1, count = 5, unlocking mutex
inc_count: thread 2, count = 6, unlocking mutex
inc_count: thread 1, count = 7, unlocking mutex
inc_count: thread 2, count = 8, unlocking mutex
inc_count: thread 1, count = 9, unlocking mutex
inc_count: thread 2, count = 10, unlocking mutex
inc_count: thread 1, count = 11, unlocking mutex
```

```
inc_count: thread 2, count = 12 Threshold reached.
inc_count: thread 2, count = 12, unlocking mutex
watch_count: thread 0 signal received.
watch_count: thread 0 count now = 137.
inc_count: thread 1, count = 138, unlocking mutex
inc_count: thread 2, count = 139, unlocking mutex
inc_count: thread 1, count = 140, unlocking mutex
inc_count: thread 2, count = 141, unlocking mutex
inc_count: thread 1, count = 142, unlocking mutex
inc_count: thread 2, count = 143, unlocking mutex
inc_count: thread 1, count = 144, unlocking mutex
inc_count: thread 2, count = 145, unlocking mutex
Main(): Waited on 3 threads. Done.
```

Semaphore

Semaphore

- Combines the functionalities of a mutex and a condition variable
- Can be shared across processes

Semaphore - features

- A combination of an unsigned integer **x**, a lock, and a condition variable
- **int sem_init(sem_t *sem, int pshared, unsigned int value);**

```
x = value;
```

- **int sem_wait(sem_t *sem);**

```
if (x == 0) {  
    wait();  
}  
x--;
```

- **int sem_post(sem_t *sem);**

```
x++;
```

x is automatically protected by locking, and wait() is automatically managed with mechanisms similar to condition variable

Semaphore - replacing mutex

- Set the initial value to 1
- Lock: **sem_wait()**
- Unlock: **sem_post()**

```
void *child(void *arg) {  
    int i;  
    for (i = 0; i < 10000000; i++) {  
        sem_wait(&mutex);  
        counter++;  
        sem_post(&mutex);  
    }  
    return NULL;  
}
```

```
int main(int argc, char *argv[]) {  
    sem_init(&mutex, 0, 1);  
    pthread_t c1, c2;  
    pthread_create(&c1, NULL, child, NULL);  
    pthread_create(&c2, NULL, child, NULL);  
    pthread_join(c1, NULL);  
    pthread_join(c2, NULL);  
    printf("result: %d (should be  
20000000)\n", counter);  
    return 0;  
}
```

Semaphore - replacing condition variable

- Set the initial value to 0
- Wait: **sem_wait()**
- Signal: **sem_post()**

```
void *child(void *arg) {  
    sleep(2);  
    printf("child\n");  
    sem_post(&s); // signal here: child  
is done  
    return NULL;  
}
```

```
int main(int argc, char *argv[]) {  
    sem_init(&s, 0, 0);  
    printf("parent: begin\n");  
    pthread_t c;  
    pthread_create(&c, NULL, child, NULL);  
    sem_wait(&s); // wait here for child  
    printf("parent: end\n");  
    return 0;  
}
```

Semaphore - throttling

- Control how many threads can access a resource / perform an action at the same time
- Initialize the value to the maximum concurrent accesses allowed

Frequently asked questions

Coding & Debugging

Why can't my code be compiled?

- Inspect the error message. (Though not easy for C++)
- Are you using the wrong compiler / standard? (Use **gcc** for C and use **g++** for C++, use flags **-std=c++11** / **-std=c++14** / **-std=c99** as needed)
- Incompatible types
 - Confusing value types, pointers and references
 - Mixing up STL string (**std::string**), C string (**char***), character (**char**) and array of C strings (**char****)
 - Stick with the language you are familiar with

Coding & Debugging

Why can't my code be linked?

- Check the command line arguments
 - Add **-c** for creating object files
 - **-o** should directly follow output file
 - One and only one source file should contain **main()**
 - Mixing C and C++
 - Use **g++** to compile C code
 - Or, disable name mangling in C++
- <https://eclass.srv.ualberta.ca/mod/page/view.php?id=3849191>

Coding & Debugging

Why do I encounter segmentation faults?

- Your code / the external function you called is accessing invalid memory
- Use **gdb** or **valgrind** to locate the cause

```
$ gcc dragonshell.c -g      (Add -g to compiler options)
$ gdb ./a.out              (Open gdb)
(gdb) run                  (run your program in gdb)
.....                    (some output)
Segmentation fault         (segmentation fault happened)
(gdb) where                 (locate the cause)
```

Coding & Debugging

Why is my code not working as expected?

- Should I use passing by value / passing by reference?
- Check variable scope
 - Anything between { and } is a block
 - Variables declared inside a blocked can not be accessed outside
- Check errors of system calls
- Add assertions


```
#include <cassert>
// .....
assert(x >= 0);
```

```
Assertion failed: (x < 0),
function main, file assert.c,
line 6.
```

Coding & Debugging

Why is my code not working as expected?


- Track the execution of your program line by line
- Or, print debugging messages 🤪
(although make them meaningful)



Coding & Debugging

Why is my code not working as expected?

- Try to explain your code to yourself, your friend, or to a lovely rubber duck
- Find out more:
<https://rubberduckdebugging.com>
- [Talk to a duck online](#)
- Don't make assumptions about how it behaves when reading through code



```
You are now in an online chat room with:

Cyberduck

- Cyberduck is a mock Rogerian psychotherapist (for debugging).
- The original concept program was described by Joseph Weizenbaum in 1966.
- Terminal themed implementation was created by Norbert Landsteiner in 2005.
- This debugging themed implementation by Ben Holland in 2017.

CYBERDUCK: Hello. Please tell me your problem. Be specific and go line by line.
YOU: I'm stuck
CYBERDUCK: Do you know anyone else who is stuck ?
YOU: No
CYBERDUCK: You are being a bit negative.
YOU: My program is strange.
CYBERDUCK: Why do you say your program is strange ?
YOU: █
```

Valgrind

Do I need to check for memory leaks?

- Yes, because
- It counts towards your marks
- It helps you resolve your bugs

Valgrind

Why can't I get rid of “still reachable” leaks?

- You (or some libraries) hold pointers to some memory that could be freed
- C++ STL / runtime will use memory pools to allocate extra memory ahead of time. They are not counted as memory leaks in our assignments. Check the Valgrind report to see where the memory was allocated.

Resources

Where can I find tutorials for beginners on pthread (and other topics)?

- [Operating Systems: Three Easy Pieces](#)
- Advanced Programming in a Unix Environment

Assignment 2

Can I modify threadpool.h?

- Yes you can freely change type definitions and include other header files
- If you change function signatures, you have to justify your modifications in the design document
- We won't test your thread pool. But your MapReduce library must use your own thread pool implementation

Assignment 2

Can I modify mapreduce.h?

- No and you may need to make sure the original **distwc.c** works

Assignment 2

Do I need to handle errors?

- You are encouraged to do so, but it won't be tested (unless mentioned in the starter code or description)
- If you have a way to handle errors, you may briefly describe them in your design documents

More questions?