

# Operating System Concepts

## Lecture 2: Interactions between OS and Hardware

Omid Ardakanian  
[oardakan@ualberta.ca](mailto:oardakan@ualberta.ca)  
University of Alberta

MWF 12:00-12:50 VVC 2 215

# Today's class

---

- A closer look at OS
  - Why do we study it?
- Interactions between Operating System and hardware
  - How does OS make it easier to use/program a computer?

# Why do we study Operating Systems?

---

- knowledge of how OS works is crucial to proper, efficient, and secure programming
  - OS code is **efficient** (runs fast using minimum resources), **reliable** (fails infrequently), **complex** (addresses various timing and concurrency issues), and **secure**

# Why do we study Operating Systems?

---

- knowledge of how OS works is crucial to proper, efficient, and secure programming
  - OS code is **efficient** (runs fast using minimum resources), **reliable** (fails infrequently), **complex** (addresses various timing and concurrency issues), and **secure**
- studying OS is studying the design of large software systems
  - Windows Vista is more than 50 million lines of code

# Why do we study Operating Systems?

---

- knowledge of how OS works is crucial to proper, efficient, and secure programming
  - OS code is **efficient** (runs fast using minimum resources), **reliable** (fails infrequently), **complex** (addresses various timing and concurrency issues), and **secure**
- studying OS is studying the design of large software systems
  - Windows Vista is more than 50 million lines of code
- studying OS makes you a better programmer!
  - teaches you how to make tradeoffs between performance and usability, performance and design simplicity, ....
  - teaches you how to manage complexity through appropriate abstractions (e.g., file system, process)

# Why do we study Operating Systems?

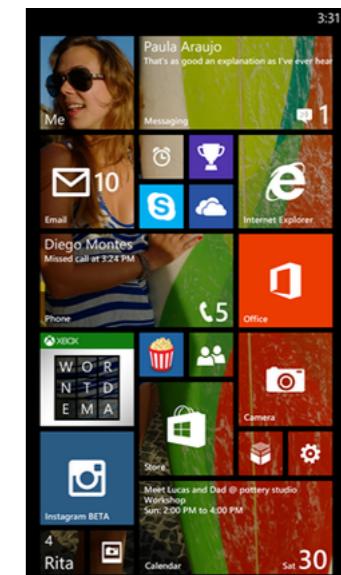
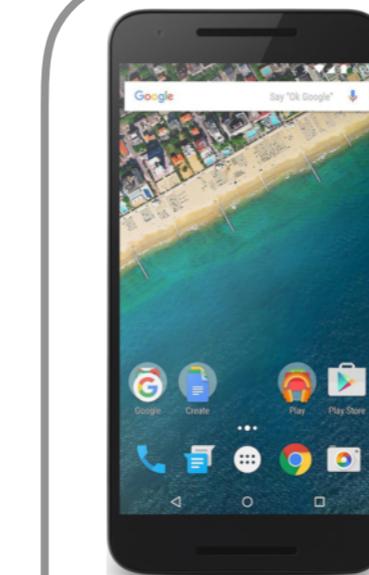
---

- knowledge of how OS works is crucial to proper, efficient, and secure programming
  - OS code is **efficient** (runs fast using minimum resources), **reliable** (fails infrequently), **complex** (addresses various timing and concurrency issues), and **secure**
- studying OS is studying the design of large software systems
  - Windows Vista is more than 50 million lines of code
- studying OS makes you a better programmer!
  - teaches you how to make tradeoffs between performance and usability, performance and design simplicity, ....
  - teaches you how to manage complexity through appropriate abstractions (e.g., file system, process)
- it's an active area of research!  **SIGOPS**
  - new application spaces: cloud, mobile systems, embedded systems, cyber-physical systems
  - SOSP 2019: 27th ACM Symposium on Operating Systems Principles
  - OSDI 2018: 13th USENIX Symposium on Operating Systems Design and Implementation

# And because Operating Systems are everywhere!



**Desktop & Laptop Computers**



**Mobile Devices**



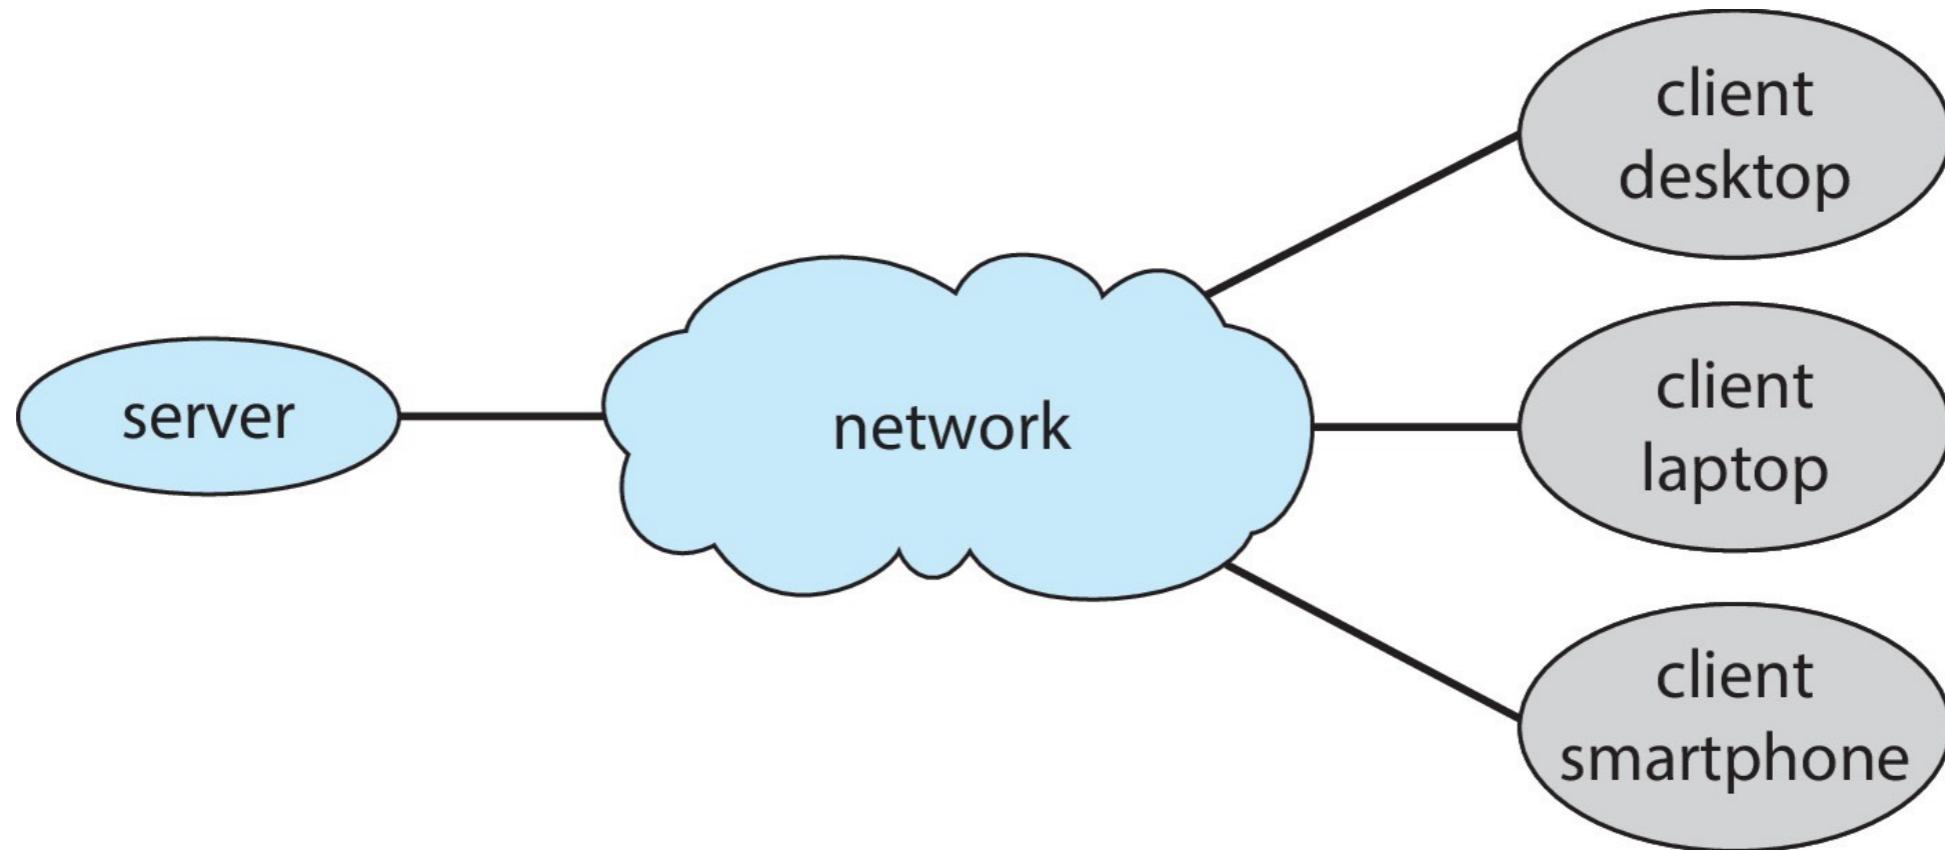
**Embedded Computers**



**Wearables**

# An example: client-server model

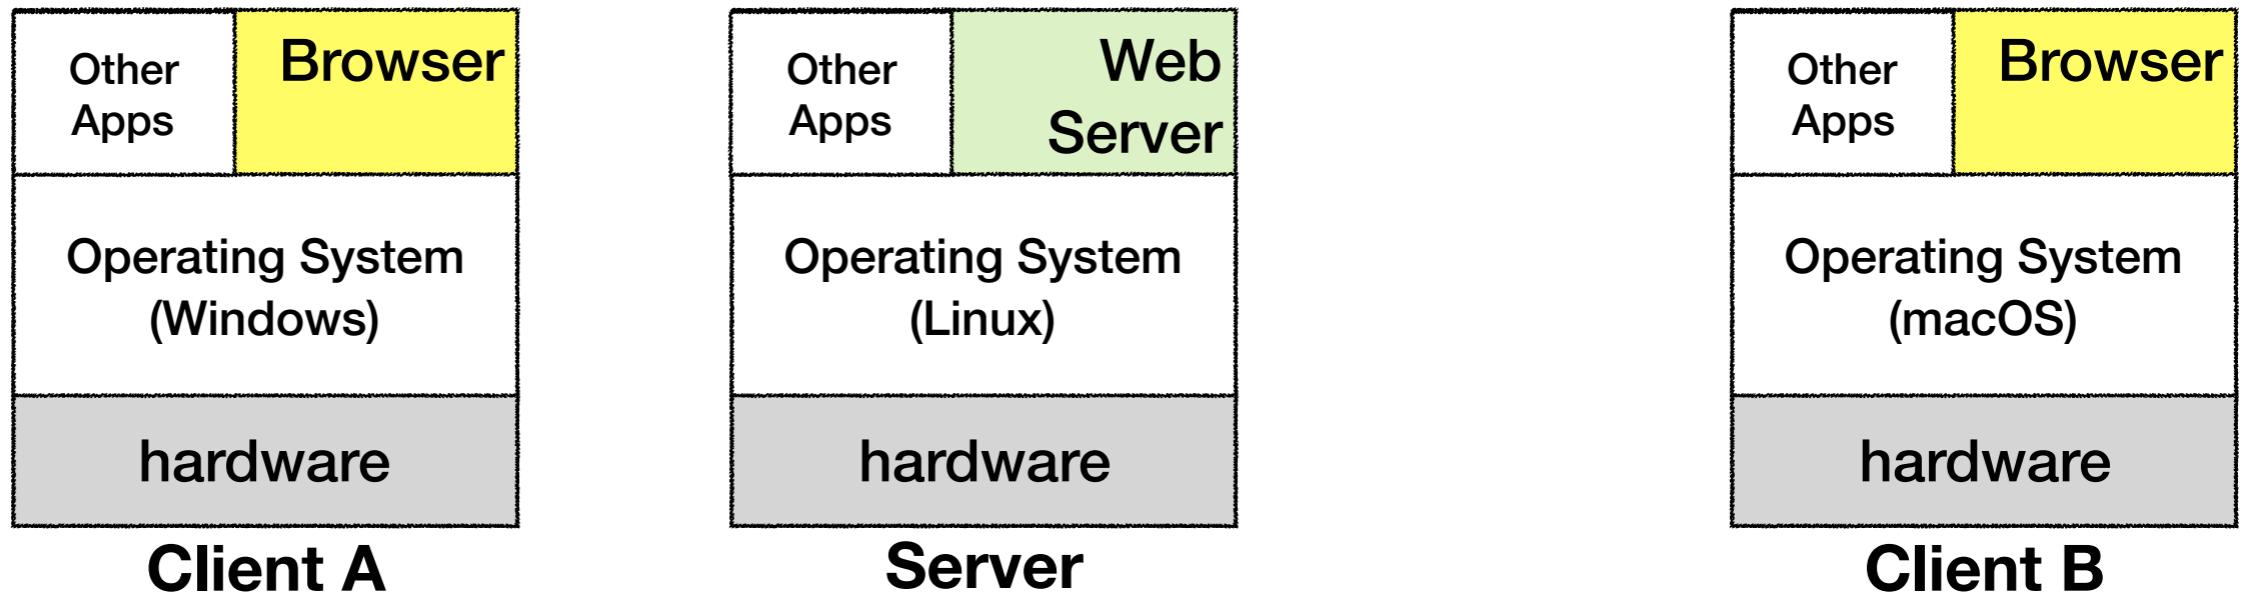
---



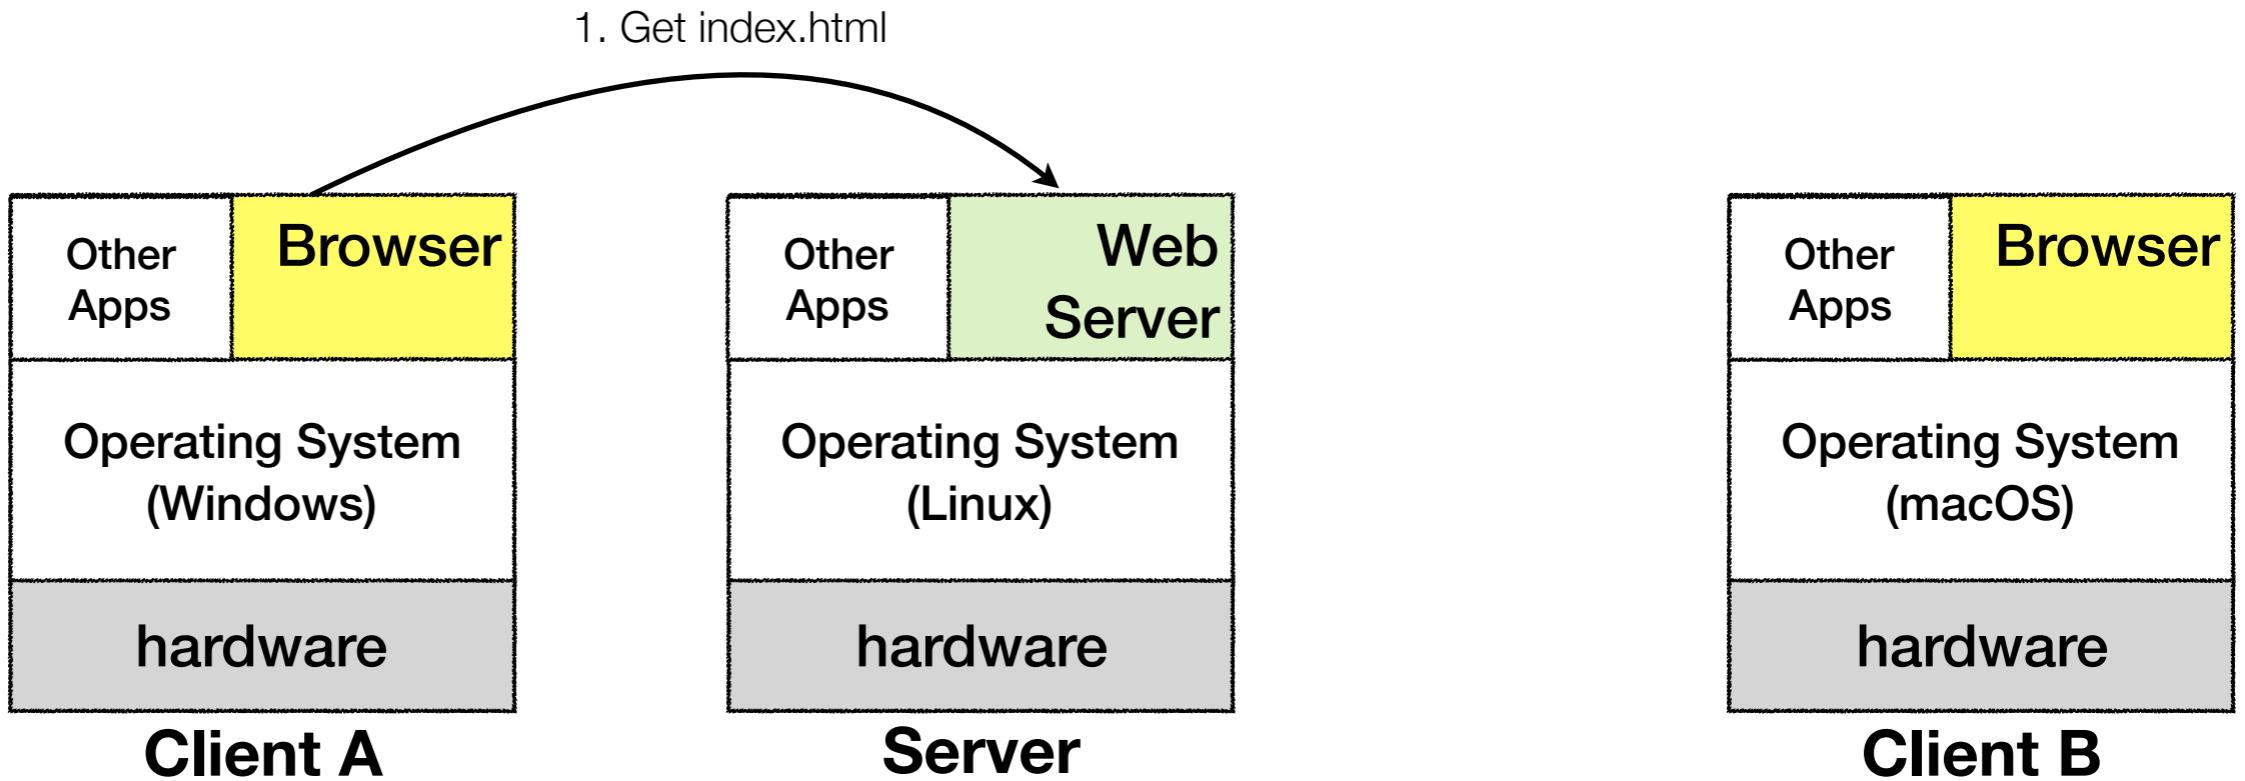
- each client can send a request to the server to perform some operation (e.g., running a query)
- the server executes the operation upon receiving the request, and sends the results back to the client

# An example: client-server model

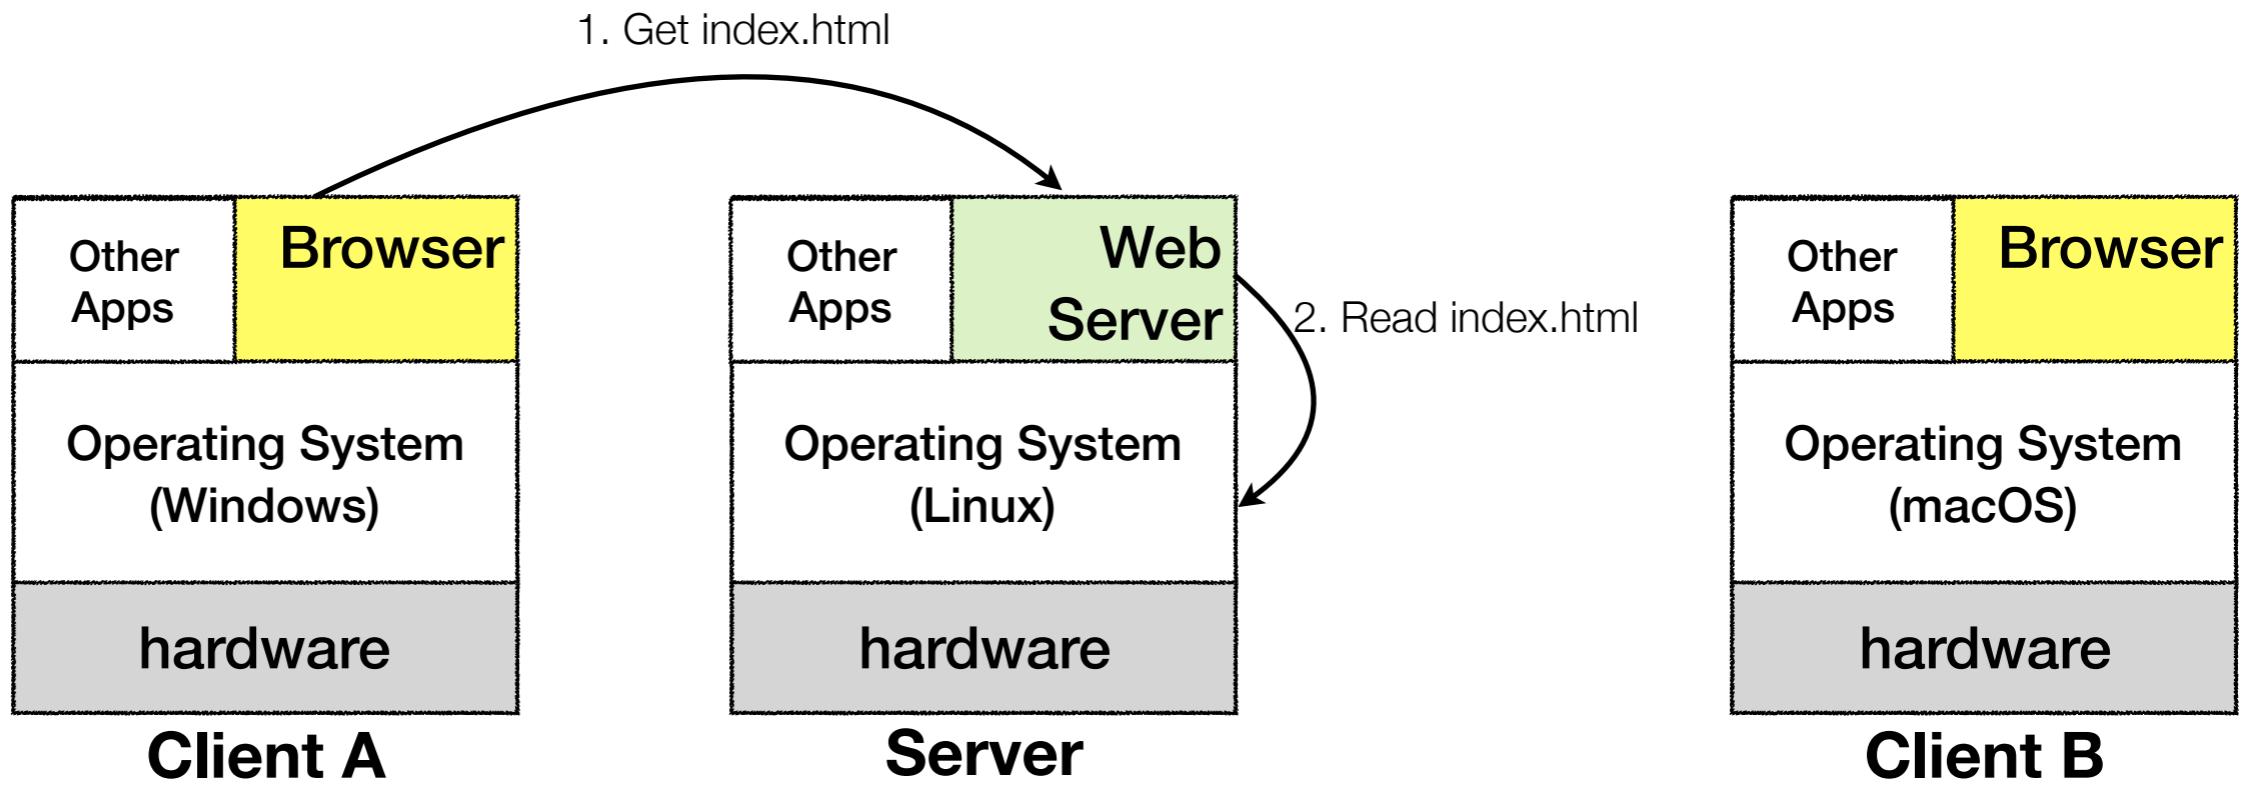
---



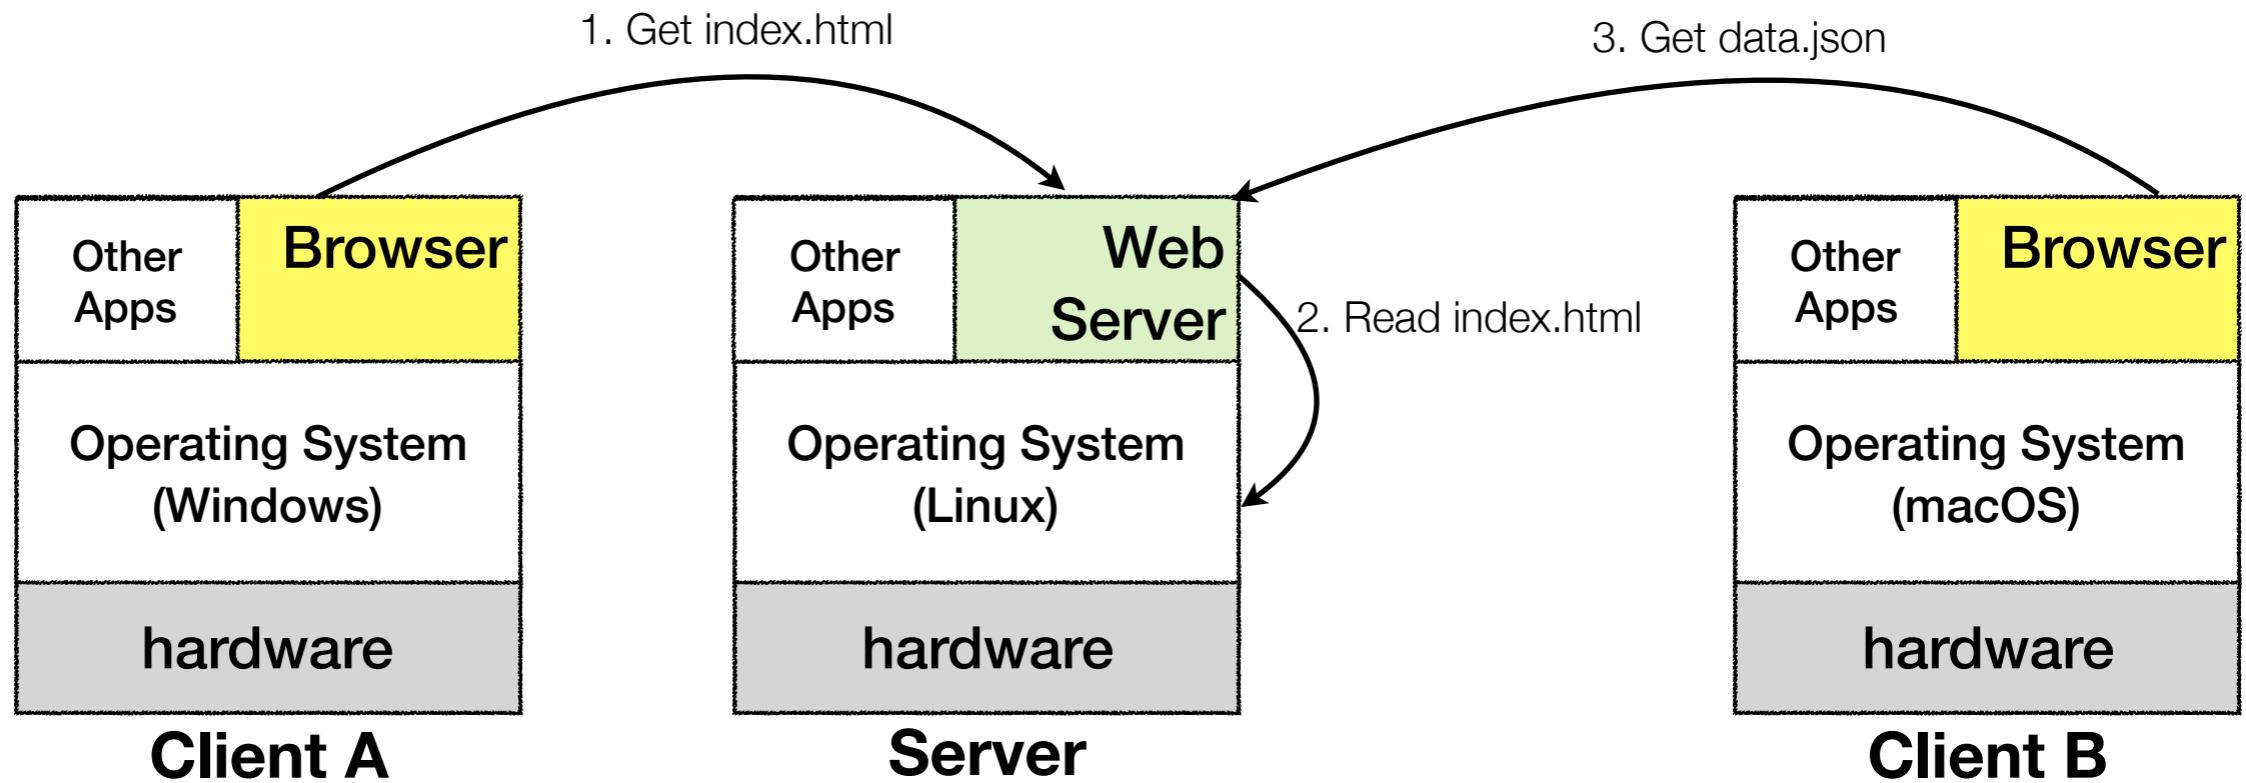
# An example: client-server model



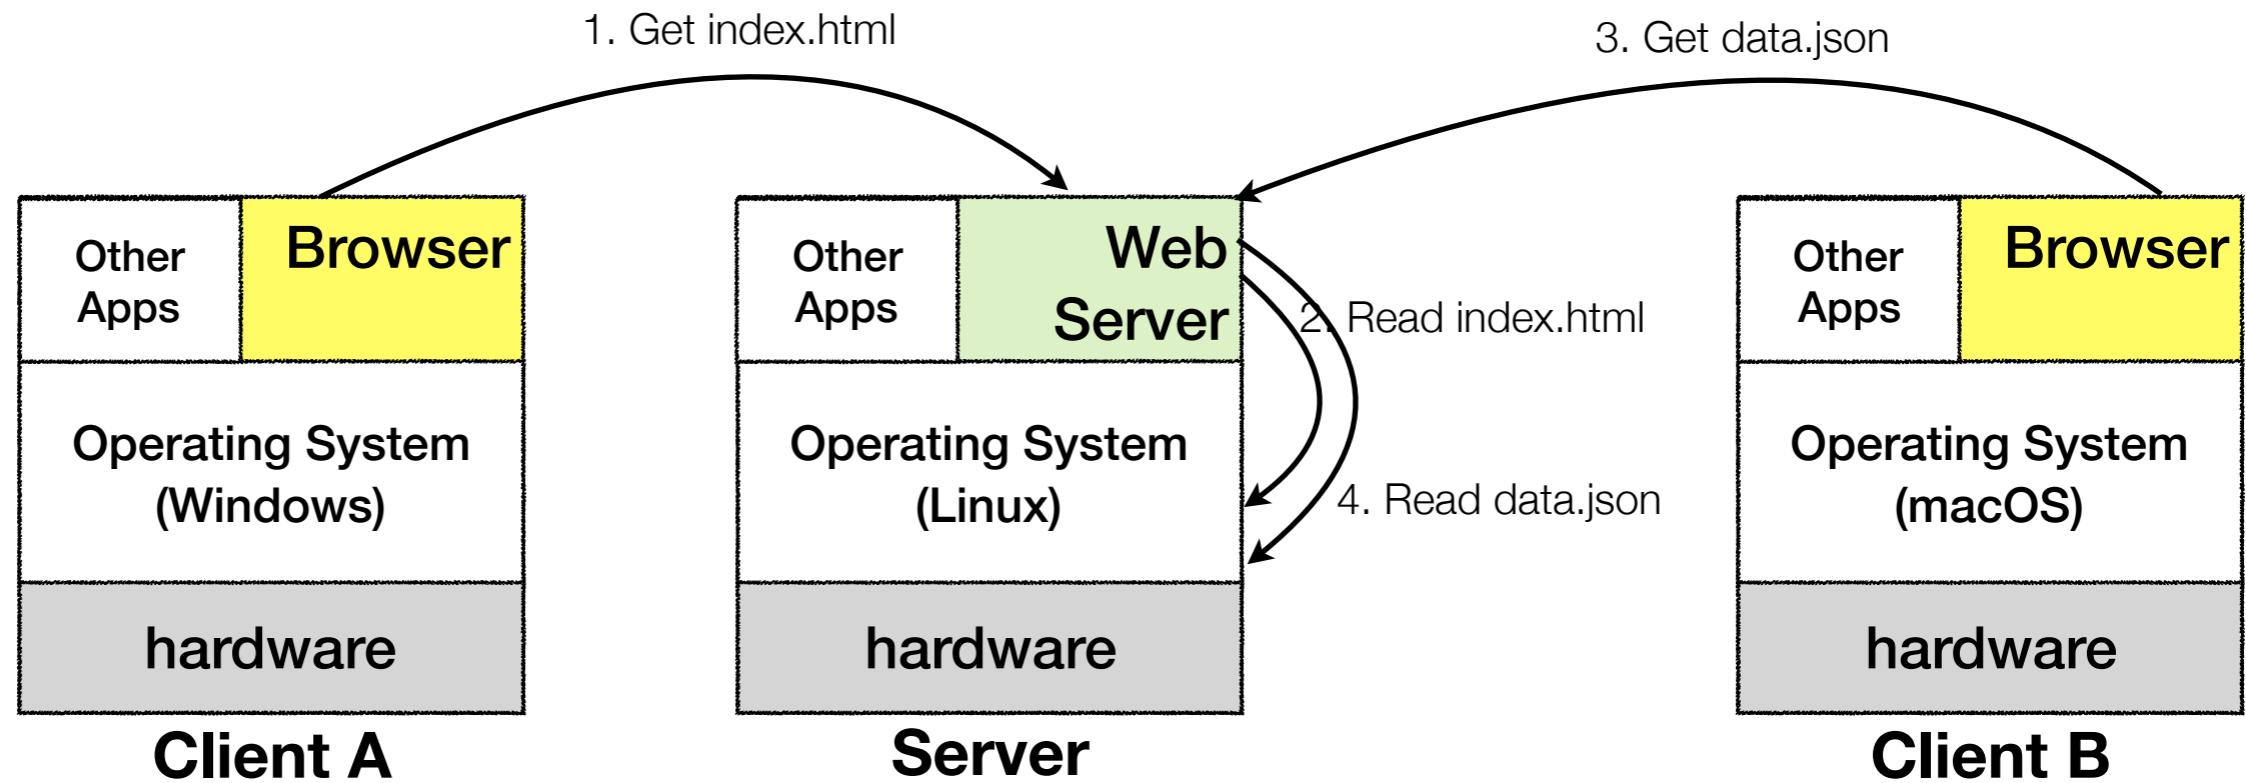
# An example: client-server model



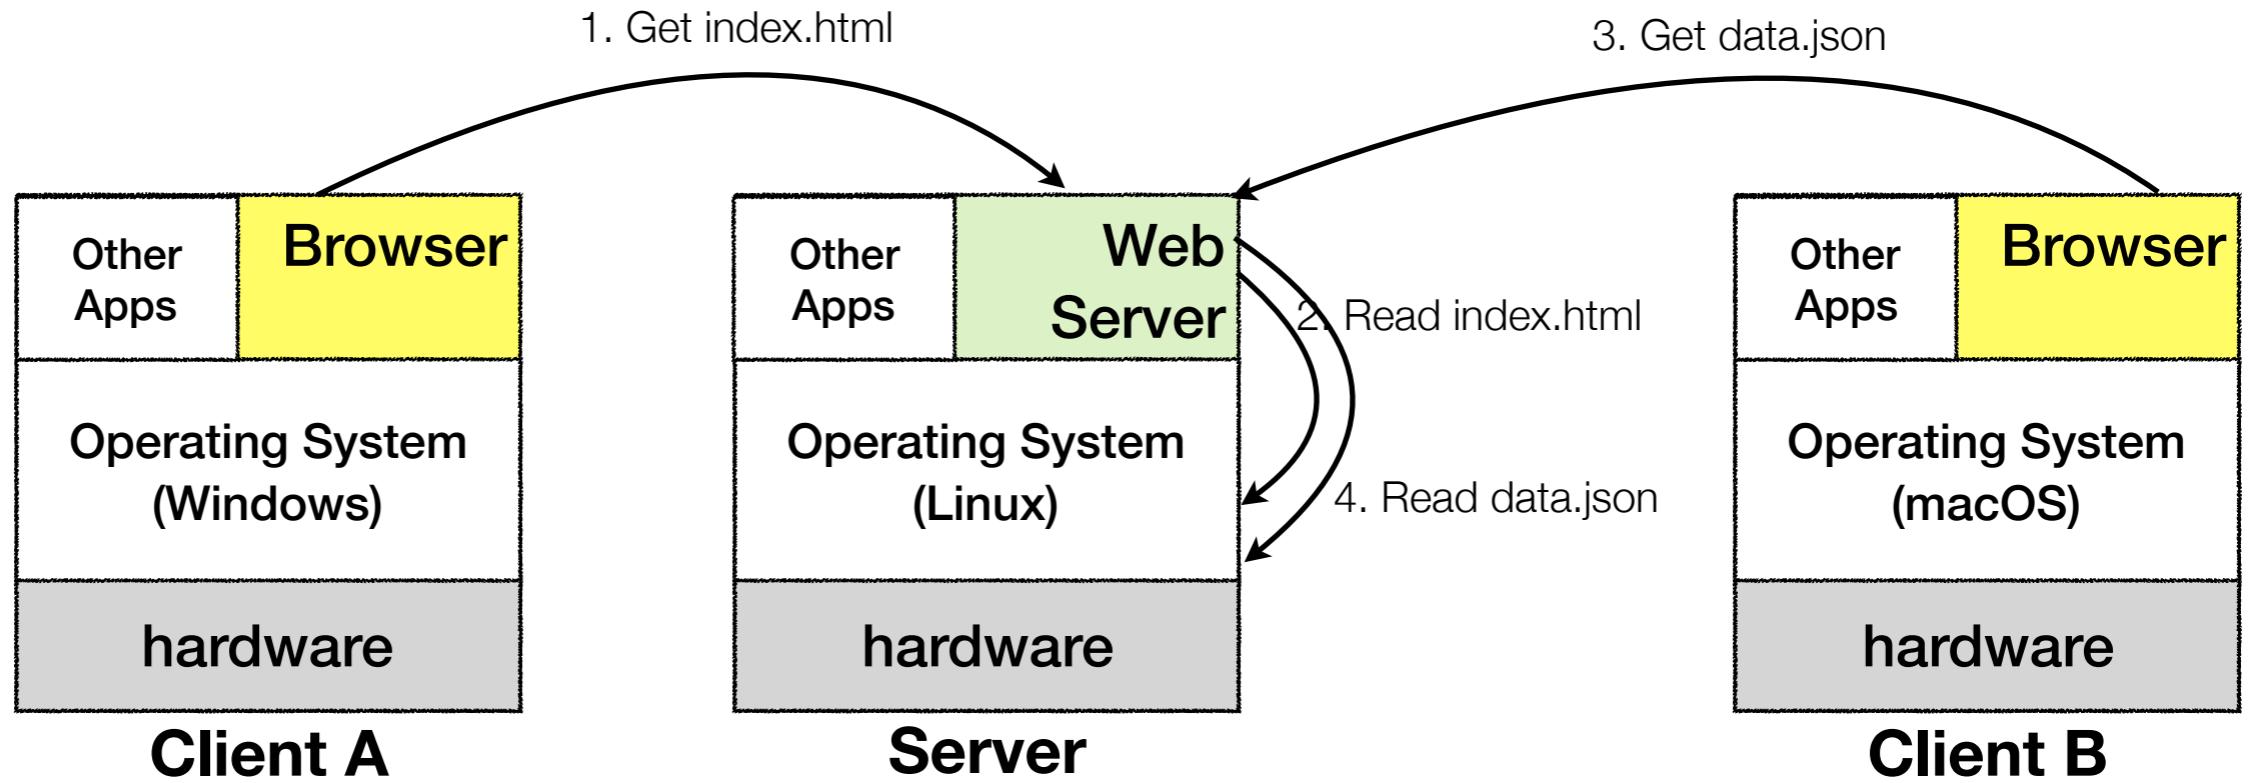
# An example: client-server model



# An example: client-server model



# An example: client-server model



- In this course you will learn about how the OS
  - allows multiple local and remote user applications to communicate with each other and share data
  - handles multiple concurrent requests
  - supports access to shared data using atomic operations
  - protects the system against malicious scripts

# Interfacing Hardware

# Recap

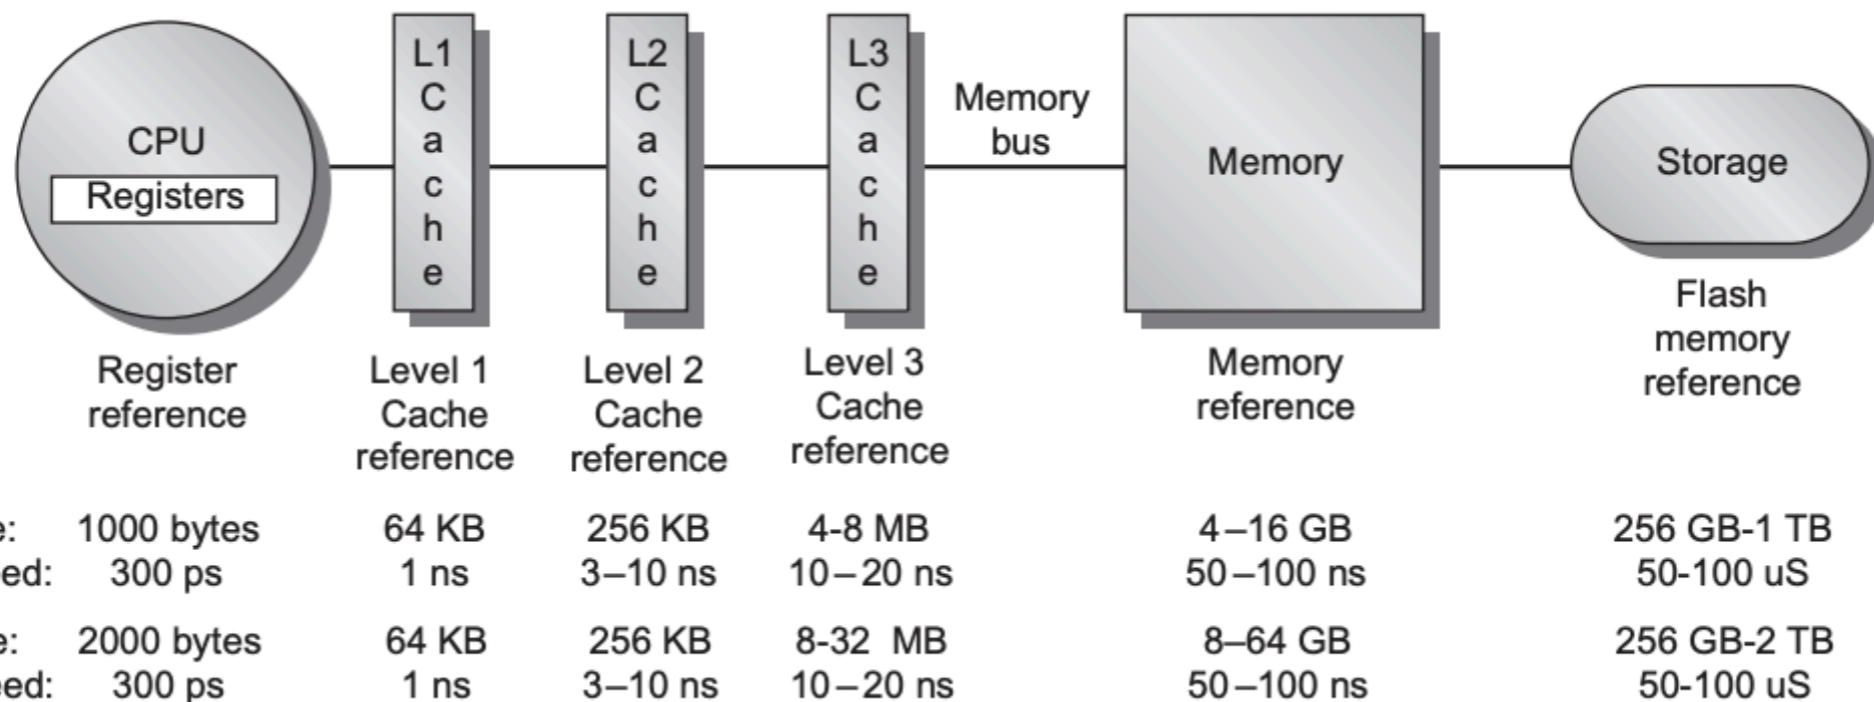
---

- CPU fetches and runs instructions one at a time

# Recap

---

- CPU fetches and runs instructions one at a time
  - Registers are CPU's workspace

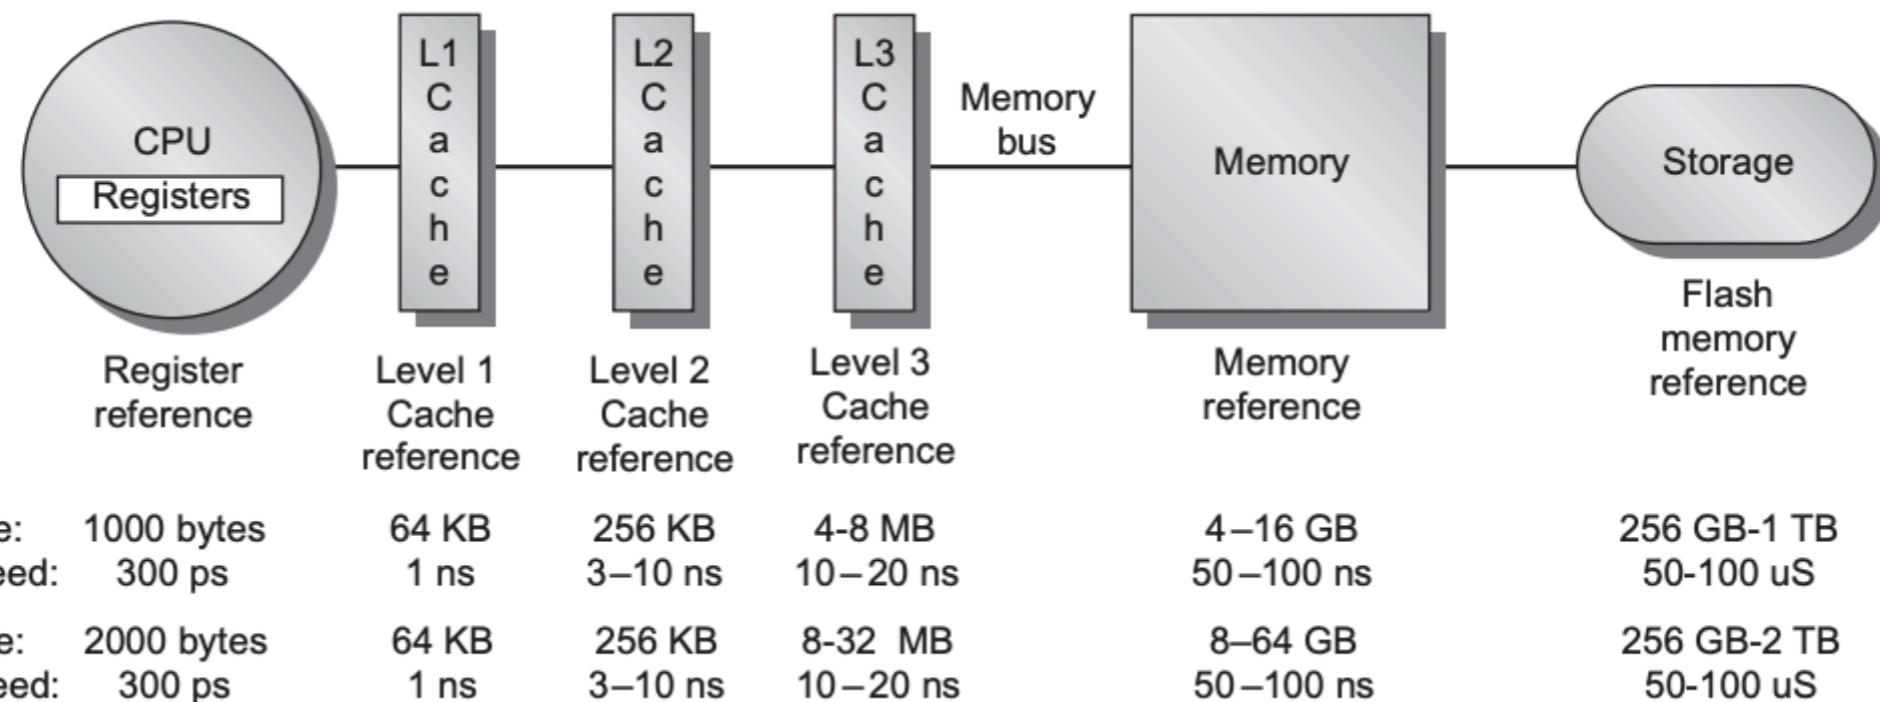


Memory hierarchy for a desktop and a laptop

# Recap

---

- CPU fetches and runs instructions one at a time
  - Registers are CPU's workspace
- Programmers desire a large amount of fast memory

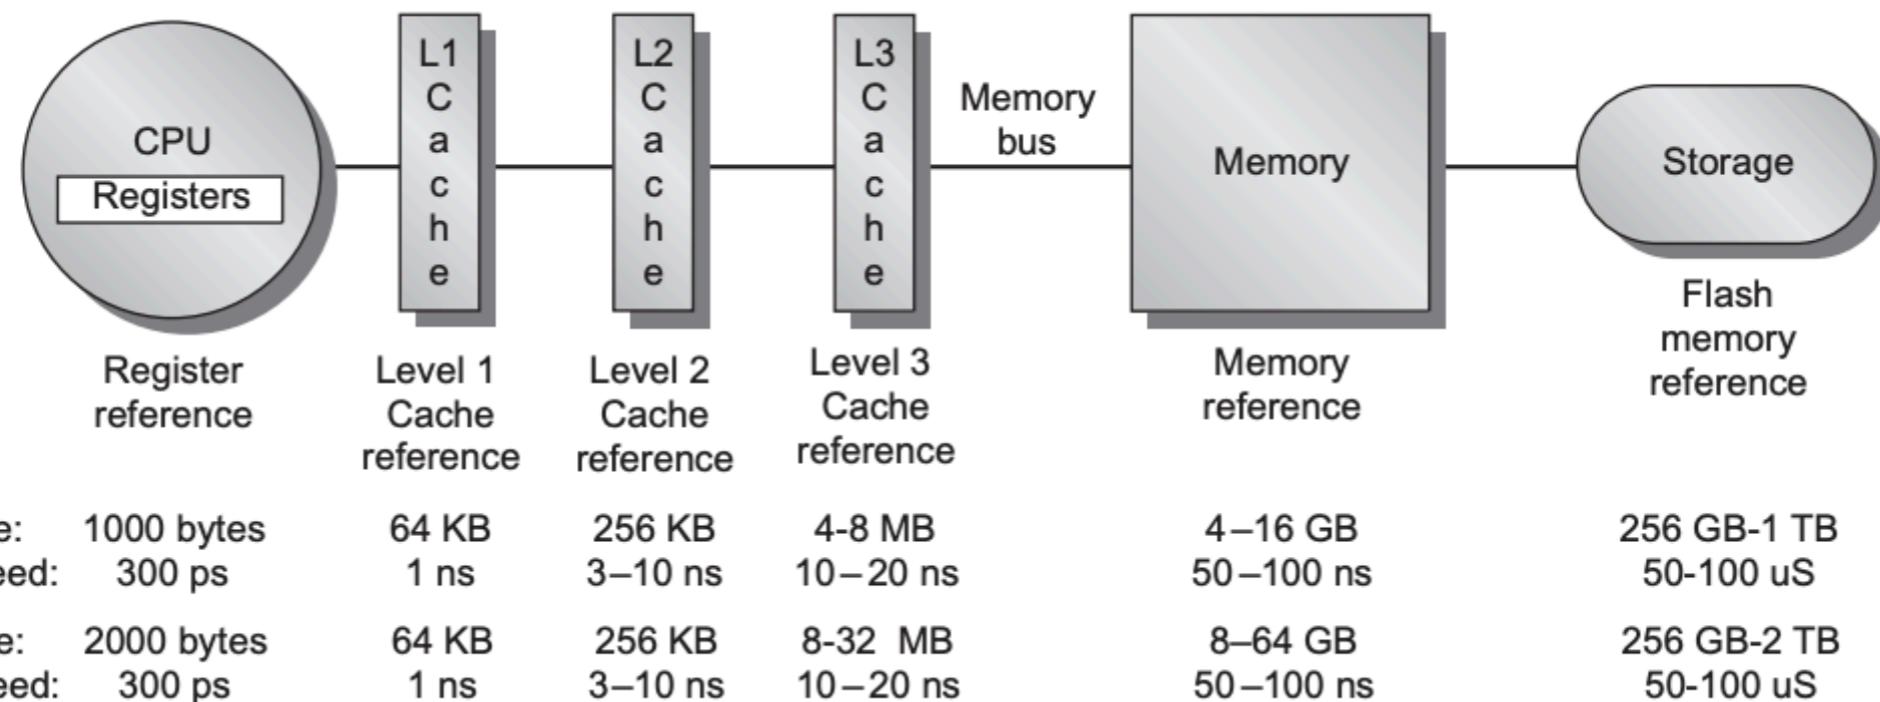


Memory hierarchy for a desktop and a laptop

# Recap

---

- CPU fetches and runs instructions one at a time
  - Registers are CPU's workspace
- Programmers desire a large amount of fast memory
  - an economical solution: a memory hierarchy organized into several levels; each level is smaller, faster, and more expensive per byte than the next

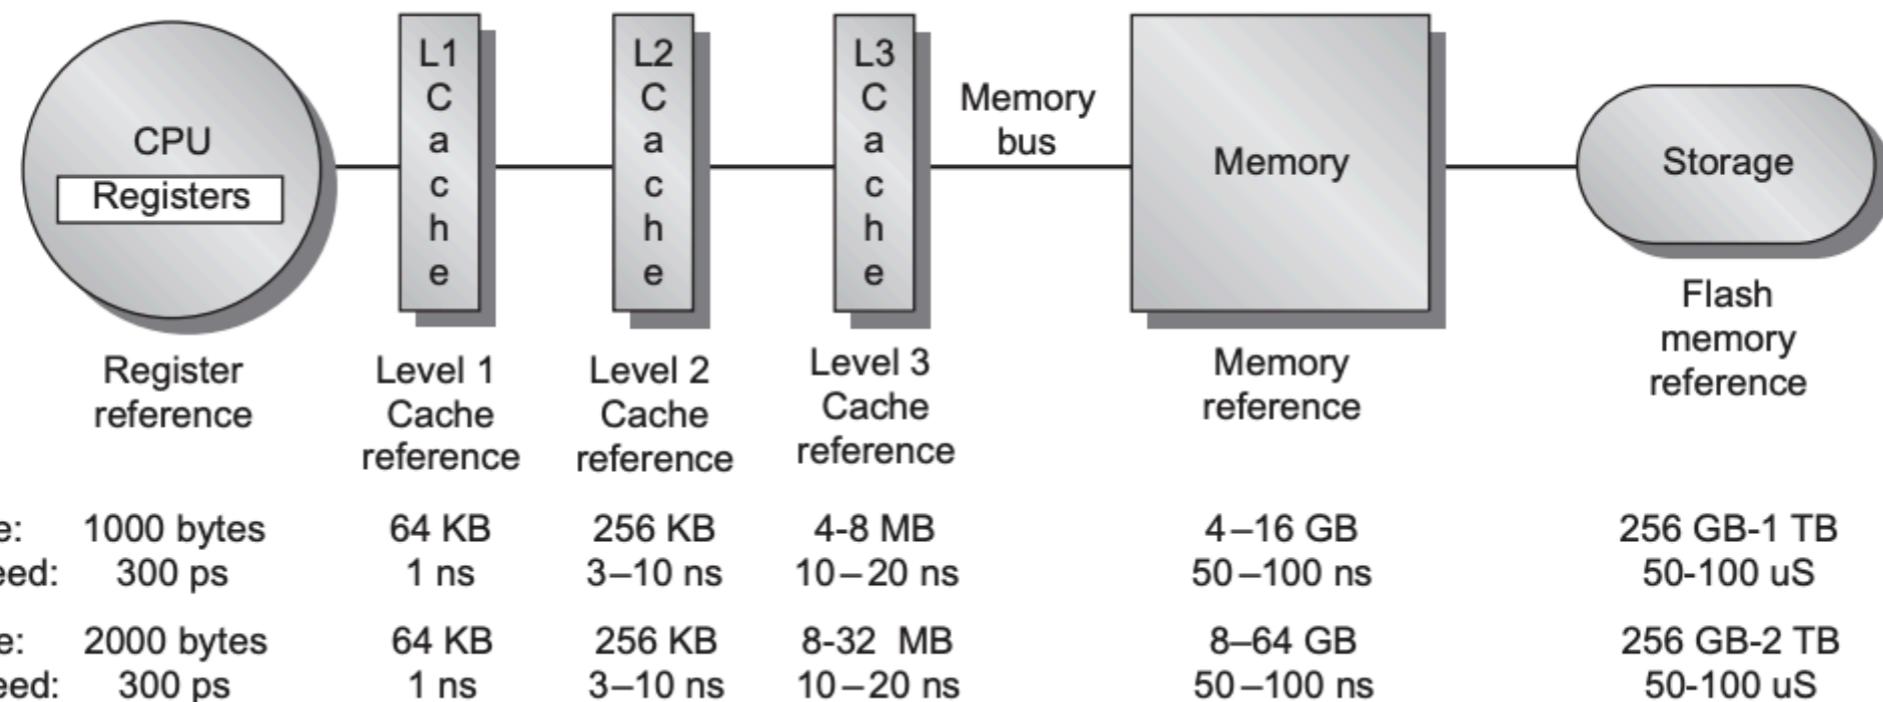


Memory hierarchy for a desktop and a laptop

# Recap

---

- CPU fetches and runs instructions one at a time
  - Registers are CPU's workspace
- Programmers desire a large amount of fast memory
  - an economical solution: a memory hierarchy organized into several levels; each level is smaller, faster, and more expensive per byte than the next
  - it takes advantage of **locality** and **cost-performance tradeoffs**



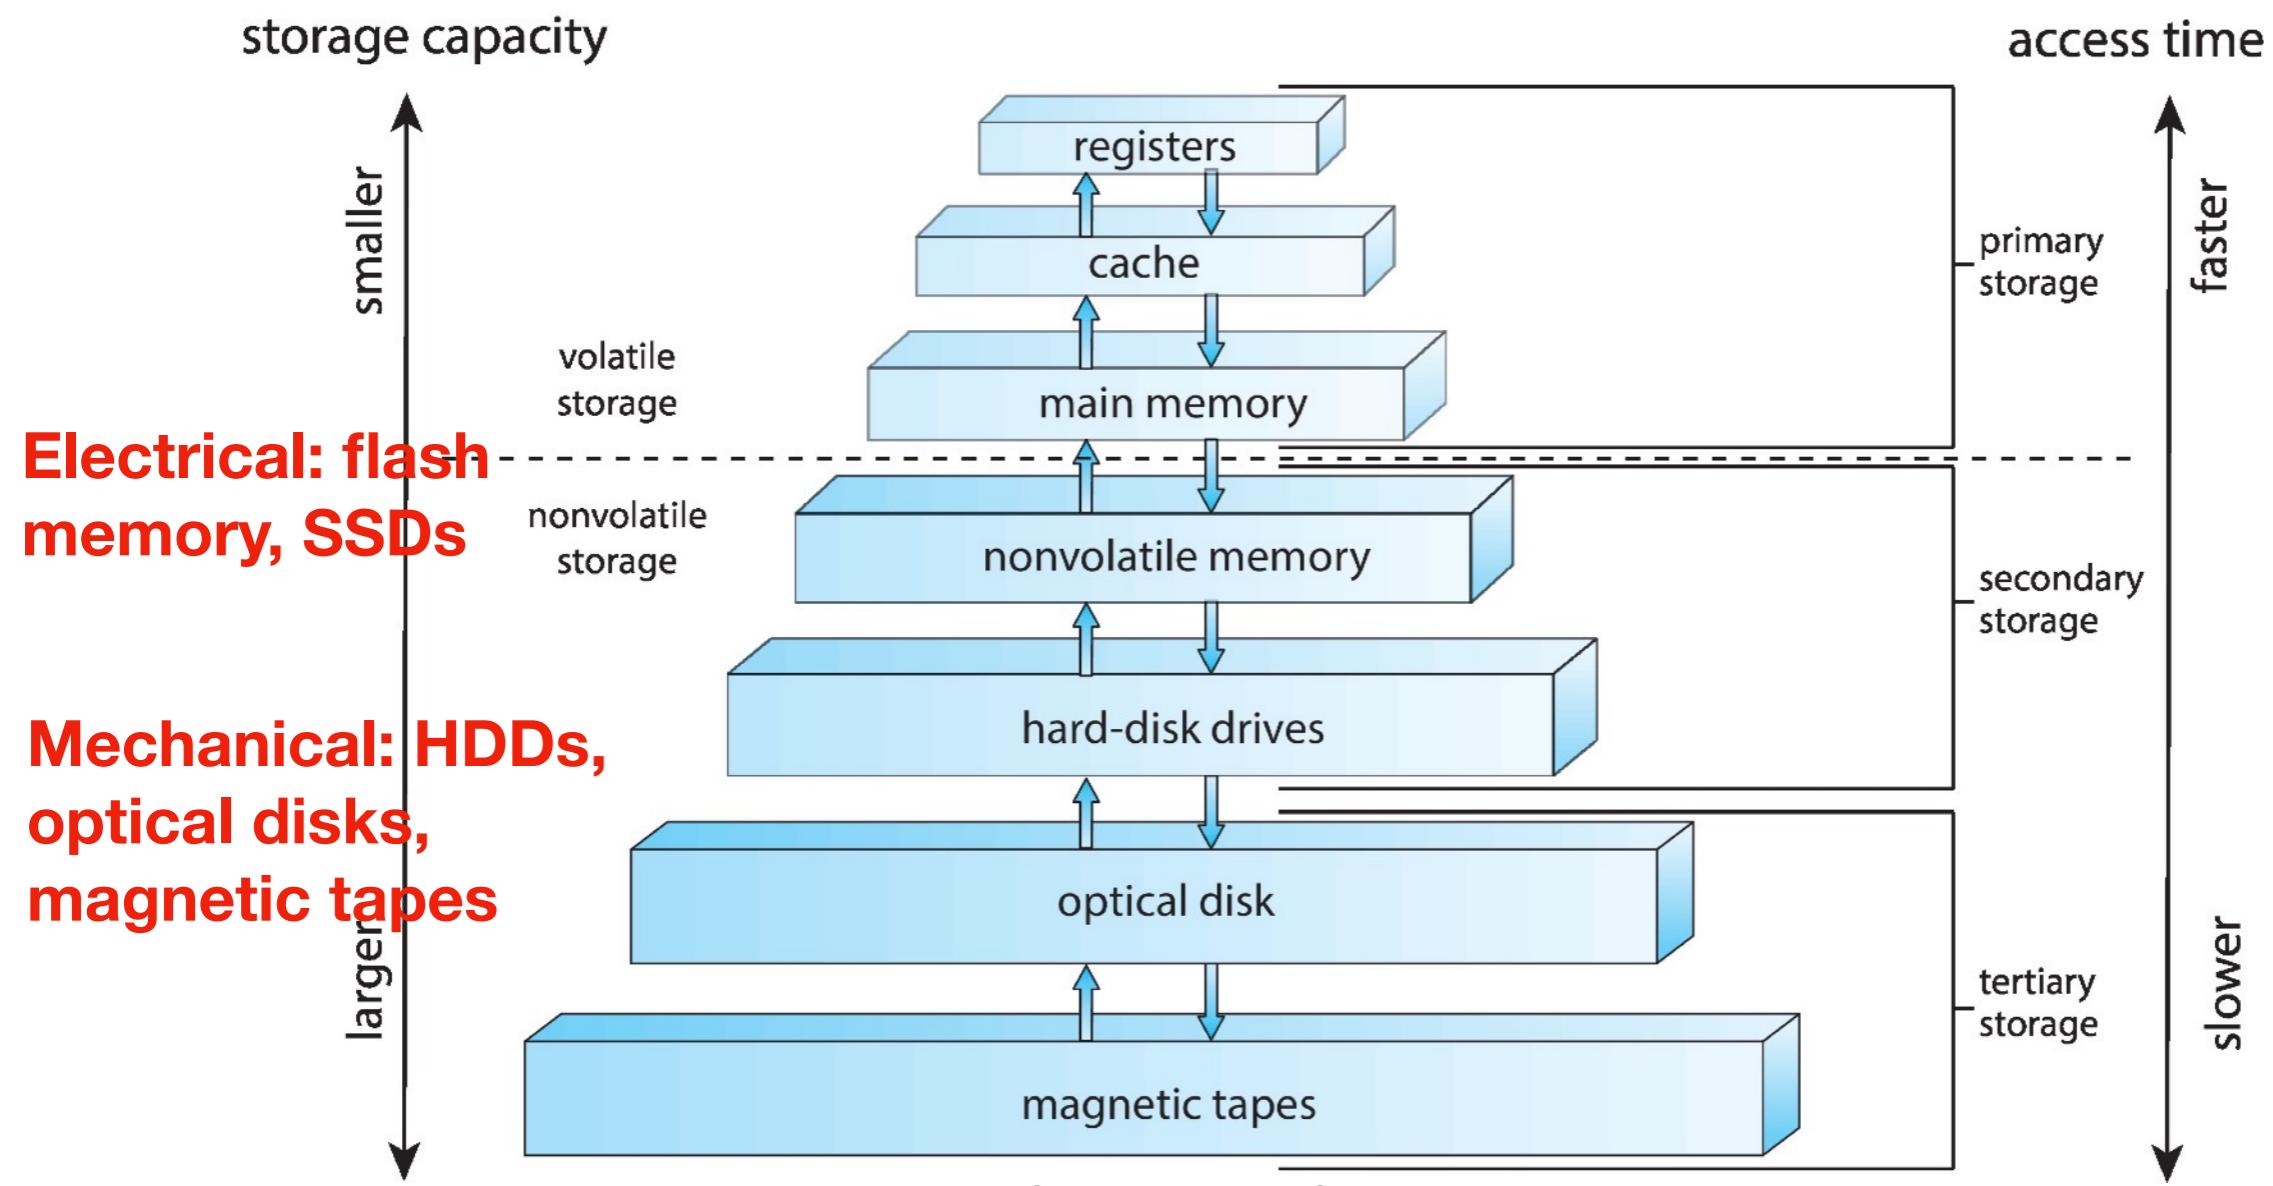
10<sup>-12</sup>

10<sup>-6</sup>

Memory hierarchy for a desktop and a laptop

# Memory hierarchy — recap

- The main differences among storage systems lie in speed (access time), size, and volatility
- Volatile storage loses its content when power is lost



# Single-processor, multiprocessor and multicore systems

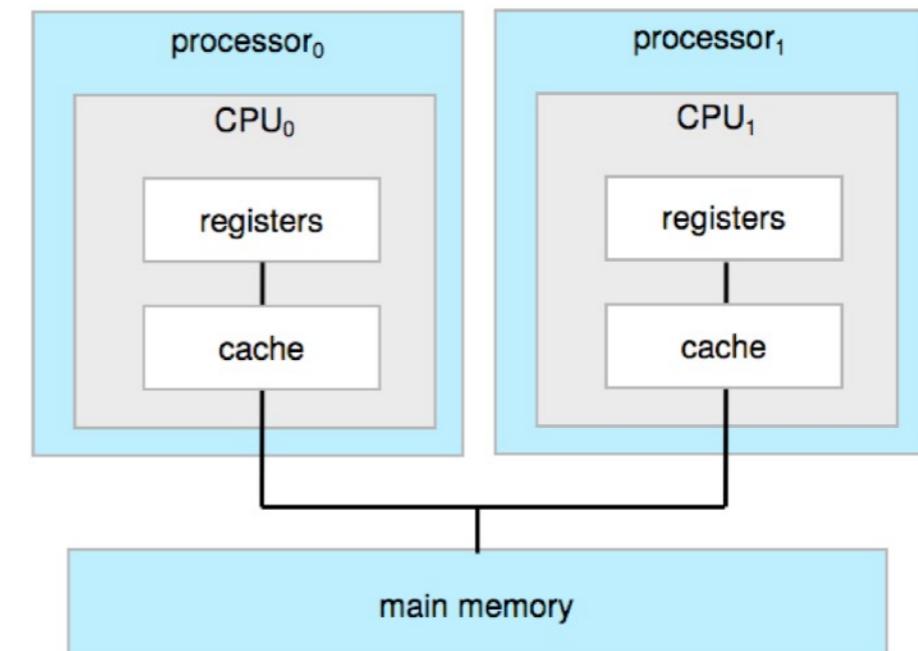
---

- a single-processor system contains a single general-purpose processor with one processing core
  - there might be several device-specific processors (for disk, keyboard, and graphics controllers) but they do not run processes

# Single-processor, multiprocessor and multicore systems

---

- a single-processor system contains a single general-purpose processor with one processing core
  - there might be several device-specific processors (for disk, keyboard, and graphics controllers) but they do not run processes
- a multiprocessor system contains two (or more) processors each with a single-core CPU
  - each processor has its own set of registers and cache but they all share the main memory
  - many processes can run simultaneously
  - **load balancing** might be necessary

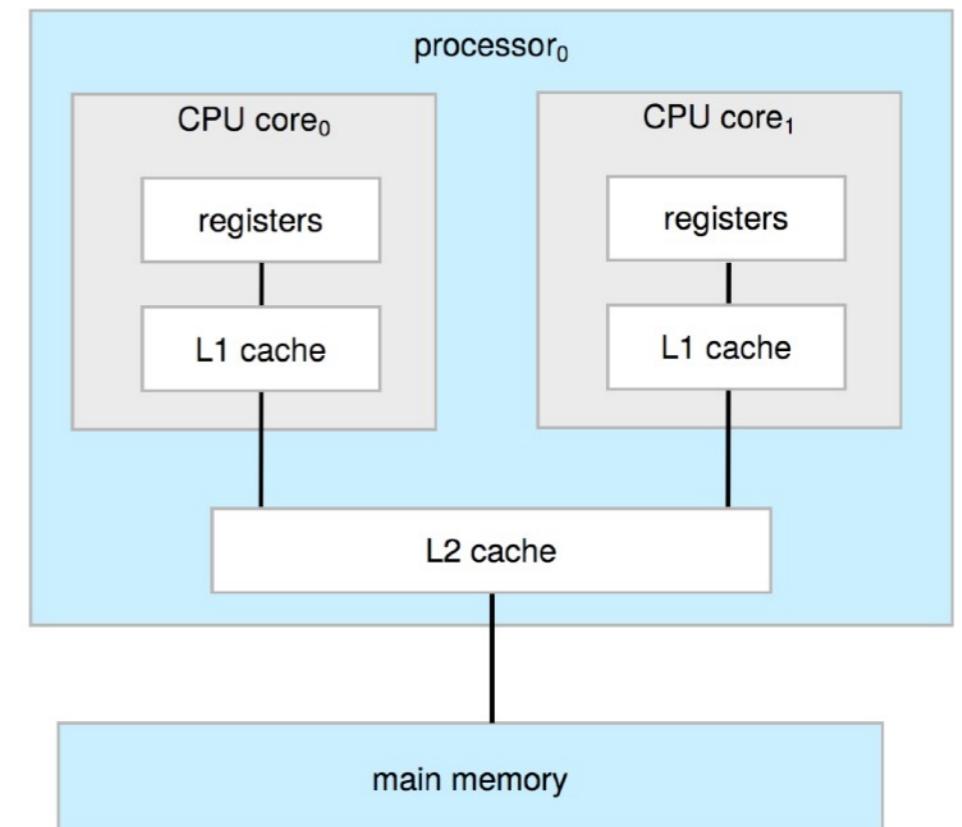


# Single-processor, multiprocessor and multicore systems

---

- a multicore system has multiple cores residing on a single processor chip
  - each core has its own registers and L1 cache; L2 cache is shared between processing cores
  - it can be more efficient than a multiprocessor system, because on-chip communication is **faster** than between-chip communication, and uses significantly **less power**

**N cores appear to the OS  
as N standard CPUs**

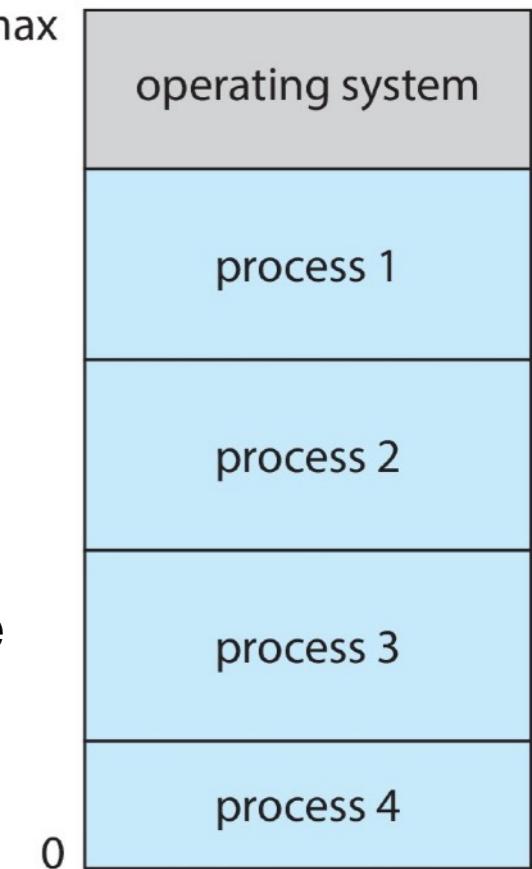


# Multiprogramming and multitasking

---

- Multiprogramming

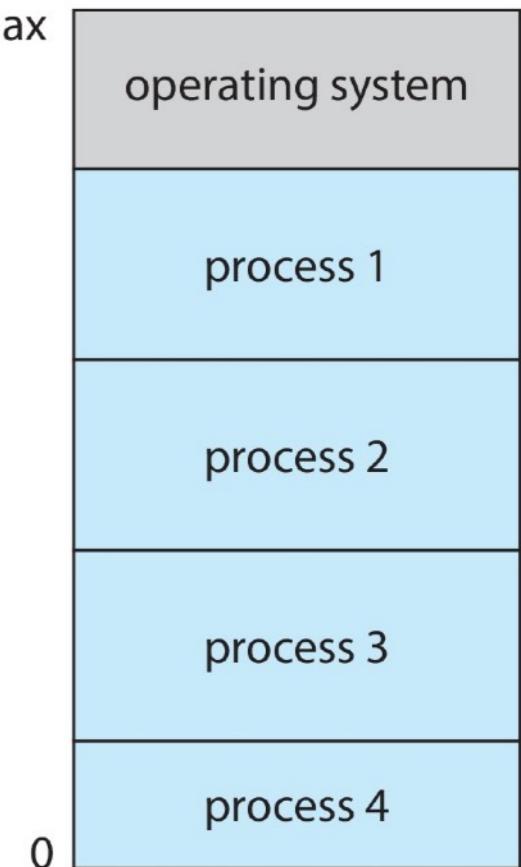
- OS keeps several processes in memory simultaneously
- OS picks one of them for execution (**scheduling**); when that process terminates or has to wait (e.g., for I/O), OS runs another process on the CPU
- Unlike a non-multiprogrammed system, the CPU wouldn't be idle as long as at least one process has to execute
- This increases **CPU utilization**



# Multiprogramming and multitasking

---

- Multiprogramming
  - OS keeps several processes in memory simultaneously
  - OS picks one of them for execution (**scheduling**); when that process terminates or has to wait (e.g., for I/O), OS runs another process on the CPU
  - Unlike a non-multiprogrammed system, the CPU wouldn't be idle as long as at least one process has to execute
  - This increases **CPU utilization**
- Multitasking is an extension of multiprogramming
  - CPU executes multiple processes by switching between them frequently to ensure reasonable response time (typically less than 1 second)
  - Choosing the process that runs next requires CPU scheduling



# How multiprogramming works?

---

- When a process waits for I/O, multiprogramming allows CPU to execute another process
  - the process waiting for I/O is added to the waiting queue of the I/O device (e.g., terminal, disks, video board, printer, network card)

# How multiprogramming works?

---

- When a process waits for I/O, multiprogramming allows CPU to execute another process
  - the process waiting for I/O is added to the waiting queue of the I/O device (e.g., terminal, disks, video board, printer, network card)
- Adds buffering
  - Data fills the buffer

# How multiprogramming works?

---

- When a process waits for I/O, multiprogramming allows CPU to execute another process
  - the process waiting for I/O is added to the waiting queue of the I/O device (e.g., terminal, disks, video board, printer, network card)
- Adds buffering
  - Data fills the buffer
- Adds interrupt handling
  - I/O events trigger a signal (known as **interrupt**)
  - control is then transferred to the OS
  - OS adds the process back into the ready queue

# Interrupt handling – basic idea

---

- Definition: a signal sent by a **device controller** to CPU's interrupt-request-line to inform that an event has occurred (e.g., I/O request completed by device driver)

# Interrupt handling – basic idea

---

- Definition: a signal sent by a **device controller** to CPU's interrupt-request-line to inform that an event has occurred (e.g., I/O request completed by device driver)
  - Each device controller has its own small processor which executes **asynchronously** with the main CPU

# Interrupt handling – basic idea

---

- Definition: a signal sent by a **device controller** to CPU's interrupt-request-line to inform that an event has occurred (e.g., I/O request completed by device driver)
  - Each device controller has its own small processor which executes **asynchronously** with the main CPU
  - CPU stops whatever it was doing, catches the interrupt, and dispatches it to the interrupt-handler routine through the in-memory **interrupt vector**, which contains the addresses of all handler routines;

0: 0x2ff080000	keyboard
1: 0x2ff100000	mouse
2: 0x2ff100480	timer
3: 0x2ff123010	Disk 1

# Interrupt handling – basic idea

---

- Definition: a signal sent by a **device controller** to CPU's interrupt-request-line to inform that an event has occurred (e.g., I/O request completed by device driver)
  - Each device controller has its own small processor which executes **asynchronously** with the main CPU
  - CPU stops whatever it was doing, catches the interrupt, and dispatches it to the interrupt-handler routine through the in-memory **interrupt vector**, which contains the addresses of all handler routines;
  - An interrupt-specific handler clears the interrupt by servicing the device

0: 0x2ff080000	keyboard
1: 0x2ff100000	mouse
2: 0x2ff100480	timer
3: 0x2ff123010	Disk 1

# Interrupt handling – basic idea

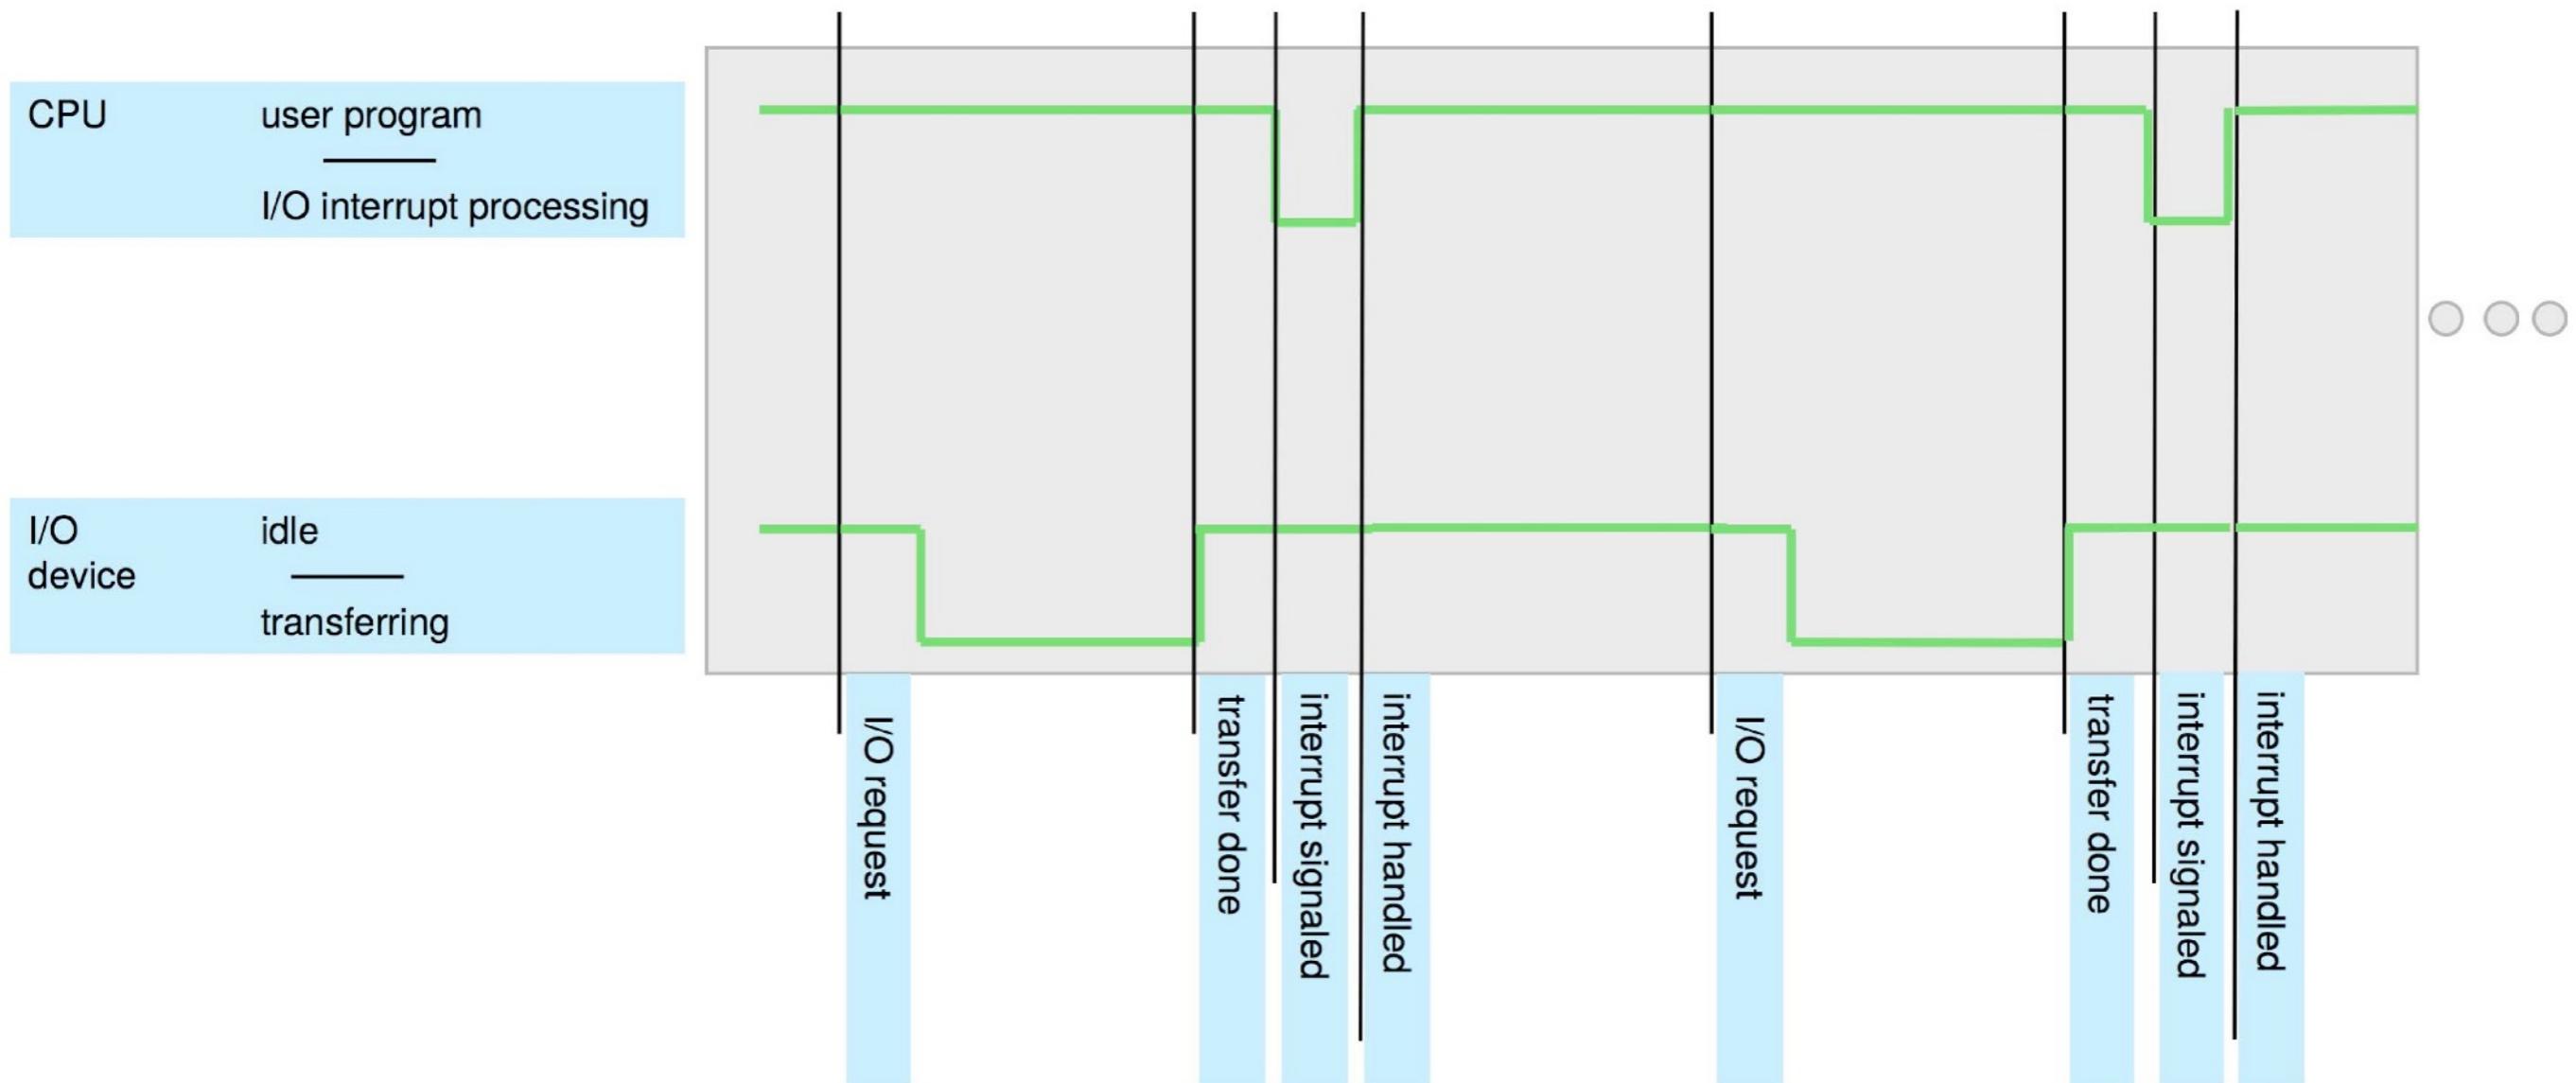
---

- Definition: a signal sent by a **device controller** to CPU's interrupt-request-line to inform that an event has occurred (e.g., I/O request completed by device driver)
  - Each device controller has its own small processor which executes **asynchronously** with the main CPU
  - CPU stops whatever it was doing, catches the interrupt, and dispatches it to the interrupt-handler routine through the in-memory **interrupt vector**, which contains the addresses of all handler routines;
  - An interrupt-specific handler clears the interrupt by servicing the device
  - The handler also preserves the state of the CPU by storing registers and the program counter; it restores the state after clearing the interrupt

0: 0x2ff080000	keyboard
1: 0x2ff100000	mouse
2: 0x2ff100480	timer
3: 0x2ff123010	Disk 1

# Interrupt-based asynchronous I/O

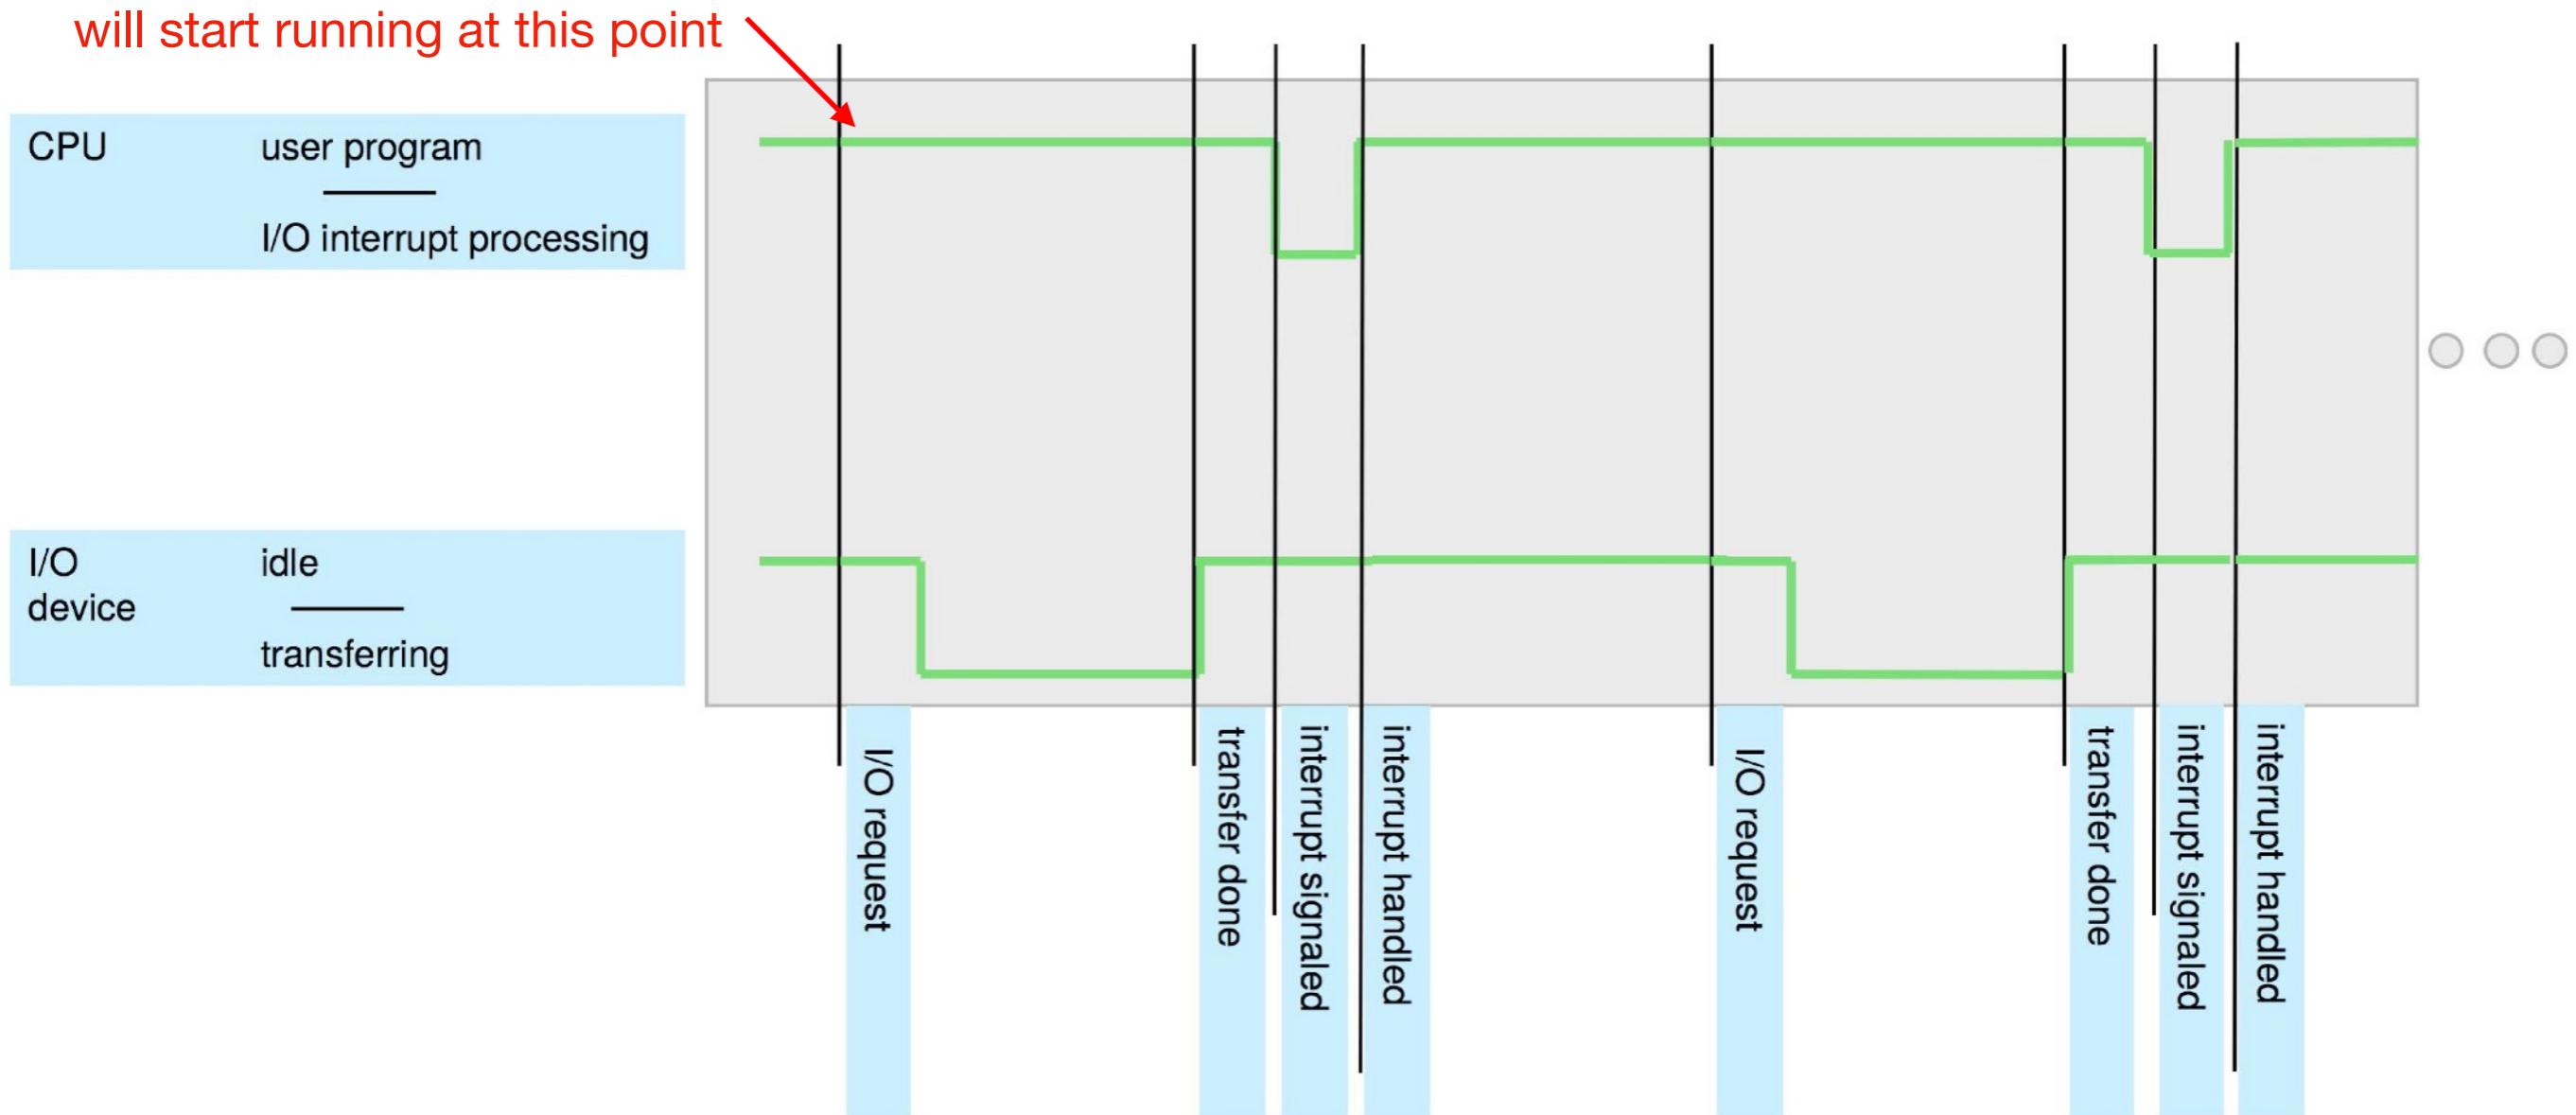
After I/O starts, control returns to user program without waiting for I/O completion



# Interrupt-based asynchronous I/O

After I/O starts, control returns to user program without waiting for I/O completion

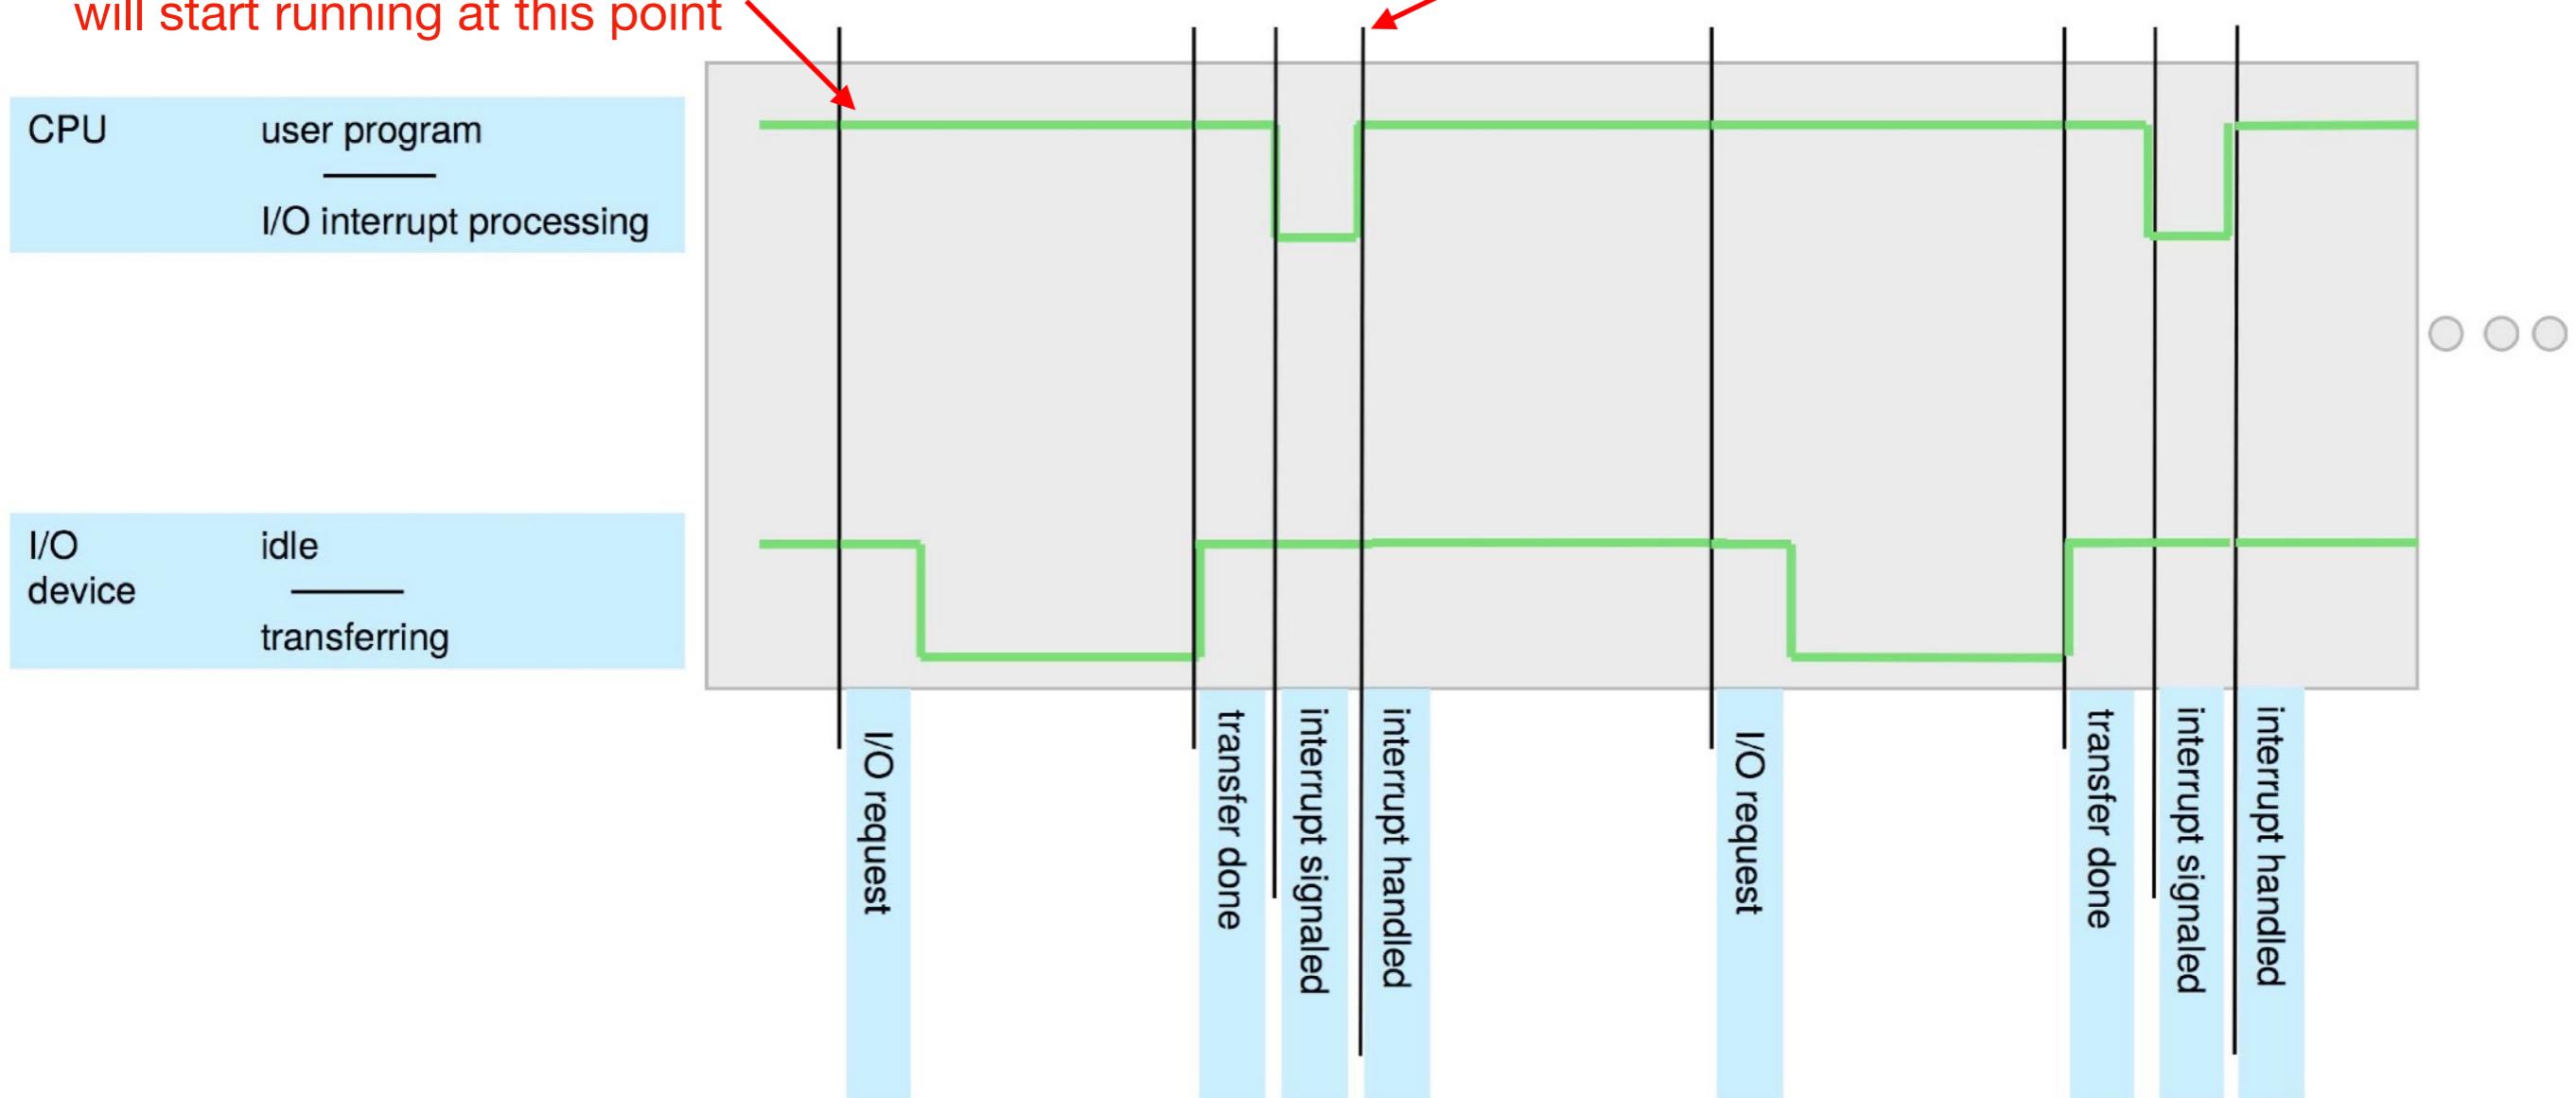
another user program  
will start running at this point



# Interrupt-based asynchronous I/O

After I/O starts, control returns to user program without waiting for I/O completion

another user program  
will start running at this point



# Interrupt handling – other considerations

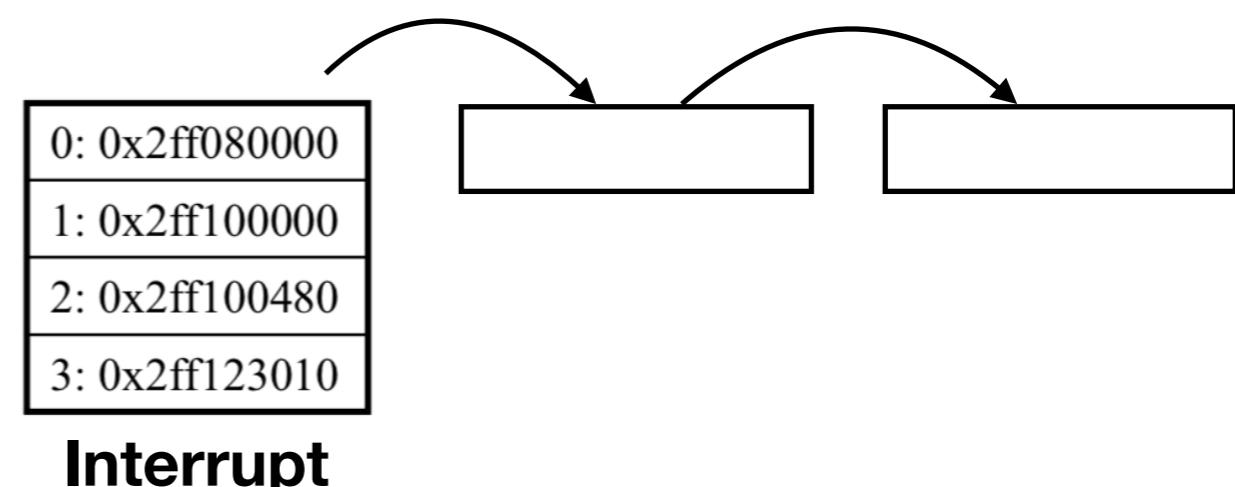
---

- How to defer interrupt handling during the execution of a critical instruction sequence? **two interrupt-request-lines**
  - CPU has two interrupt-request-lines: non-maskable and maskable
  - device controllers raise an interrupt through the maskable line
  - the non-maskable line is reserved for hardware faults such as unrecoverable memory errors

# Interrupt handling – other considerations

---

- How to defer interrupt handling during the execution of a critical instruction sequence? **two interrupt-request-lines**
  - CPU has two interrupt-request-lines: non-maskable and maskable
  - device controllers raise an interrupt through the maskable line
  - the non-maskable line is reserved for hardware faults such as unrecoverable memory errors
- What if there are more interrupt handlers than addresses available in the interrupt vector? **interrupt chaining**
  - each element in the interrupt vector points to the head of a list of interrupt handlers; handlers in the corresponding list are called one by one until the interrupt can be cleared by one handler



# Interrupt handling – other considerations

---

- How to defer interrupt handling during the execution of a critical instruction sequence? **two interrupt-request-lines**
  - CPU has two interrupt-request-lines: non-maskable and maskable
  - device controllers raise an interrupt through the maskable line
  - the non-maskable line is reserved for hardware faults such as unrecoverable memory errors
- What if there are more interrupt handlers than addresses available in the interrupt vector? **interrupt chaining**
  - each element in the interrupt vector points to the head of a list of interrupt handlers; handlers in the corresponding list are called one by one until the interrupt can be cleared by one handler
- How to differentiate urgent events from non-urgent ones? **interrupt priority levels**
  - CPU defers the handling of low-level interrupts to handle a high-level interrupt

# Interrupt handling – basic idea

---

- Definition: a trap or an exception is a **software-generated interrupt** caused either by an error (e.g., division by zero) or a user request for OS service (known as a **system call**)

0: 0x00080000	Illegal address
1: 0x00100000	Memory violation
2: 0x00100480	Illegal instruction
3: 0x00123010	System call

**Trap Vector**

# Protection

---

- Hardware has a status bit that indicates the current mode (user or kernel)
  - there could be more than two operation modes  
e.g., ARMv8 systems have seven modes
  - user applications run in the user mode
  - kernel code runs in the kernel model with the **full privileges** of the hardware

# Protection

---

- Hardware has a status bit that indicates the current mode (user or kernel)
  - there could be more than two operation modes  
e.g., ARMv8 systems have seven modes
  - user applications run in the user mode
  - kernel code runs in the kernel model with the **full privileges** of the hardware
- Operation modes provide the means for protecting OS from errant users
  - the code could be buggy or malicious

# Protection

---

- Hardware has a status bit that indicates the current mode (user or kernel)
  - there could be more than two operation modes  
e.g., ARMv8 systems have seven modes
  - user applications run in the user mode
  - kernel code runs in the kernel model with the **full privileges** of the hardware
- Operation modes provide the means for protecting OS from errant users
  - the code could be buggy or malicious
- examples of privileged instructions are I/O control, timer management, interrupt management; they can only run in the kernel mode

# Protection

- Hardware has a status bit that indicates the current mode (user or kernel)
  - there could be more than two operation modes  
e.g., ARMv8 systems have seven modes
  - user applications run in the user mode
  - kernel code runs in the kernel mode with the **full privileges** of the hardware
- Operation modes provide the means for protecting OS from errant users
  - the code could be buggy or malicious
- examples of privileged instructions are I/O control, timer management, interrupt management; they can only run in the kernel mode
- invoking a system call allows a user program to run privileged instructions

