

Operating System Concepts

Lecture 3: System Calls, Linking and Loading

Omid Ardakanian
oardakan@ualberta.ca
University of Alberta

MWF 12:00-12:50 VVC 2 215

Today's class

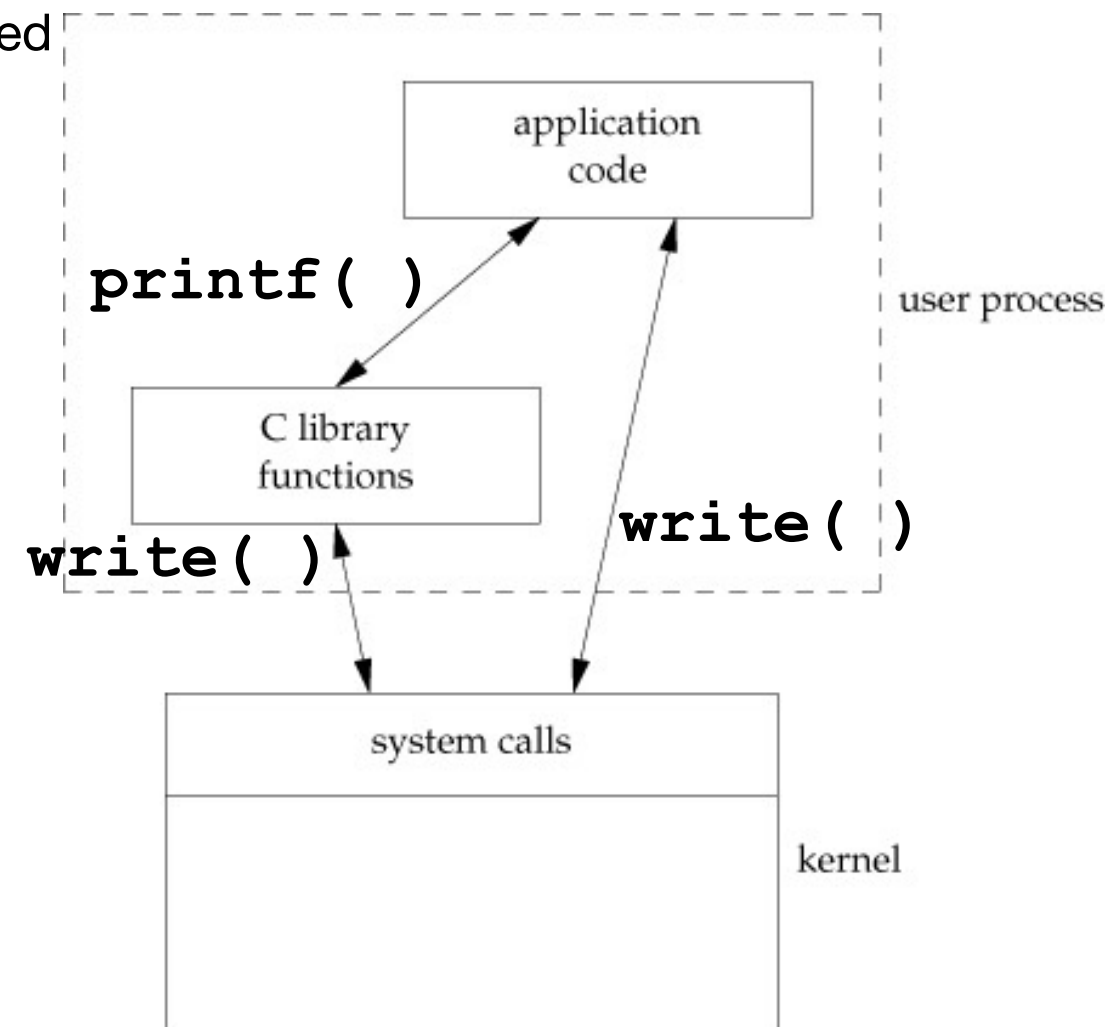
- Hardware support
 - System calls
 - Protection
- Basics of compiling, linking, and loading
- OS structure
 - Examples

System calls... what are they?

- Definition: services provided by the kernel to application programs
 - There are instructions that only the OS can run
 - A typical OS exports a few hundred (~300-400) system calls

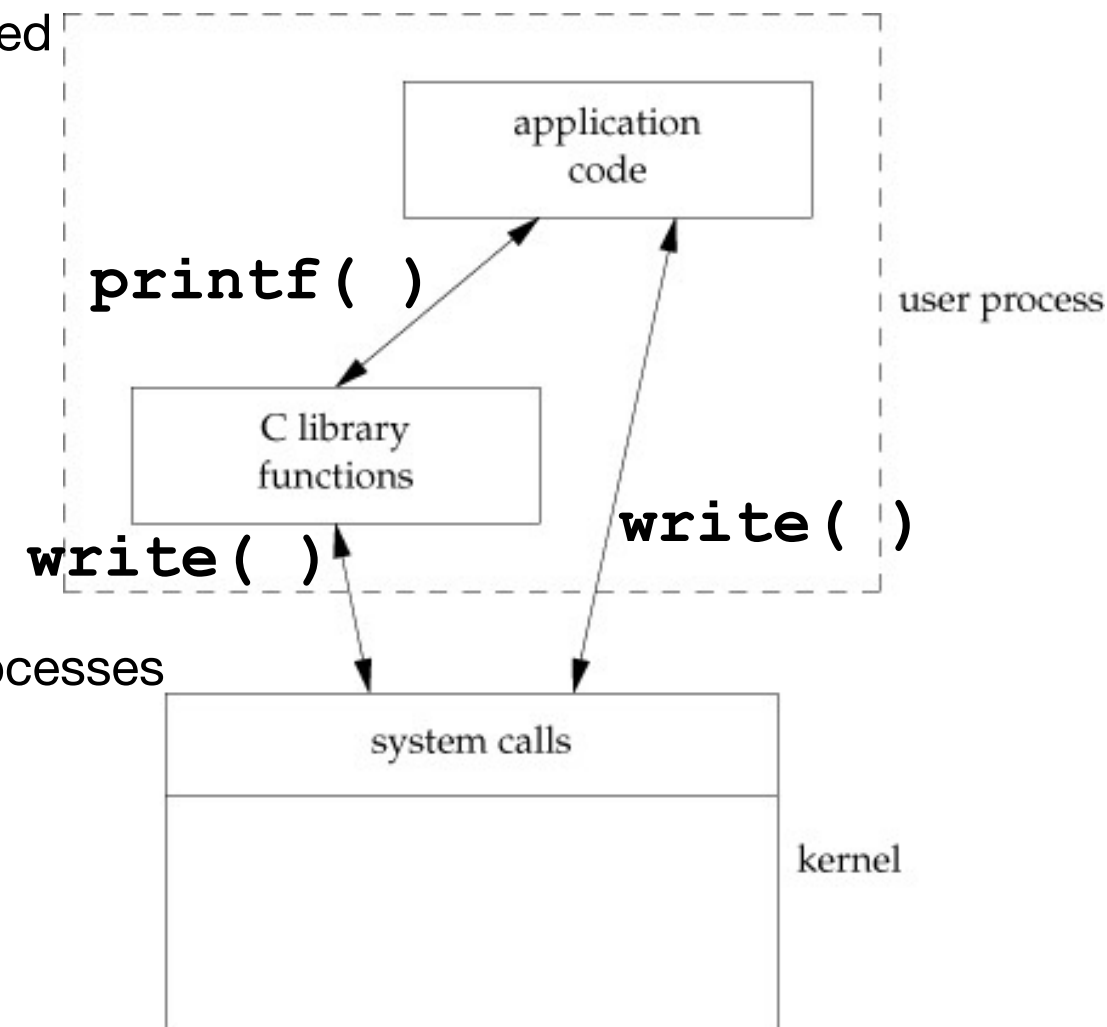
System calls... what are they?

- Definition: services provided by the kernel to application programs
 - There are instructions that only the OS can run
 - A typical OS exports a few hundred (~300-400) system calls
- Each system call is typically defined as a function in the standard C library
 - In UNIX-based systems most system calls are declared in `unistd.h` and some are declared in `fcntl.h`



System calls... what are they?

- Definition: services provided by the kernel to application programs
 - There are instructions that only the OS can run
 - A typical OS exports a few hundred (~300-400) system calls
- Each system call is typically defined as a function in the standard C library
 - In UNIX-based systems most system calls are declared in `unistd.h` and some are declared in `fcntl.h`
- System calls are necessary to
 - Access devices and files
 - Request memory
 - Set access permissions
 - Stop and start processes, and communicate with processes
 - Set a timer



UNIX system calls

- Process management
 - `getpid()`, `fork()`, `exec()`, `wait()`, `exit()`, ...

UNIX system calls

- Process management

- `getpid()`, `fork()`, `exec()`, `wait()`, `exit()`, ...

- File manipulation

- `open()`, `close()`, `read()`, `write()`, `stat()`, `lseek()`, `link()`, ...

UNIX system calls

- Process management

- `getpid()`, `fork()`, `exec()`, `wait()`, `exit()`, ...

- File manipulation

- `open()`, `close()`, `read()`, `write()`, `stat()`, `lseek()`, `link()`, ...

- Protection

- `chmod()`, `chown()`, ...

UNIX system calls

- Process management
 - `getpid()`, `fork()`, `exec()`, `wait()`, `exit()`, ...
- File manipulation
 - `open()`, `close()`, `read()`, `write()`, `stat()`, `lseek()`, `link()`, ...
- Protection
 - `chmod()`, `chown()`, ...
- Communications
 - `pipe()`, `shm_open()`, `mmap()`, `socket()`, `accept()`, `send()`, `recv()`, ...

UNIX system calls

- Process management
 - `getpid()`, `fork()`, `exec()`, `wait()`, `exit()`, ...
- File manipulation
 - `open()`, `close()`, `read()`, `write()`, `stat()`, `lseek()`, `link()`, ...
- Protection
 - `chmod()`, `chown()`, ...
- Communications
 - `pipe()`, `shm_open()`, `mmap()`, `socket()`, `accept()`, `send()`, `recv()`, ...

See Linux system calls here:

<http://man7.org/linux/man-pages/man2/syscalls.2.html>

- ```
[candakar@us01: ~] $ strings pud
```

[illegible]

# POSIX Application Programming Interface (API)

---

- Application programmers usually use functions from an API rather than invoking system calls directly

# POSIX Application Programming Interface (API)

---

- Application programmers usually use functions from an API rather than invoking system calls directly
  - **Potability**: UNIX-based operating systems could have different declarations and implementations for their system calls, but they all support the same **POSIX API**; hence if you use the POSIX API you can expect that your program compiles and runs on any system that supports it
  - **Ease of use**: system calls are often more detailed and more difficult to use than their corresponding API functions

# POSIX Application Programming Interface (API)

---

- Application programmers usually use functions from an API rather than invoking system calls directly
  - **Potability**: UNIX-based operating systems could have different declarations and implementations for their system calls, but they all support the same **POSIX API**; hence if you use the POSIX API you can expect that your program compiles and runs on any system that supports it
  - **Ease of use**: system calls are often more detailed and more difficult to use than their corresponding API functions
- POSIX (Portable Operating System Interface) is a family of standards implemented primarily for UNIX-based operating systems



# POSIX Application Programming Interface (API)

---

- Application programmers usually use functions from an API rather than invoking system calls directly
  - **Potability**: UNIX-based operating systems could have different declarations and implementations for their system calls, but they all support the same **POSIX API**; hence if you use the POSIX API you can expect that your program compiles and runs on any system that supports it
  - **Ease of use**: system calls are often more detailed and more difficult to use than their corresponding API functions
- POSIX (Portable Operating System Interface) is a family of standards implemented primarily for UNIX-based operating systems
  - POSIX compliant systems like Linux and macOS must implement the POSIX core standard (POSIX.1)

# POSIX Application Programming Interface (API)

---

- Application programmers usually use functions from an API rather than invoking system calls directly
  - **Potability**: UNIX-based operating systems could have different declarations and implementations for their system calls, but they all support the same **POSIX API**; hence if you use the POSIX API you can expect that your program compiles and runs on any system that supports it
  - **Ease of use**: system calls are often more detailed and more difficult to use than their corresponding API functions
- POSIX (Portable Operating System Interface) is a family of standards implemented primarily for UNIX-based operating systems
  - POSIX compliant systems like Linux and macOS must implement the POSIX core standard (POSIX.1)
  - POSIX thread library (aka pthreads) is defined in an extension of this API known as POSIX.1c



# POSIX Application Programming Interface (API)

---

- Application programmers usually use functions from an API rather than invoking system calls directly
  - **Potability**: UNIX-based operating systems could have different declarations and implementations for their system calls, but they all support the same **POSIX API**; hence if you use the POSIX API you can expect that your program compiles and runs on any system that supports it
  - **Ease of use**: system calls are often more detailed and more difficult to use than their corresponding API functions
- POSIX (Portable Operating System Interface) is a family of standards implemented primarily for UNIX-based operating systems
  - POSIX compliant systems like Linux and macOS must implement the POSIX core standard (POSIX.1)
  - POSIX thread library (aka pthreads) is defined in an extension of this API known as POSIX.1c
- **Win32 API** is the standard API for Windows operating systems

# Protection

---

- Hardware has a status bit that indicates the current operation mode
  - there could be more than two operation modes  
e.g., ARMv8 systems have seven modes
  - user applications run in the **user mode**
  - kernel code runs in the **kernel mode** with the **full privileges** of the hardware
  - examples of privileged instructions are I/O control, timer management, interrupt management; they can only run in the kernel mode

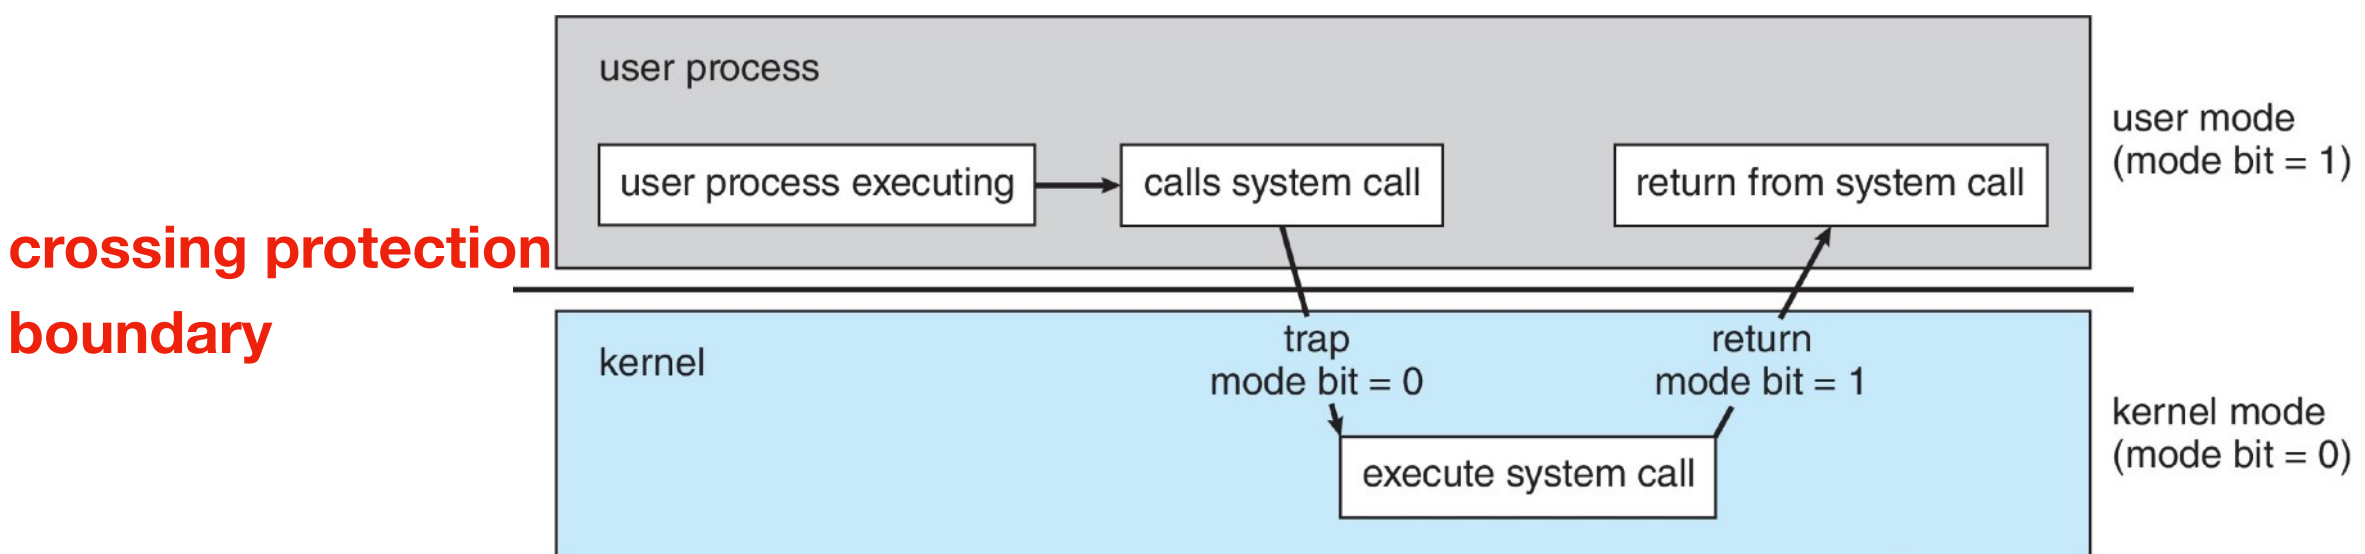
# Protection

---

- Hardware has a status bit that indicates the current operation mode
  - there could be more than two operation modes  
e.g., ARMv8 systems have seven modes
  - user applications run in the **user mode**
  - kernel code runs in the **kernel mode** with the **full privileges** of the hardware
  - examples of privileged instructions are I/O control, timer management, interrupt management; they can only run in the kernel mode
- invoking a system call allows a user program to run privileged instructions

# Protection

- Hardware has a status bit that indicates the current operation mode
  - there could be more than two operation modes  
e.g., ARMv8 systems have seven modes
  - user applications run in the **user mode**
  - kernel code runs in the **kernel model** with the **full privileges** of the hardware
  - examples of privileged instructions are I/O control, timer management, interrupt management; they can only run in the kernel mode
- invoking a system call allows a user program to run privileged instructions
- operation modes provide the means for protecting OS from errant users
  - the code could be buggy or malicious



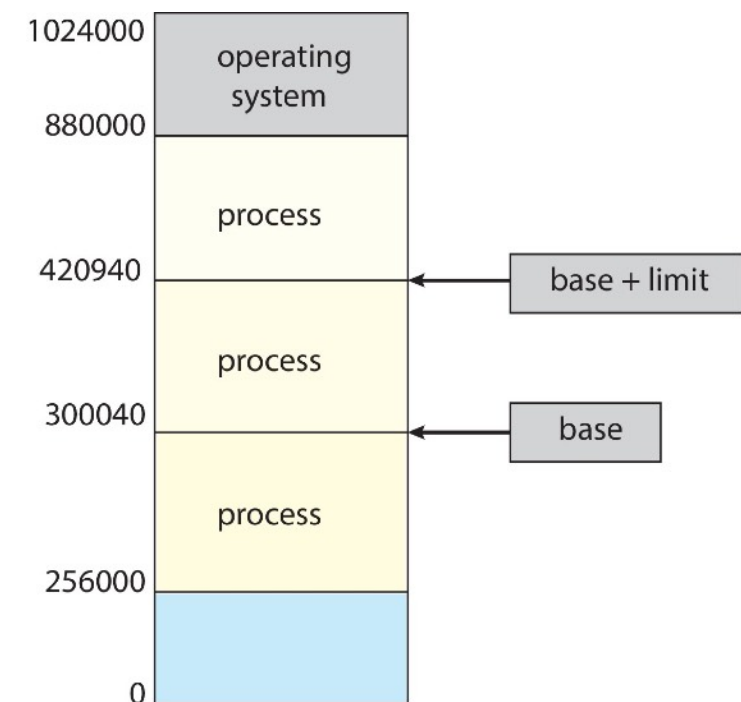
# Memory protection

---

- hardware must provide support so that the OS can
  - protect user programs from each other
  - protect the OS from user programs

# Memory protection

- hardware must provide support so that the OS can
  - protect user programs from each other
  - protect the OS from user programs
- the simplest technique is to use base and limit registers
  - base and limit registers are loaded by the OS before creating a process
  - base register holds the smallest legal physical memory address of the process
  - limit register holds the size of the memory allocated to a process
  - CPU checks each reference in user mode (instruction and data addresses) to ensure that it falls between base and base+limit



# Timer interrupt

---

**What if the process running on the CPU does not wait for I/O or signals?**

# Timer interrupt

---

**What if the process running on the CPU does not wait for I/O or signals?**

- the kernel must periodically get back control



# Timer interrupt

---

**What if the process running on the CPU does not wait for I/O or signals?**

- the kernel must periodically get back control
- CPU is protected from being hogged using timer interrupts that occur at regular intervals (e.g., every 100 microseconds)
  - frequency is set by the kernel
  - yet another protection mechanism

|                |          |
|----------------|----------|
| 0: 0x2ff080000 | keyboard |
| 1: 0x2ff100000 | mouse    |
| 2: 0x2ff100480 | timer    |
| 3: 0x2ff123010 | Disk 1   |

**Interrupt Vector**

# Timer interrupt

---

**What if the process running on the CPU does not wait for I/O or signals?**

- the kernel must periodically get back control
- CPU is protected from being hogged using timer interrupts that occur at regular intervals (e.g., every 100 microseconds)
  - frequency is set by the kernel
  - yet another protection mechanism
- at each timer interrupt, the CPU chooses a new process to execute

|                |          |
|----------------|----------|
| 0: 0x2ff080000 | keyboard |
| 1: 0x2ff100000 | mouse    |
| 2: 0x2ff100480 | timer    |
| 3: 0x2ff123010 | Disk 1   |

**Interrupt Vector**

# Timer interrupt

---

**What if the process running on the CPU does not wait for I/O or signals?**

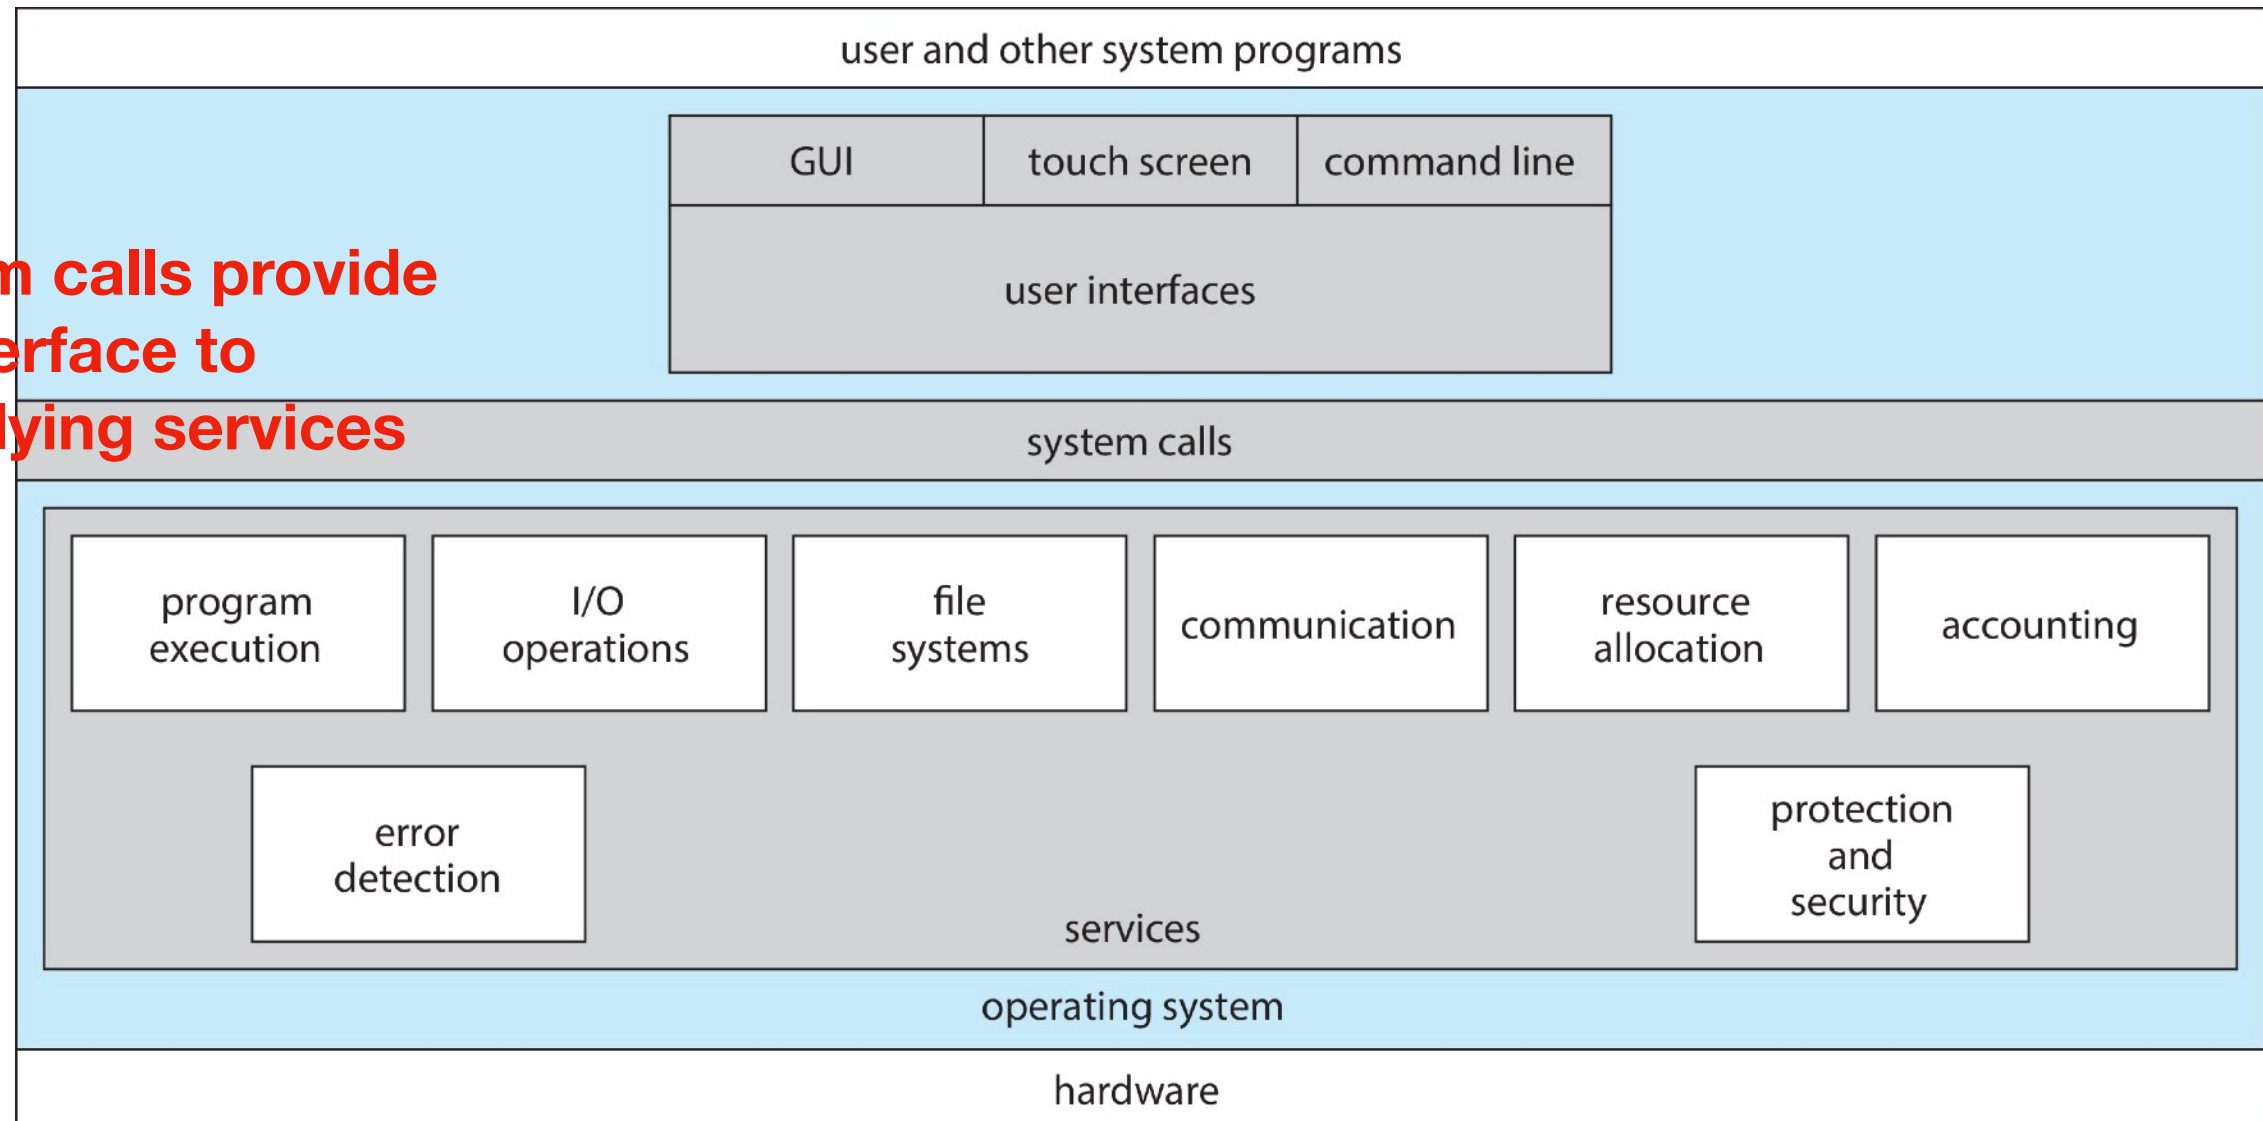
- the kernel must periodically get back control
- CPU is protected from being hogged using timer interrupts that occur at regular intervals (e.g., every 100 microseconds)
  - frequency is set by the kernel
  - yet another protection mechanism
- at each timer interrupt, the CPU chooses a new process to execute
- interrupts can be temporarily deferred
  - it is crucial for implementing mutual exclusion

|                |          |
|----------------|----------|
| 0: 0x2ff080000 | keyboard |
| 1: 0x2ff100000 | mouse    |
| 2: 0x2ff100480 | timer    |
| 3: 0x2ff123010 | Disk 1   |

**Interrupt Vector**

# Operating System services

system calls provide  
an interface to  
underlying services



- making the programming task easier and increasing the convenience of users
- ensuring the efficient operation of the system (i.e., resource allocation, logging, protection and security)

# Summary

---

## OS Services

## Hardware Support

---

Interrupts

interrupt request lines, interrupt vector

---

System calls

trap vector

---

I/O

interrupt and **memory mapping**

---

Protection

operation modes, privileged instructions, base and limit registers, timer

---

Scheduling & accounting

timer

---

**Synchronization**

**atomic operations**

---


**Virtual memory**

**MMU, translation look-aside buffer (TLB)**

# Running a user program

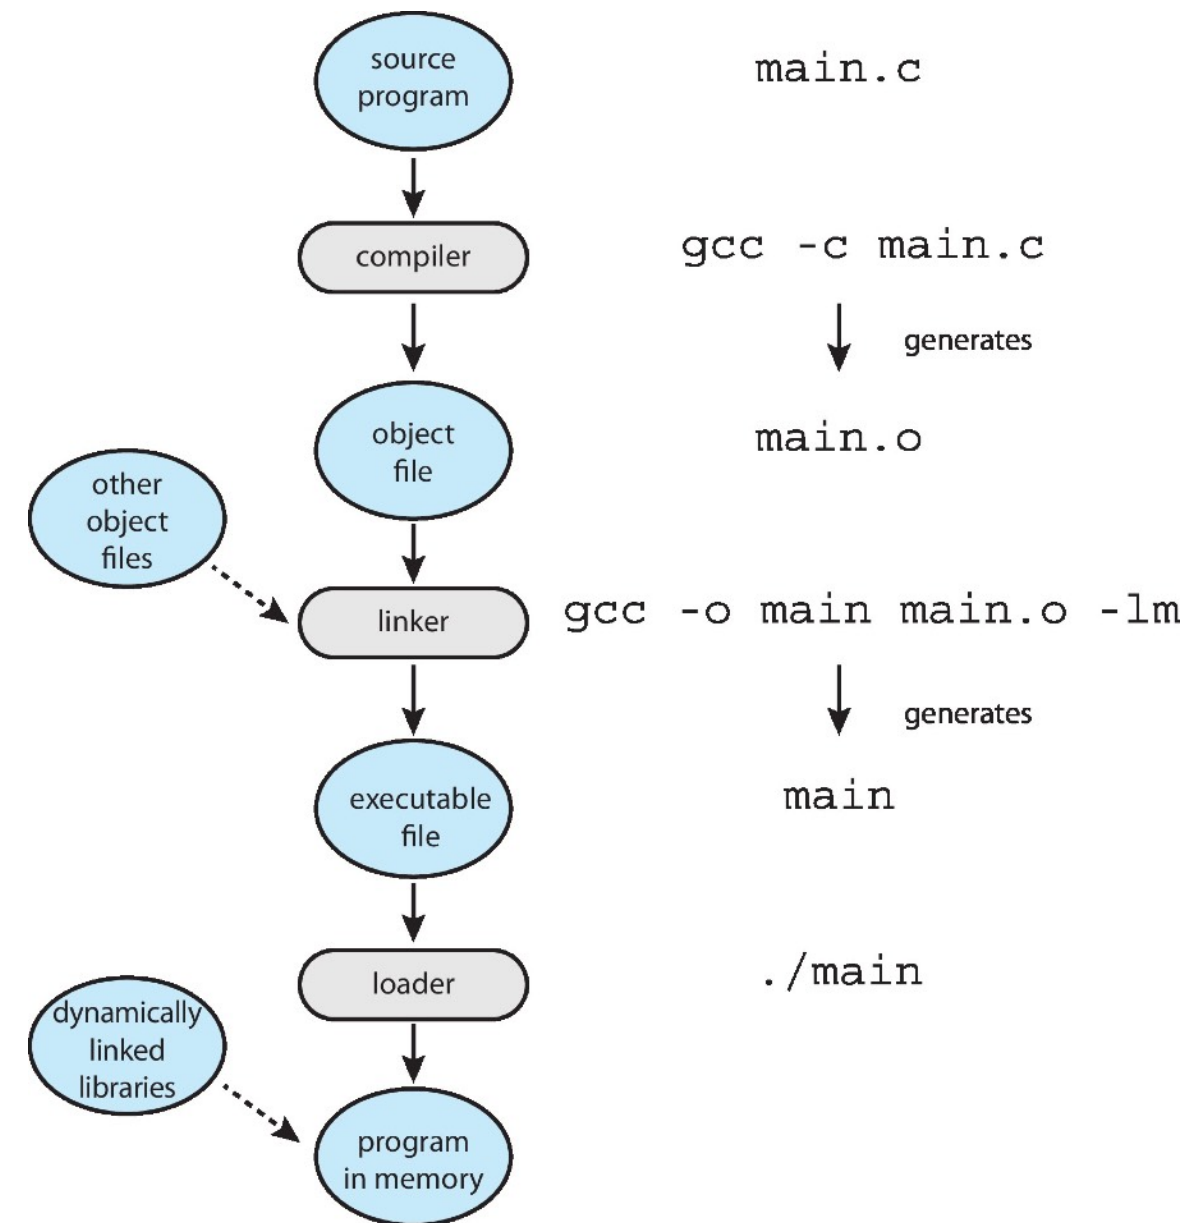
# Program to Process

- Compiling is the process of converting a source file (ASCII code) into an object file
  - object files miss certain information, e.g., functions called from other files and libraries



# Program to Process

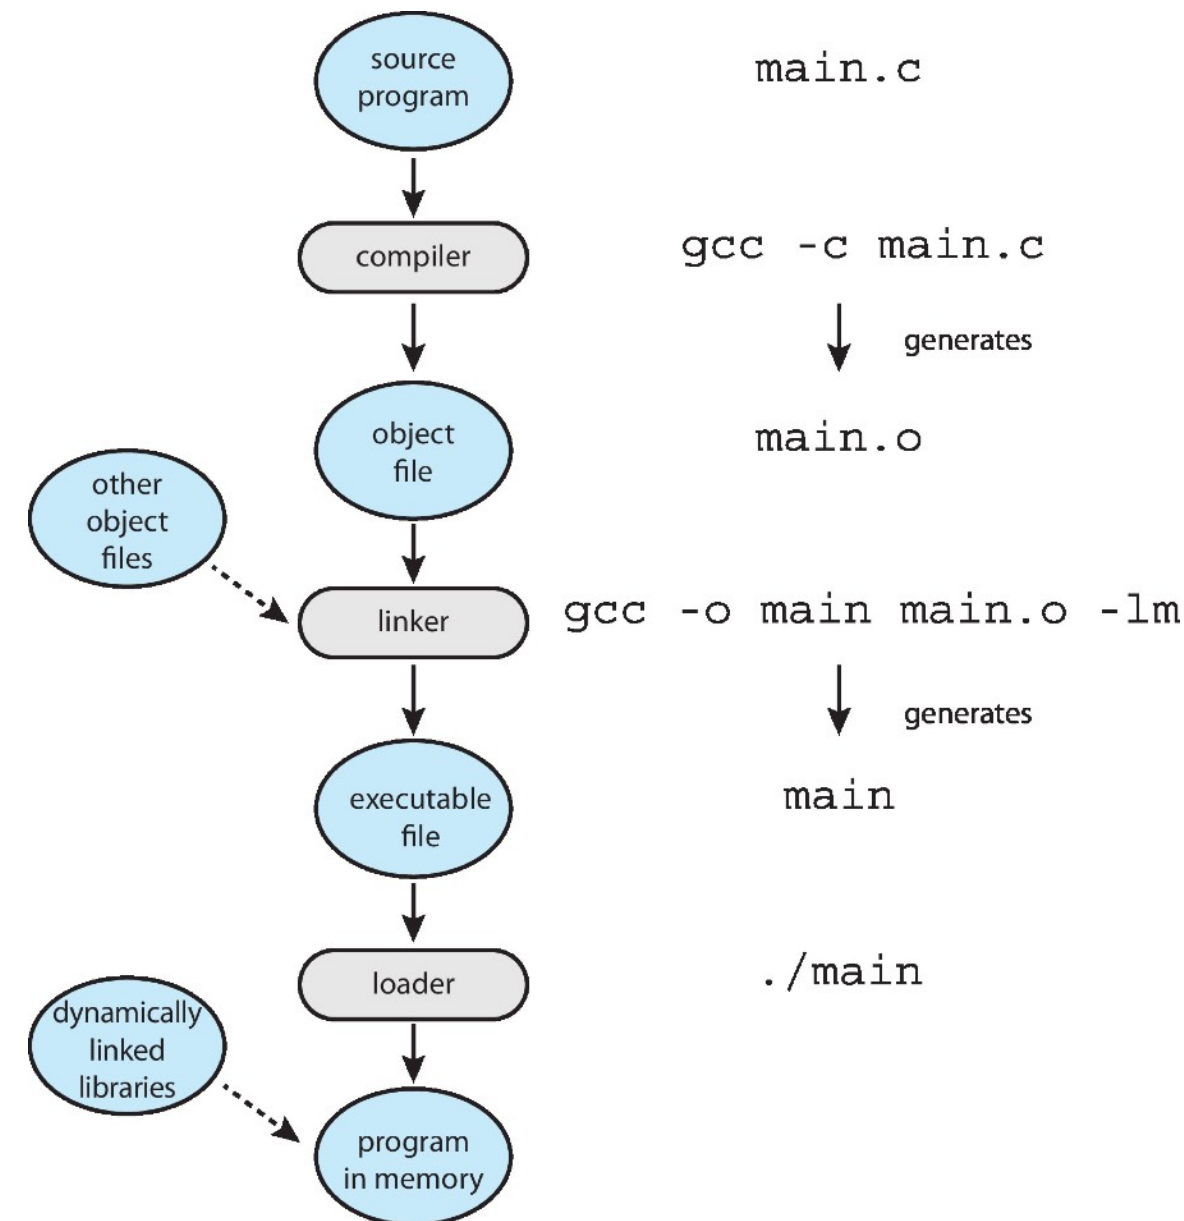
- Compiling is the process of converting a source file (ASCII code) into an object file
  - object files miss certain information, e.g., functions called from other files and libraries
- Linking is the process of combining relocatable object files and specific libraries into a single binary executable





# Program to Process

- Compiling is the process of converting a source file (ASCII code) into an object file
  - object files miss certain information, e.g., functions called from other files and libraries
- Linking is the process of combining relocatable object files and specific libraries into a single binary executable
- Loading is the process of bringing this executable file into memory
  - the loader is executed when you enter the name of the executable file on the command line; this is done using `execve( )` system call
  - the loader sets up the process memory to contain code and data from executable
  - a library can be conditionally linked and loaded if it is required during the run time; this can be done using dynamically linked libraries (DLLs)



# Executable and Linkable Format (ELF)

---

- Executable files must have a standard format, so that the OS knows how to load and start them

# Executable and Linkable Format (ELF)

---

- Executable files must have a standard format, so that the OS knows how to load and start them
- ELF is a common standard across modern UNIX operating systems, such as Linux

# Executable and Linkable Format (ELF)

---

- Executable files must have a standard format, so that the OS knows how to load and start them
- ELF is a common standard across modern UNIX operating systems, such as Linux
- **ELF relocatable file** contains the compiled machine code and a symbol table containing metadata about functions and variables referenced in the program
  - the symbol table is necessary for linking the relocatable object file with other object files

# Executable and Linkable Format (ELF)

---

- Executable files must have a standard format, so that the OS knows how to load and start them
- ELF is a common standard across modern UNIX operating systems, such as Linux
- **ELF relocatable file** contains the compiled machine code and a symbol table containing metadata about functions and variables referenced in the program
  - the symbol table is necessary for linking the relocatable object file with other object files
- **ELF executable file** contains the address of the first instruction of the program (program's entry point: `_start` function)
  - this file can be loaded into memory

# Executable and Linkable Format (ELF)

---

- Executable files must have a standard format, so that the OS knows how to load and start them
- ELF is a common standard across modern UNIX operating systems, such as Linux
- **ELF relocatable file** contains the compiled machine code and a symbol table containing metadata about functions and variables referenced in the program
  - the symbol table is necessary for linking the relocatable object file with other object files
- **ELF executable file** contains the address of the first instruction of the program (program's entry point: `_start` function)
  - this file can be loaded into memory
- the ELF header starts with a 4-byte magic string:
  - 0x7F followed by 0x45 0x4C 0x46 (ELF in ASCII)

# Executable and Linkable Format (ELF)

---

An ELF file is divided into sections

# Executable and Linkable Format (ELF)

---

An ELF file is divided into sections

- **.text** contains the machine code
  - In UNIX-based systems, see the content of this section using `objdump -drS objfile`



# Executable and Linkable Format (ELF)

---

An ELF file is divided into sections

- **.text** contains the machine code
  - In UNIX-based systems, see the content of this section using `objdump -drS objfile`
- **.data** contains initialized global variables

# Executable and Linkable Format (ELF)

---

An ELF file is divided into sections

- **.text** contains the machine code
  - In UNIX-based systems, see the content of this section using `objdump -drS objfile`
- **.data** contains initialized global variables
- **.bss** (block storage start) contains uninitialized global variables
  - occupies no space in the object file actually
  - there is no point storing more zeroes on your disk

# Executable and Linkable Format (ELF)

---

An ELF file is divided into sections

- **.text** contains the machine code
  - In UNIX-based systems, see the content of this section using `objdump -drS objfile`
- **.data** contains initialized global variables
- **.bss** (block storage start) contains uninitialized global variables
  - occupies no space in the object file actually
  - there is no point storing more zeroes on your disk
- **.rodata** contains read-only data such as constant strings
  - e.g., the format strings in `printf` statements

# Executable and Linkable Format (ELF)

---

An ELF file is divided into sections

- **.text** contains the machine code
  - In UNIX-based systems, see the content of this section using `objdump -drS objfile`
- **.data** contains initialized global variables
- **.bss** (block storage start) contains uninitialized global variables
  - occupies no space in the object file actually
  - there is no point storing more zeroes on your disk
- **.rodata** contains read-only data such as constant strings
  - e.g., the format strings in `printf` statements
- **.symtab** contains the symbol table
  - information about functions and global variables defined and referenced in the program

# Executable and Linkable Format (ELF)

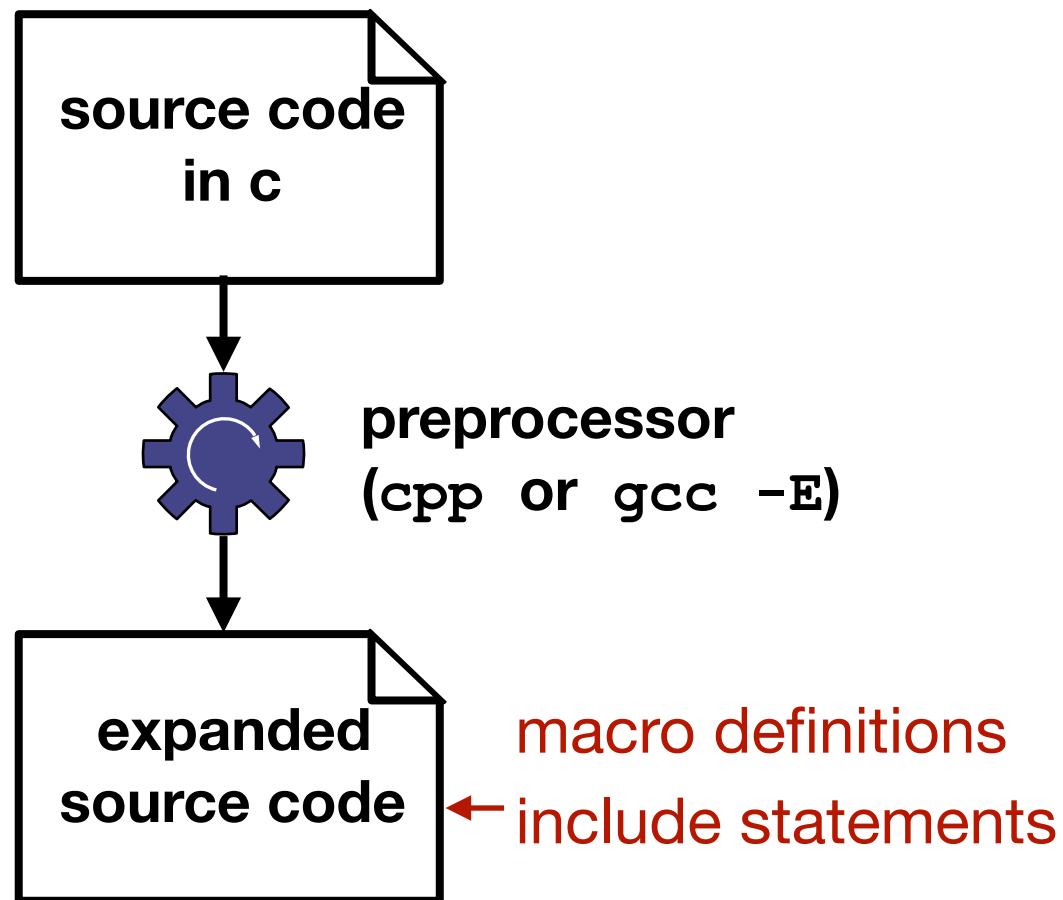
---

An ELF file is divided into sections

- **.text** contains the machine code
  - In UNIX-based systems, see the content of this section using `objdump -drS objfile`
- **.data** contains initialized global variables
- **.bss** (block storage start) contains uninitialized global variables
  - occupies no space in the object file actually
  - there is no point storing more zeroes on your disk
- **.rodata** contains read-only data such as constant strings
  - e.g., the format strings in `printf` statements
- **.symtab** contains the symbol table
  - information about functions and global variables defined and referenced in the program
- **.rel.text** and **.rel.data** contain relocation information for functions and global variables that are referenced but not defined (external references)
  - linker modifies this section when combining the object files, resolving external references

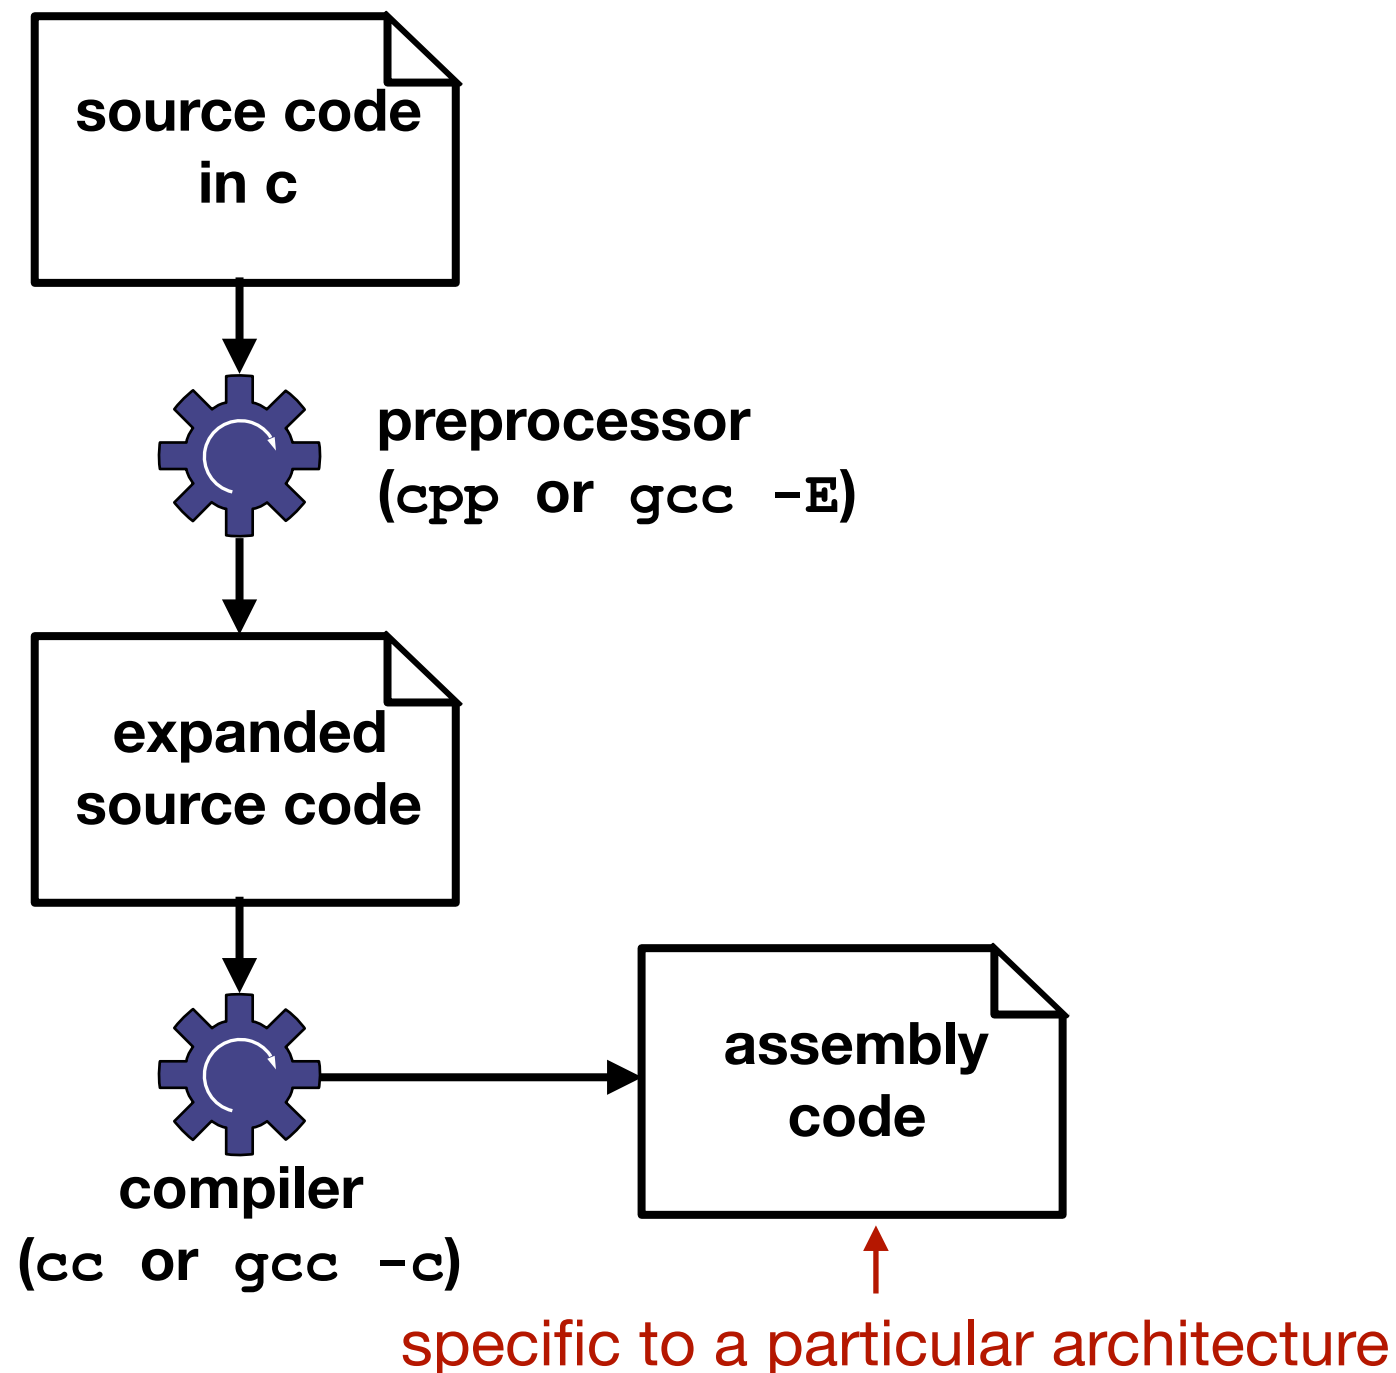
# Compiler, Linker and Loader in action

---



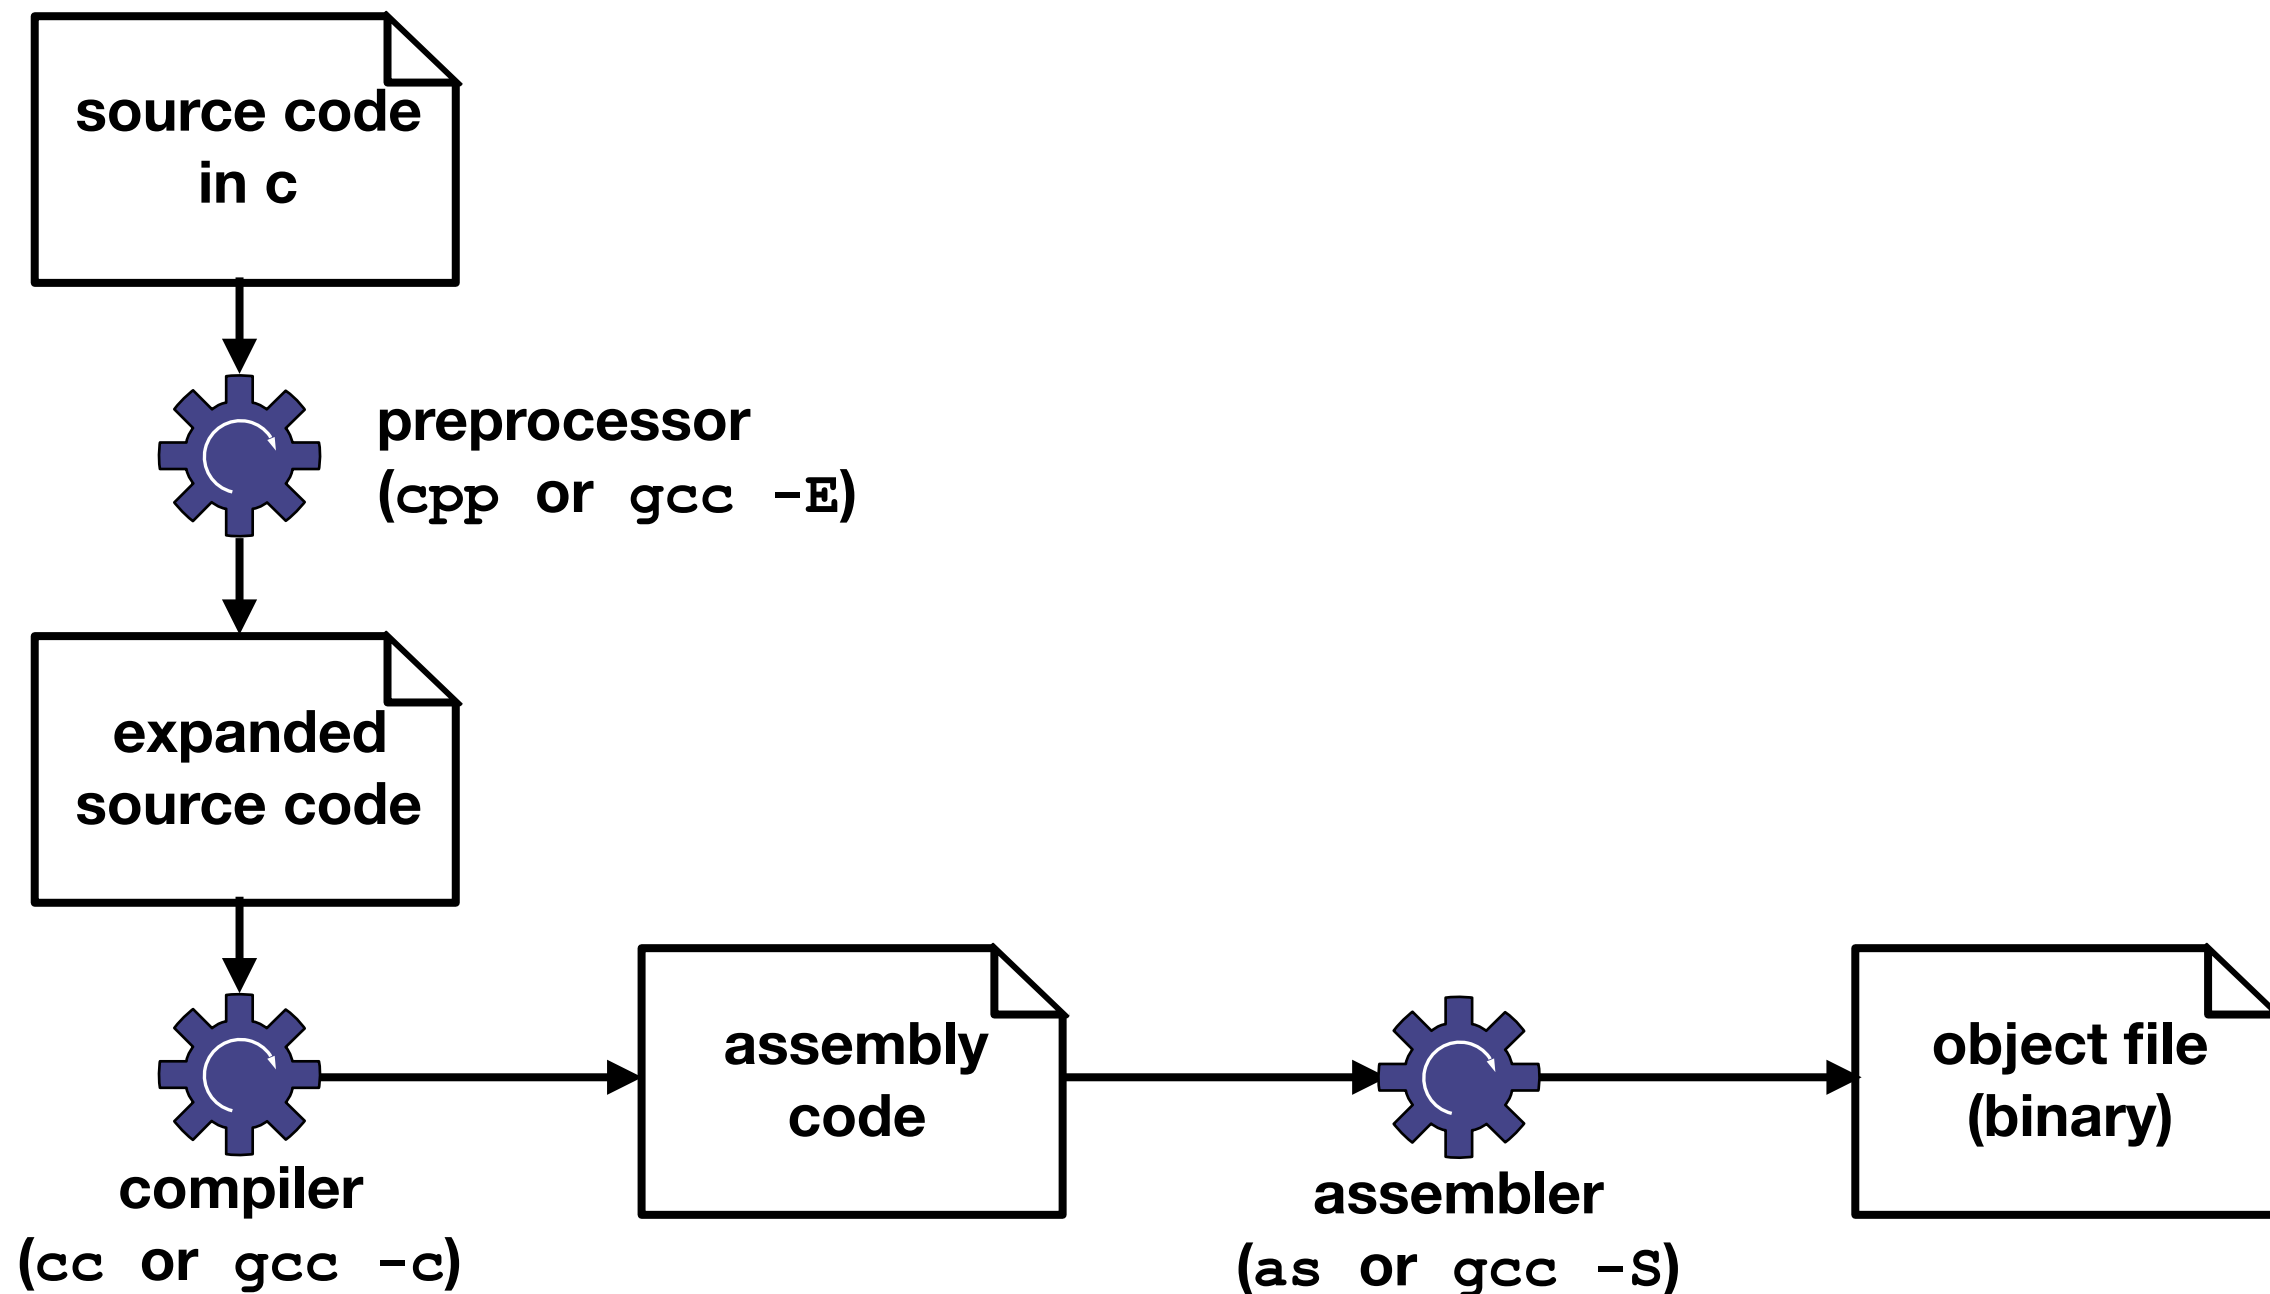
# Compiler, Linker and Loader in action

---



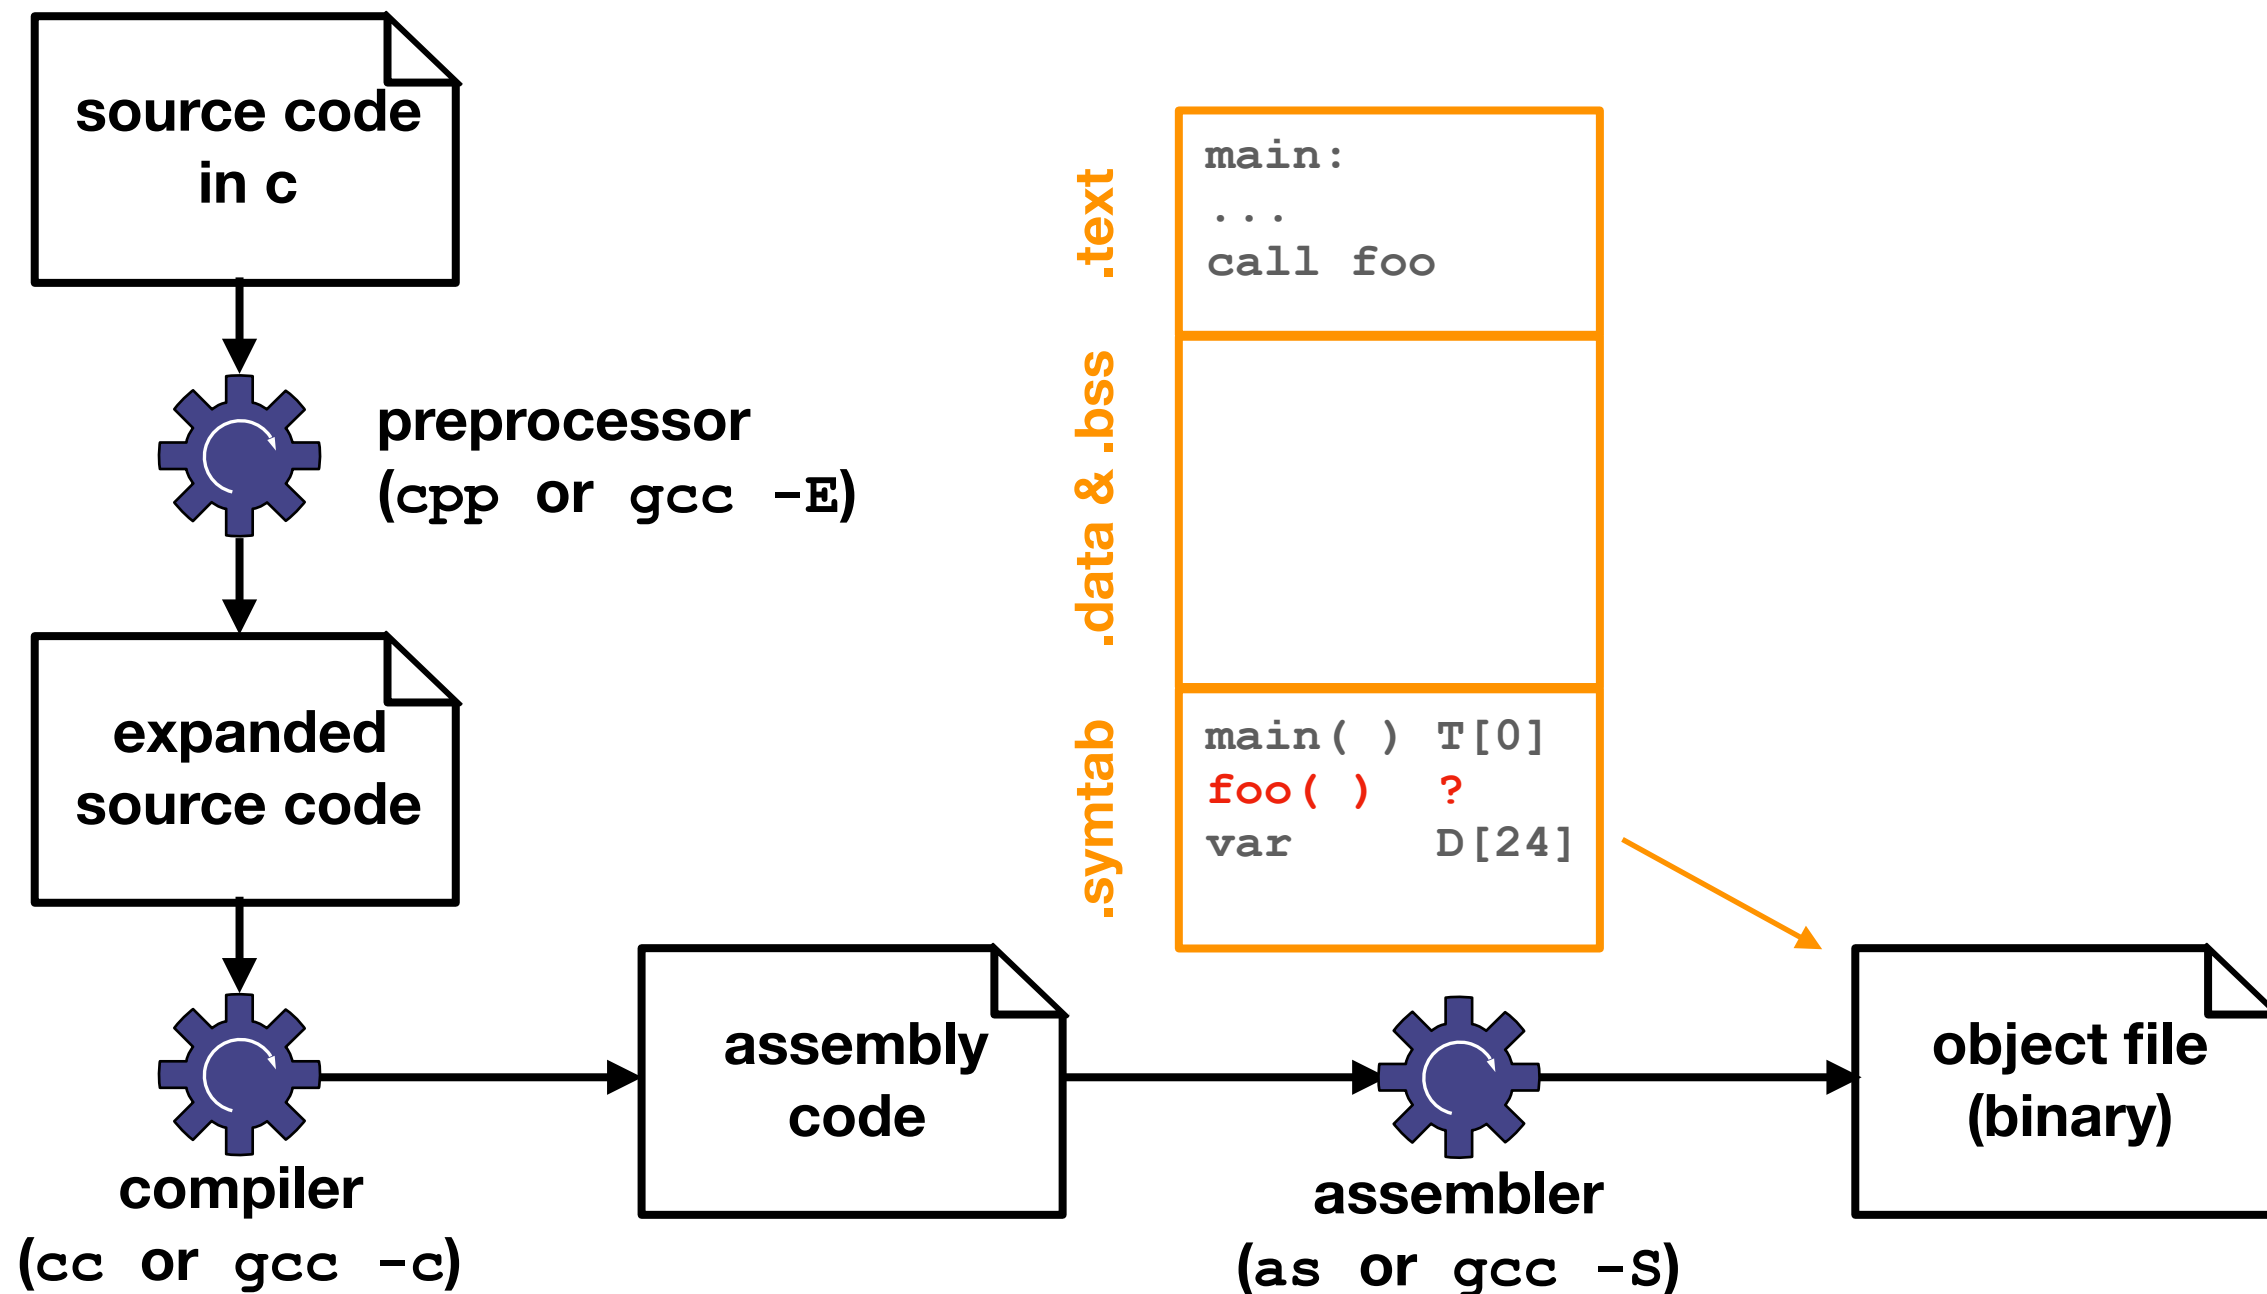
# Compiler, Linker and Loader in action

---



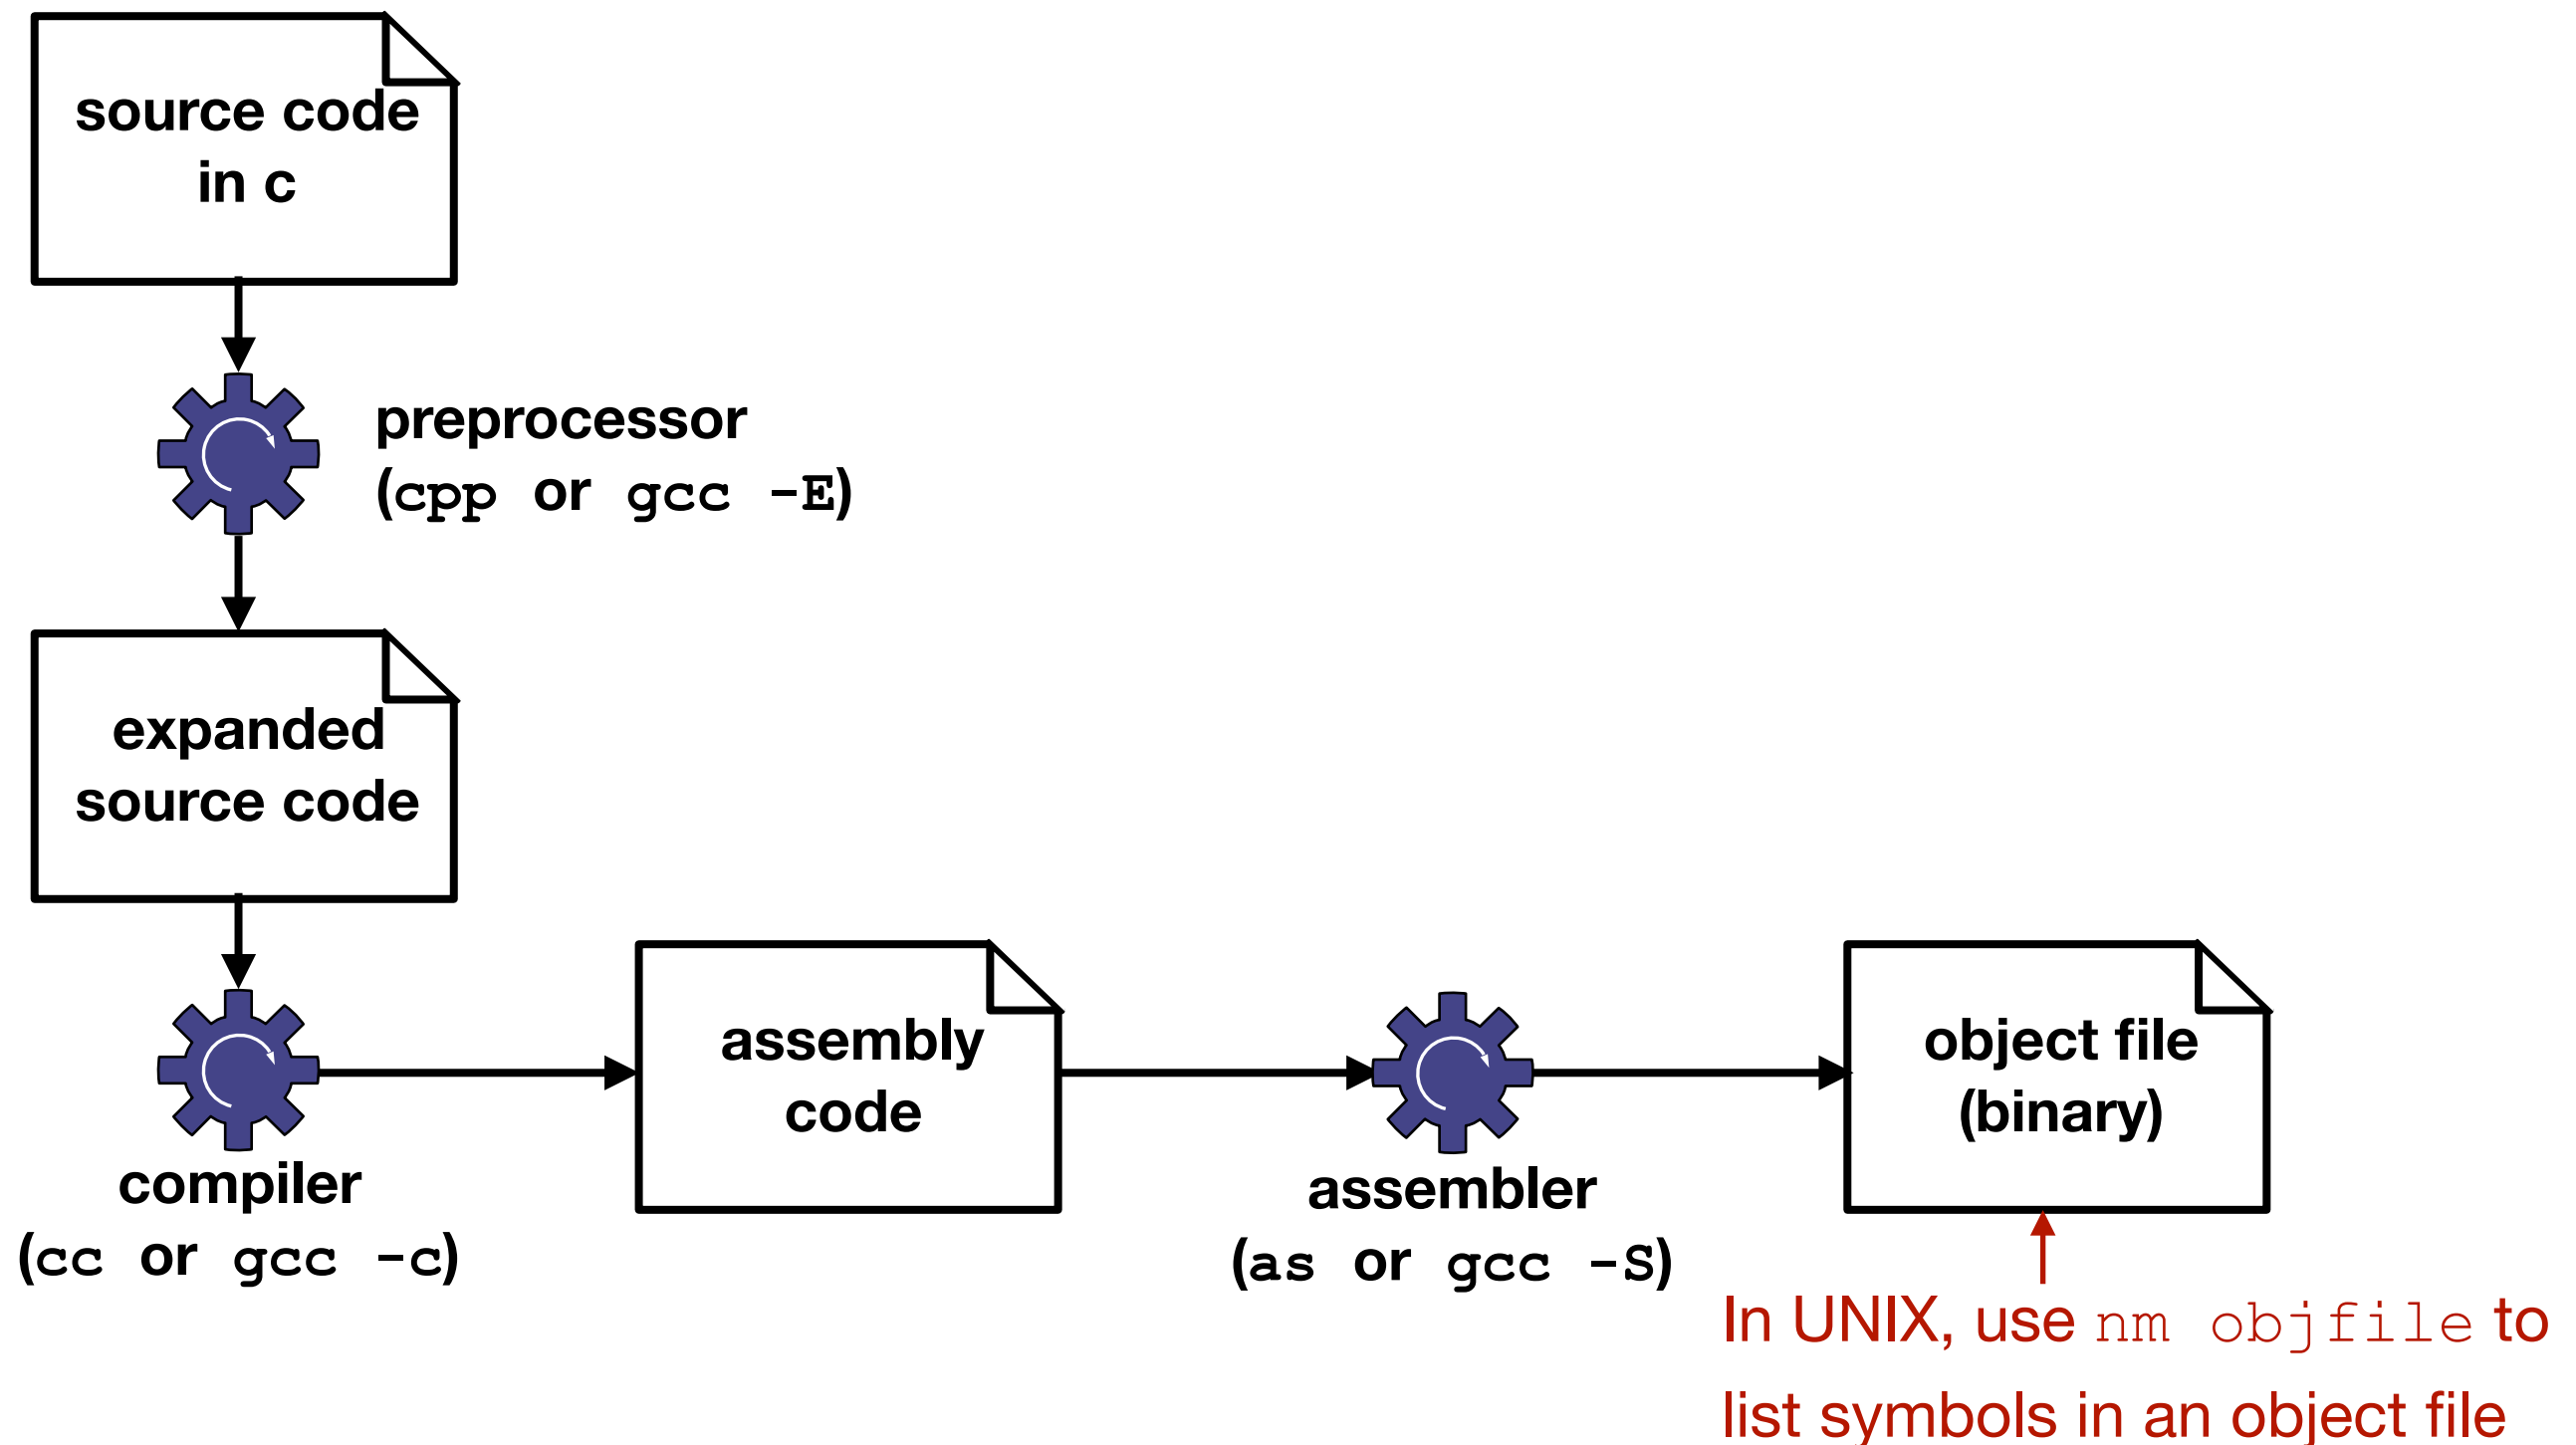


# Compiler, Linker and Loader in action



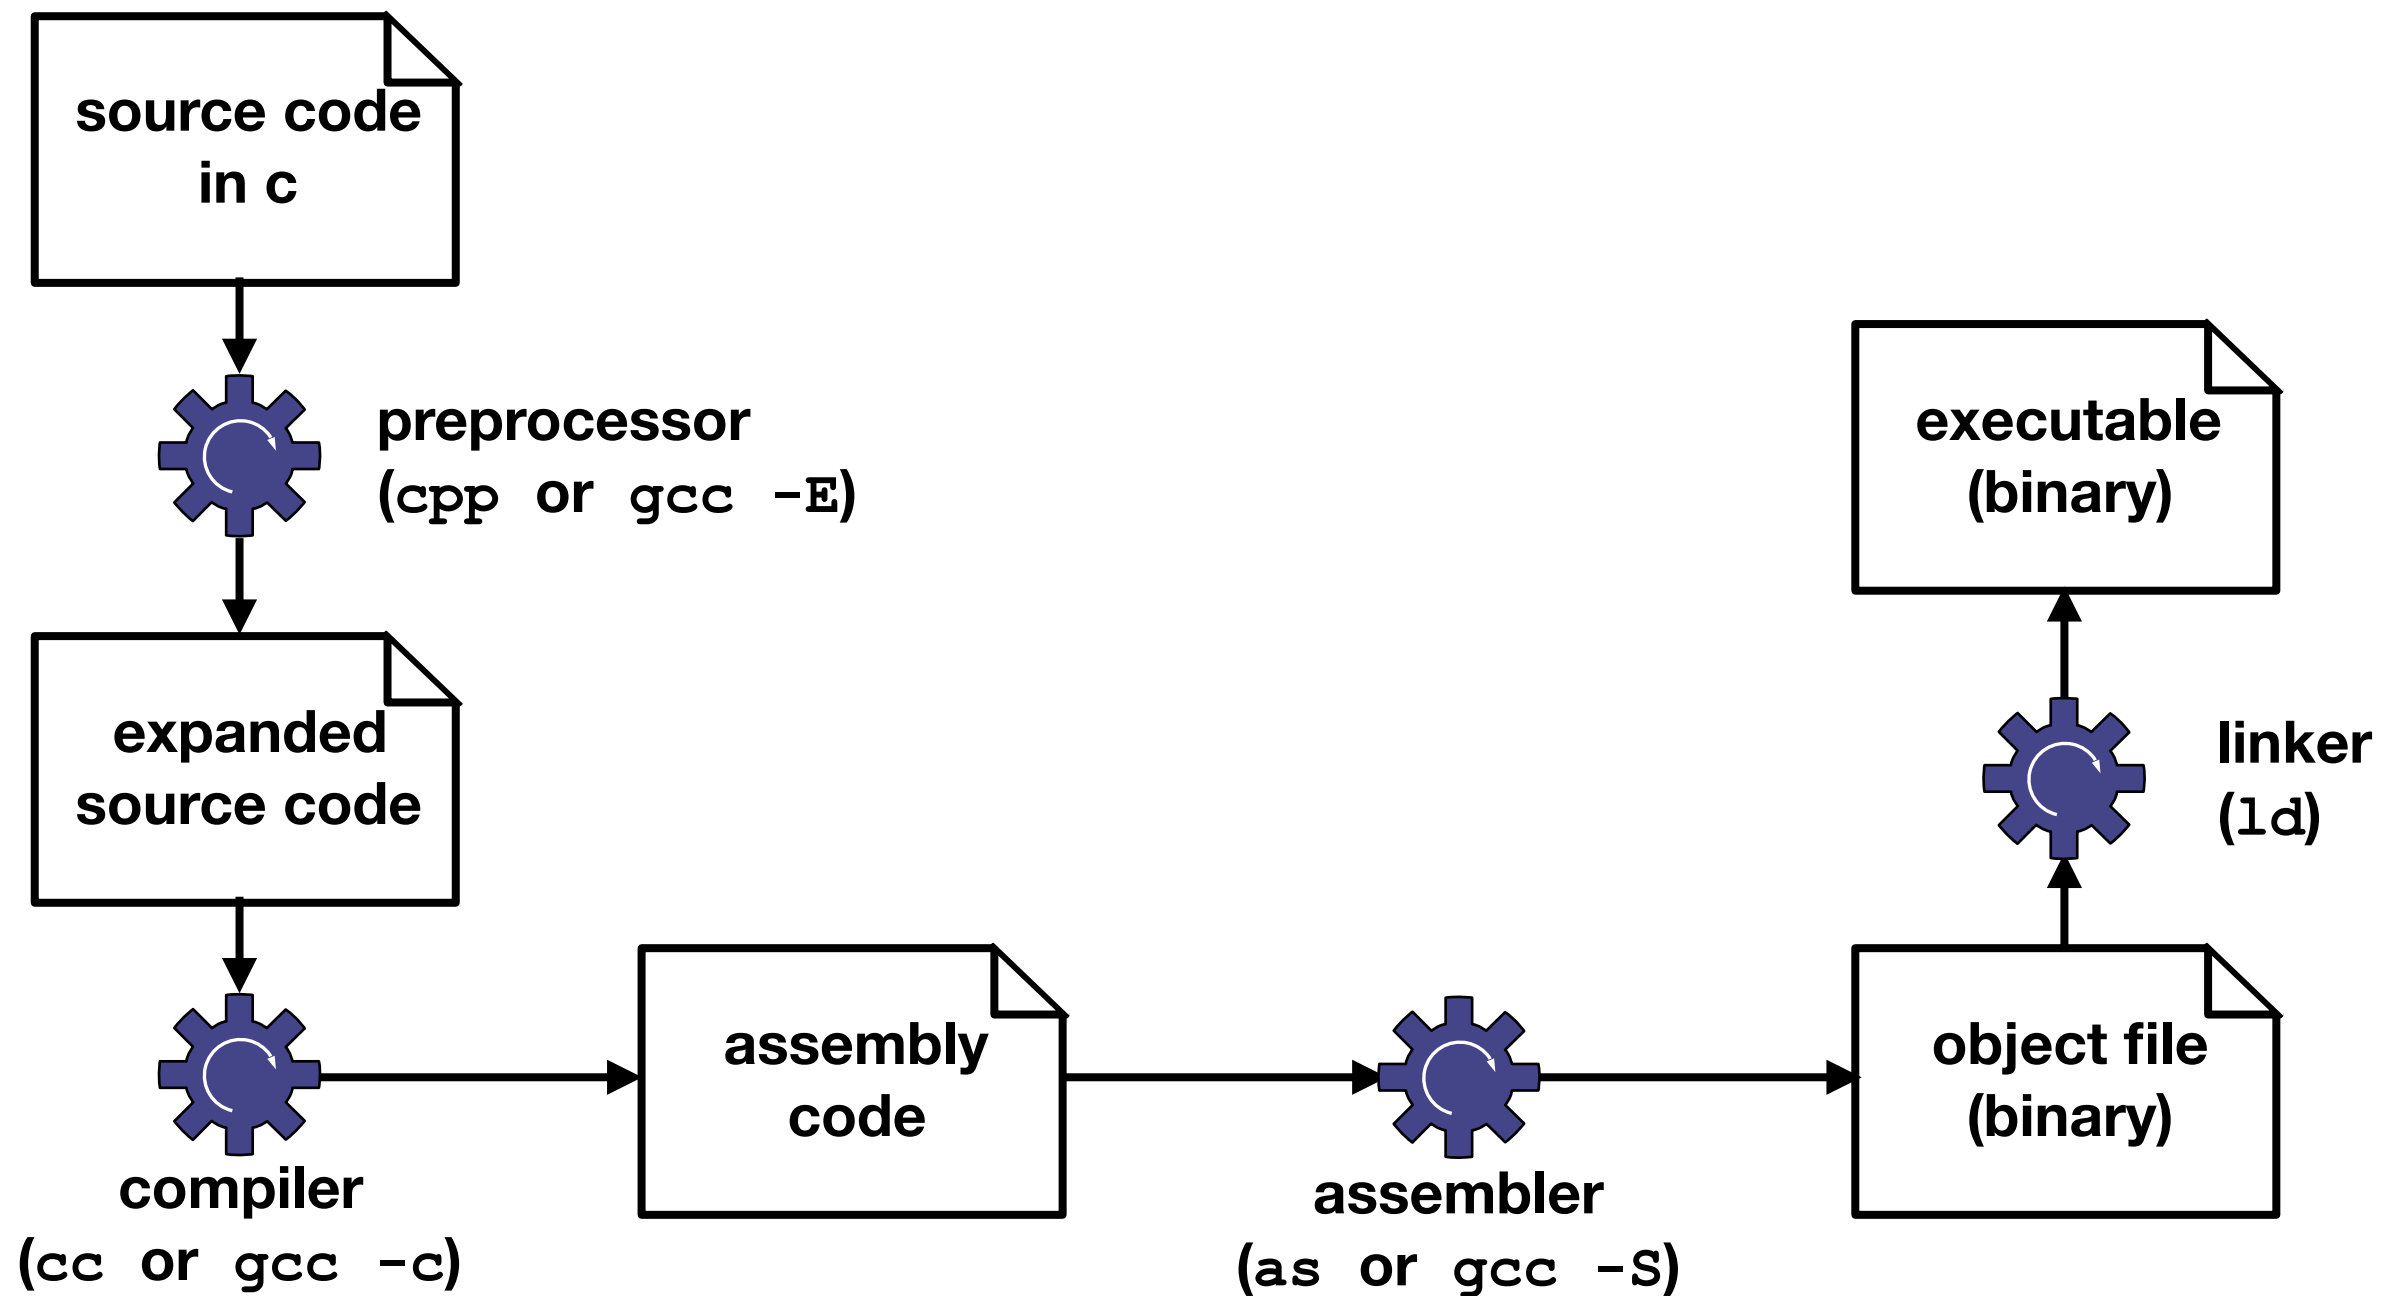
# Compiler, Linker and Loader in action

---

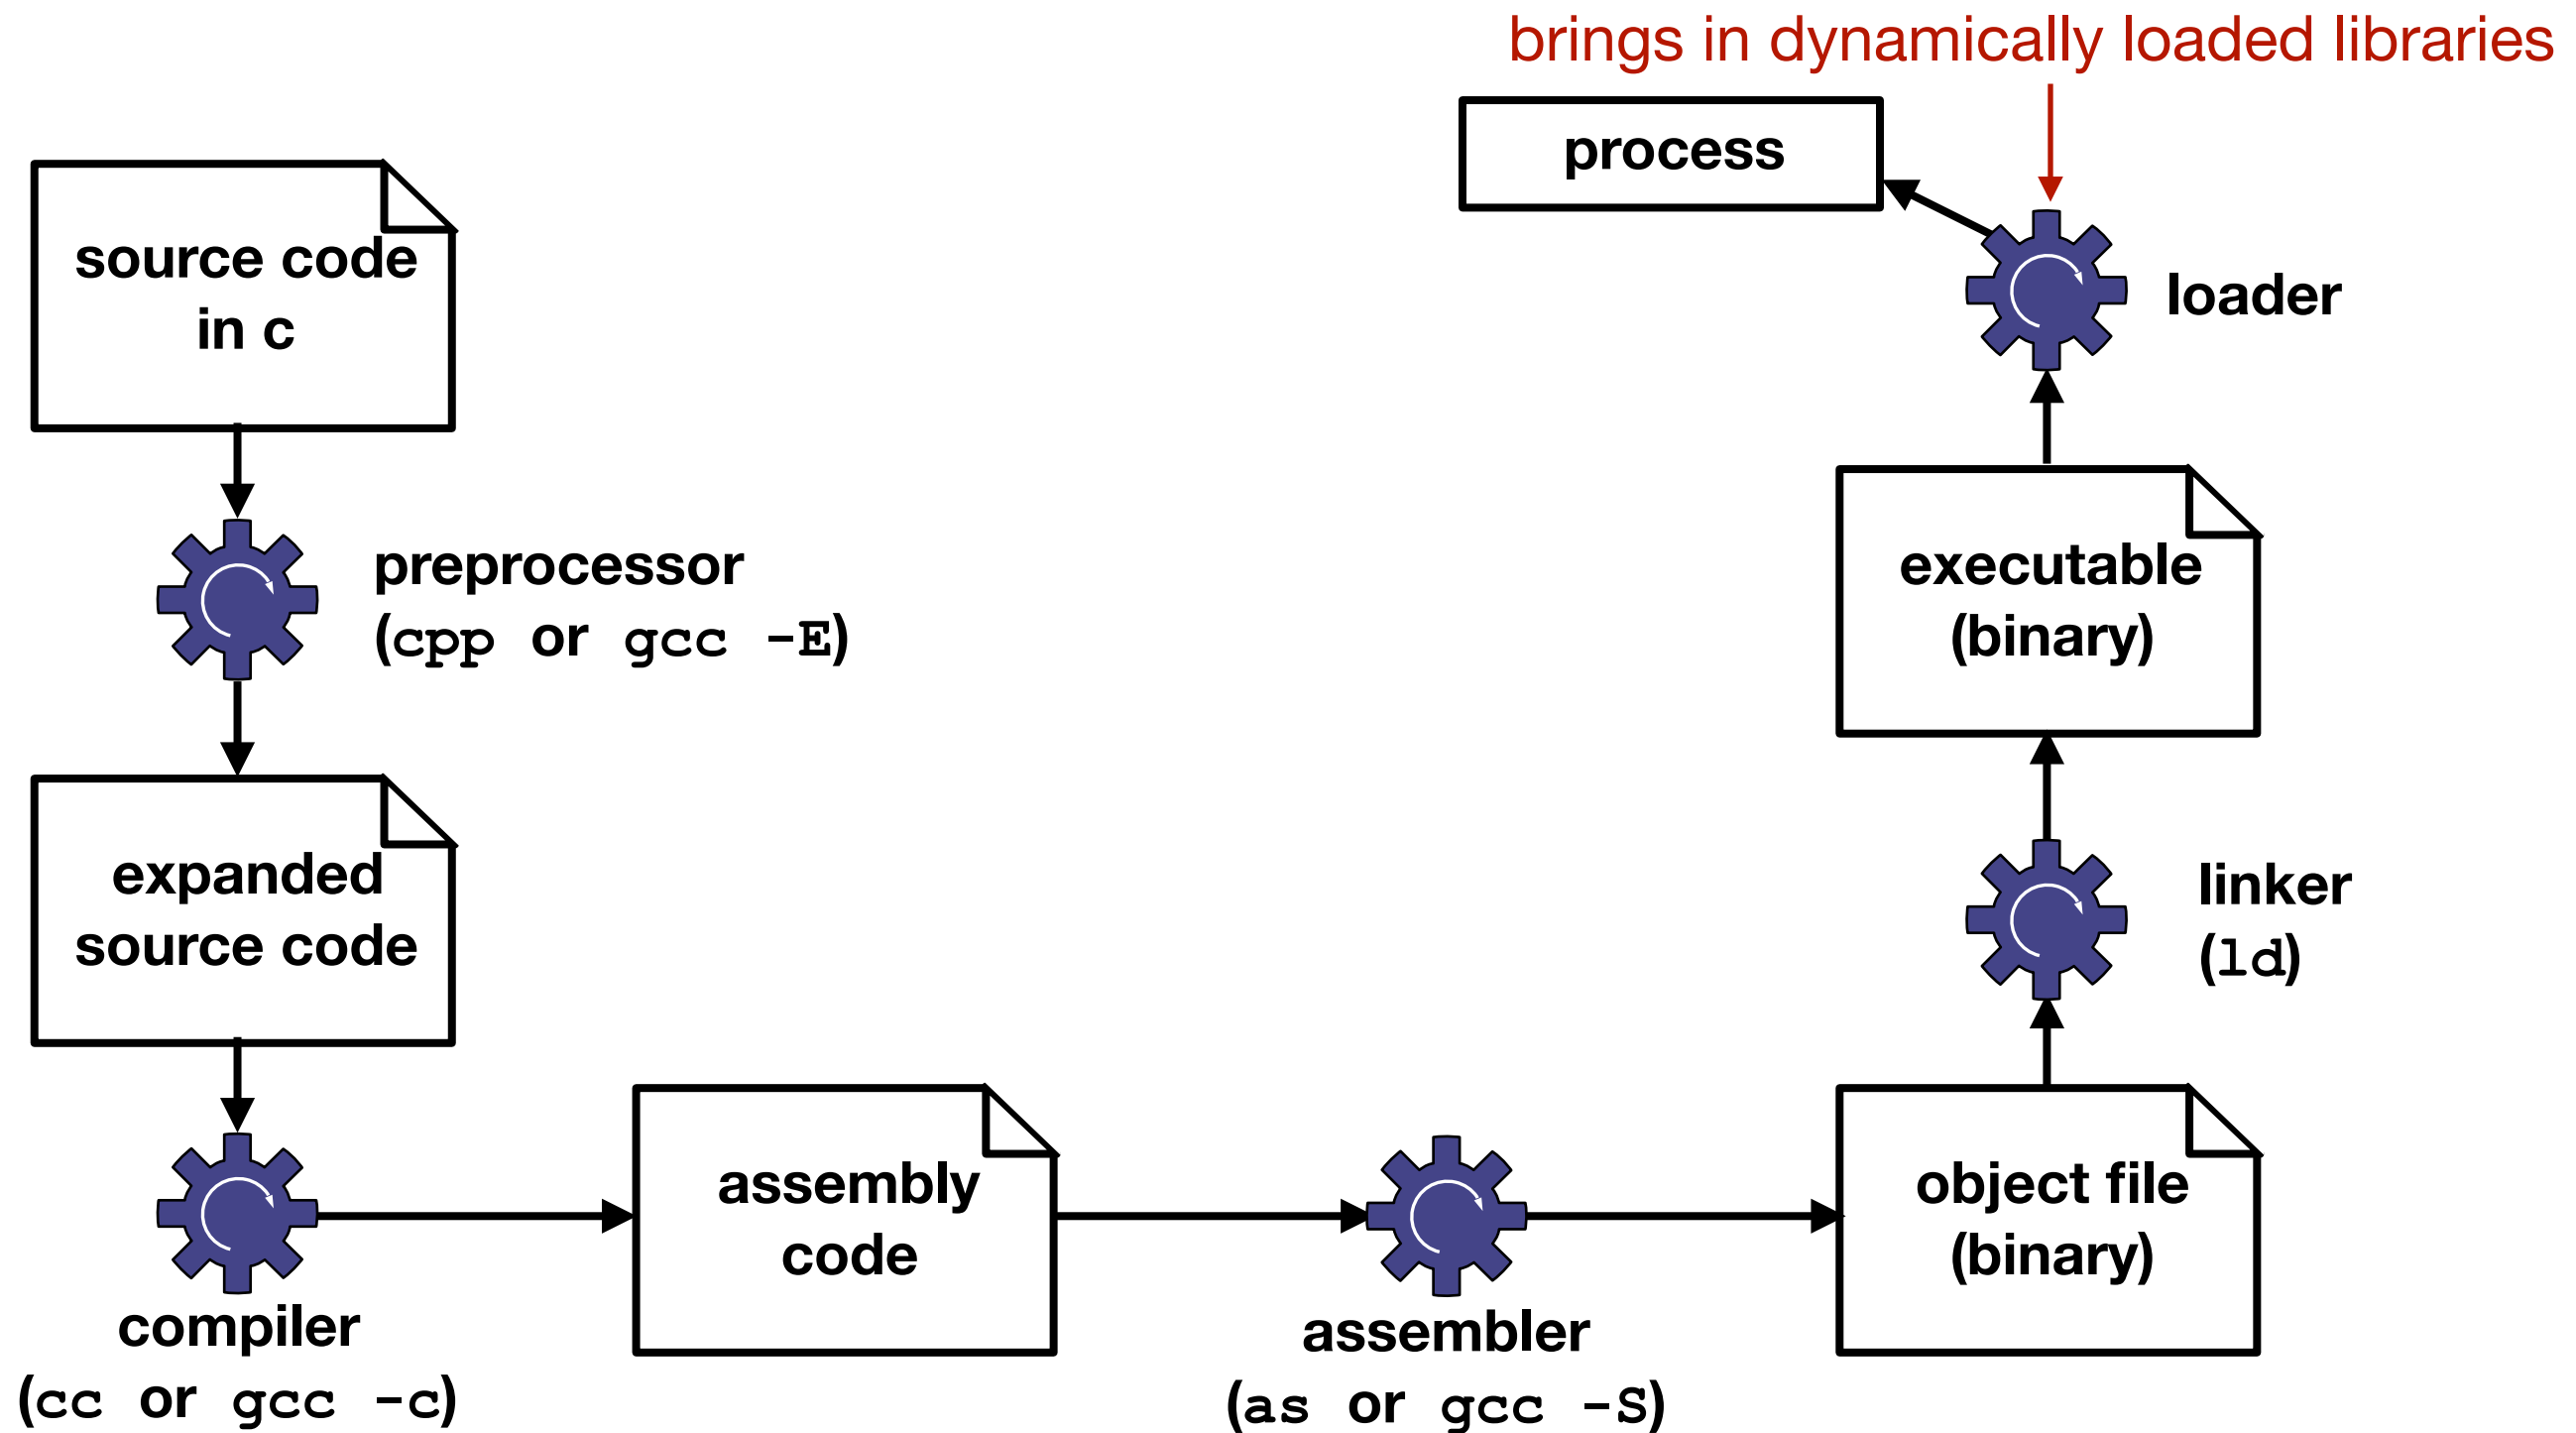


# Compiler, Linker and Loader in action


---



# Compiler, Linker and Loader in action




# Linking



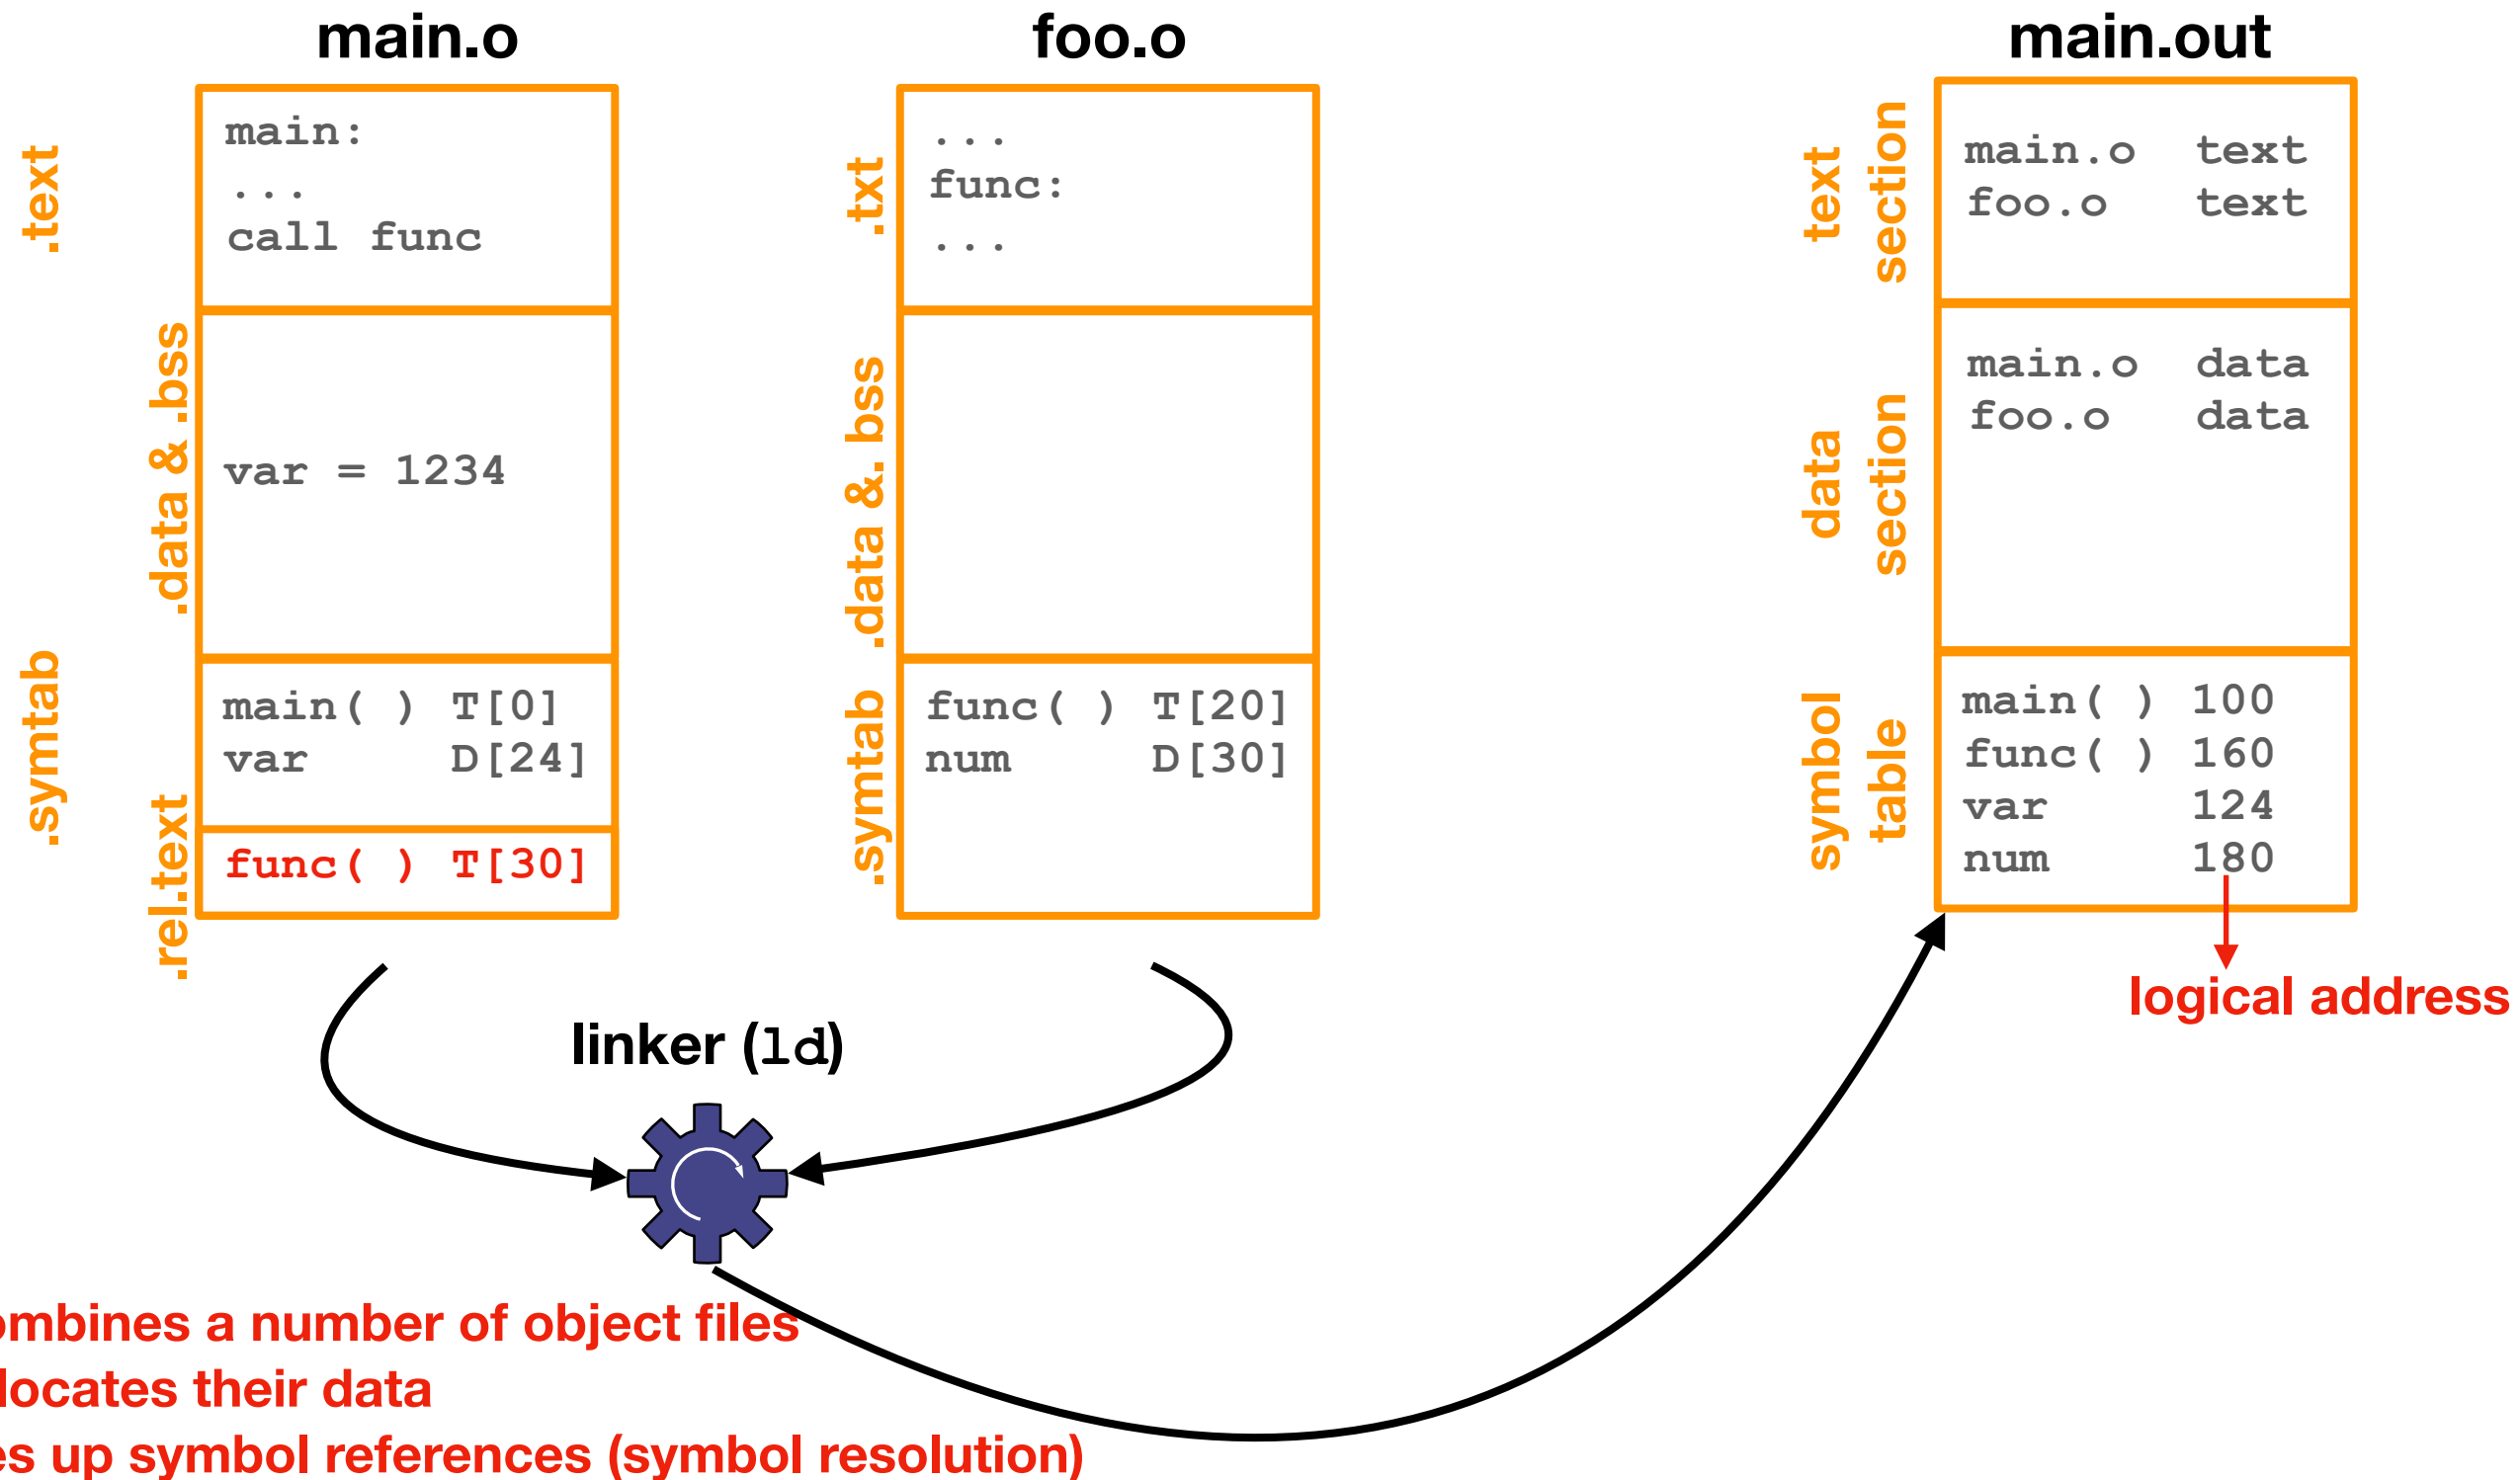
1. combines a number of object files
2. relocates their data
3. ties up symbol references (symbol resolution)

# Linking



1. combines a number of object files
2. relocates their data
3. ties up symbol references (symbol resolution)

# Linking



# Homework

---

- Familiarize yourself with Linux and Windows system calls
- Read the Linux Journal article about linking and loading