

Operating System Concepts

Lecture 5: Process Management

Omid Ardakanian
oardakan@ualberta.ca
University of Alberta

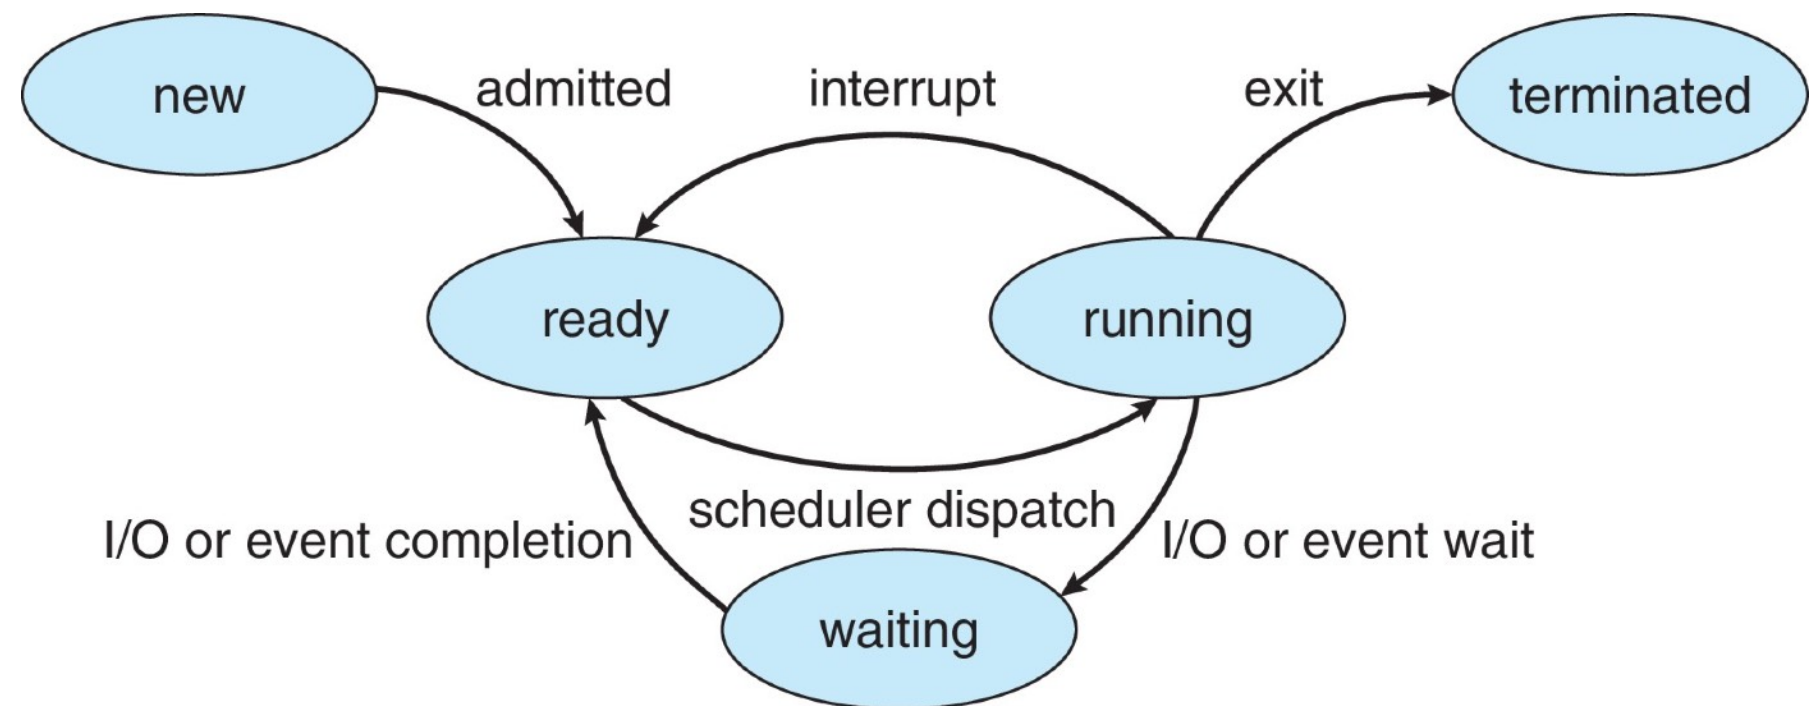
MWF 12:00-12:50 VVC 2 215

Today's class

- Process Abstraction
 - What are the roles of the scheduler and the dispatcher?
 - What happens during context switching?
- Process Control
 - How to create a new process?

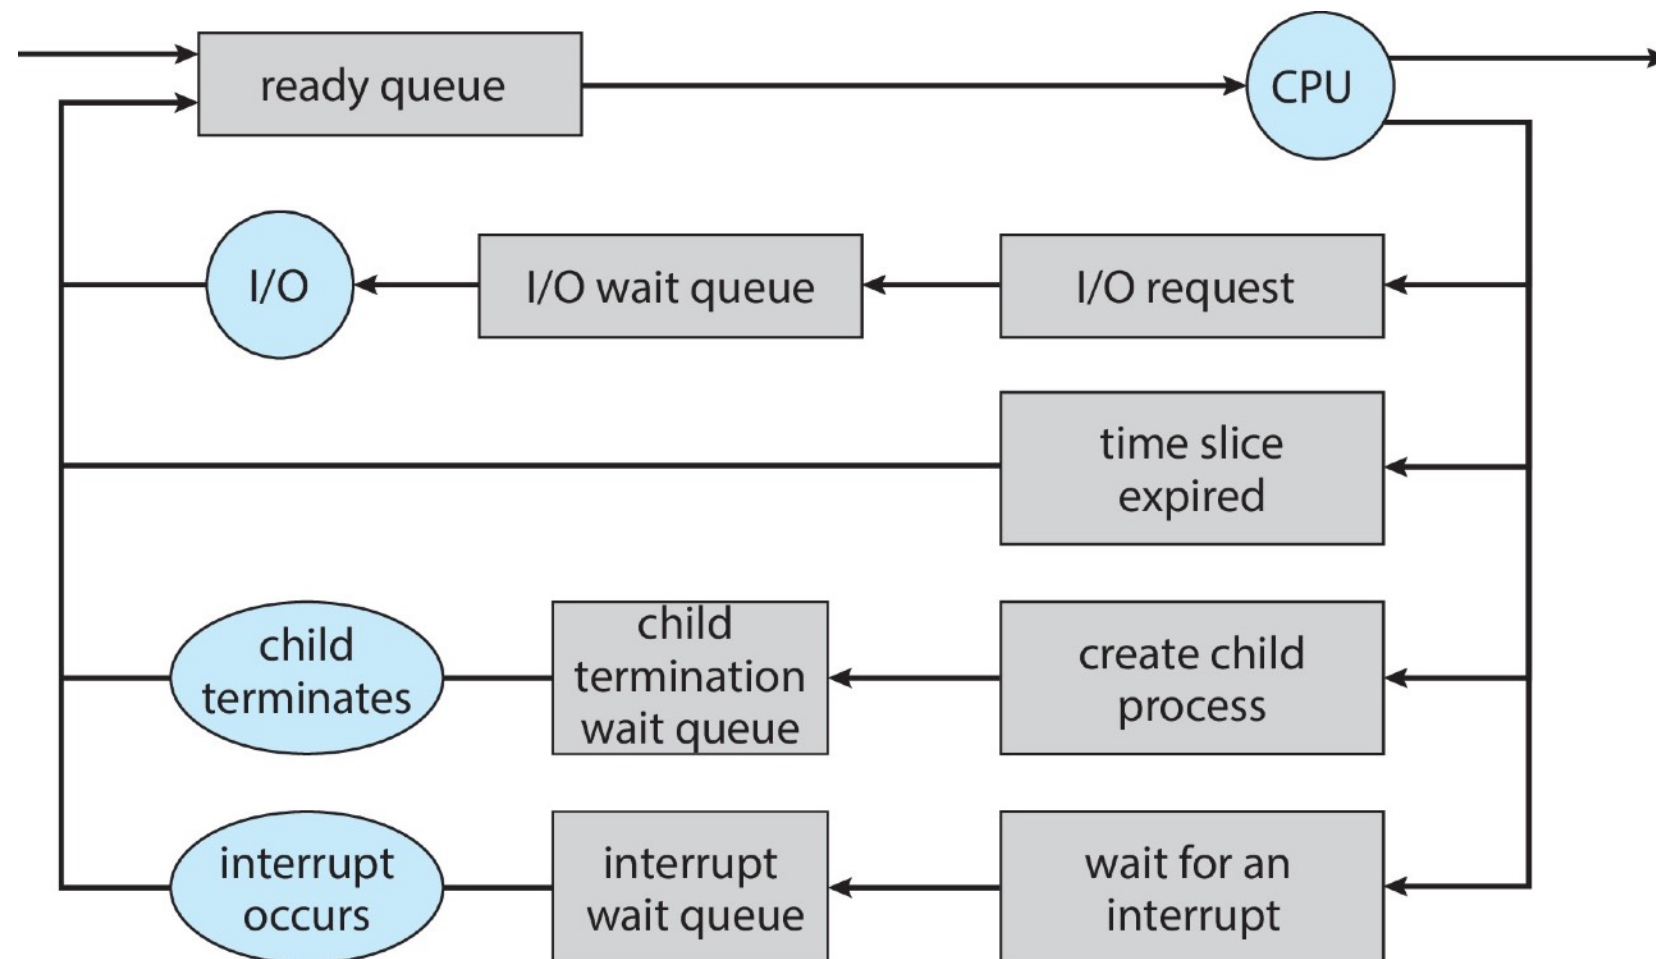
Process lifecycle

- New: being created
- Ready: waiting to be assigned a CPU core
- Running: instructions are running on the CPU core
- Waiting: waiting for some event (e.g., I/O)
- Terminated: finished or aborted




Multiprogramming

- only one process is active at a time (per CPU core)




Multiprogramming

- only one process is active at a time (per CPU core)
- OS gives out CPU to different processes in the ready queue
 - it also gives out other resources to them



Scheduling

- the scheduler maintains a data structure (e.g., a **doubly linked list**) of PCBs
 - PCBs are moved from one queue to another queue
 - job queue -> every process in the system
 - ready queue -> processes residing in the main memory, waiting to run on CPU
 - device queue -> processes waiting for a particular I/O device, each device has its own queue



Scheduling

- the scheduler maintains a data structure (e.g., a **doubly linked list**) of PCBs
 - PCBs are moved from one queue to another queue
 - job queue -> every process in the system
 - ready queue -> processes residing in the main memory, waiting to run on CPU
 - device queue -> processes waiting for a particular I/O device, each device has its own queue
- the scheduler selects a process from the ready queue to run
 - the selected process runs on CPU
 - if no process left in the ready queue, the CPU runs an idle process
 - scheduling can be performed for fairness, minimum latency, providing real-time guarantees, etc.

Scheduling

- the scheduler maintains a data structure (e.g., a **doubly linked list**) of PCBs
 - PCBs are moved from one queue to another queue
 - job queue -> every process in the system
 - ready queue -> processes residing in the main memory, waiting to run on CPU
 - device queue -> processes waiting for a particular I/O device, each device has its own queue
- the scheduler selects a process from the ready queue to run
 - the selected process runs on CPU
 - if no process left in the ready queue, the CPU runs an idle process
 - scheduling can be performed for fairness, minimum latency, providing real-time guarantees, etc.
- the (short-term) scheduler makes a decision very fast (<10 ms) and is called very often (e.g., every 100ms)

Context switching

- a context switch is stopping a process and starting another one
 - it is a relatively expensive operation
 - time-sharing systems may do hundreds of context switches per second
 - the **context** includes the value of CPU registers, the process state, and memory management information

Context switching

- a context switch is stopping a process and starting another one
 - it is a relatively expensive operation
 - time-sharing systems may do hundreds of context switches per second
 - the **context** includes the value of CPU registers, the process state, and memory management information
- OS starts executing a process in the ready state by loading hardware registers (PC, SP, etc) from its PCB
 - while a process is running, the CPU modifies the Program Counter (PC), Stack Pointer (SP), registers, etc.

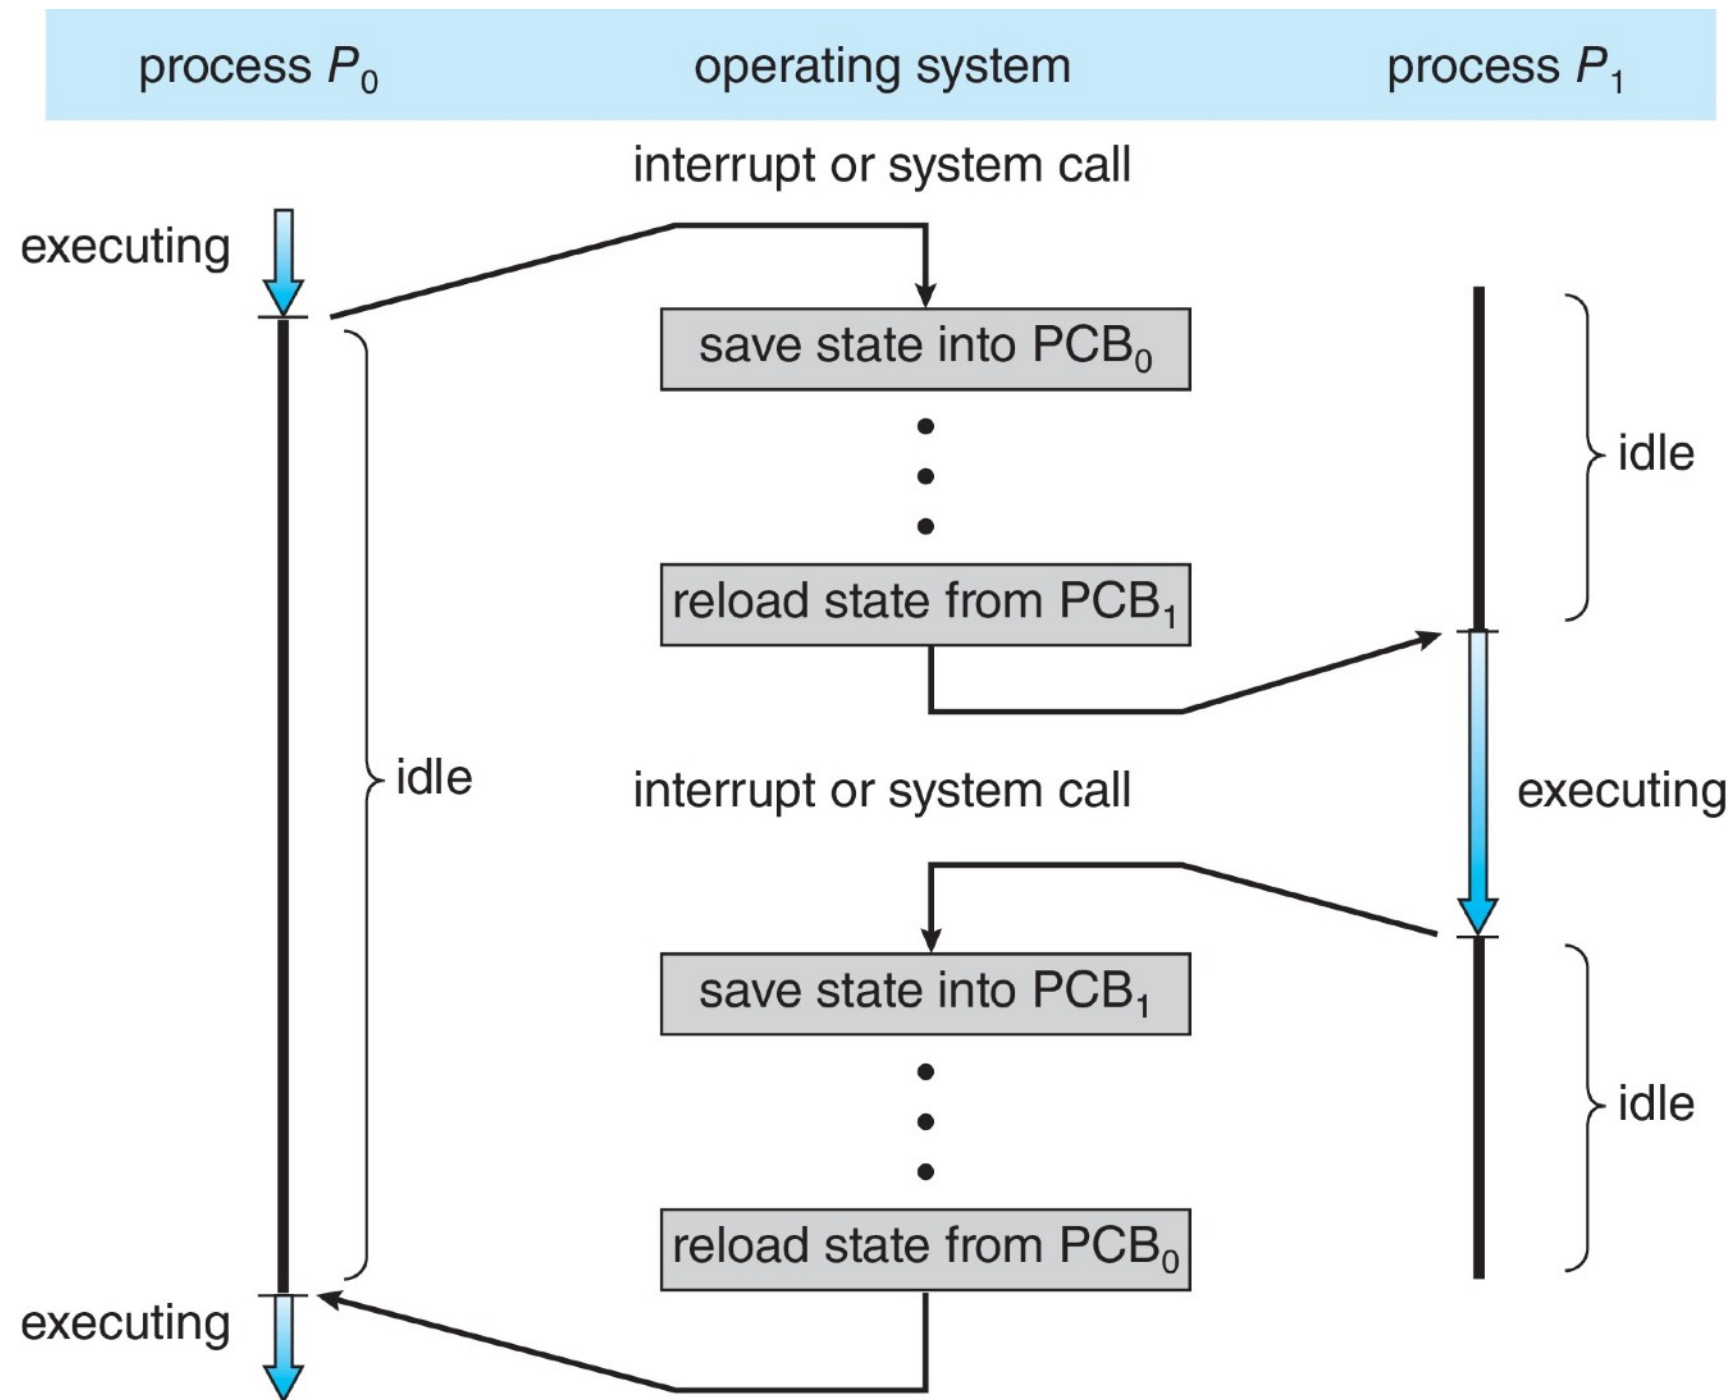
Context switching

- a context switch is stopping a process and starting another one
 - it is a relatively expensive operation
 - time-sharing systems may do hundreds of context switches per second
 - the **context** includes the value of CPU registers, the process state, and memory management information
- OS starts executing a process in the ready state by loading hardware registers (PC, SP, etc) from its PCB
 - while a process is running, the CPU modifies the Program Counter (PC), Stack Pointer (SP), registers, etc.
- when OS stops executing a process, it saves the values of the registers (PC, SP, etc.) into its PCB
 - so that they can be restored the next time the process is selected for running on the CPU

Context switching

- a context switch is stopping a process and starting another one
 - it is a relatively expensive operation
 - time-sharing systems may do hundreds of context switches per second
 - the **context** includes the value of CPU registers, the process state, and memory management information
- OS starts executing a process in the ready state by loading hardware registers (PC, SP, etc) from its PCB
 - while a process is running, the CPU modifies the Program Counter (PC), Stack Pointer (SP), registers, etc.
- when OS stops executing a process, it saves the values of the registers (PC, SP, etc.) into its PCB
 - so that they can be restored the next time the process is selected for running on the CPU
- context switching is pure **overhead**, i.e., the system does no useful work
 - the cost of a context switch and the time between switches are closely related
 - but fast context switching is necessary for responsiveness
 - OS must balance the context switch frequency with the scheduling requirement

Context switching



Scheduler and dispatcher

Scheduler and dispatcher

- the scheduler selects a process to run next

Scheduler and dispatcher

- the scheduler selects a process to run next
- the dispatcher makes it happen
 - performs context switching
 - switches to the user mode
 - jumps to the proper location in the user program

Process Management

Process management system calls in UNIX

Process management system calls in UNIX

- `getpid()` returns the current PID

Process management system calls in UNIX

- `getpid()` returns the current PID
- `fork()` copies the current process (a new PID is assigned to the child process)

Process management system calls in UNIX

- `getpid()` returns the current PID
- `fork()` copies the current process (a new PID is assigned to the child process)
- `exec()` loads a new binary file into memory (without changing its PID)

Process management system calls in UNIX

- `getpid()` returns the current PID
- `fork()` copies the current process (a new PID is assigned to the child process)
- `exec()` loads a new binary file into memory (without changing its PID)
- `wait()` waits until one of its child processes terminates

Process management system calls in UNIX

- `getpid()` returns the current PID
- `fork()` copies the current process (a new PID is assigned to the child process)
- `exec()` loads a new binary file into memory (without changing its PID)
- `wait()` waits until one of its child processes terminates
- `waitpid()` waits until the specified child process terminates

Process management system calls in UNIX

- `getpid()` returns the current PID
- `fork()` copies the current process (a new PID is assigned to the child process)
- `exec()` loads a new binary file into memory (without changing its PID)
- `wait()` waits until one of its child processes terminates
- `waitpid()` waits until the specified child process terminates
- `exit()` terminates a process

Process management system calls in UNIX

- `getpid()` returns the current PID
- `fork()` copies the current process (a new PID is assigned to the child process)
- `exec()` loads a new binary file into memory (without changing its PID)
- `wait()` waits until one of its child processes terminates
- `waitpid()` waits until the specified child process terminates
- `exit()` terminates a process
- `kill()` sends a signal (interrupt-like notification) to another process

Process management system calls in UNIX

- `getpid()` returns the current PID
- `fork()` copies the current process (a new PID is assigned to the child process)
- `exec()` loads a new binary file into memory (without changing its PID)
- `wait()` waits until one of its child processes terminates
- `waitpid()` waits until the specified child process terminates
- `exit()` terminates a process
- `kill()` sends a signal (interrupt-like notification) to another process
- `pause()` causes the calling process to sleep until a signal is delivered that either terminates the process or causes the invocation of a signal-catching function

Process management system calls in UNIX

- `getpid()` returns the current PID
- `fork()` copies the current process (a new PID is assigned to the child process)
- `exec()` loads a new binary file into memory (without changing its PID)
- `wait()` waits until one of its child processes terminates
- `waitpid()` waits until the specified child process terminates
- `exit()` terminates a process
- `kill()` sends a signal (interrupt-like notification) to another process
- `pause()` causes the calling process to sleep until a signal is delivered that either terminates the process or causes the invocation of a signal-catching function
- `nanosleep()` suspends execution of a process for at least the specified time (can be interrupted by a signal that triggers the invocation of a handler)

Process management system calls in UNIX

- `getpid()` returns the current PID
- `fork()` copies the current process (a new PID is assigned to the child process)
- `exec()` loads a new binary file into memory (without changing its PID)
- `wait()` waits until one of its child processes terminates
- `waitpid()` waits until the specified child process terminates
- `exit()` terminates a process
- `kill()` sends a signal (interrupt-like notification) to another process
- `pause()` causes the calling process to sleep until a signal is delivered that either terminates the process or causes the invocation of a signal-catching function
- `nanosleep()` suspends execution of a process for at least the specified time (can be interrupted by a signal that triggers the invocation of a handler)
- `sigaction()` sets handlers for signals


Creating a process with the fork system call

- the `fork()` system call creates a **child** process which inherits a copy of its **parent**'s memory, file descriptors, CPU registers, etc.

Creating a process with the fork system call

- the `fork()` system call creates a **child** process which inherits a copy of its **parent**'s memory, file descriptors, CPU registers, etc.

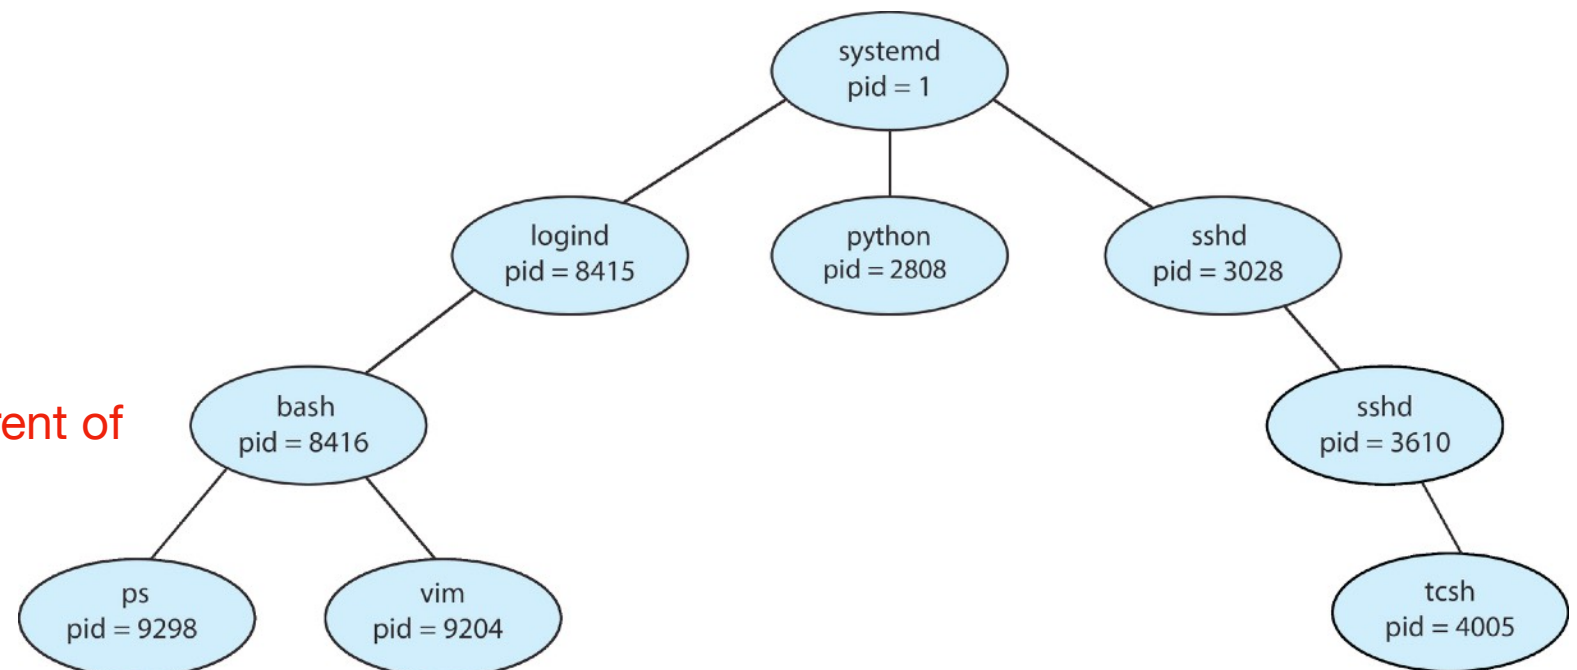
the `systemd` process is the root of this tree (parent of all user processes created when the system has booted); it has PID = 1



Creating a process with the fork system call

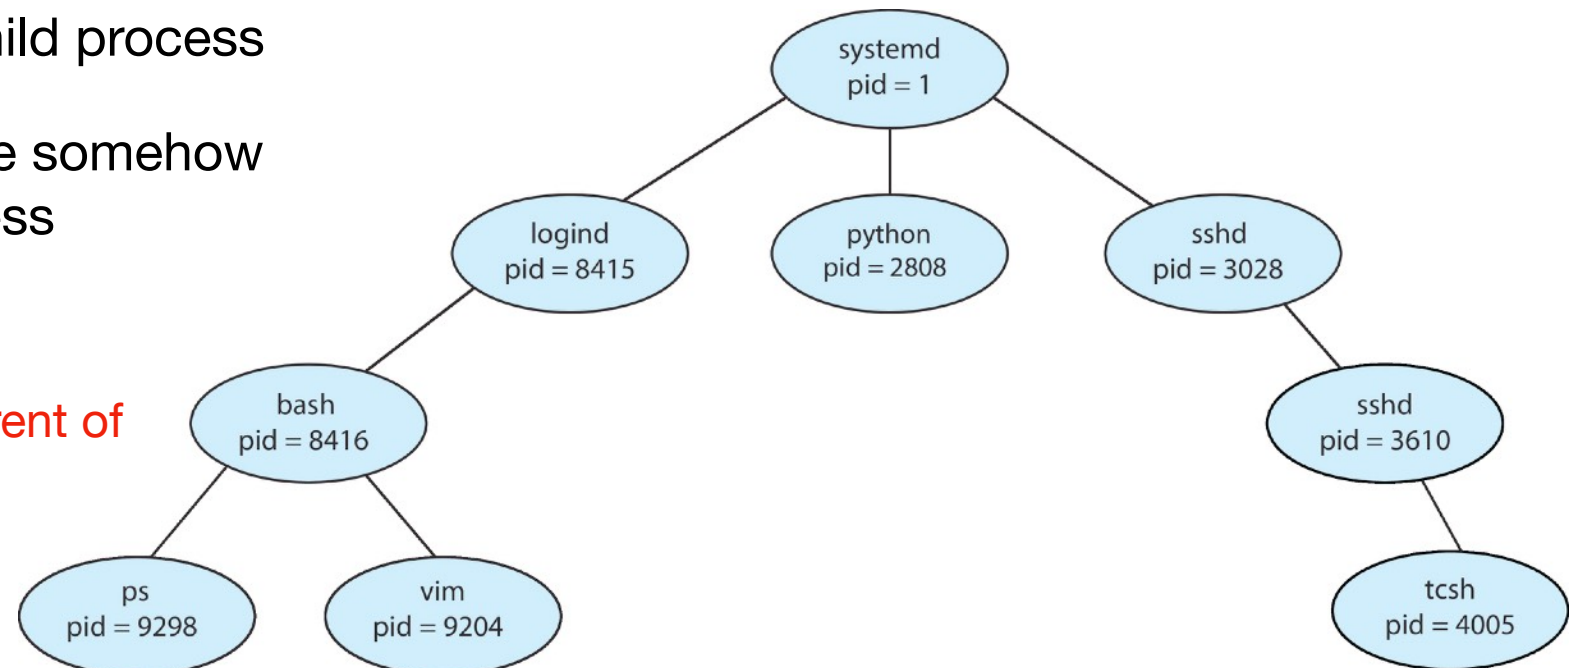
- the `fork()` system call creates a **child** process which inherits a copy of its **parent**'s memory, file descriptors, CPU registers, etc.
- both parent and child processes execute from the instruction following `fork()`
 - either the child or the parent might run first

the `systemd` process is the root of this tree (parent of all user processes created when the system has booted); it has PID = 1



Creating a process with the fork system call


- the `fork()` system call creates a **child** process which inherits a copy of its **parent**'s memory, file descriptors, CPU registers, etc.
- both parent and child processes execute from the instruction following `fork()`
 - either the child or the parent might run first
- the return value from `fork()` is of type `pid_t` (like an integer)
 - when > 0 : running in (original) parent process and the return value is pid of new child
 - when $= 0$: running in new child process
 - when < 0 : error! must handle somehow still running in original process



the `systemd` process is the root of this tree (parent of all user processes created when the system has booted); it has PID = 1

What happens on a fork?

```
pid_t cid = fork( );
```



Common problems with fork

Common problems with fork

the `fork()` system call is

Common problems with fork

the `fork()` system call is

- inefficient and slow
 - the cost of copying the entire address space of a process is high

Common problems with fork

the `fork()` system call is

- inefficient and slow
 - the cost of copying the entire address space of a process is high
- insecure
 - the parent process must explicitly remove states that the child process does not need (scrubbing secrets from memory)

Common problems with fork

the `fork()` system call is

- inefficient and slow
 - the cost of copying the entire address space of a process is high
- insecure
 - the parent process must explicitly remove states that the child process does not need (scrubbing secrets from memory)
- not thread-safe
 - the child process created by the fork system call will have a single thread only (a copy of the calling thread)
 - **Problem:** one thread doing memory allocation and holding a **heap lock**, while another thread forks. Any attempt to allocate memory in the child (and thus acquire the same lock) will immediately **deadlock** waiting for an unlock operation that will never happen
 - **Solution:** not using fork in a multithreaded process, or calling `exec` immediately afterwards

Program loading with the exec system call

- the `exec ()` system call allows a process to load a different program and start execution at main

Program loading with the exec system call

- the `exec ()` system call allows a process to load a different program and start execution at main
- it allows a process to specify the number of arguments (`argc`) and the string argument array (`argv`)

Program loading with the exec system call

- the `exec ()` system call allows a process to load a different program and start execution at main
- it allows a process to specify the number of arguments (`argc`) and the string argument array (`argv`)
- if the call is successful, the same process runs a different program

Program loading with the exec system call

- the `exec ()` system call allows a process to load a different program and start execution at main
- it allows a process to specify the number of arguments (`argc`) and the string argument array (`argv`)
- if the call is successful, the same process runs a different program
- code, data, stack and heap sections are overwritten

Program loading with the exec system call

- the `exec ()` system call allows a process to load a different program and start execution at main
- it allows a process to specify the number of arguments (`argc`) and the string argument array (`argv`)
- if the call is successful, the same process runs a different program
- code, data, stack and heap sections are overwritten
- in most cases we call `exec ()` after calling `fork ()`
 - hence, the memory copied during `fork ()` is useless
 - the `vfork ()` system call allows for creating a process without creating an identical memory image; in this case child process must call `exec ()` immediately

Waiting for the child process to terminate

Waiting for the child process to terminate

- the parent can execute concurrently with its children or can wait until some or all of them terminate

Waiting for the child process to terminate

- the parent can execute concurrently with its children or can wait until some or all of them terminate
- the `wait()` system call enables the parent process to wait for the child process to terminate and receive its return value

Waiting for the child process to terminate

- the parent can execute concurrently with its children or can wait until some or all of them terminate
- the `wait()` system call enables the parent process to wait for the child process to terminate and receive its return value
 - it puts the parent to sleep waiting for a child's result

Waiting for the child process to terminate

- the parent can execute concurrently with its children or can wait until some or all of them terminate
- the `wait()` system call enables the parent process to wait for the child process to terminate and receive its return value
 - it puts the parent to sleep waiting for a child's result
 - when a child calls `exit()`, the OS unblocks the parent and returns the value passed by `exit()` as a result of the wait call (along with the pid of the child)

Waiting for the child process to terminate

- the parent can execute concurrently with its children or can wait until some or all of them terminate
- the `wait()` system call enables the parent process to wait for the child process to terminate and receive its return value
 - it puts the parent to sleep waiting for a child's result
 - when a child calls `exit()`, the OS unblocks the parent and returns the value passed by `exit()` as a result of the wait call (along with the pid of the child)
 - if there are no children alive, `wait()` returns immediately

Waiting for the child process to terminate

- the parent can execute concurrently with its children or can wait until some or all of them terminate
- the `wait()` system call enables the parent process to wait for the child process to terminate and receive its return value
 - it puts the parent to sleep waiting for a child's result
 - when a child calls `exit()`, the OS unblocks the parent and returns the value passed by `exit()` as a result of the wait call (along with the pid of the child)
 - if there are no children alive, `wait()` returns immediately
 - also, if there are zombies waiting for their parents, `wait()` returns one of the values immediately (and deallocates the zombie)

Waiting for the child process to terminate

- the parent can execute concurrently with its children or can wait until some or all of them terminate
- the `wait()` system call enables the parent process to wait for the child process to terminate and receive its return value
 - it puts the parent to sleep waiting for a child's result
 - when a child calls `exit()`, the OS unblocks the parent and returns the value passed by `exit()` as a result of the wait call (along with the pid of the child)
 - if there are no children alive, `wait()` returns immediately
 - also, if there are zombies waiting for their parents, `wait()` returns one of the values immediately (and deallocates the zombie)

