

Operating System Concepts

Lecture 6: Process Control

Omid Ardakanian
oardakan@ualberta.ca
University of Alberta


MWF 12:00-12:50 VVC 2 215

Today's class

- Process Control
 - How to create a new process?
 - How to terminate a process?
- Examples

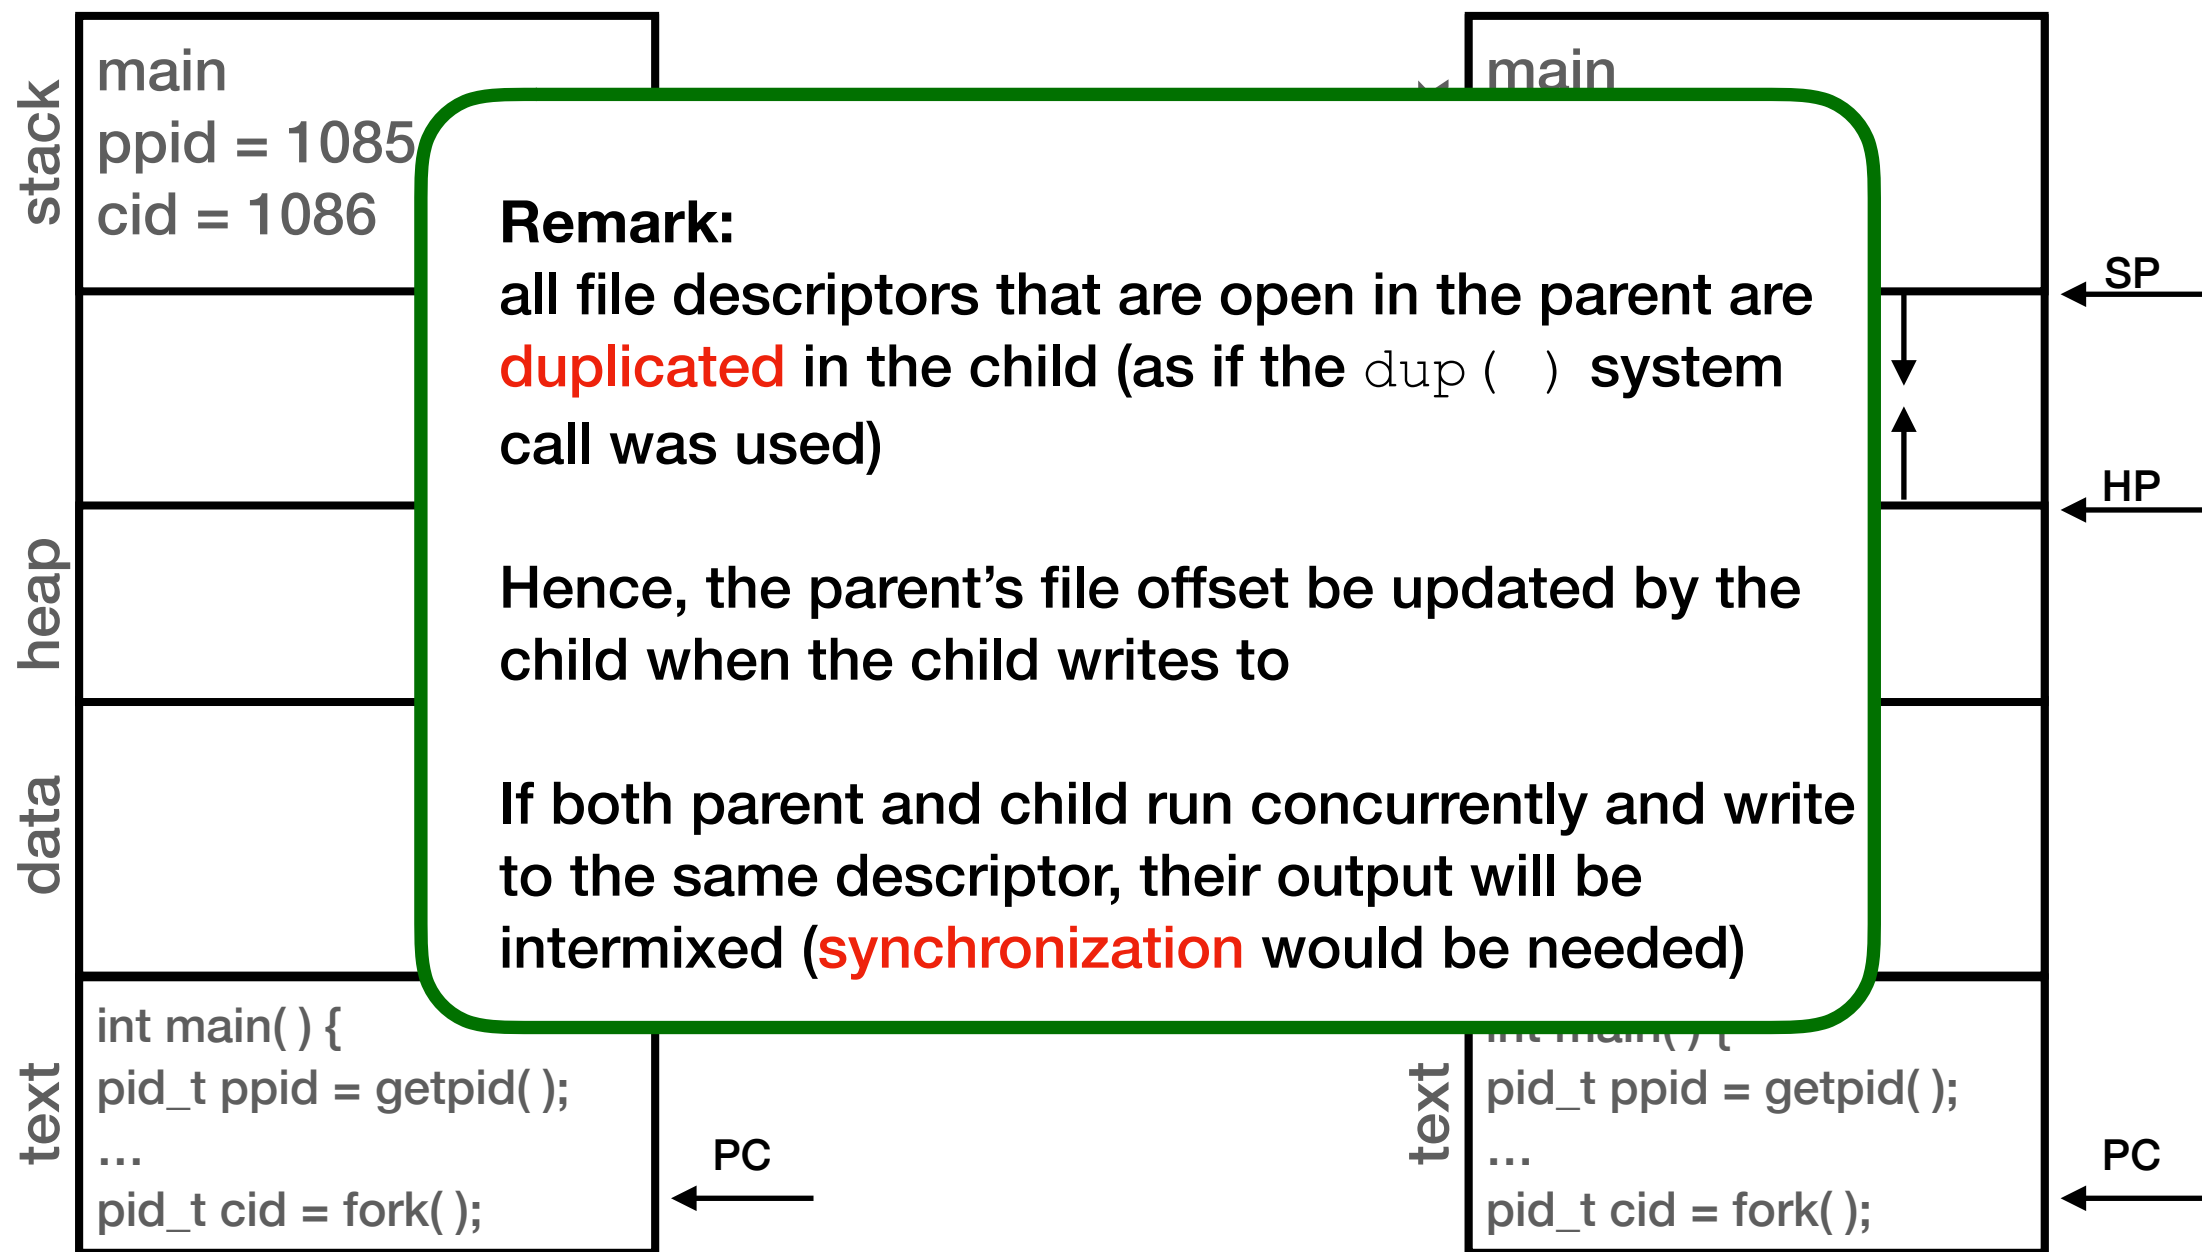
What happens on a fork?

```
pid_t cid = fork( );
```



What happens on a fork?

```
pid_t cid = fork( );
```



Terminating a process

- Process termination is the ultimate resource reclamation by the OS
 - closes all open files, connections, etc.
 - deallocates memory and most of the OS structures supporting the process
 - checks if parent is alive
 - if so, holds the **exit status** until parent requests it; in this case, process does not really die, but it enters the zombie state (**WHY?**)
 - If not, it deallocates all data structures; the process is dead at this point
 - cleans up all waiting zombies

Normal and abnormal termination

- a process can terminate normally by returning from `main`, or directly calling the standard C library function `exit()` or the system call `_exit()`;

When the `main` function returns, `exit()` is called indirectly

- `exit()` calls all exit handlers that have been registered using `atexit()` — a glibc function
- `_exit()` doesn't call exit handlers

Normal and abnormal termination

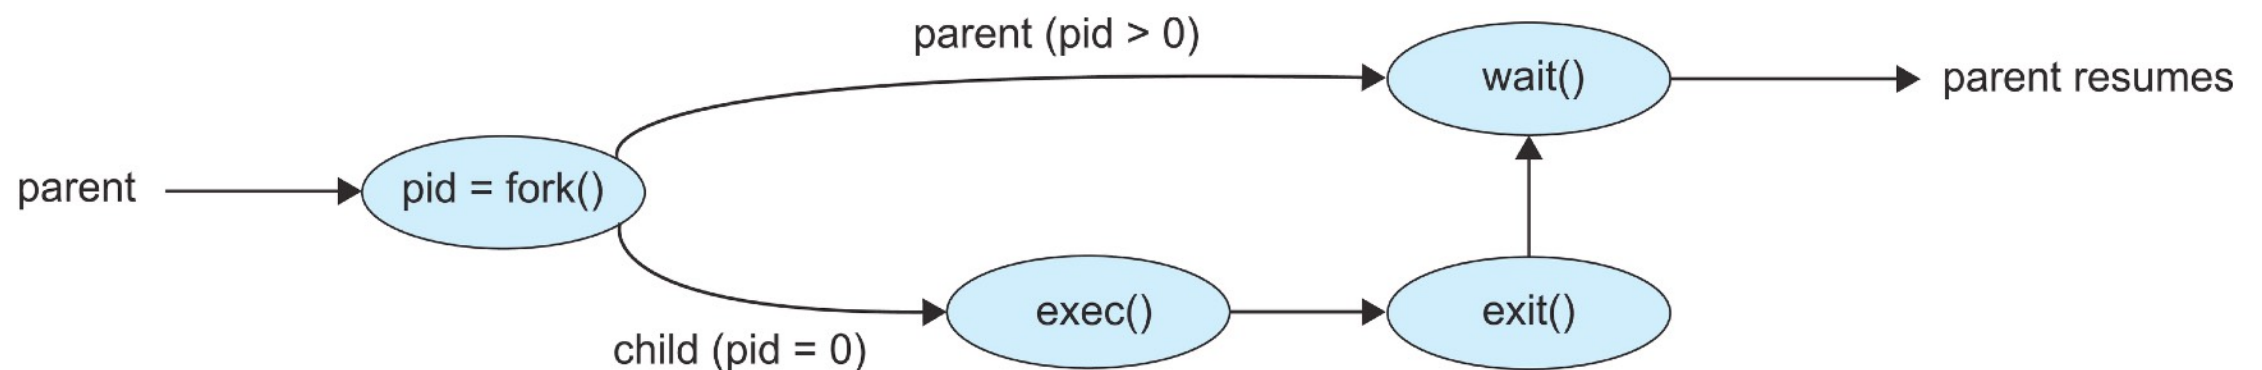
- a process can terminate normally by returning from `main`, or directly calling the standard C library function `exit()` or the system call `_exit()`;
When the `main` function returns, `exit()` is called indirectly
 - `exit()` calls all exit handlers that have been registered using `atexit()` — a glibc function
 - `_exit()` doesn't call exit handlers
- for abnormal termination of a process we can call `abort()` which generates `SIGABRT`
 - functions registered using `atexit()` are not called
 - it may not close open files or flush stream buffers

Normal and abnormal termination

- a process can terminate normally by returning from `main`, or directly calling the standard C library function `exit()` or the system call `_exit()`;
When the `main` function returns, `exit()` is called indirectly
 - `exit()` calls all exit handlers that have been registered using `atexit()` — a glibc function
 - `_exit()` doesn't call exit handlers
- for abnormal termination of a process we can call `abort()` which generates `SIGABRT`
 - functions registered using `atexit()` are not called
 - it may not close open files or flush stream buffers
- a process can terminate a child using the `kill()` system call
 - `kill(cid, SIGKILL)`


Waiting for the child process to terminate

- the parent can execute concurrently with its children or can wait until some or all of them terminate




Waiting for the child process to terminate

- the parent can execute concurrently with its children or can wait until some or all of them terminate
- the `wait()` system call enables the parent process to wait for the child process to terminate and receive its return value



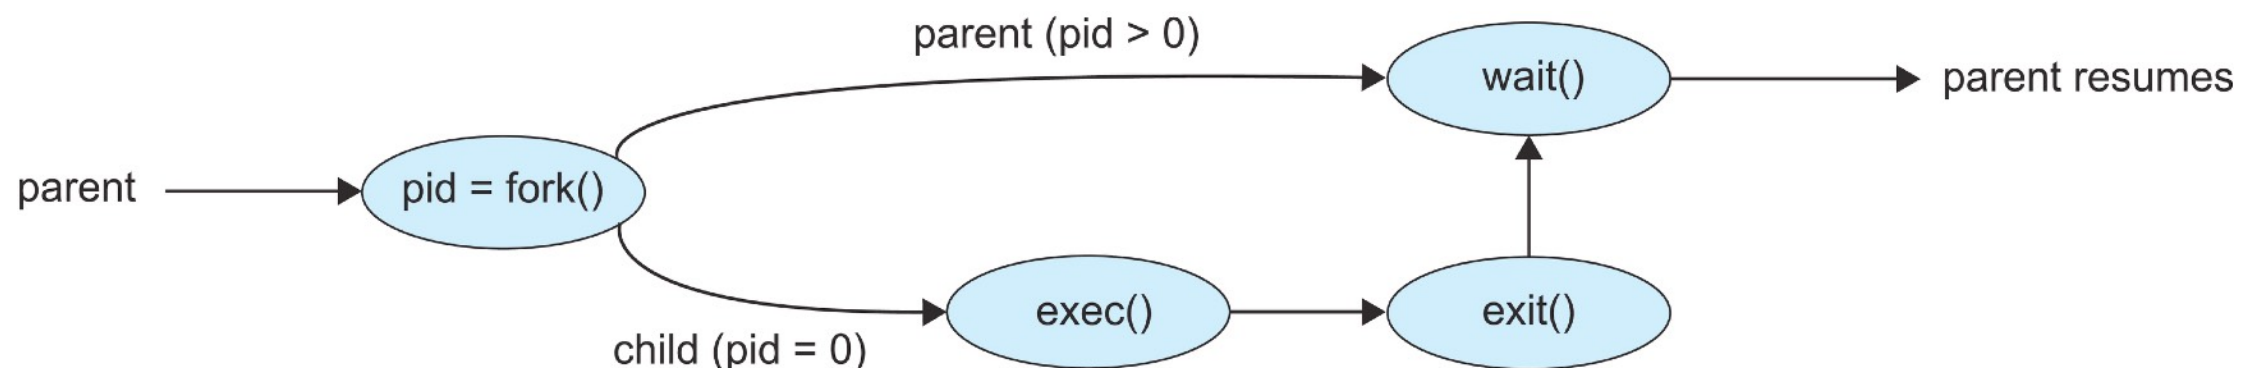
Waiting for the child process to terminate

- the parent can execute concurrently with its children or can wait until some or all of them terminate
- the `wait()` system call enables the parent process to wait for the child process to terminate and receive its return value
 - it puts the parent to sleep waiting for a child's result



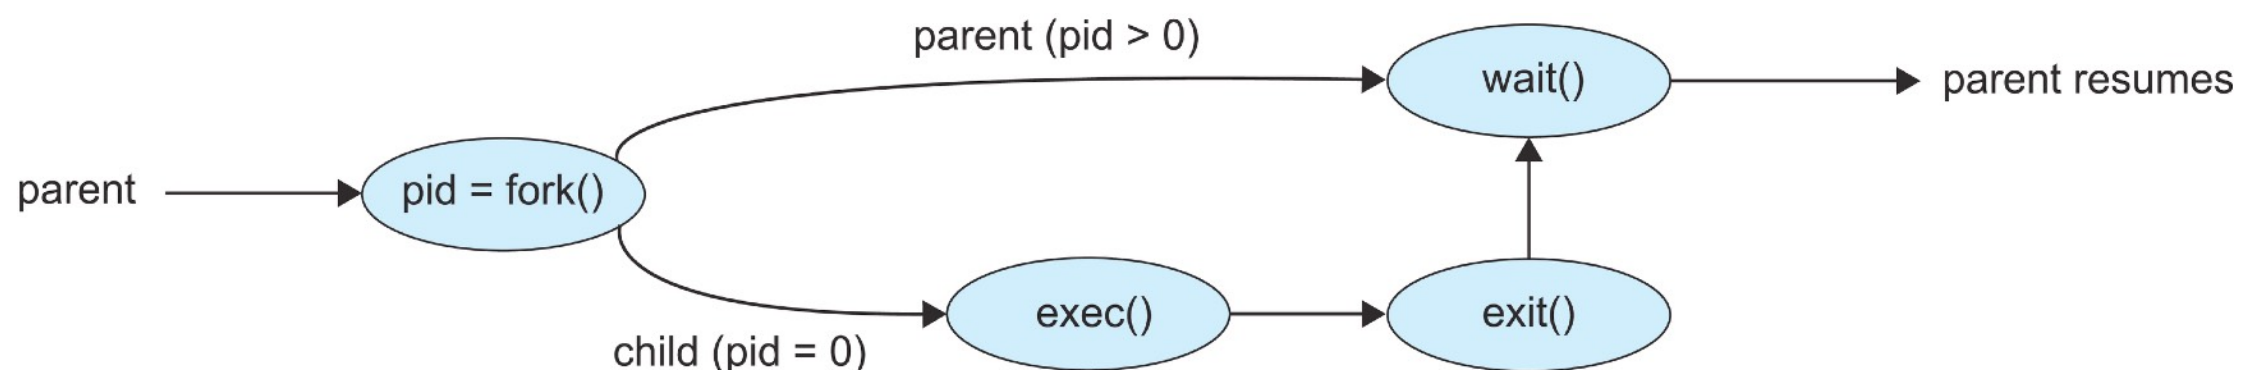
Waiting for the child process to terminate

- the parent can execute concurrently with its children or can wait until some or all of them terminate
- the `wait()` system call enables the parent process to wait for the child process to terminate and receive its return value
 - it puts the parent to sleep waiting for a child's result
 - when a child calls `exit()`, the kernel notifies the parent by sending the `SIGCHLD` signal to the parent; this unblocks the parent and returns the child's exit value along with the child's PID



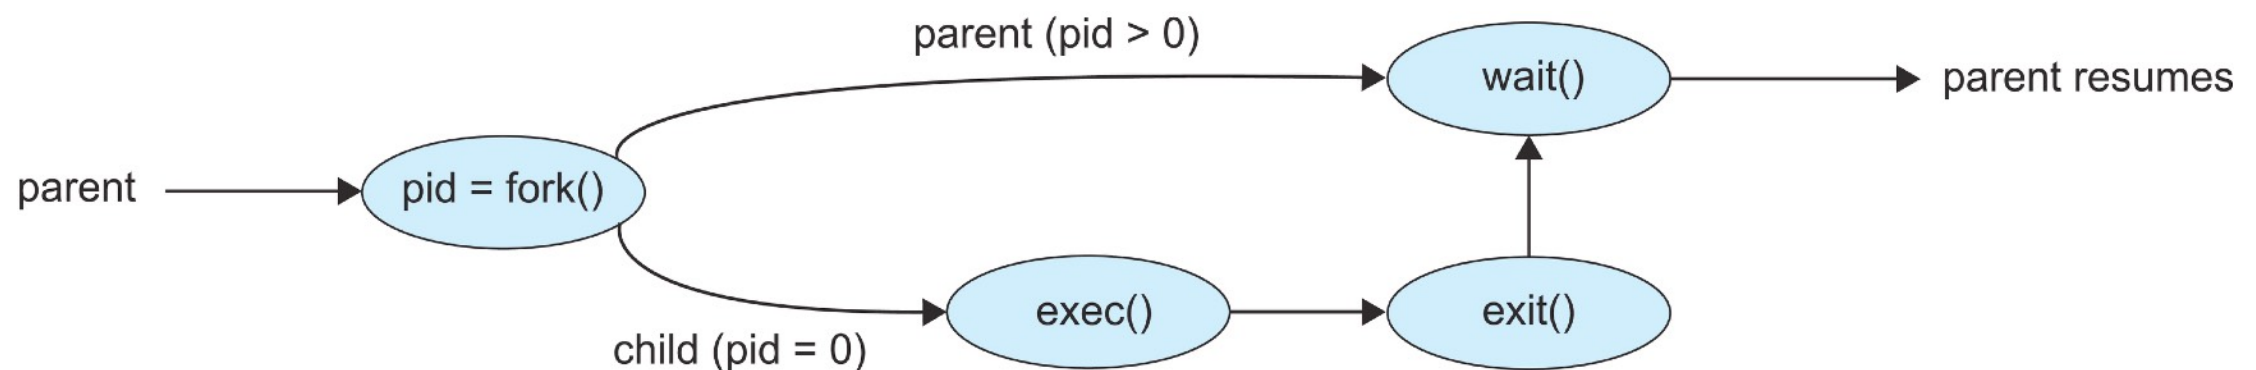
Waiting for the child process to terminate

- the parent can execute concurrently with its children or can wait until some or all of them terminate
- the `wait()` system call enables the parent process to wait for the child process to terminate and receive its return value
 - it puts the parent to sleep waiting for a child's result
 - when a child calls `exit()`, the kernel notifies the parent by sending the `SIGCHLD` signal to the parent; this unblocks the parent and returns the child's exit value along with the child's PID
 - if there are no children alive, `wait()` returns immediately



Waiting for the child process to terminate

- the parent can execute concurrently with its children or can wait until some or all of them terminate
- the `wait()` system call enables the parent process to wait for the child process to terminate and receive its return value
 - it puts the parent to sleep waiting for a child's result
 - when a child calls `exit()`, the kernel notifies the parent by sending the `SIGCHLD` signal to the parent; this unblocks the parent and returns the child's exit value along with the child's PID
 - if there are no children alive, `wait()` returns immediately
 - also, if there are zombies waiting for their parents, `wait()` returns one of the exit statuses immediately (and deallocates the zombie)



Zombie and orphan processes

- a process that has terminated, but its parent has not (yet) called `wait ()` becomes **zombie**

Zombie and orphan processes

- a process that has terminated, but its parent has not (yet) called `wait()` becomes **zombie**
 - the `ps` command prints the state of a zombie process as Z

Zombie and orphan processes

- a process that has terminated, but its parent has not (yet) called `wait()` becomes **zombie**
 - the `ps` command prints the state of a zombie process as Z
- a process becomes **orphan** when its parent terminates while it is still running

Zombie and orphan processes

- a process that has terminated, but its parent has not (yet) called `wait()` becomes **zombie**
 - the `ps` command prints the state of a zombie process as Z
- a process becomes **orphan** when its parent terminates while it is still running
 - hence the parent terminates without invoking the `wait()` system call

Zombie and orphan processes

- a process that has terminated, but its parent has not (yet) called `wait()` becomes **zombie**
 - the `ps` command prints the state of a zombie process as Z
- a process becomes **orphan** when its parent terminates while it is still running
 - hence the parent terminates without invoking the `wait()` system call
 - UNIX and Linux systems assign the `init` process (PID=1) as the new parent of the orphan process (aka **reparenting**)

Zombie and orphan processes

- a process that has terminated, but its parent has not (yet) called `wait()` becomes **zombie**
 - the `ps` command prints the state of a zombie process as `Z`
- a process becomes **orphan** when its parent terminates while it is still running
 - hence the parent terminates without invoking the `wait()` system call
 - UNIX and Linux systems assign the `init` process (PID=1) as the new parent of the orphan process (aka **reparenting**)
 - you can check this `if (getppid() == 1)`

Zombie and orphan processes

- a process that has terminated, but its parent has not (yet) called `wait()` becomes **zombie**
 - the `ps` command prints the state of a zombie process as `Z`
- a process becomes **orphan** when its parent terminates while it is still running
 - hence the parent terminates without invoking the `wait()` system call
 - UNIX and Linux systems assign the `init` process (PID=1) as the new parent of the orphan process (aka **reparenting**)
 - you can check this `if (getppid() == 1)`
 - the `init` process periodically calls `wait()` allowing the exit status of any orphaned process to be collected and their process table entries be deleted

Fork example

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>

int main() {
    ...
    pid_t ppid = getpid();           // store parent's pid
    pid_t pid = fork();              // create a child
    if(pid == 0){                    // child continues here
        printf("Child pid: [%d]\n", getpid());
        ...
    } else if (pid > 0) {             // parent continues here
        printf("Parent pid: [%d] Child pid: [%d]\n", ppid, pid);
        ...
    } else {
        perror("fork failed!");
        exit(1);
    }
    ...
}
```

Fork example

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>

int main() {
    ...
    pid_t ppid = getpid();           // store parent's pid
    pid_t pid = fork();              // create a child
    if(pid == 0){                    // child continues here
        printf("Child pid: [%d]\n", getpid());
        ...
    } else if (pid > 0) {             // parent continues here
        printf("Parent pid: [%d] Child pid: [%d]\n", ppid, pid);
        ...
    } else {
        perror("fork failed!");
        exit(1);
    }
    ...
}
```

Run the `ps` command to
check the processes' IDs

Combining fork and wait

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>

int main() {
    ...
    pid_t ppid = getpid();           // store parent's pid
    pid_t pid = fork();              // create a child
    if(pid == 0){                    // child continues here
        ...
    } else if (pid > 0) {            // parent continues here
        ...
        pid_t cpid = wait(&child_status);
    } else {
        perror("fork failed!");
    }
    ...
}
```


Combining fork and exec

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>

main() {
    ...
    pid_t ppid = getpid();           // store parent's pid
    pid_t pid = fork();              // create a child
    if(pid == 0) {                   // child continues here
        status = execl("/bin/ls", arg0, arg1, ...); // mark the end with a null pointer
        /* exec doesn't return when it works.
           so if we got here, it must have failed! */
        perror("exec failed!");
    } else if (pid > 0) {            // parent continues here
        ...
        cpid = wait(&status);        // pass NULL if not interested in exit status
        if (WIFEXITED(status))
            printf("child exit status was %d\n", WEXITSTATUS(status));
    } else {
        perror("fork failed!");
    }
    ...
}
```

Parent can kill its child!

```
#include <signal.h>
#include <unistd.h>
#include <stdio.h>

main() {
    ...
    int ppid = getpid();           // store parent's pid
    int pid = fork();              // create a child
    if(pid == 0){                  // child continues here
        sleep(10);                // child sleeps for 10 seconds
        ...
        exit(0);
    }
    else {                         // parent continues here
        ...
        printf( "Type any character to kill the child.\n" );
        char answer[10];
        gets(answer);
        if ( !kill(pid, SIGKILL) ) {
            printf("Killed the child.\n");
        }
    }
}
```

Other system calls for process control

- OS must include calls to enable special control of a process:

Other system calls for process control

- OS must include calls to enable special control of a process:
 - Priority manipulation:

Other system calls for process control

- OS must include calls to enable special control of a process:
 - Priority manipulation:
 - the `nice(incr)` system call adjusts the process priority by adding `incr` to its nice value

Other system calls for process control

- OS must include calls to enable special control of a process:
 - Priority manipulation:
 - the `nice(incr)` system call adjusts the process priority by adding `incr` to its nice value
 - lower nice values have higher scheduling priority

Other system calls for process control

- OS must include calls to enable special control of a process:
 - Priority manipulation:
 - the `nice(incr)` system call adjusts the process priority by adding `incr` to its nice value
 - lower nice values have higher scheduling priority
 - a process could be “nice” and reduce its share of the CPU by adjusting its nice value

Other system calls for process control

- OS must include calls to enable special control of a process:
 - Priority manipulation:
 - the `nice(incr)` system call adjusts the process priority by adding `incr` to its nice value
 - lower nice values have higher scheduling priority
 - a process could be “nice” and reduce its share of the CPU by adjusting its nice value
 - Debugging support:

Other system calls for process control

- OS must include calls to enable special control of a process:
 - Priority manipulation:
 - the `nice(incr)` system call adjusts the process priority by adding `incr` to its nice value
 - lower nice values have higher scheduling priority
 - a process could be “nice” and reduce its share of the CPU by adjusting its nice value
 - Debugging support:
 - the `ptrace()` system call allows a process to be put under control of another process

Other system calls for process control

- OS must include calls to enable special control of a process:
 - Priority manipulation:
 - the `nice(incr)` system call adjusts the process priority by adding `incr` to its nice value
 - lower nice values have higher scheduling priority
 - a process could be “nice” and reduce its share of the CPU by adjusting its nice value
 - Debugging support:
 - the `ptrace()` system call allows a process to be put under control of another process
 - the other process can check the arguments of the system call made by the process being traced, set breakpoints, examine registers, etc.

Other system calls for process control

- OS must include calls to enable special control of a process:
 - Priority manipulation:
 - the `nice(incr)` system call adjusts the process priority by adding `incr` to its nice value
 - lower nice values have higher scheduling priority
 - a process could be “nice” and reduce its share of the CPU by adjusting its nice value
 - Debugging support:
 - the `ptrace()` system call allows a process to be put under control of another process
 - the other process can check the arguments of the system call made by the process being traced, set breakpoints, examine registers, etc.
 - Alarms and time:

Other system calls for process control

- OS must include calls to enable special control of a process:
 - Priority manipulation:
 - the `nice(incr)` system call adjusts the process priority by adding `incr` to its nice value
 - lower nice values have higher scheduling priority
 - a process could be “nice” and reduce its share of the CPU by adjusting its nice value
 - Debugging support:
 - the `ptrace()` system call allows a process to be put under control of another process
 - the other process can check the arguments of the system call made by the process being traced, set breakpoints, examine registers, etc.
 - Alarms and time:
 - the `sleep()` system call puts a process on a timer queue waiting for some number of seconds, supporting alarm functionality

Process termination in UNIX systems

- the `kill` command sends an arbitrary signal to processes based on PID

Process termination in UNIX systems

- the `kill` command sends an arbitrary signal to processes based on PID
 - it sends a `SIGTERM` signal by default

Process termination in UNIX systems

- the `kill` command sends an arbitrary signal to processes based on PID
 - it sends a `SIGTERM` signal by default
- the `killall` command sends an arbitrary signal to processes based on process name

Process monitoring in UNIX systems

- `ps` displays information about a selection of the active processes

Process monitoring in UNIX systems

- `ps` displays information about a selection of the active processes
 - `ps -el` lists complete information about all processes that are currently active in the system

Process monitoring in UNIX systems

- `ps` displays information about a selection of the active processes
 - `ps -el` lists complete information about all processes that are currently active in the system
 - `ps -u [username]` lists all processes created by a specific user

Process monitoring in UNIX systems

- `ps` displays information about a selection of the active processes
 - `ps -el` lists complete information about all processes that are currently active in the system
 - `ps -u [username]` lists all processes created by a specific user
- `top` provides a dynamic real-time view of a running system (repetitive update on active processes)

Process monitoring in UNIX systems

- `ps` displays information about a selection of the active processes
 - `ps -el` lists complete information about all processes that are currently active in the system
 - `ps -u [username]` lists all processes created by a specific user
- `top` provides a dynamic real-time view of a running system (repetitive update on active processes)
- `pstree` displays a tree of processes

Shell

- act as a process control system
 - allows programmer to create and manage a set of processes to do some tasks
 - Windows, Linux, MacOS have shells

Shell

- act as a process control system
 - allows programmer to create and manage a set of processes to do some tasks
 - Windows, Linux, MacOS have shells
- when you log in to a machine running UNIX, you create a shell process

Shell

- act as a process control system
 - allows programmer to create and manage a set of processes to do some tasks
 - Windows, Linux, MacOS have shells
- when you log in to a machine running UNIX, you create a shell process
- every command launched in the shell is a child of the shell process and is an implicit `fork()` and `exec()` pair

Shell

- act as a process control system
 - allows programmer to create and manage a set of processes to do some tasks
 - Windows, Linux, MacOS have shells
- when you log in to a machine running UNIX, you create a shell process
- every command launched in the shell is a child of the shell process and is an implicit `fork()` and `exec()` pair
- the separation of `fork()` and `exec()` enables features like input/output redirection, pipes, etc.
 - the shell runs code after the call to `fork()` and before the call to `exec()`

Summary

- OS creates, deletes, suspends, and resumes processes
- OS allocates resources to active processes
 - memory, I/O devices, files
- OS schedules processes
 - context switches between them
- OS supports **interprocess communication** and provides **synchronization** mechanisms

Homework

- Compile and run the examples available on eClass