

# Operating System Concepts

## Lecture 8: Interprocess Communication

Omid Ardakanian  
[oardakan@ualberta.ca](mailto:oardakan@ualberta.ca)  
University of Alberta

MWF 12:00-12:50 VVC 2 215

# Today's class

---

- Interprocess communication
  - Shared memory
  - Message Passing

# Cooperating processes

---

- any two processes are either independent or cooperating
  - a process is cooperating if it can affect or be affected by other processes

# Cooperating processes

---

- any two processes are either independent or cooperating
  - a process is cooperating if it can affect or be affected by other processes
- cooperating processes work with each other to accomplish a single task
  - improve performance
    - parallel execution results in computation speedup on multicore systems
  - improve program structure/modularity
    - each cooperating process is smaller than a single monolithic program

# Cooperating processes

---

- any two processes are either independent or cooperating
  - a process is cooperating if it can affect or be affected by other processes
- cooperating processes work with each other to accomplish a single task
  - improve performance
    - parallel execution results in computation speedup on multicore systems
  - improve program structure/modularity
    - each cooperating process is smaller than a single monolithic program
- cooperating processes need to share information
  - OS makes this happen using interprocess communication mechanisms

# What are IPC mechanisms in UNIX?

---

IPC type	SUS	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10
half-duplex pipes FIFOs	• •	(full) •	• •	• •	(full) •
full-duplex pipes named full-duplex pipes	allowed obsolescent	•, UDS UDS	UDS UDS	UDS UDS	•, UDS •, UDS
XSI message queues XSI semaphores XSI shared memory	XSI XSI XSI	• • •	• • •	• • •	• • •
message queues (real-time) semaphores shared memory (real-time)	MSG option • SHM option	• • •	• • •	• • •	• • •
sockets STREAMS	• obsolescent	•	•	•	• •

# Two fundamental approaches

---

- Message passing
  - kernel intervention (through system calls) is required in every send/receive operation
  - extensible to communication in distributed systems
  - **positive aspects:** all sharing is explicit; less chance for error
  - **negative aspects:** high overhead. data copying, cross protection domains

# Two fundamental approaches


---

- Message passing
  - kernel intervention (through system calls) is required in every send/receive operation
  - extensible to communication in distributed systems
  - **positive aspects:** all sharing is explicit; less chance for error
  - **negative aspects:** high overhead. data copying, cross protection domains
- Shared memory: processes (or threads) can read and write a set of shared memory locations
  - system calls are only required to establish shared memory regions
  - processes are responsible for ensuring that they do not access a location concurrently (synchronization is explicit)
  - difficult to provide across machine boundaries
  - **positive aspects:** faster: set up shared memory once, then access w/o crossing protection domains
  - **negative aspects:** things might change behind your back; error prone

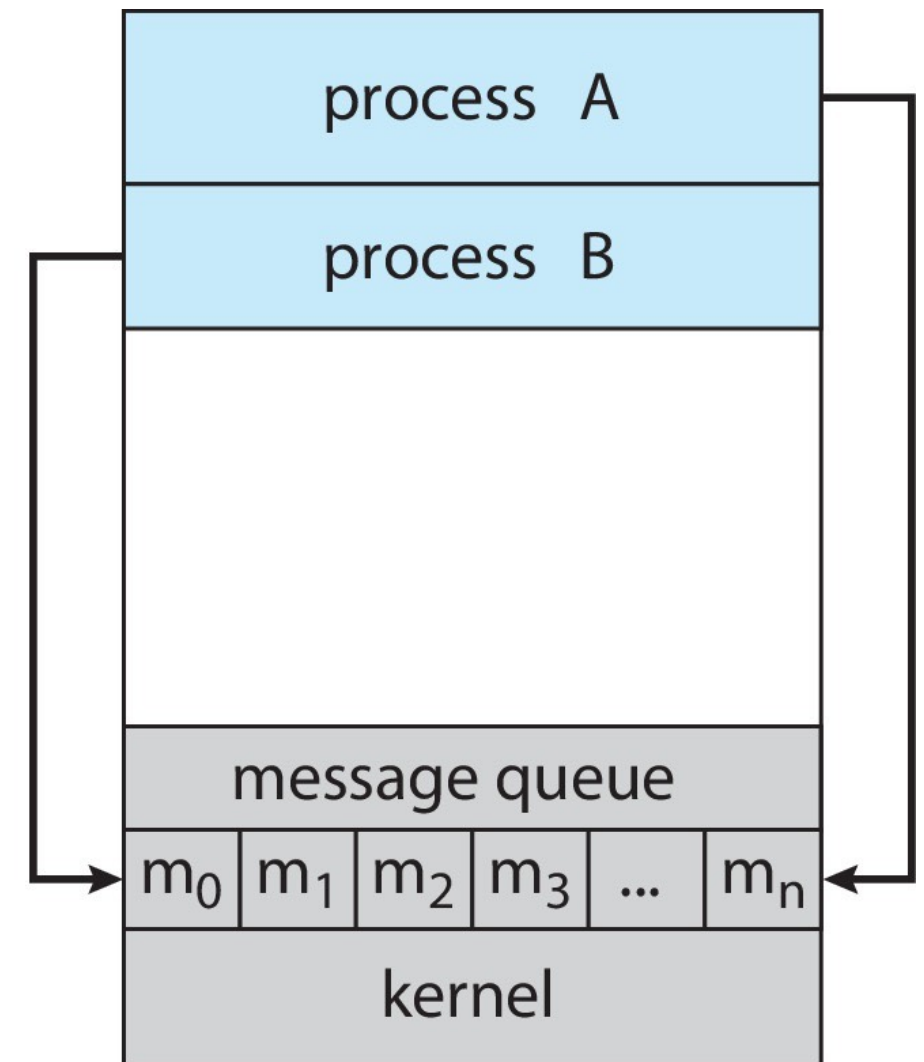


# Two fundamental approaches

processes read/write data  
from/to this shared  
memory region established  
by the mmap system call



(a)



(b)

# Communication using message passing

---

- distributed systems typically communicate using message passing
  - each process needs to be able to name the other process or name the mailbox/message queue (POSIX implementation)

# Communication using message passing

---

- distributed systems typically communicate using message passing
  - each process needs to be able to name the other process or name the mailbox/message queue (POSIX implementation)
- a common system message queue is
  - a linked list of messages stored within the kernel and identified by a message queue identifier
  - this identifier is shared between the processes to access to the queue (**HOW?**)

# Communication using message passing

---

- distributed systems typically communicate using message passing
  - each process needs to be able to name the other process or name the mailbox/message queue (POSIX implementation)
- a common system message queue is
  - a linked list of messages stored within the kernel and identified by a message queue identifier
  - this identifier is shared between the processes to access to the queue (**HOW?**)
- OS keeps track of messages
  - copies them, notifies receiving process, etc.

# Properties of message passing systems

---

- Direction
  - simplex (one-way)
  - half-duplex (two-way, but only one-way at a time)
  - full-duplex (two-way)

# Properties of message passing systems

---

- Direction
  - simplex (one-way)
  - half-duplex (two-way, but only one-way at a time)
  - full-duplex (two-way)
- Message boundaries
  - datagram model: message boundaries
  - byte stream model: no boundaries

# Properties of message passing systems

---

- Direction
  - simplex (one-way)
  - half-duplex (two-way, but only one-way at a time)
  - full-duplex (two-way)
- Message boundaries
  - datagram model: message boundaries
  - byte stream model: no boundaries
- Connection model
  - connection-oriented model: recipient is specified at connection time, hence it does not need be specified for individual send operations
  - connectionless models: recipient is specified as a parameter to each send operation

# Properties of message passing systems

---

- Direction
  - simplex (one-way)
  - half-duplex (two-way, but only one-way at a time)
  - full-duplex (two-way)
- Message boundaries
  - datagram model: message boundaries
  - byte stream model: no boundaries
- Connection model
  - connection-oriented model: recipient is specified at connection time, hence it does not need be specified for individual send operations
  - connectionless models: recipient is specified as a parameter to each send operation
- Reliability
  - messages can get lost, corrupted, or reordered



# Message passing issues

---


**Does a send/receive operation block?**

# Message passing issues

---

## Does a send/receive operation block?

- blocking operations




# Message passing issues

---

## Does a send/receive operation block?

- blocking operations
  - sender has to wait until its message is received
  - receiver has to wait if no message is available





# Message passing issues

---

## Does a send/receive operation block?

- blocking operations
  - sender has to wait until its message is received
  - receiver has to wait if no message is available
- non-blocking operations




# Message passing issues

---

## Does a send/receive operation block?


- blocking operations

- sender has to wait until its message is received
- receiver has to wait if no message is available



- non-blocking operations

- send operation returns immediately
- receive operation returns if no message is available




# Message passing issues

---

## Does a send/receive operation block?


- blocking operations

- sender has to wait until its message is received
- receiver has to wait if no message is available



- non-blocking operations

- send operation returns immediately
- receive operation returns if no message is available



- partially blocking/non-blocking


# Message passing issues

---

## Does a send/receive operation block?


- blocking operations

- sender has to wait until its message is received
- receiver has to wait if no message is available



- non-blocking operations

- send operation returns immediately
- receive operation returns if no message is available



- partially blocking/non-blocking

- send( )/receive( ) with timeout

# Communication using shared memory

---

- need to establish a mapping between the process's address space to a named memory object that may be shared across processes



# Communication using shared memory

---

- need to establish a mapping between the process's address space to a named memory object that may be shared across processes
- the `mmap ( )` systems call performs this function

# Communication using shared memory


---

- need to establish a mapping between the process's address space to a named memory object that may be shared across processes
- the `mmap ( )` systems call performs this function
- can fork processes that need to share the data structure so that they know the name of the shared object
  - any other solution?

# Producer-Consumer problem

---


- Definition: producers puts information into a shared buffer; consumers takes it out



# Producer-Consumer problem

---


- Definition: producers puts information into a shared buffer; consumers takes it out
- it is also known as the bounded-buffer problem
  - examples:
    - web servers (server:producer & client:consumer)
    - compiling your code w/ gcc: `cpp | cc1 | as | ld`



# Producer-Consumer problem

---


- Definition: producers puts information into a shared buffer; consumers takes it out
- it is also known as the bounded-buffer problem
  - examples:
    - web servers (server:producer & client:consumer)
    - compiling your code w/ gcc: `cpp | cc1 | as | ld`
- For implementation, **interprocess communication** is necessary



# Producer-Consumer problem

---


- Definition: producers puts information into a shared buffer; consumers takes it out
- it is also known as the bounded-buffer problem
  - examples:
    - web servers (server:producer & client:consumer)
    - compiling your code w/ gcc: `cpp | cc1 | as | ld`
- For implementation, **interprocess communication** is necessary
- What should the producer do when the buffer is full?



# Producer-Consumer problem

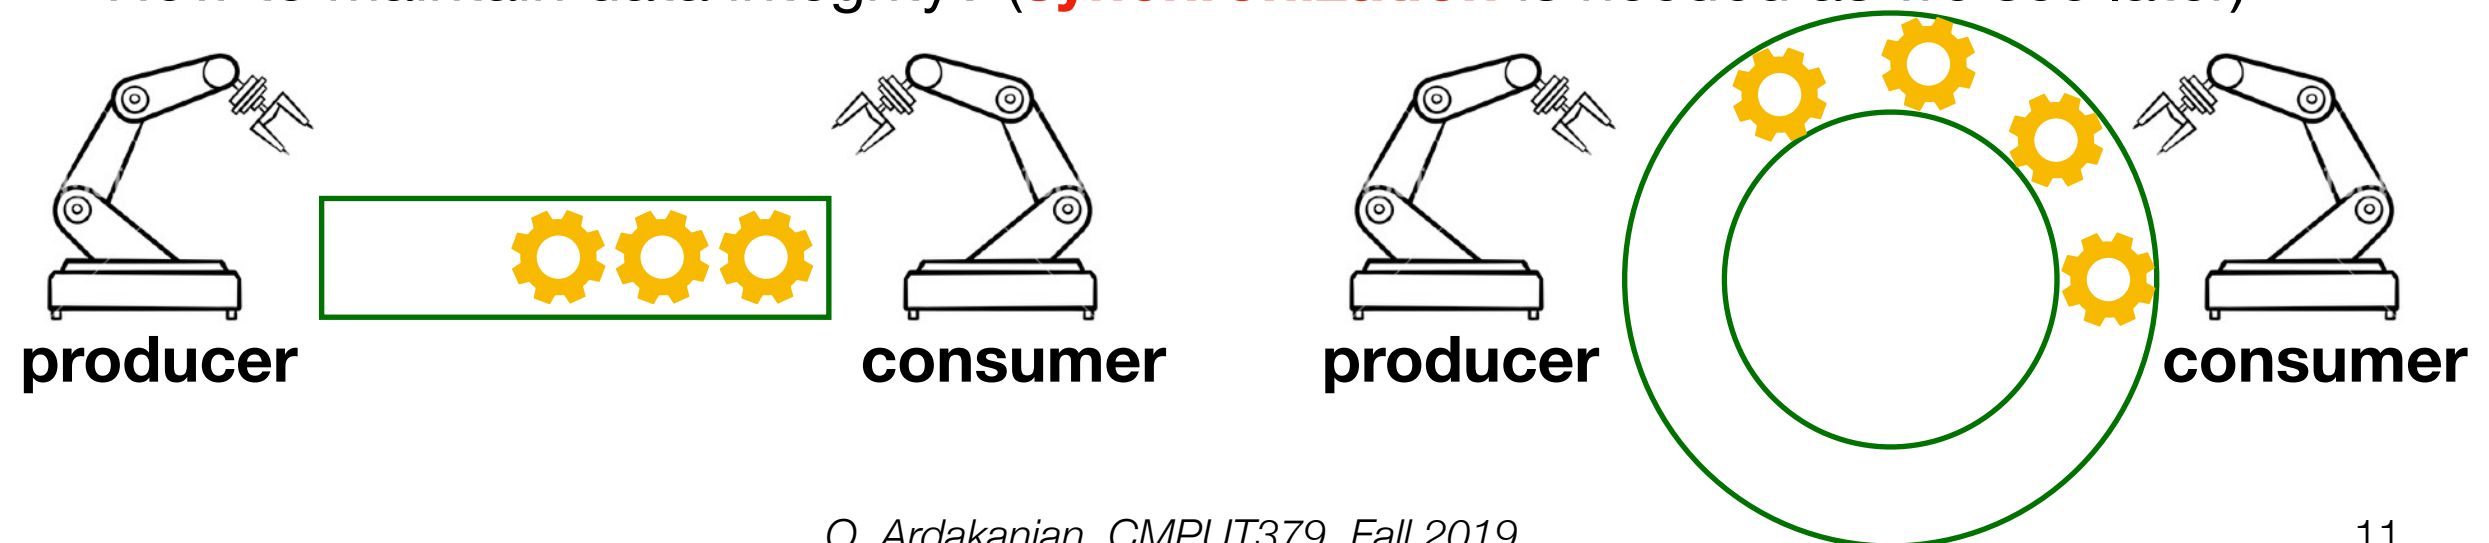
---

- Definition: producers puts information into a shared buffer; consumers takes it out
- it is also known as the bounded-buffer problem
  - examples:
    - web servers (server:producer & client:consumer)
    - compiling your code w/ gcc: `cpp | cc1 | as | ld`
- For implementation, **interprocess communication** is necessary
- What should the producer do when the buffer is full?
- What should the consumer do when the buffer is empty?



# Producer-Consumer problem

- Definition: producers puts information into a shared buffer; consumers takes it out
- it is also known as the bounded-buffer problem
  - examples:
    - web servers (server:producer & client:consumer)
    - compiling your code w/ gcc: `cpp | cc1 | as | ld`
- For implementation, **interprocess communication** is necessary
- What should the producer do when the buffer is full?
- What should the consumer do when the buffer is empty?
- How to maintain data integrity? (**synchronization** is needed as we see later)





# Producer-Consumer problem

---

- implementation based on message passing
  - using `send(c_pid, nextp)` and `receive(p_pid, nextc)`
  - each process needs to be able to name the other process
  - the kernel manages the buffer (message queue)

# Producer-Consumer problem

---

- implementation based on message passing
  - using `send(c_pid, nextp)` and `receive(p_pid, nextc)`
  - each process needs to be able to name the other process
  - the kernel manages the buffer (message queue)

## Main

```
int main() {  
    ...  
    if (fork() != 0) producer();  
    else consumer();  
    ...  
}
```

# Producer-Consumer problem

---

- implementation based on message passing
  - using send(c\_pid, nextp) and receive(p\_pid, nextc)
  - each process needs to be able to name the other process
  - the kernel manages the buffer (message queue)

## Main

```
int main( ) {  
    ...  
    if (fork( ) != 0) producer( );  
    else consumer( );  
    ...  
}
```

## Producer

```
int producer( ) {  
    ...  
    while(true) {  
        ...  
        nextp = produced item  
        send(C_pid, nextp)  
        ...  
    }  
}
```

# Producer-Consumer problem

---

- implementation based on message passing
  - using send(c\_pid, nextp) and receive(p\_pid, nextc)
  - each process needs to be able to name the other process
  - the kernel manages the buffer (message queue)

## Main

```
int main() {  
    ...  
    if (fork() != 0) producer();  
    else consumer();  
    ...  
}
```

## Producer

```
int producer() {  
    ...  
    while(true) {  
        ...  
        nextp = produced item  
        send(C_pid, nextp)  
        ...  
    }  
}
```

## Consumer

```
int consumer() {  
    ...  
    while(true) {  
        ...  
        receive(P_pid, &nextc)  
        consume nextc  
        ...  
    }  
}
```

# Producer-Consumer problem

---

- implementation based on shared memory
  - `n` is the size of the buffer
  - `in` points to the next free location
  - `out` points to the first full location
  - `in` and `out` are shared between producer and consumer
  - this way we can have at most  $n - 1$  items in the buffer (**WHY?**)

# Producer-Consumer problem

---

- implementation based on shared memory
  - `n` is the size of the buffer
  - `in` points to the next free location
  - `out` points to the first full location
  - `in` and `out` are shared between producer and consumer
  - this way we can have at most  $n - 1$  items in the buffer (**WHY?**)

## Main

```
int main() {  
    ...  
    mmap(..., in, out, PROT_WRITE, PROT_SHARED,  
    ...);  
    in = 0  
    out = 0  
    if (fork() != 0) producer();  
    else consumer();  
    ...  
}
```

# Producer-Consumer problem

---

- implementation based on shared memory
  - `n` is the size of the buffer
  - `in` points to the next free location
  - `out` points to the first full location
  - `in` and `out` are shared between producer and consumer
  - this way we can have at most  $n - 1$  items in the buffer (**WHY?**)

## Main

```
int main() {  
    ...  
    mmap(..., in, out, PROT_WRITE, PROT_SHARED,  
    ...);  
    in = 0  
    out = 0  
    if (fork() != 0) producer();  
    else consumer();  
    ...  
}
```

## Producer

```
int producer() {  
    ...  
    while(true) {  
        ...  
        nextp = produce item  
        while (in+1 mod n == out) {}  
        buffer[in] = nextp  
        in = in+1 mod n  
    }  
}
```

# Producer-Consumer problem

---

- implementation based on shared memory
  - `n` is the size of the buffer
  - `in` points to the next free location
  - `out` points to the first full location
  - `in` and `out` are shared between producer and consumer
  - this way we can have at most  $n - 1$  items in the buffer (**WHY?**)

## Main

```
int main() {  
    ...  
    mmap(..., in, out, PROT_WRITE, PROT_SHARED,  
    ...);  
    in = 0  
    out = 0  
    if (fork() != 0) producer();  
    else consumer();  
    ...  
}
```

## Producer

```
int producer() {  
    ...  
    while(true) {  
        ...  
        nextp = produce item  
        while (in+1 mod n == out) {}  
        buffer[in] = nextp  
        in = in+1 mod n  
    }  
}
```

## Consumer

```
int consumer() {  
    ...  
    while(true) {  
        ...  
        while (in == out) {}  
        nextc = buffer[out]  
        out = out+1 mod n  
        consume nextc  
    }  
}
```



# POSIX shared memory support

---

# POSIX shared memory support

---

POSIX shared memory is organized using memory-mapped files, i.e., associating the region of shared memory with a file

# POSIX shared memory support

---

POSIX shared memory is organized using memory-mapped files, i.e., associating the region of shared memory with a file

- the `shm_open( )` and `ftruncate( )` system calls are used to create a shared memory object with a specified name and to set the size of this object, respectively

# POSIX shared memory support

---

POSIX shared memory is organized using memory-mapped files, i.e., associating the region of shared memory with a file

- the `shm_open( )` and `ftruncate( )` system calls are used to create a shared memory object with a specified name and to set the size of this object, respectively
  - **shared memory object**: a handle which can be used by unrelated processes to memory map the same region of shared memory

# POSIX shared memory support

---

POSIX shared memory is organized using memory-mapped files, i.e., associating the region of shared memory with a file

- the `shm_open( )` and `ftruncate( )` system calls are used to create a shared memory object with a specified name and to set the size of this object, respectively
  - **shared memory object**: a handle which can be used by unrelated processes to memory map the same region of shared memory
  - `shm_open( )` returns a new file descriptor referring to the newly created shared memory object

# POSIX shared memory support

---

POSIX shared memory is organized using memory-mapped files, i.e., associating the region of shared memory with a file

- the `shm_open( )` and `ftruncate( )` system calls are used to create a shared memory object with a specified name and to set the size of this object, respectively
  - **shared memory object**: a handle which can be used by unrelated processes to memory map the same region of shared memory
  - `shm_open( )` returns a new file descriptor referring to the newly created shared memory object
- the `mmap( )` system call establishes a memory-mapped file containing this object and returns a pointer to this file;

# POSIX shared memory support

---

POSIX shared memory is organized using memory-mapped files, i.e., associating the region of shared memory with a file

- the `shm_open( )` and `ftruncate( )` system calls are used to create a shared memory object with a specified name and to set the size of this object, respectively
  - **shared memory object**: a handle which can be used by unrelated processes to memory map the same region of shared memory
  - `shm_open( )` returns a new file descriptor referring to the newly created shared memory object
- the `mmap( )` system call establishes a memory-mapped file containing this object and returns a pointer to this file;
  - the file pointer is used to write to or read from this shared memory object

# POSIX shared memory support

---

POSIX shared memory is organized using memory-mapped files, i.e., associating the region of shared memory with a file

- the `shm_open( )` and `ftruncate( )` system calls are used to create a shared memory object with a specified name and to set the size of this object, respectively
  - **shared memory object**: a handle which can be used by unrelated processes to memory map the same region of shared memory
  - `shm_open( )` returns a new file descriptor referring to the newly created shared memory object
- the `mmap( )` system call establishes a memory-mapped file containing this object and returns a pointer to this file;
  - the file pointer is used to write to or read from this shared memory object
- the `shm_unlink( )` system call removes the shared memory object



# POSIX shared memory support

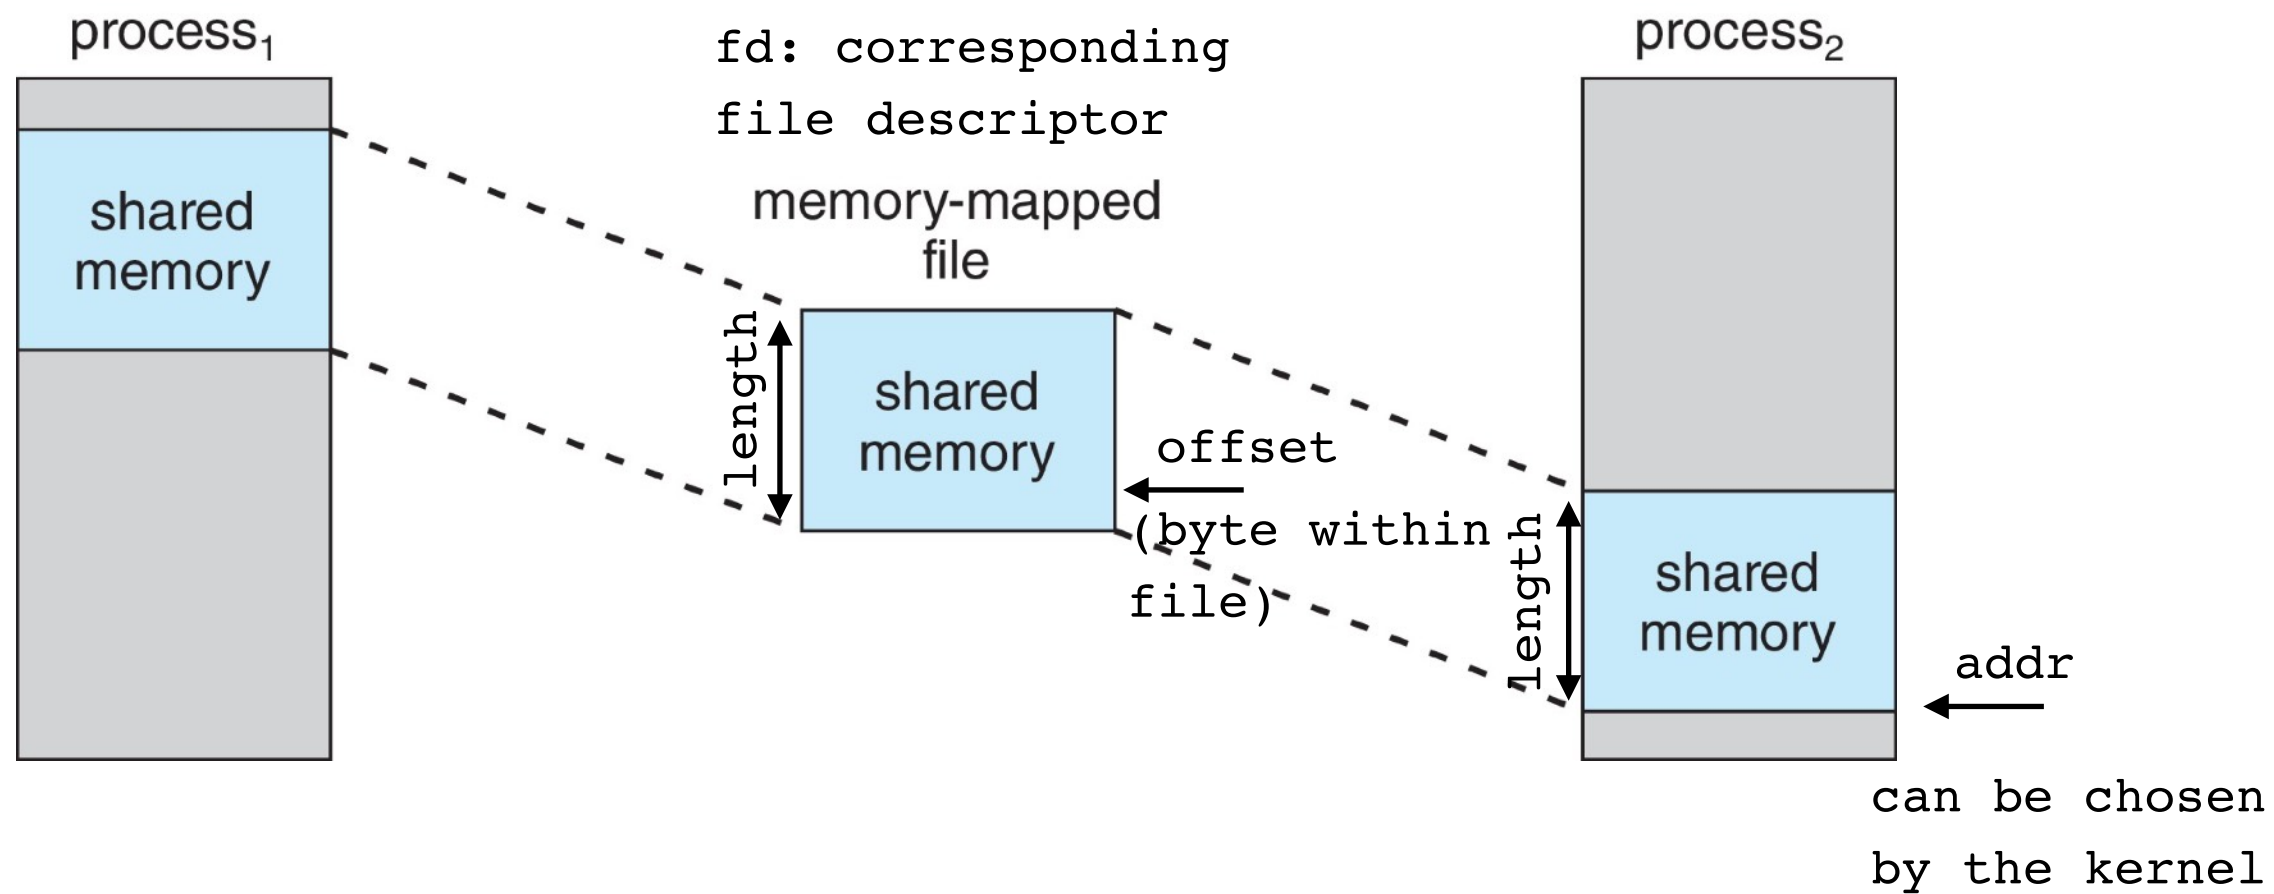
---

POSIX shared memory is organized using memory-mapped files, i.e., associating the region of shared memory with a file

- the `shm_open( )` and `ftruncate( )` system calls are used to create a shared memory object with a specified name and to set the size of this object, respectively
  - **shared memory object**: a handle which can be used by unrelated processes to memory map the same region of shared memory
  - `shm_open( )` returns a new file descriptor referring to the newly created shared memory object
- the `mmap( )` system call establishes a memory-mapped file containing this object and returns a pointer to this file;
  - the file pointer is used to write to or read from this shared memory object
- the `shm_unlink( )` system call removes the shared memory object
  - once all processes have unmapped the object, it de-allocates and destroys the contents of the associated memory region

# the mmap system call

```
void *mmap(void *addr, size_t length, int prot, int flags, int fd, off_t offset)
```



# POSIX producer

---

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
#include <sys/mman.h>

int main() {
    const int SIZE = 4096;
    const char* name = "/prog-shm";

    const char* message_0 = "Hello";
    const char* message_1 = "World!";

    int shm_fd;
    void* ptr;

    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);
    ftruncate(shm_fd, SIZE);
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    sprintf(ptr, "%s", message_0);
    ptr += strlen(message_0);
    sprintf(ptr, "%s", message_1);
    ptr += strlen(message_1);

    return 0;
}
```

/\* size (B) of shared memory object \*/  
/\* name of the shared memory object \*/

/\* shared memory file descriptor \*/  
/\* pointer to shared memory object \*/

/\* create the shared memory object \*/  
/\* configure size of the shared memory object \*/  
/\* memory map the shared memory object \*/

/\* write to the shared memory object \*/

# POSIX consumer

---

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
#include <sys/mman.h>

int main() {
    const int SIZE = 4096;                /* size (B) of shared memory object */
    const char* name = "/prog-shm";      /* name of the shared memory object */

    int shm_fd;                          /* shared memory file descriptor */
    void* ptr;                           /* pointer to shared memory object */

    shm_fd = shm_open(name, O_RDONLY, 0666); /* open the shared memory object */

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    printf("%s", (char*)ptr);            /* read from the shared memory object */
    shm_unlink(name);                    /* remove the shared memory object */

    return 0;
}
```

# System V message passing support

---

- `ftok( )` generates a unique key

# System V message passing support

---

- `ftok ( )` generates a unique key
- `msgget ( )` either returns the message queue identifier for a newly created message queue or returns the identifiers for a queue which exists with the same key value

# System V message passing support

---

- `ftok ( )` generates a unique key
- `msgget ( )` either returns the message queue identifier for a newly created message queue or returns the identifiers for a queue which exists with the same key value
- `msgsnd ( )` places data onto a message queue

# System V message passing support

---

- `ftok ( )` generates a unique key
- `msgget ( )` either returns the message queue identifier for a newly created message queue or returns the identifiers for a queue which exists with the same key value
- `msgsnd ( )` places data onto a message queue
- `msgrcv ( )` retrieves messages from a queue



# System V message passing support

---

- `ftok ( )` generates a unique key
- `msgget ( )` either returns the message queue identifier for a newly created message queue or returns the identifiers for a queue which exists with the same key value
- `msgsnd ( )` places data onto a message queue
- `msgrcv ( )` retrieves messages from a queue
- `msgctl ( )` performs various operations on a queue.
  - is generally used to destroy message queue

# Example — writer.c

---

```
#include <stdio.h>
#include <sys/ipc.h>
#include <sys/msg.h>

// structure for message queue
struct mesg_buffer {
    long mesg_type;
    char mesg_text[100];
} message;

int main()
{
    key_t key;
    int msgid;

    key = ftok("progfile", 65); // generates a unique key
    msgid = msgget(key, 0666 | IPC_CREAT); // creates a message queue
    message.mesg_type = 1;

    printf("Write Data: ");
    scanf("%s", message.mesg_text);

    msgsnd(msgid, &message, sizeof(message), 0); // sends the message

    printf("Data sent is: %s \n", message.mesg_text);

    return 0;
}
```

# Example — reader.c

---

```
#include <stdio.h>
#include <sys/ipc.h>
#include <sys/msg.h>

// structure for message queue
struct mesg_buffer {
    long mesg_type;
    char mesg_text[100];
} message;

int main()
{
    key_t key;
    int msgid;

    key = ftok("progfile", 65);                // generates a unique key
    msgid = msgget(key, 0666 | IPC_CREAT);       // creates a message queue
    msgrcv(msgid, &message, sizeof(message), 1, 0); // receives the message

    printf("Data Received is : %s \n", message.mesg_text);

    msgctl(msgid, IPC_RMID, NULL);              // destroys the message queue

    return 0;
}
```