

# Operating System Concepts

## Lecture 10: Distributed Systems

Omid Ardakanian  
[oardakan@ualberta.ca](mailto:oardakan@ualberta.ca)  
University of Alberta

MWF 12:00-12:50 VVC 2 215

# Today's class

---

- Distributed systems
  - Motivation
  - Design issues
- Communication basics
  - Network structure
  - Network protocol

# Distributed systems


---

- Definition: a set of physically separate, **loosely coupled** nodes connected by a communication network (high-speed buses or the Internet)
  - each node has an independent OS along with its own memory and other resources
  - nodes are variously called processors, machines, computers, and hosts

# Distributed systems

---

- Definition: a set of physically separate, **loosely coupled** nodes connected by a communication network (high-speed buses or the Internet)
  - each node has an independent OS along with its own memory and other resources
  - nodes are variously called processors, machines, computers, and hosts
- communication over a network occurs through **message passing**
  - inherently unreliable: messages sometimes do not reach their destination



# Distributed systems

---

- Definition: a set of physically separate, **loosely coupled** nodes connected by a communication network (high-speed buses or the Internet)
  - each node has an independent OS along with its own memory and other resources
  - nodes are variously called processors, machines, computers, and hosts
- communication over a network occurs through **message passing**
  - inherently unreliable: messages sometimes do not reach their destination
- nodes may exist in a client-server, peer-to-peer, or hybrid configuration
  - in client server model the server has a resource that the client (at a different site) wants to use
  - in peer-to-peer model each node shares equal responsibilities and can act as both clients and servers

# Distributed systems

---

- Definition: a set of physically separate, **loosely coupled** nodes connected by a communication network (high-speed buses or the Internet)
  - each node has an independent OS along with its own memory and other resources
  - nodes are variously called processors, machines, computers, and hosts
- communication over a network occurs through **message passing**
  - inherently unreliable: messages sometimes do not reach their destination
- nodes may exist in a client-server, peer-to-peer, or hybrid configuration
  - in client server model the server has a resource that the client (at a different site) wants to use
  - in peer-to-peer model each node shares equal responsibilities and can act as both clients and servers
- nearly all systems today are distributed in some way
  - complex services are built from a large collection of machines cooperating to provide the particular service of the site
  - networks hook them together

# Why distributed systems are popular?

---

- resource sharing
  - resources need not be replicated at each processor, e.g., shared files
  - expensive and/or scarce resources can be shared, e.g., printers, GPUs
  - user can use specialized and licensed software on another machine

# Why distributed systems are popular?

---

- resource sharing
  - resources need not be replicated at each processor, e.g., shared files
  - expensive and/or scarce resources can be shared, e.g., printers, GPUs
  - user can use specialized and licensed software on another machine
- computational speedup
  - distribute tasks among various sites to run concurrently
    - they do not compete with each other for a single CPU core
  - load balancing (moving jobs to more lightly-loaded sites) would help
  - coordination and communication are necessary yet quite challenging



# Why distributed systems are popular?

---

- resource sharing
  - resources need not be replicated at each processor, e.g., shared files
  - expensive and/or scarce resources can be shared, e.g., printers, GPUs
  - user can use specialized and licensed software on another machine
- computational speedup
  - distribute tasks among various sites to run concurrently
    - they do not compete with each other for a single CPU core
  - load balancing (moving jobs to more lightly-loaded sites) would help
  - coordination and communication are necessary yet quite challenging
- reliability and availability
  - resource replication yields fault tolerance (no single point of failure)
    - performance will degrade but system remains operational
  - must be able to detect and recover from site failure (function transfer, reintegrate failed sites, etc.)

# Design issues of distributed systems

---

- robustness — the system should withstand failures
  - including failure of a link, failure of a site, and loss of messages
  - a fault-tolerant system can tolerate a certain level of failure
  - the degree of fault tolerance depends on the system design and the specific fault
  - detecting hardware failure is difficult (can use a heartbeat protocol)

# Design issues of distributed systems

---

- robustness — the system should withstand failures
  - including failure of a link, failure of a site, and loss of messages
  - a fault-tolerant system can tolerate a certain level of failure
  - the degree of fault tolerance depends on the system design and the specific fault
  - detecting hardware failure is difficult (can use a heartbeat protocol)
- transparency to user — in terms of where files are stored and user mobility
  - the distributed system should appear as a conventional, centralized system to the user
    - user interface should not distinguish between local and remote resources
    - user mobility allows users to log into any machine in the environment and see his/her environment

# Design issues of distributed systems

---

- scalability — the system should react gracefully to increased load (by accepting new resources)
  - data **compression** and **deduplication** can cut down on storage and network resources used

# Design issues of distributed systems

---

- scalability — the system should react gracefully to increased load (by accepting new resources)
  - data **compression** and **deduplication** can cut down on storage and network resources used
- consistency — the cached copy of the data must be consistent with the master copy
  - consistency checks must be performed periodically
  - nodes must keep track of the cached data to detect potential inconsistencies

# Network OS

---

- all general-purpose operating systems today are network operating systems
  - allow the user to login remotely and transfer **files**
  - users are aware of multiplicity of machines

# Network OS

---

- all general-purpose operating systems today are network operating systems
  - allow the user to login remotely and transfer **files**
  - users are aware of multiplicity of machines
- access to resources of various machines is done explicitly by
  - remote login via Secure Shell (SSH) protocol  
`ssh ohaton.cs.ualberta.edu`
  - remote file transfer via File Transfer Protocol (FTP) and Secure File Transfer Protocol (SFTP)  
`get, put, ls, cd` commands are used for locating and transferring files

# Network OS

---

- all general-purpose operating systems today are network operating systems
  - allow the user to login remotely and transfer **files**
  - users are aware of multiplicity of machines
- access to resources of various machines is done explicitly by
  - remote login via Secure Shell (SSH) protocol  
`ssh ohaton.cs.ualberta.edu`
  - remote file transfer via File Transfer Protocol (FTP) and Secure File Transfer Protocol (SFTP)  
`get, put, ls, cd` commands are used for locating and transferring files
- users must change paradigms – establish a **session**, give network-based commands (might be different from the normal OS commands)
  - session is a complete round of communication



# Distributed OS

---

- users are not aware of multiplicity of machines
  - access to remote resources similar to access to local resources
  - data migration and process migration are handled **seamlessly** by the distributed operating system
    - data translation may be required (when they have different character-code representations, different order of bits)

# Distributed OS

---

- users are not aware of multiplicity of machines
  - access to remote resources similar to access to local resources
  - data migration and process migration are handled **seamlessly** by the distributed operating system
    - data translation may be required (when they have different character-code representations, different order of bits)
- data migration
  - transfer data by transferring the entire file, or transferring only those portions of the file necessary for the immediate task

# Distributed OS

---

- users are not aware of multiplicity of machines
  - access to remote resources similar to access to local resources
  - data migration and process migration are handled **seamlessly** by the distributed operating system
    - data translation may be required (when they have different character-code representations, different order of bits)
- data migration
  - transfer data by transferring the entire file, or transferring only those portions of the file necessary for the immediate task
- computation migration
  - transfer the computation, rather than the data, across the system via remote procedure calls (RPCs) for example

# Distributed OS

---


- users are not aware of multiplicity of machines
  - access to remote resources similar to access to local resources
  - data migration and process migration are handled **seamlessly** by the distributed operating system
    - data translation may be required (when they have different character-code representations, different order of bits)
- data migration
  - transfer data by transferring the entire file, or transferring only those portions of the file necessary for the immediate task
- computation migration
  - transfer the computation, rather than the data, across the system via remote procedure calls (RPCs) for example
- process migration
  - execute the entire process (computation and data), or parts of it, at different sites

# Communication Basics

# Definitions

---


- network: one or more communication links allowing two computers to communicate
  - each computer has a network address



# Definitions

---


- network: one or more communication links allowing two computers to communicate
  - each computer has a network address
- network interface: computer's interface to the network
  - each network interface card (NIC) has a unique hardware address



# Definitions

---

- network: one or more communication links allowing two computers to communicate
  - each computer has a network address
- network interface: computer's interface to the network
  - each network interface card (NIC) has a unique hardware address
- packet: network's basic transmission unit; a sequence of bits






# Definitions

---


- network: one or more communication links allowing two computers to communicate
  - each computer has a network address
- network interface: computer's interface to the network
  - each network interface card (NIC) has a unique hardware address
- packet: network's basic transmission unit; a sequence of bits



# Definitions

---

- network: one or more communication links allowing two computers to communicate
  - each computer has a network address
- network interface: computer's interface to the network
  - each network interface card (NIC) has a unique hardware address
- packet: network's basic transmission unit; a sequence of bits
- protocol: a set of rules for communication that are agreed to by all parties



# Network structure

---

- Local Area Network (LAN) covers a small geographical area (e.g., a building)
  - consists of multiple computers, peripherals (printers, storage arrays), and routers providing access to other networks
  - needs to be fast and reliable

# Network structure

---

- Local Area Network (LAN) covers a small geographical area (e.g., a building)
  - consists of multiple computers, peripherals (printers, storage arrays), and routers providing access to other networks
  - needs to be fast and reliable
- Ethernet and/or Wireless (WiFi) most common way to construct LANs
  - Ethernet defined by IEEE 802.3 standard with speeds typically varying from 10Mbps to over 10Gbps
    - everyone taps into a single wire
    - everyone gets packets and discards them if it is not the target
  - WiFi defined by IEEE 802.11 standard with speeds typically varying from 11Mbps to over 400Mbps

# Network structure

---

- Wide Area Network links geographically separated sites (across the country or planet)
  - WAN is slower and less reliable than LAN
  - Internet WAN enables hosts world wide to communicate

# Network structure

---

- Wide Area Network links geographically separated sites (across the country or planet)
  - WAN is slower and less reliable than LAN
  - Internet WAN enables hosts world wide to communicate
- point-to-point connections via links
  - telephone lines, leased (dedicated data) lines, optical cable, microwave links, radio waves, and satellite channels
  - speeds vary
    - many backbone providers have speeds at 40-100Gbps
    - local Internet Service Providers (ISPs) may be slower

# Network structure

---

- Wide Area Network links geographically separated sites (across the country or planet)
  - WAN is slower and less reliable than LAN
  - Internet WAN enables hosts world wide to communicate
- point-to-point connections via links
  - telephone lines, leased (dedicated data) lines, optical cable, microwave links, radio waves, and satellite channels
  - speeds vary
    - many backbone providers have speeds at 40-100Gbps
    - local Internet Service Providers (ISPs) may be slower
- WANs and LANs interconnect

# Principles of network communication

---

- data sent through the network is chopped into **packets**



# Principles of network communication

---

- data sent through the network is chopped into **packets**
- computers at the switching points control the packet flow
  - analogy: cars/road/police - packets/network link/computer

# Principles of network communication

---

- data sent through the network is chopped into **packets**
- computers at the switching points control the packet flow
  - analogy: cars/road/police - packets/network link/computer
- shared resources can lead to contention (just like traffic jams)


# Principles of network communication

---

- data sent through the network is chopped into **packets**
- computers at the switching points control the packet flow
  - analogy: cars/road/police - packets/network link/computer
- shared resources can lead to contention (just like traffic jams)
- the destination computer is interrupted when a packet arrives

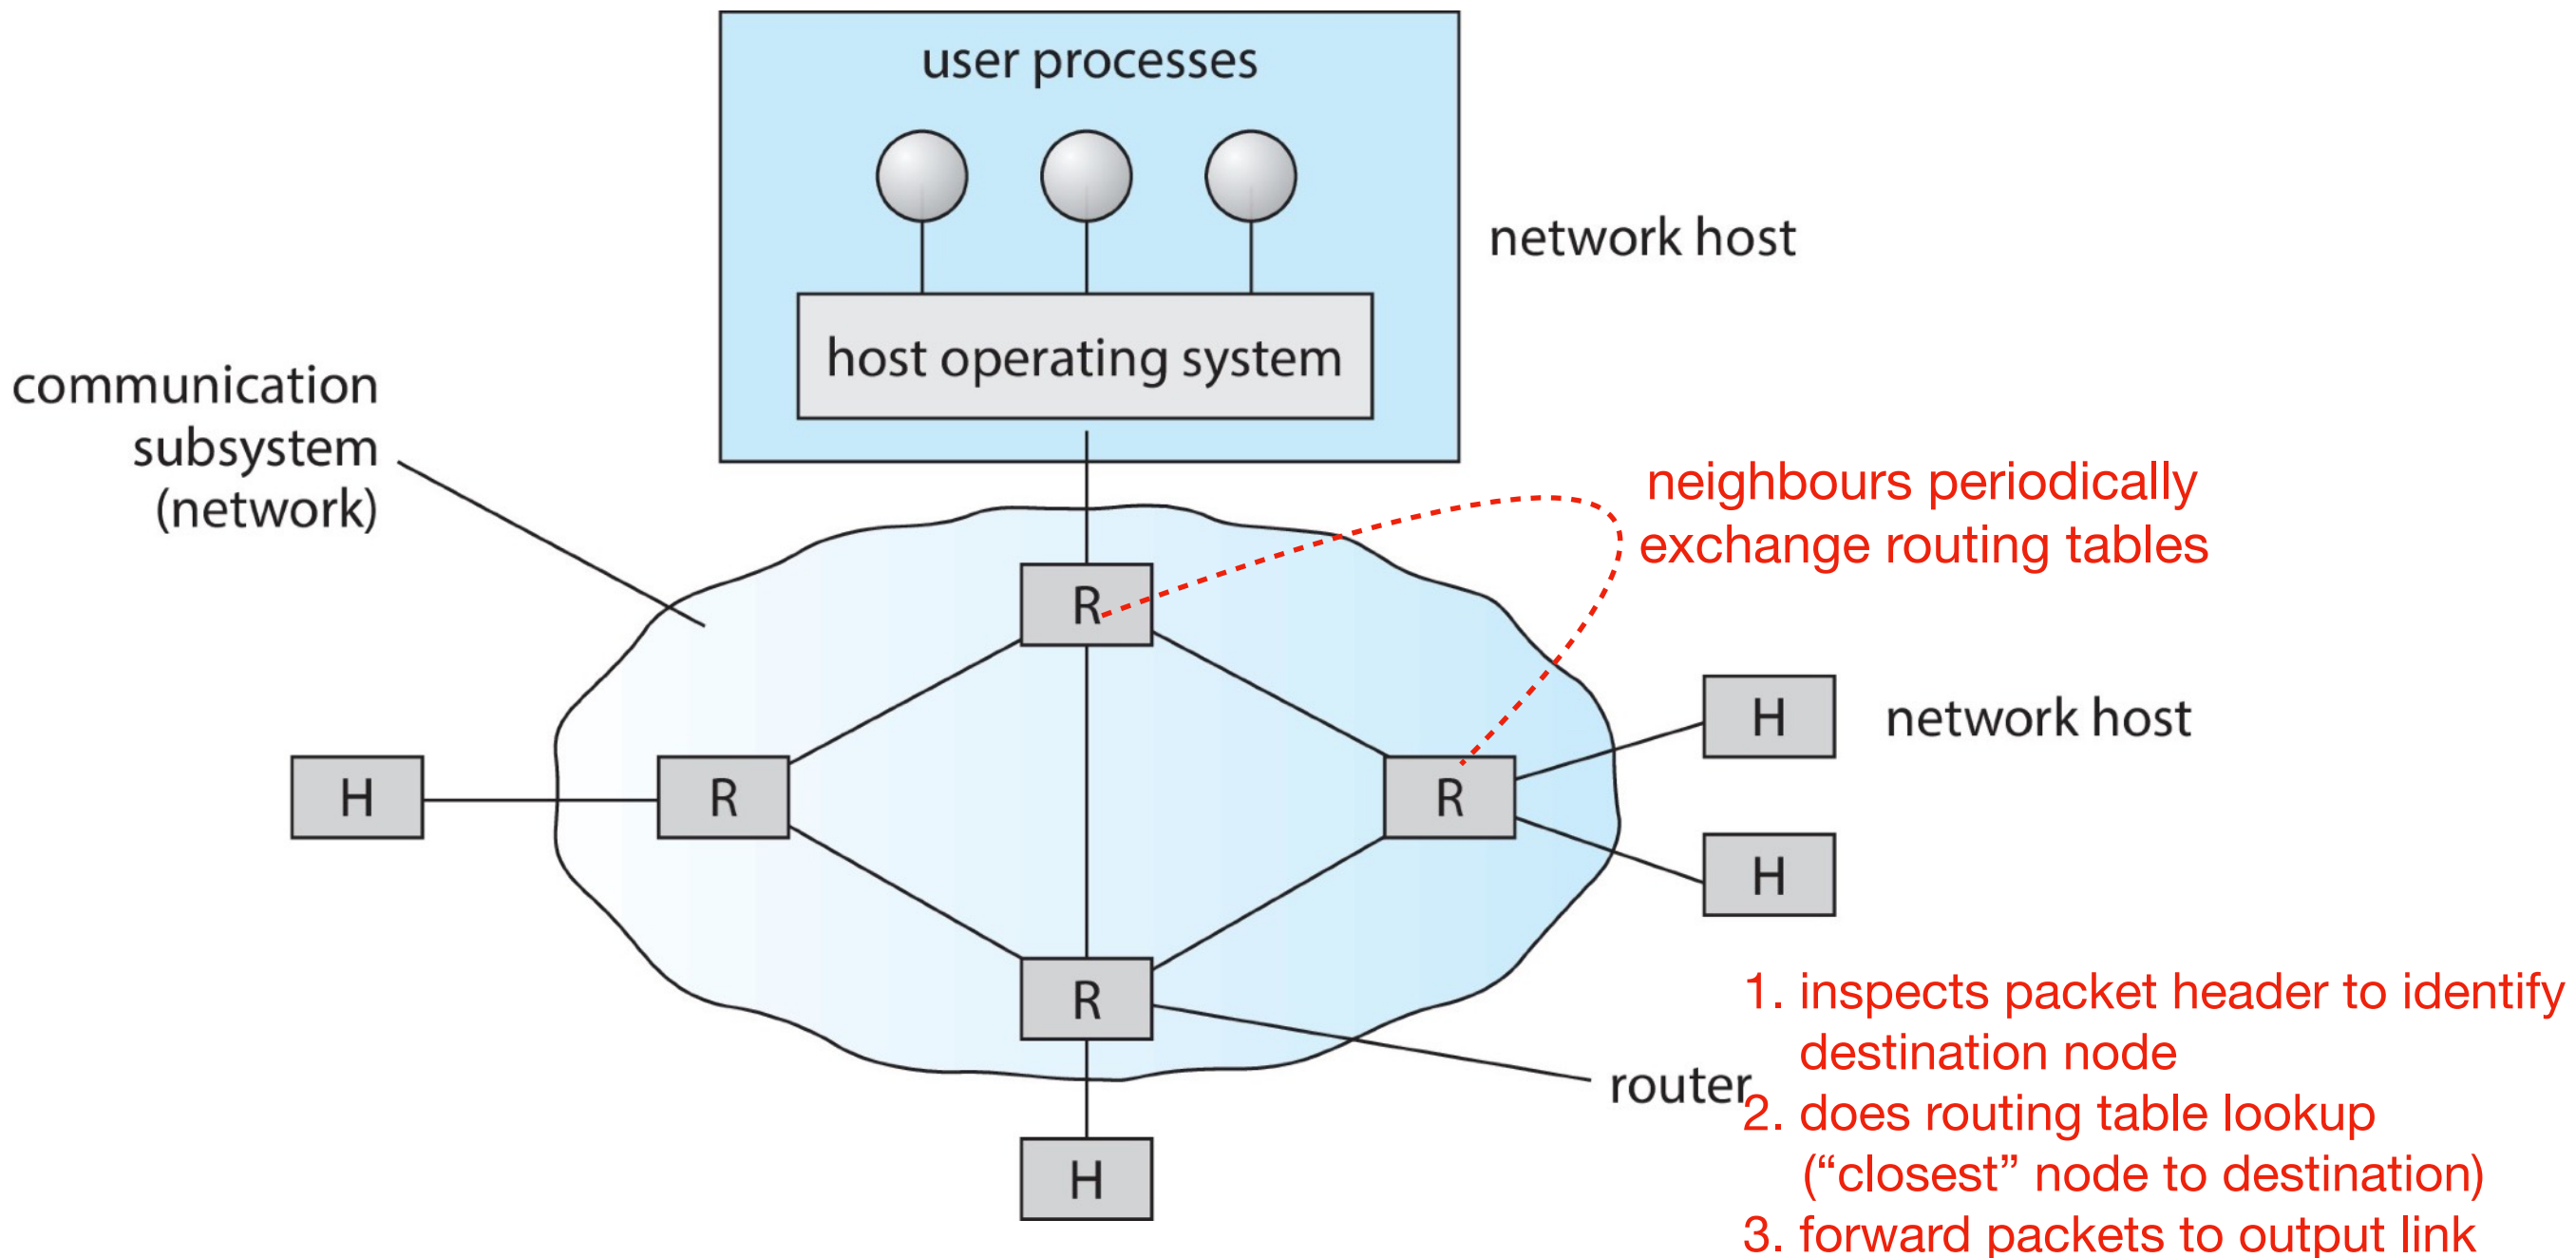
# How a packet is delivered to destination if the hosts are located on separate networks?

---



WAN is implemented via routers to direct traffic from one network to another

# How a packet is delivered to destination if the hosts are located on separate networks?



WAN is implemented via routers to direct traffic from one network to another

# How to identify the destination system/process?

---

- a process on a remote system is identified by  
`<host-name, identifier>` pair

# How to identify the destination system/process?

---

- a process on a remote system is identified by `<host-name, identifier>` pair
- each process in a given system has a unique name (process-id)

# How to identify the destination system/process?

---

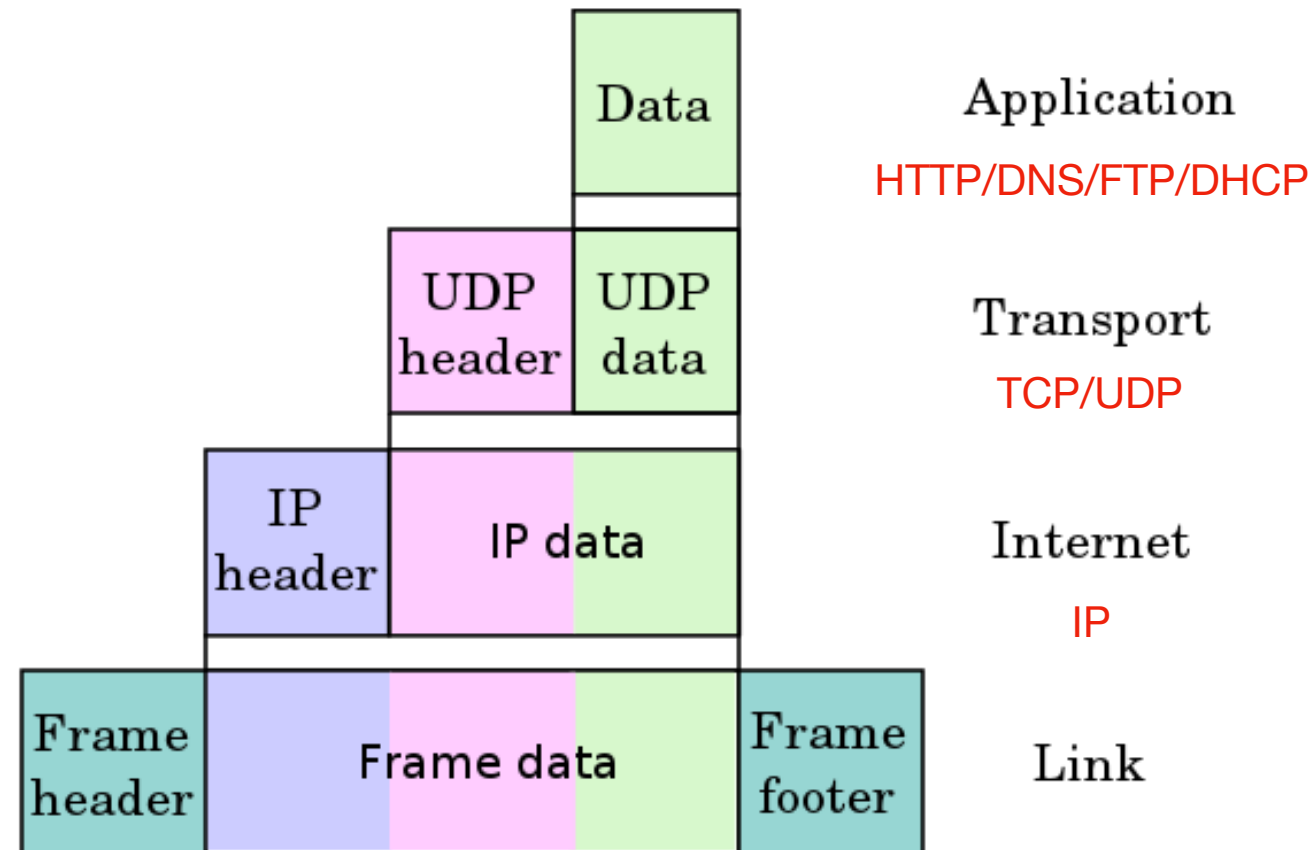
- a process on a remote system is identified by `<host-name, identifier>` pair
- each process in a given system has a unique name (process-id)
- each computer system in the network has a unique name
  - Domain Name System (DNS): a global distributed database system for resolving hostname-IP address mappings
    - 32-bit (IPv4) address: e.g., 129.128.5.180
    - humane readable names: e.g., gpu.srv.ualberta.ca



# TCP/IP protocol stack (Internet protocol suite)

---

enables end-to-end communications using functions organized into four abstraction layers




**routers only have the last two layers**

# TCP/IP protocol stack (Internet protocol suite)

---

enables end-to-end communications using functions organized into four abstraction layers

- Application layer
  - provides process-to-process data exchange for applications



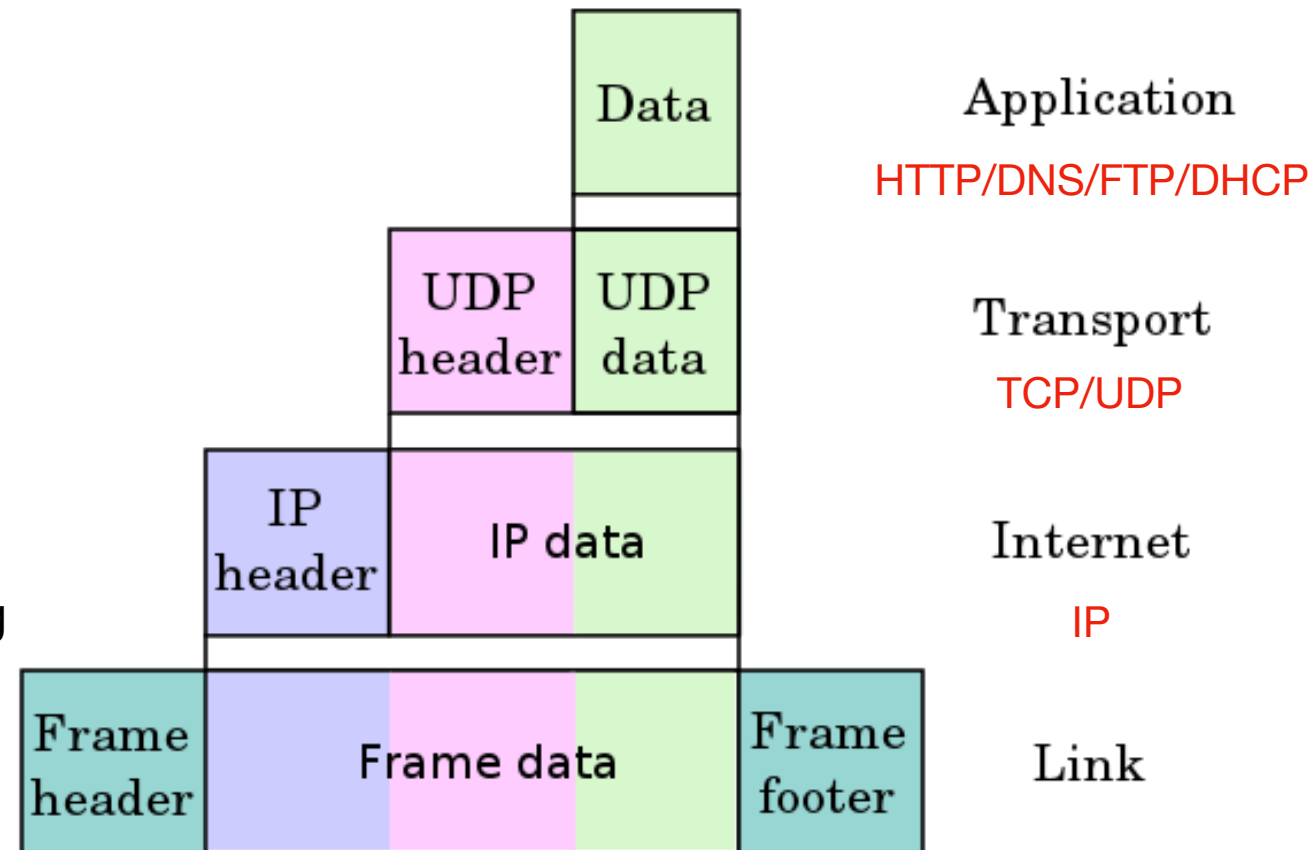
**routers only have the last two layers**

# TCP/IP protocol stack (Internet protocol suite)

---

enables end-to-end communications using functions organized into four abstraction layers

- Application layer
  - provides process-to-process data exchange for applications
- Transport layer
  - responsible for message transfer between hosts, including partitioning messages into TCP/UDP packets, maintaining packet order, and controlling flow




**routers only have the last two layers**

# TCP/IP protocol stack (Internet protocol suite)

---

enables end-to-end communications using functions organized into four abstraction layers

- Application layer
  - provides process-to-process data exchange for applications
- Transport layer
  - responsible for message transfer between hosts, including partitioning messages into TCP/UDP packets, maintaining packet order, and controlling flow
- Internet layer
  - responsible for routing IP packets through the network, and encoding/decoding addresses




**routers only have the last two layers**

# TCP/IP protocol stack (Internet protocol suite)

---

enables end-to-end communications using functions organized into four abstraction layers

- Application layer
  - provides process-to-process data exchange for applications
- Transport layer
  - responsible for message transfer between hosts, including partitioning messages into TCP/UDP packets, maintaining packet order, and controlling flow
- Internet layer
  - responsible for routing IP packets through the network, and encoding/decoding addresses
- Link layer
  - handles the frames, or fixed-length parts of packets, including transmission over a physical medium, and any error detection and recovery that occurred in the physical layer



**routers only have the last two layers**

# Media Access Control (MAC) address

---

How does a packet move from sender (host or router) to receiver on the same LAN?

# Media Access Control (MAC) address

---

How does a packet move from sender (host or router) to receiver on the same LAN?

- every Ethernet/WiFi device has a unique medium access control (MAC) address

# Media Access Control (MAC) address

---

How does a packet move from sender (host or router) to receiver on the same LAN?

- every Ethernet/WiFi device has a unique medium access control (MAC) address
- if a system wants to send data to another system, it needs to perform the IP to MAC address mapping
  - using **address resolution protocol (ARP)**
  - run `arp -a` to see the content of your arp table



# Internet Protocol (IP) address

---

- every host has a name and an associated **IP address**
  - 32-bit (IPv4) address: approximately 4.3 billion addresses
  - 128-bit (IPv6) address: approximately  $3.4 \times 10^{38}$  addresses
  - a special address is reserved for local host: 127.0.0.1

# Internet Protocol (IP) address

---

- every host has a name and an associated **IP address**
  - 32-bit (IPv4) address: approximately 4.3 billion addresses
  - 128-bit (IPv6) address: approximately  $3.4 \times 10^{38}$  addresses
  - a special address is reserved for local host: 127.0.0.1
- the sending system checks routing tables and locates a router to send packet

# Internet Protocol (IP) address

---

- every host has a name and an associated **IP address**
  - 32-bit (IPv4) address: approximately 4.3 billion addresses
  - 128-bit (IPv6) address: approximately  $3.4 \times 10^{38}$  addresses
  - a special address is reserved for local host: 127.0.0.1
- the sending system checks routing tables and locates a router to send packet
- each router uses the network part of host-id to determine where to transfer packet

# Internet Protocol (IP) address

---

- every host has a name and an associated **IP address**
  - 32-bit (IPv4) address: approximately 4.3 billion addresses
  - 128-bit (IPv6) address: approximately  $3.4 \times 10^{38}$  addresses
  - a special address is reserved for local host: 127.0.0.1
- the sending system checks routing tables and locates a router to send packet
- each router uses the network part of host-id to determine where to transfer packet
- the destination system receives the packet
  - it may be complete message, or it may need to be reassembled into larger message spanning multiple packets

# Transport layer address

---

- once a host with a specific IP address receives a packet, it must somehow pass it to the correct waiting process

# Transport layer address

---

- once a host with a specific IP address receives a packet, it must somehow pass it to the correct waiting process
- transport protocols, TCP and UDP, identify receiving and sending processes through the use of a port number (16 bits)
  - allows a host with a single IP address to have multiple processes sending and receiving packets
  - system or well-known ports are used to implement **standard** services: 0 through 1023
    - FTP – port 21/TCP; ssh – port 22/TCP; SMTP – port 25/TCP; HTTP – port 80/TCP; NTP – port 123/UDP
  - registered ports: 1024 through 49151
    - registered for specific services
  - dynamic or private ports: 49152 to 65535
    - available for use by any application communicating using TCP/UDP

# Transport layer address


---

- once a host with a specific IP address receives a packet, it must somehow pass it to the correct waiting process
- transport protocols, TCP and UDP, identify receiving and sending processes through the use of a port number (16 bits)
  - allows a host with a single IP address to have multiple processes sending and receiving packets
  - system or well-known ports are used to implement **standard** services: 0 through 1023
    - FTP – port 21/TCP; ssh – port 22/TCP; SMTP – port 25/TCP; HTTP – port 80/TCP; NTP – port 123/UDP
  - registered ports: 1024 through 49151
    - registered for specific services
  - dynamic or private ports: 49152 to 65535
    - available for use by any application communicating using TCP/UDP
- transport protocol can be simple or can add reliability to network packet stream

# User Datagram Protocol (UDP)

---

- UDP packets are also called **datagrams**
  - **messages** up to some maximum size






# User Datagram Protocol (UDP)

---


- UDP packets are also called **datagrams**
  - **messages** up to some maximum size
- UDP is unreliable
  - bare-bones extension to IP with the addition of port number
  - packets may be lost or received out-of-order



# User Datagram Protocol (UDP)

---

- UDP packets are also called **datagrams**
  - **messages** up to some maximum size
- UDP is unreliable
  - bare-bones extension to IP with the addition of port number
  - packets may be lost or received out-of-order
- UDP is also connectionless
  - no connection setup at the beginning of the transmission to set up state
  - also no connection tear-down at the end of transmission




# User Datagram Protocol (UDP)

- UDP packets are also called **datagrams**
  - **messages** up to some maximum size
- UDP is unreliable
  - bare-bones extension to IP with the addition of port number
  - packets may be lost or received out-of-order
- UDP is also connectionless
  - no connection setup at the beginning of the transmission to set up state
  - also no connection tear-down at the end of transmission

## Why to use unreliable communication?

many applications simply want to send data to a destination and not worry about packet loss



# Transmission Control Protocol (TCP)

---

- TCP is both reliable and connection-oriented

# Transmission Control Protocol (TCP)

---

- TCP is both reliable and connection-oriented
- in addition to port number, TCP provides abstraction to allow in-order, uninterrupted byte-stream across an unreliable network
  - whenever host sends packet, the receiver must send an acknowledgement packet (ACK)
  - if ACK is not received before a timer expires, the sender will timeout and retransmit the packet
    - requires keeping a copy of messages sent and not yet acknowledged
  - sequence counters in TCP header allow the receiver to put packets in order and notice duplicate packets (ack was lost)

# Transmission Control Protocol (TCP)

---

- TCP is both reliable and connection-oriented
- in addition to port number, TCP provides abstraction to allow in-order, uninterrupted byte-stream across an unreliable network
  - whenever host sends packet, the receiver must send an acknowledgement packet (ACK)
  - if ACK is not received before a timer expires, the sender will timeout and retransmit the packet
    - requires keeping a copy of messages sent and not yet acknowledged
  - sequence counters in TCP header allow the receiver to put packets in order and notice duplicate packets (ack was lost)
- connections are initiated with series of control packets
  - three-way handshake (SYN, SYN+ACK, ACK)

# Transmission Control Protocol (TCP)

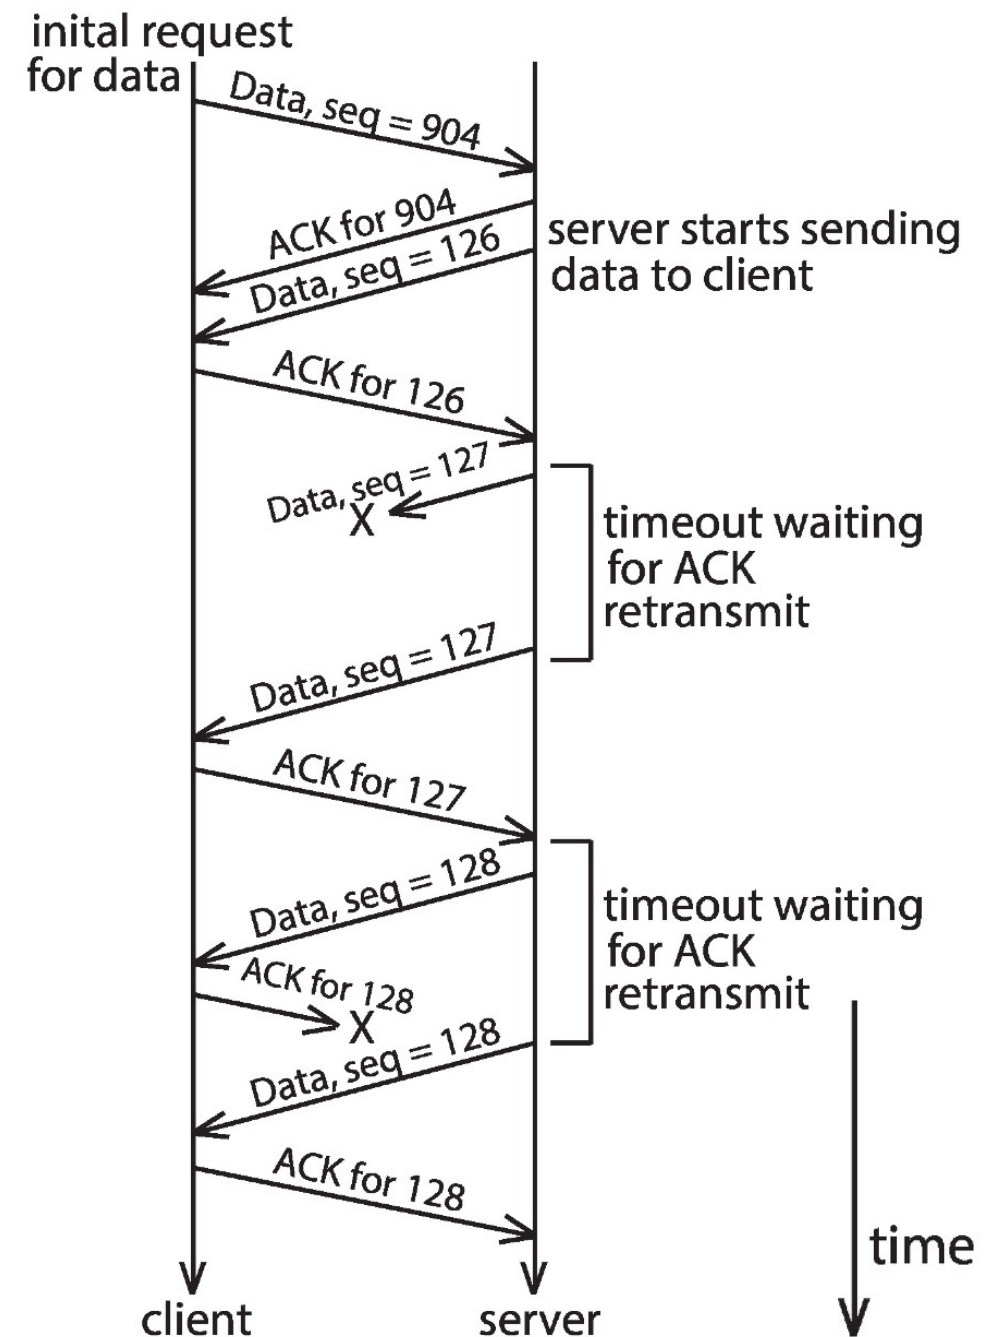
---

- TCP is both reliable and connection-oriented
- in addition to port number, TCP provides abstraction to allow in-order, uninterrupted byte-stream across an unreliable network
  - whenever host sends packet, the receiver must send an acknowledgement packet (ACK)
  - if ACK is not received before a timer expires, the sender will timeout and retransmit the packet
    - requires keeping a copy of messages sent and not yet acknowledged
  - sequence counters in TCP header allow the receiver to put packets in order and notice duplicate packets (ack was lost)
- connections are initiated with series of control packets
  - three-way handshake (SYN, SYN+ACK, ACK)
- connections also closed with series of control packets

# TCP data transfer

- receiver can send a cumulative ACK to acknowledge series of packets

- server can also send multiple packets before waiting for ACKs
- takes advantage of network throughput





# TCP data transfer

- receiver can send a cumulative ACK to acknowledge series of packets
  - server can also send multiple packets before waiting for ACKs
  - takes advantage of network throughput
- flow of packets regulated through flow control and congestion control
  - **flow control** – prevents sender from overrunning capacity of receiver
  - **congestion control** – approximates congestion of the network to slow down or speed up packet sending rate

