

Operating System Concepts

Lecture 16: Concurrency Problems

Omid Ardakanian
oardakan@ualberta.ca
University of Alberta

MWF 12:00-12:50 VVC 2 215

Today's class

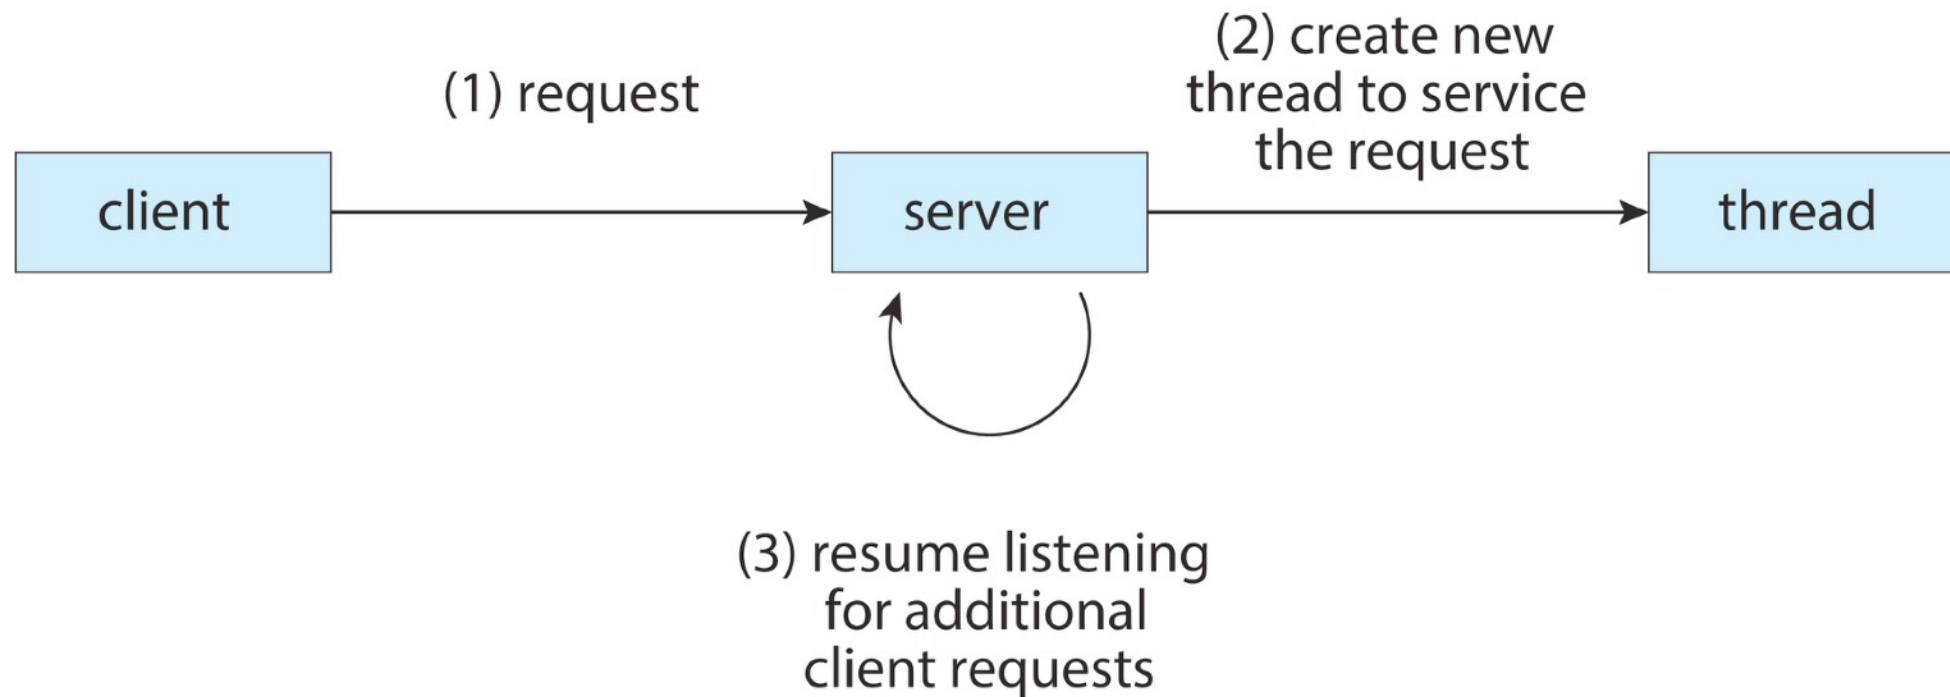
- Race condition
- Critical sections
- Example: too much milk
 - Three attempts to solve the problem

Recap

- cooperating processes (or threads) need to share information
 - sharing eliminates the need to copy data (increasing performance)
 - processes share data using IPC mechanisms
 - threads share global variables
- concurrent access to shared data may result in data inconsistency
 - they may not have a consistent view of the world
 - programmers are responsible for synchronizing access (**protecting**) globally shared data
- maintaining data consistency requires mechanisms to ensure orderly execution of cooperating processes

Sharing memory among threads can increase performance but can lead to problems...

- recall that in the concurrent (web) server example, each request is handled by a different thread



Sharing memory among threads can increase performance but can lead to problems...

- recall that in the concurrent (web) server example, each request is handled by a different thread
- there might be a **global** variable of type integer that keeps track of the number hits the index page of the cs.ualberta.ca website gets in a day
 - each thread has to increment this **shared** variable (`hits = hits + 1`) if the request is to access the index page
 - `hits = hits + 1` is not a single operation; it involves reading `hits` from memory into a register (loading), incrementing the register, and storing the register back to memory:

```
register1 <- hits
register1 <- register1 + 1
hits <- register1
```
 - thread execution can be interleaved because of time-slicing

Race condition

- Definition: when the outcome depends on the particular order in which threads that access some shared data are scheduled
 - the order of thread execution is non-deterministic
 - the last thread that changes a value wins

Race condition

- Definition: when the outcome depends on the particular order in which threads that access some shared data are scheduled
 - the order of thread execution is non-deterministic
 - the last thread that changes a value wins
- consider the following example
 - suppose $n = \text{STACK_SIZE} - 1$ before the two threads start and it is a global variable

Thread A runs:

```
if (n == STACK_SIZE) // A1
    return -1;        // A2
stack[n] = value;     // A3
n = n + 1;            // A4
```

Thread B runs:

```
if (n == STACK_SIZE) // B1
    return -1;        // B2
stack[n] = value;     // B3
n = n + 1;            // B4
```

Race condition

- Definition: when the outcome depends on the particular order in which threads that access some shared data are scheduled
 - the order of thread execution is non-deterministic
 - the last thread that changes a value wins
- consider the following example
 - suppose `n = STACK_SIZE - 1` before the two threads start and it is a global variable

Thread A runs:

```
if (n == STACK_SIZE) // A1
    return -1;        // A2
stack[n] = value;     // A3
n = n + 1;            // A4
```

Thread B runs:

```
if (n == STACK_SIZE) // B1
    return -1;        // B2
stack[n] = value;     // B3
n = n + 1;            // B4
```

- below are three possible scheduling orders:
 - A1,A3,A4,B1,B2 → thread B does not write the value (thread A wins)
 - A1,B1,A3,B3,A4,B4 → thread B overwrites the value written by thread A (thread B wins)
 - A1,B1,A3,A4,B3,B4 → thread B attempts to write the value at `stack[STACK_SIZE]` (**overflowing**)

Race condition (what's the value of x?)

Thread A runs:

```
int x = 0;
int y = 0;
...

funcA() {
    x = y + 1;    // A1
}
```

Thread B runs:

```
int x = 0;
int y = 0;
...

funcB() {
    y = 2;        // B1
    y = y * 2;    // B2
}
```

Race condition (what's the value of x?)

Thread A runs:

```
int x = 0;
int y = 0;
...

funcA() {
    x = y + 1;    // A1
}
```

Thread B runs:

```
int x = 0;
int y = 0;
...

funcB() {
    y = 2;        // B1
    y = y * 2;    // B2
}
```

- below are three possible scheduling orders:
 - A1,B1,B2 \rightarrow x = 1
 - B1,B2,A1 \rightarrow x = 5
 - B1,A1,B2 \rightarrow x = 3
 - we say that thread A races against thread B

Race condition (what's the value of x?)

Thread A runs:

```
int x = 0;
int y = 0;
...

funcA() {
    x = y + 1;    // A1
}
```

Thread B runs:

```
int x = 0;
int y = 0;
...

funcB() {
    y = 2;        // B1
    y = y * 2;    // B2
}
```

- below are three possible scheduling orders:
 - A1,B1,B2 \rightarrow x = 1
 - B1,B2,A1 \rightarrow x = 5
 - B1,A1,B2 \rightarrow x = 3
 - we say that thread A races against thread B
- bugs are intermittent and difficult to catch: a small change in this code (e.g., adding a print statement) can hide the bug
 - correct results are produced for some interleavings

Race condition in the producer-consumer problem

Process A runs:

```
while(1) {  
    while (counter == BUFFER_SIZE)  
        ; /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++; // num items in the buffer  
}
```

Process B runs:

```
while(1) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--; // num items in the buffer  
}
```

- the two processes use shared memory for communication
 - counter is a variable in the shared memory
 - one item is produced and one time is consumed so we expected to have counter = 5

Race condition in the producer-consumer problem

Process A runs:

```
while(1) {  
    while (counter == BUFFER_SIZE)  
        ; /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++; // num items in the buffer  
}
```

Process B runs:

```
while(1) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--; // num items in the buffer  
}
```

- the two processes use shared memory for communication
 - counter is a variable in the shared memory
 - one item is produced and one time is consumed so we expected to have counter = 5
- consider this execution interleaving with counter = 5 initially:
S0: producer execute register1 = counter {register1 = 5}
S1: producer execute register1 = register1 + 1 {register1 = 6}
S2: consumer execute register2 = counter {register2 = 5}
S3: consumer execute register2 = register2 - 1 {register2 = 4}
S4: producer execute counter = register1 {counter = 6}
S5: consumer execute counter = register2 {counter = 4}

Why race conditions exist?

- the order of thread execution is non-deterministic
 - multiprocessing: a system may contain multiple processors: cooperating threads/processes can execute simultaneously
 - multi-programming: thread/process execution can be interleaved because of time-slicing
- operations are not typically atomic
 - for example $x = x + 1$ is not a single operation

Why race conditions exist?

- the order of thread execution is non-deterministic
 - multiprocessing: a system may contain multiple processors: cooperating threads/processes can execute simultaneously
 - multi-programming: thread/process execution can be interleaved because of time-slicing
- operations are not typically atomic
 - for example $x = x + 1$ is not a single operation
- Goal: ensuring that your concurrent program produces the correct output under all possible interleavings
 - an operation must run to completion or not at all (atomicity), and an instruction sequence must be guaranteed to execute indivisibly
 - one thread executes an instruction sequence at a time

Critical sections

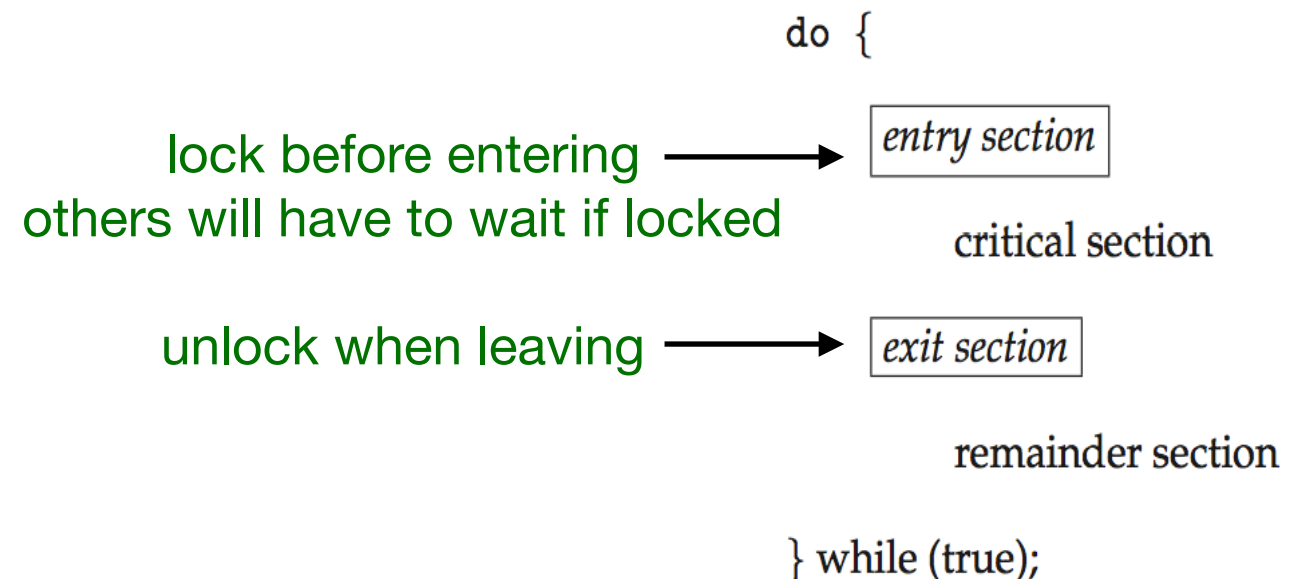
- critical section is a block of code (a number of consecutive program instructions) that cannot be executed in parallel by multiple threads (to avoid a race condition)
 - one process/thread running a critical section excludes the other ones (i.e., mutual exclusion). For example, two threads adding to a linked list can corrupt it
 - it is typically the code that accesses and/or modifies the values of shared variables (files, data structures, etc.)

Critical sections

- critical section is a block of code (a number of consecutive program instructions) that cannot be executed in parallel by multiple threads (to avoid a race condition)
 - one process/thread running a critical section excludes the other ones (i.e., mutual exclusion). For example, two threads adding to a linked list can corrupt it
 - it is typically the code that accesses and/or modifies the values of shared variables (files, data structures, etc.)
- there can be multiple critical sections in a program

Critical sections

- critical section is a block of code (a number of consecutive program instructions) that cannot be executed in parallel by multiple threads (to avoid a race condition)
 - one process/thread running a critical section excludes the other ones (i.e., mutual exclusion). For example, two threads adding to a linked list can corrupt it
 - it is typically the code that accesses and/or modifies the values of shared variables (files, data structures, etc.)
- there can be multiple critical sections in a program
- synchronization primitives (such as locks) are required to ensure that only one thread/process runs in the critical section at a time



Requirements

- a critical section implementation must be:
 - correct: the system behaves as if only one thread can execute in the critical section at any given time
 - efficient: getting in and getting out of a critical section must be fast critical sections should be as short as possible
 - flexible: must have as few restrictions as practically possible
 - supporting high concurrency: allows maximum concurrency while preserving correctness

A real life example

Roommate A

3:00 Arrive home: no milk in fridge

3:05 Leave for store

3:10 Arrive at store

3:15 Leave store

3:20 Return home, put milk away

3:25

3:30



A real life example

Roommate A

3:00 Arrive home: no milk in fridge
3:05 Leave for store
3:10 Arrive at store
3:15 Leave store
3:20 Return home, put milk away
3:25
3:30

Roommate B

Arrive home: no milk in fridge
Leave for store
Arrive at store
Leave store
Return home: **too much milk!**



milk is the
shared data
structure



A real life example

Roommate A

3:00 Arrive home: no milk in fridge

3:05 Leave for store

3:10 Arrive at store

3:15 Leave store

3:20 Return home, put milk away

3:25

3:30

Roommate B

Arrive home: no milk in fridge

Leave for store

Arrive at store

Leave store

Return home: **too much milk!**



milk is the
shared data
structure



Requirements (correctness properties):

1. someone buys milk if there is no milk in the fridge, otherwise starvation (**liveness**)
2. only one person buys milk, otherwise it spoils (**safety**)

First attempt: just leave a note!

```
while(1){  
    if(milk == 0) {  
        if(note == 0) {  
            note = 1;  
            buy_milk();  
            note = 0;  
        }  
    }  
}
```

each roommate should check if there is a note before buying milk (**waiting**)

each roommate should leave a note when going out to buy milk (**locking**)

add some value to milk

each roommate should remove the note after buying milk (**unlocking**)

Assumption: load and store operations are atomic
no hardware support is required

Failed attempt — threads can get context-switched at any time

Thread A

```
if(milk == 0) {  
    if(note == 0) {
```

```
        note = 1;  
        buy_milk();  
        note = 0;
```

```
    }
```

```
}
```

Thread B

```
if(milk == 0) {  
    if(note == 0) {  
        note = 1;  
        buy_milk();  
        note = 0;  
    }  
}
```

still too much milk


Second attempt

Thread A

```
note[0] = 1;
if(note[1] == 0) {
    if(milk == 0) {
        buy_milk();
    }
}
note[0] = 0;
```

Thread B

```
note[1] = 1;
if(note[0] == 0) {
    if(milk == 0) {
        buy_milk();
    }
}
note[1] = 0;
```



each roommate will leave a
labelled note before looking in
fridge
`boolean note[2];`

Failed attempt

Thread A

```
note[0] = 1;

if(note[1] == 0) {
    if(milk == 0) {
        buy_milk();
    }
}
note[0] = 0;
```

Thread B

```
note[1] = 1;

if(note[0] == 0) {
    if(milk == 0) {
        buy_milk();
    }
}
note[1] = 0;
```

this time we got no milk (starvation)

Third attempt

Thread A

```
note[0] = 1;
if(note[1] == 0) {
    if(milk == 0) {
        buy_milk();
    }
}
note[0] = 0;
```

Thread B

```
note[1] = 1;
while(note[0] == 1) {
    ; // spin
}
if(milk == 0) {
    buy_milk();
}
note[1] = 0;
```

Third attempt - scenario 1

Thread A

```
note[0] = 1;

if(note[1] == 0) {
    if(milk == 0) {
        buy_milk();
    }
}
note[0] = 0;
```

**only Thread B will execute
buy_milk()**

Thread B

```
note[1] = 1;

while(note[0] == 1) {
    ; // spin
}
```

```
if(milk == 0)
    buy_milk();
note[1] = 0;
```



Third attempt - scenario 2

Thread A

```
note[0] = 1;  
if(note[1] == 0) {
```

```
    if(milk == 0) {  
        buy_milk();  
    }  
}  
note[0] = 0;
```

**only Thread A will execute
buy_milk()**



Thread B

```
note[1] = 1;  
while(note[0] == 1) {  
    ; // spin  
}
```

```
if(milk == 0)  
    buy_milk();  
note[1] = 0;
```

Third attempt - scenario 3


Thread A

```
note[0] = 1;
if(note[1] == 0) {
    if(milk == 0) {
        buy_milk();
    }
}
note[0] = 0;
```

**only Thread B will execute
buy_milk()**

Thread B

```
note[1] = 1;
while(note[0] == 1) {
    ; // spin
}
if(milk == 0)
```

 buy_milk();
note[1] = 0;

Correctness of the third attempt

Thread A

```
note[0] = 1;
X: if(note[1] == 0) {
    if(milk == 0) {
        buy_milk();
    }
}
note[0] = 0;
```

Thread B

```
note[1] = 1;
Y: while(note[0] == 1) {
    ; // spin
}
if(milk == 0) {
    buy_milk();
}
note[1] = 0;
```

- at point X either there is a note left by Thread B or not
 - if there is a note, then B is either checking and buying milk as needed, or is waiting for A to remove the note. So in both cases, A must remove its note ASAP
 - if not, B has either bought milk or hasn't started yet. In both cases, A can safely check if milk is needed and buy

Correctness of the third attempt

Thread A

```
note[0] = 1;
X: if(note[1] == 0) {
    if(milk == 0) {
        buy_milk();
    }
}
note[0] = 0;
```

Thread B

```
note[1] = 1;
Y: while(note[0] == 1) {
    ; // spin
}
if(milk == 0) {
    buy_milk();
}
note[1] = 0;
```

- at point Y either there is a note left by Thread A or not
 - if there is a note, then A must be checking B's note or buying milk as needed. So B has to wait until there is no longer a note left by A; Once this happens B either finds milk that A bought or buys it if needed
 - if not, it is safe for B to buy milk as needed since A has either not yet started or has quit

Is this a good solution though?

- relies on load and store operations being atomic ✓
- is too complicated — it was hard to convince ourselves this solution works
- is asymmetrical — operations executed by Threads A and B are different
 - adding more threads would require different code for each new thread and modifications to existing threads
- requires **busy waiting** — Thread B is consuming CPU resources despite the fact that it is not doing any useful work

Homework

- write a program that finds the minimum of a list of integers using K threads
 - the list is partitioned equally among the threads
 - the minimum value found so far is stored in a global variable
- will there be race condition?