

Operating System Concepts

Lecture 18: Synchronization Primitives

Omid Ardakanian
oardakan@ualberta.ca
University of Alberta

MWF 12:00-12:50 VVC 2 215

Today's class

- Synchronization primitives
 - Mutex locks
 - Condition variables
 - Semaphores

Programming abstractions for synchronization

- With low-level hardware support, programming languages can provide atomic operations for synchronization
 - locks: can be held by at most one process/thread at a time; it is grabbed before critical section and released afterward
 - conditional variables: provide conditional synchronization
 - semaphores: more general version of locks
 - monitors: connect shared data to synchronized primitives

Higher-level API	locks	condition variables	semaphores	monitors
Hardware Support	atomic load and store	Interrupt disable	test_and_set	compare_and_swap

Mutex locks

- Definition: a high-level programming abstraction; an object that only one thread can hold at a time
 - can be implemented as a **spin-lock** which does busy waiting

```
while(test_and_set(&lock))  
    ; /* spin */  
/* critical section */  
lock = false;
```

- can be implemented as a blocking lock (next slide)

Blocking implementation of locks

```
class Lock {  
public:  
    void Acquire(); ← waits until lock is free and then grabs it  
    void Release(); ← releases the lock and wakes up any waiters  
private:  
    int locked;  
    Queue Q;  
}
```

Blocking implementation of locks

```
class Lock {  
public:  
    void Acquire();  
    void Release();  
private:  
    int locked;  
    Queue Q;  
}
```

- mutual exclusion can be implemented using locks (symmetric solution)

```
Lock milklock;   
...  
milklock.Acquire( )  
if(milk == 0)  
    buy_milk();  
milklock.Release( )
```

← initially free (not held by any process)

← acquired before
accessing shared data

← critical section

← released after
accessing shared data

Example of using locks (loose syntax)

```
void *malloc(size_t size) {  
    heaplock.acquire();  
    p = allocate memory of the specified size  
    heaplock.release();  
    return p;  
}  
  
void free(void *p) {  
    heaplock.acquire();  
    deallocate memory & put it back on free list  
    heaplock.release();  
}
```

threads of a process share the heap

How to implement locks on uniprocessors?

```
Lock::Acquire(Thread T) {
    intr_disable();
    if (locked == 0) {    // lock is free
        locked = 1;
    } else {              // lock is held by another thread
        queue_add(Q, T);
        thread_block(T); // put to sleep
    }
    intr_enable();
}

Lock::Release() {
    intr_disable();
    if(queue_empty(Q)) {
        locked = 0;           // release the lock
    } else {
        thread_unblock(queue_remove(Q)); // put on ready queue
    }
    intr_enable();
}
```

CLI and STI (privileged) instructions are used to clear and set interrupts respectively

How to implement locks on multiprocessors?

- a thread/process executing a CLI instruction does not disable interrupts on other processors!
- so we have to use other hardware support to implement acquire and release methods
 - test_and_set
 - compare_and_swap

Compare_and_swap

- test the value against some constant
- if the test returns true, set value in memory to different value
 - if [addr] == r1 then [addr] = r2;
- report the result of the test in a flag

```
Lock::Lock {  
    locked = 0;  
}
```

```
Lock::Acquire(Thread T) {  
    while(compare_and_swap(&locked, 0, 1) != 0)  
        ; // if busy, do nothing  
}
```

```
Lock::Release() {  
    value = 0;  
}
```

Test_and_set

- if lock is free (value = 0), test&set reads 0, sets value to 1, and returns 0
 - the Lock is now busy: the test in the while fails (Acquire is complete)
- if lock is busy (value = 1), test&set reads 1, sets value to 1, and returns 1
 - continues to loop until a Release is executed

```
Lock::Lock {  
    locked = 0;  
}
```

```
Lock::Acquire(Thread T) {  
    while (test_and_set(&locked) == 1) {  
        ; // if busy, do nothing  
    }  
}
```

```
Lock::Release() {  
    locked = 0;  
}
```

Can we build test_and_set locks without busy-waiting?

- we can't eliminate busy waiting entirely but we can minimize the busy-waiting time
 - instead of busy waiting until lock is free, we busy wait to atomically check the lock value and give up the CPU if we find that the lock is busy

```
class Lock {  
public:  
    void Acquire(Thread T);  
    void Release();  
private:  
    int locked;  
    int guard;  
    Queue Q;  
}
```

```
Lock::Lock {  
    locked = 0; // lock is free initially  
    guard = 0;  
}
```

Test_and_set — minimal waiting

```
Lock::Acquire(Thread T) {
    while(test_and_set(guard) == 1)
        ;
    if(locked != 0) {                // lock is busy
        queue_add(Q, T);
        thread_block(T, guard); // set guard to 0 before blocking thread
    } else {                        // lock is free
        locked = 1;
        guard = 0;
    }
}

Lock::Release() {
    while(test_and_set(guard) == 1)
        ;
    if(!queue_empty(Q)) {
        // take thread off wait queue and place on ready queue
        thread_unblock(queue_remove(Q));
    } else {
        locked = 0;
    }
    guard = 0;
}
```

Comparing to “interrupt disable” solution

```
Lock::Acquire(Thread T) {  
    intr_disable();  
    if (locked == 0) {  
        locked = 1;  
    } else {  
        queue_add(Q, T);  
        thread_block(T);  
    }  
    intr_enable();  
}
```

```
Lock::Release() {  
    intr_disable();  
    if(queue_empty(Q)) {  
        locked = 0;  
    } else {  
        thread_unblock(queue_remove(Q));  
    }  
    intr_enable();  
}
```

Replace

- `intr_disable()` with `while(test&set(guard))`;
- `intr_enable()` with `guard = 0`

Going beyond locks

- locks provide mutual exclusion but sometimes a thread has to wait only if a certain condition is true (synchronizing on a condition)
- Example: producer puts things in a fixed-size buffer, consumer takes them out
what are the constraints for bounded buffer?
 1. only one thread can manipulate buffer queue at a time (*mutual exclusion*)
 2. consumer must wait for producer to fill buffers **if all empty** (*scheduling constraint*)
 3. producer must wait for consumer to empty buffers **if all full** (*scheduling constraint*)

Condition variables

- Definition: an abstraction that supports **conditional synchronization**
 - a queue of threads waiting for a specific **event** inside a critical section
 - free memory is getting low, run the garbage collector
 - new data has arrived in the I/O port, process it
 - the condition of the condition variable depends on data protected by mutex lock

Condition variables

- Definition: an abstraction that supports **conditional synchronization**
 - a queue of threads waiting for a specific **event** inside a critical section
 - free memory is getting low, run the garbage collector
 - new data has arrived in the I/O port, process it
 - the condition of the condition variable depends on data protected by mutex lock
- provide three operations
 - `Wait()`
 - atomically release lock and go to sleep (block the thread until signalled)
 - reacquire lock upon waking up
 - `Notify()` — historically called `Signal()`
 - wake up a waiting thread, if any
 - `NotifyAll()` — historically called `Broadcast()`
 - wake up all waiting threads

Condition variables

- Definition: an abstraction that supports **conditional synchronization**
 - a queue of threads waiting for a specific **event** inside a critical section
 - free memory is getting low, run the garbage collector
 - new data has arrived in the I/O port, process it
 - the condition of the condition variable depends on data protected by mutex lock
- provide three operations
 - `Wait()`
 - atomically release lock and go to sleep (block the thread until signalled)
 - reacquire lock upon waking up
 - `Notify()` — historically called `Signal()`
 - wake up a waiting thread, if any
 - `NotifyAll()` — historically called `Broadcast()`
 - wake up all waiting threads
- thread must hold the lock when doing these condition variable operations
 - 1st reason: these operations may update the state
 - 2nd reason: to ensure signal and wait operations are not interleaved (by two threads)

Protocol for using condition variables

- acquire the lock to enter the critical section
- check condition inside the critical section
 - if condition is true: block the thread and release the lock
 - if condition is false: only release the lock

Example: the coke machine

Condition variables are used with a mutex lock and in a loop (to check the condition)

```
Class CokeMachine{
    ...
    storage for cokes (buffer)
    Lock lock;
    int count = 0;
    Condition notFull, notEmpty;
}

CokeMachine::Deposit(){
    lock->acquire( );
    while(count == n)
        notFull.wait(&lock); //release lock before blocking; reacquire when waking up
    add coke to the machine;
    count++;
    notEmpty.notify();
    lock->release();
}

CokeMachine::Remove(){
    lock->acquire();
    while(count == 0)
        notEmpty.wait(&lock); //release lock before blocking; reacquire when waking up
    remove coke from the machine;
    count--;
    notFull.notify();
    lock->release();
}
```

Semaphores

- semaphores are generalized locks invented by Dijkstra in 1965
 - they are a (non-negative) integer variable that supports two **atomic** operations: `wait` and `signal` (or `down` and `up`)

Semaphores

- semaphores are generalized locks invented by Dijkstra in 1965
 - they are a (non-negative) integer variable that supports two **atomic** operations: `wait` and `signal` (or `down` and `up`)
- **binary semaphore** (or mutex lock): used for mutual exclusion
 - guarantees mutually exclusive access to a resource (i.e., only one process is in the critical section at a time)
 - can vary from 0 to 1 – It is initialized to free (value = 1)

Semaphores

- semaphores are generalized locks invented by Dijkstra in 1965
 - they are a (non-negative) integer variable that supports two **atomic** operations: `wait` and `signal` (or `down` and `up`)
- **binary semaphore** (or mutex lock): used for mutual exclusion
 - guarantees mutually exclusive access to a resource (i.e., only one process is in the critical section at a time)
 - can vary from 0 to 1 – It is initialized to free (value = 1)
- **counting semaphore**: used for conditional synchronization
 - useful when multiple units of a specific resource are available
 - the initial count to which the semaphore is initialized is usually the number of resources
 - a process can acquire access so long as at least one unit of the resource is available

Atomic operations with semaphores

- each semaphore supports a queue of processes that are waiting to access the critical section (e.g., to check and buy milk)
- if a thread executes `Wait()` and the semaphore is free (non-zero), it continues executing after decrementing the semaphore's variable. But if it is not free, the OS puts the thread on the wait queue for that semaphore
- `Signal()` unblocks one thread on the semaphore's wait queue
- semaphores can be used to implement mutual exclusion:

Atomic operations with semaphores

- each semaphore supports a queue of processes that are waiting to access the critical section (e.g., to check and buy milk)
- if a thread executes `Wait()` and the semaphore is free (non-zero), it continues executing after decrementing the semaphore's variable. But if it is not free, the OS puts the thread on the wait queue for that semaphore
- `Signal()` unblocks one thread on the semaphore's wait queue
- semaphores can be used to implement mutual exclusion:

```
Semaphore milksemaphore; // suppose value is set to 1
```

```
milksemaphore.Wait( )    ← acquired before  
<critical section>      ← accessing shared data  
milksemaphore.Signal( ) ← released after  
                        ← accessing shared data
```

Implementing signal and wait by disabling interrupts

```
class Semaphore {
public:
    void Wait(Thread T);
    void Signal();
private:
    int value;
    Queue Q;
}

Semaphore::Semaphore(int val) {
    value = val; // initialized to the number of available resources
}

Semaphore::Wait(Thread T) {
    intr_disable();
    value = value - 1;
    if(value < 0) {        // |value| is the number of waiting threads
        queue_add(Q, T);
        thread_block(T);
    }
    intr_enable();
}

Semaphore::Signal() {
    intr_disable();
    value = value + 1;
    if(value <= 0)        // if there is a waiting thread
        thread_unblock(queue_remove(Q));
    intr_enable();
}
```

Implementing signal and wait by disabling interrupts

```
class Semaphore {
public:
    void Wait(Thread T);
    void Signal();
private:
    int value;
    Queue Q;
}

Semaphore::Semaphore(int val) {
    value = val; // initialized to the number of available resources
}

Semaphore::Wait(Thread T) {
    intr_disable();
    value = value - 1;
    if(value < 0) { // |value| is the number of waiting threads
        queue_add(Q, T);
        thread_block(T);
    }
    intr_enable();
}

Semaphore::Signal() {
    intr_disable();
    value = value + 1;
    if(value <= 0) // if there is a waiting thread
        thread_unblock(queue_remove(Q));
    intr_enable();
}
```

can you implement these atomic operations using test_and_set?