

Operating System Concepts

Lecture 20: Synchronization Examples

Omid Ardakanian
oardakan@ualberta.ca
University of Alberta

MWF 12:00-12:50 VVC 2 215

Today's class

- Classic synchronization problems
- POSIX synchronization

CLASSIC PROBLEMS OF SYNCHRONIZATION

The Bounded-Buffer Problem

- shared data: a finite buffer
- requirements
 - mutual exclusion: no lost items & no corrupt writes
 - producers must wait if buffer is full
 - consumers must wait if buffer is empty
- real world example: a web server

The Bounded-Buffer Problem

- shared data: a finite buffer
- requirements
 - mutual exclusion: no lost items & no corrupt writes
 - producers must wait if buffer is full
 - consumers must wait if buffer is empty
- real world example: a web server

```
class BoundedBufferMonitor{
    public:
        void P(Item item);
        void C(Item &item);
    private:
        Item buffer[N];
        int last, count;
        Condition full, empty;
        Lock lock;
}
```

```
BoundedBufferMonitor::BoundedBufferMonitor(){
    count = 0; // we need to keep track of the used spaces
    last = 0;
}
```

The Bounded-Buffer Problem

- shared data: a finite buffer
- requirements
 - mutual exclusion: no lost items & no corrupt writes
 - producers must wait if buffer is full
 - consumers must wait if buffer is empty
- real world example: a web server

```
class BoundedBufferMonitor{
    public:
        void P(Item item);
        void C(Item &item);
    private:
        Item buffer[N];
        int last, count;
        Condition full, empty;
        Lock lock;
}
```

**You can also write it with
binary and counting semaphores**

```
BoundedBufferMonitor::BoundedBufferMonitor(){
    count = 0; // we need to keep track of the used spaces
    last = 0;
}
```

The Bounded-Buffer Problem

```
void BoundedBufferMonitor::P(Item item){
    lock.Acquire();
    while(count == N)
        empty.Wait(lock);
    buffer[last] = item;
    last = (last + 1)%N;
    count++;
    full.Signal();
    lock.Release();
}
```

```
void BoundedBufferMonitor::C(Item &item){
    lock.Acquire();
    while(count == 0)
        full.Wait(lock);
    item = buffer[(last-count)%N]; // first=(last-count) mod N
    count--;
    empty.Signal();
    lock.Release();
}
```

The Readers-Writers Problem

- shared data: a buffer or database
- two types of threads: readers and writers
 - readers only read the shared data without modifying it
 - writers read the data and modify it

The Readers-Writers Problem

- shared data: a buffer or database
- two types of threads: readers and writers
 - readers only read the shared data without modifying it
 - writers read the data and modify it
- requirements
 - writers must have exclusive access to the shared data
 - multiple readers should be allowed to access shared data simultaneously as long as there is no writer

The Readers-Writers Problem

- shared data: a buffer or database
- two types of threads: readers and writers
 - readers only read the shared data without modifying it
 - writers read the data and modify it
- requirements
 - writers must have exclusive access to the shared data
 - multiple readers should be allowed to access shared data simultaneously as long as there is no writer
- real world example: account accesses and updates in a banking application

The Readers-Writers Problem

- shared data: a buffer or database
- two types of threads: readers and writers
 - readers only read the shared data without modifying it
 - writers read the data and modify it
- requirements
 - writers must have exclusive access to the shared data
 - multiple readers should be allowed to access shared data simultaneously as long as there is no writer
- real world example: account accesses and updates in a banking application
- why do we use different logics for readers and writers?
 - using a single lock is overly restrictive
 - helps a lot if
 - there are many readers but only a few writers
 - we can easily differentiate readers from writers (i.e., a thread does not alternate between reading and writing)

The Readers-Writers Problem

```
class ReadWrite {
    public:
        void Read();
        void Write();
    private:
        int readers;          // number of readers (shared between readers)
        Semaphore mutex;      // controls access to readers
        Semaphore wrt;         // controls entry to first writer or reader
}
```

```
ReadWrite::ReadWrite {
    readers = 0;
    mutex.value = 1;
    wrt.value = 1;
}
```

The Readers-Writers Problem

```
void ReadWrite::Write(){
    wrt.Wait();    // any writers or readers in the critical section
    <perform write>
    wrt.Signal();
}

void ReadWrite::Read(){
    mutex.Wait(); // mutual exclusion for data shared between readers
    readers += 1;  // increment the number of readers
    if(readers == 1) // it's the first reader
        wrt.Wait(); // block writers or this reader
    mutex.Signal();

    <perform read>

    mutex.Wait(); // mutual exclusion for data shared between readers
    readers -= 1;  // reader is done
    if (readers == 0) // it's the last reader
        wrt.Signal(); // unblock writers
    mutex.Signal();
}
```

The Readers-Writers Problem

```
void ReadWrite::Write(){
    wrt.Wait();    // any writers or readers in the critical section
    <perform write>
    wrt.Signal();
}

void ReadWrite::Read(){
    mutex.Wait(); // mutual exclusion for data shared between readers
    readers += 1;    // increment the number of readers
    if(readers == 1) // it's the first reader
        wrt.Wait();    // block writers or this reader
    mutex.Signal();

    <perform read>

    mutex.Wait(); // mutual exclusion for data shared between readers
    readers -= 1;    // reader is done
    if (readers == 0) // it's the last reader
        wrt.Signal();    // unblock writers
    mutex.Signal();
}
```

first critical
section

second critical
section

Scenario 1

Reader Thread 1

Reader Thread 2

Writer Thread 1

`ReadWrite::Read()`



`ReadWrite::Read()`



`calls wrt.Signal();`

`ReadWrite::Write()`



Scenario 2

Writer Thread 1

`ReadWrite::Write()`
calls `wrt.Wait()`;

Reader Thread 1

`ReadWrite::Read()`
blocks at `wrt.Wait()`;

unblocks to perform read

Reader Thread 2

`ReadWrite::Read()`
blocks at `mutes.Wait()`;

unlocks to perform read

Scenario 3

Reader Thread 1

`ReadWrite::Read()`



Writer Thread 1

`ReadWrite::Write()`
`blocks at wrt.Wait();`

`unblocks to perform write`



Reader Thread 2

`ReadWrite::Read()`



`calls wrt.Signal();`

Scenario 3

Reader Thread 1

`ReadWrite::Read()`



Writer Thread 1

`ReadWrite::Write()`
`blocks at wrt.Wait();`

`unblocks to perform write`



**Writer has to wait until all
readers are finished!**

Reader Thread 2

`ReadWrite::Read()`



`calls wrt.Signal();`

Discussions

- implementation notes
 - the first reader blocks (inside the first critical section) if there is a writer; any other readers who try to enter block on mutex
 - the last reader to exit signals a waiting writer
 - when a writer exits, if there is both a reader and a writer waiting, the scheduler determines which one goes next
 - if a reader goes next, then all readers that are waiting will fall through (at least one is waiting on wrt and zero or more can be waiting on mutex)
 - if a writer goes next, then readers will continue waiting

Discussions

- implementation notes
 - the first reader blocks (inside the first critical section) if there is a writer; any other readers who try to enter block on mutex
 - the last reader to exit signals a waiting writer
 - when a writer exits, if there is both a reader and a writer waiting, the scheduler determines which one goes next
 - if a reader goes next, then all readers that are waiting will fall through (at least one is waiting on wrt and zero or more can be waiting on mutex)
 - if a writer goes next, then readers will continue waiting
- alternative desirable semantics
 - let a writer enter its critical section as soon as possible

Favouring writers

to block readers as soon as a writer enters, we need to keep track of the number of writers and use a writer mutex lock to update this number

Favouring writers

to block readers as soon as a writer enters, we need to keep track of the number of writers and use a writer mutex lock to update this number

```
void ReadWrite::Write(){
    write_mutex.Wait(); // ensure mutual exclusion
    writers += 1;        // another pending writer
    if(writers == 1)     // first writer blocks readers
        read_block.Wait();
    write_mutex.Signal();

    write_block.Wait(); // ensure mutual exclusion
    <perform write>
    write_block.Signal();

    write_mutex.Wait(); // ensure mutual exclusion
    writers -= 1;        // writer done
    if(writers == 0)     // enable readers
        read_block.Signal();
    write_mutex.Signal();
}
```

Favouring writers

to block readers as soon as a writer enters, we need to keep track of the number of writers and use a writer mutex lock to update this number

```
void ReadWrite::Write(){
    write_mutex.Wait(); // ensure mutual exclusion
    writers += 1;       // another pending writer
    if(writers == 1)    // first writer blocks readers
        read_block.Wait();
    write_mutex.Signal();

    write_block.Wait(); // ensure mutual exclusion
    <perform write>
    write_block.Signal();

    write_mutex.Wait(); // ensure mutual exclusion
    writers -= 1;       // writer done
    if(writers == 0)    // enable readers
        read_block.Signal();
    write_mutex.Signal();
}
```

first critical
section

second critical
section

Favouring writers

how to modify the Read () method?


- it's your homework

Read/Write locks

- pthreads support read/write locks
 - a thread can acquire a read lock or a write lock
 - many threads can hold the same read lock concurrently
 - only one thread can hold a write lock
- use pthread functions:
`pthread_rwlock_init()`
`pthread_rwlock_rdlock()`
`pthread_rwlock_wrlock()`
`pthread_rwlock_unlock()`


The Dining-Philosophers Problem

- shared data: a fixed number of chopsticks
- setting:
 - philosophers sit around a circular table
 - chopsticks are placed between philosophers
 - hence, $\text{\#philosophers} = \text{\#chopsticks}$




The Dining-Philosophers Problem

- shared data: a fixed number of chopsticks
- setting:
 - philosophers sit around a circular table
 - chopsticks are placed between philosophers
 - hence, $\text{\#philosophers} = \text{\#chopsticks}$
- requirements
 - philosophers spend their lives alternating thinking and eating (**what a boring life!**)
 - they occasionally get hungry and decide to eat!
 - they don't like to talk to each other and coordinate
 - philosophers need two chopsticks to eat from bowl
 - but can pick up one at a time
 - they can only pick up the two chopsticks that are closest to them
 - philosophers put down chopsticks when finished
 - they seek a deadlock-free and starvation-free solution



The Dining-Philosophers Problem

- shared data: a fixed number of chopsticks
- setting:
 - philosophers sit around a circular table
 - chopsticks are placed between philosophers
 - hence, $\# \text{philosophers} = \# \text{chopsticks}$
- requirements
 - philosophers spend their lives alternating thinking and eating (**what a boring life!**)
 - they occasionally get hungry and decide to eat!
 - they don't like to talk to each other and coordinate
 - philosophers need two chopsticks to eat from bowl
 - but can pick up one at a time
 - they can only pick up the two chopsticks that are closest to them
 - philosophers put down chopsticks when finished
 - they seek a deadlock-free and starvation-free solution
- real world example:
 - cooperating processes shared constrained resources
 - travel reservation: hotel, airline, car rental



First attempt

```
#define N    5                // number of philosophers
#define L    i                // index of i's left neighbour
#define R    (i+1)%N          // index of i's right neighbour

semaphore chopstick[N]; // one semaphore per chopstick

void philosopher(int i) { // i is the index of the philosopher
    while(true) {
        chopstick[L].wait();
        chopstick[R].wait();
        /* eat for awhile */
        chopstick[L].signal();
        chopstick[R].signal();
        /* think for awhile */
    }
}
```

First attempt

```
#define N    5                // number of philosophers
#define L    i                // index of i's left neighbour
#define R    (i+1)%N          // index of i's right neighbour

semaphore chopstick[N]; // one semaphore per chopstick

void philosopher(int i) { // i is the index of the philosopher
    while(true) {
        chopstick[L].wait();
        chopstick[R].wait();
        /* eat for awhile */
        chopstick[L].signal();
        chopstick[R].signal();
        /* think for awhile */
    }
}
```

What's the problem with this solution? could create a deadlock

Second attempt

```
#define N    5                // number of philosophers
#define L    (i+N-1)%N        // index of i's left neighbour
#define R    (i+1)%N          // index of i's right neighbour
#define THINKING  0
#define HUNGRY    1
#define EATING    2

int state[N];                // array that keeps everyone's state
semaphore mutex = 1;          // binary semaphore (mutual exclusion)
semaphore s[N];              // one semaphore per philosopher

void philosopher(int i) { // i is the index of the philosopher
    while(true) {
        think();            // return when philosopher gets hungry
        takeForks(i);
        eat();               // eating can take a while...
        releaseForks(i);
    }
}
```

Second attempt

```
void takeForks(int i) { // i is the index of the philosopher
    mutex.wait();
    state[i] = HUNGRY;    // update the state
    test(i);              // pick up the two chopsticks and eat, if possible
    mutex.signal();
    s[i].wait();
}

void releaseForks(int i) { // i is the index of the philosopher
    mutex.wait();
    state[i] = THINKING; // update the state
    test(L);             // let the left neighbour eat, if possible
    test(R);             // let the right neighbour eat, if possible
    mutex.signal();
}

void test(int i) { // i is the index of the philosopher
    if(state[i] == HUNGRY && state[L] != EATING && state[R] != EATING) {
        // can pick up two chopsticks and eat
        state[i] = EATING;
        s[i].signal();
    }
}
```


Second attempt

```
void takeForks(int i) { // i is the index of the philosopher
    mutex.wait();
    state[i] = HUNGRY;    // update the state
    test(i);              // pick up the two chopsticks and eat, if possible
    mutex.signal();
    s[i].wait();
}

void releaseForks(int i) { // i is the index of the philosopher
    mutex.wait();
    state[i] = THINKING; // update the state
    test(L);             // let the left neighbour eat, if possible
    test(R);             // let the right neighbour eat, if possible
    mutex.signal();
}

void test(int i) { // i is the index of the philosopher
    if(state[i] == HUNGRY && state[L] != EATING && state[R] != EATING) {
        // can pick up two chopsticks and eat
        state[i] = EATING;
        s[i].signal();
    }
}
```

Is it a good solution?

POSIX SYNCHRONIZATION

POSIX synchronization

- POSIX API provides
 - mutex locks
 - semaphores
 - condition variables
- they are widely used on UNIX, Linux, and macOS

POSIX mutex locks

- creating and initializing the lock

```
#include <pthread.h>

pthread_mutex_t mutex;

/* create and initialize the mutex lock */
pthread_mutex_init(&mutex, NULL);
```

- acquiring and releasing the lock

```
/* acquire the mutex lock */
pthread_mutex_lock(&mutex);

/* critical section */

/* release the mutex lock */
pthread_mutex_unlock(&mutex);
```

POSIX semaphores

- POSIX provides two versions — named and unnamed
- named semaphores can be used by unrelated processes, unnamed semaphores cannot be used in this case
- see <https://linux.die.net/man/7/sem> overview

POSIX named semaphores

- creating an initializing the semaphore:

```
#include <semaphore.h>
sem_t *sem;

/* create the semaphore and initialize it to 1 */
see = sem_open("SEM_NAME", O_CREAT, 0666, 1);
```

- another process can access the semaphore by referring to its name SEM
- acquiring and releasing the semaphore:

```
/* acquire the semaphore */
sem_wait(sem);

/* critical section */

/* release the semaphore */
sem_post(sem);
```

POSIX unnamed semaphores

- creating an initializing the semaphore:

```
#include <semaphore.h>
sem_t sem;
```

```
/* create the semaphore and initialize it to 1 */
sem_init(&sem, 0, 1); // second arg determines if sem is shared
                    // between threads of a process or between
                    // processes
```

- acquiring and releasing the semaphore:

```
/* acquire the semaphore */
sem_wait(&sem);
```

```
/* critical section */
```

```
/* release the semaphore */
sem_post(&sem);
```

POSIX condition variables

- POSIX condition variables are associated with a POSIX mutex lock to provide mutual exclusion
- creating and initializing the condition variable:

```
pthread_mutex_t mutex;  
pthread_cond_t cond_var;
```

```
pthread_mutex_init(&mutex, NULL);  
pthread_cond_init(&cond_var, NULL);
```


POSIX condition variables

- thread waiting for the condition `a == b` to become true:

```
pthread_mutex_lock(&mutex);  
while (a != b)  
    pthread_cond_wait(&cond_var, &mutex);  
pthread_mutex_unlock(&mutex);
```

- thread signalling another thread waiting on the condition variable:

```
pthread_mutex_lock(&mutex);  
a = b;  
pthread_cond_signal(&cond_var);  
pthread_mutex_unlock(&mutex);
```

Linux synchronization

- prior to kernel version 2.6
 - disables interrupts (kernel preemption) to implement short critical sections
- version 2.6 and later
 - fully preemptive
- Linux provides:
 - semaphores
 - atomic integers
 - spinlocks
 - reader-writer versions of both