

Operating System Concepts

Lecture 22: Deadlock — Part 2

Omid Ardakanian
oardakan@ualberta.ca
University of Alberta

MWF 12:00-12:50 VVC 2 215


Imposing a total ordering of resources

- eliminate circular waiting by ordering all locks (or semaphores, or resources)
 - define a one-to-one function $F : R \rightarrow N$
 - all code grabs locks in a predefined order
 - if several instances of the same resource are needed, a single request is issued
- problems?
 - maintaining global order is difficult, especially in a large project
 - programmers may not follow the ordering...
 - the global order can force the programmer to grab a lock earlier than it would like, tying up a resource for too long

Imposing a total ordering of resources

- in the dining philosopher problem
 - number chopsticks from 1 to 5
 - pick up the lower number chopstick first
 - this will break the circular wait condition


```
void philosopher(int i) {  
    while(true) {  
        chopstick[LowerNum(i)].wait();  
        chopstick[HigherNum(i)].wait();  
        ...  
    }  
}
```



Deadlock detection

if there is a single instance of each resource type

$T_i \rightarrow R_j$ and $R_j \rightarrow T_k$ yields $T_i \rightarrow T_k$



Deadlock detection (single instance of each resource type)

- scan the resource allocation graph for cycles
 - detecting cycles is $O(|E|+|V|)$
- when should we execute this algorithm?
 - whenever a resource request can't be filled
 - on a regular schedule (hourly or ...)
 - when CPU utilization drops below some threshold
 - just before granting a resource, check if granting it would lead to a cycle (it is deadlock avoidance rather than detection in this case)

Deadlock detection (single instance of each resource type)

- scan the resource allocation graph for cycles
 - detecting cycles is $O(|E|+|V|)$
- when should we execute this algorithm?
 - whenever a resource request can't be filled
 - on a regular schedule (hourly or ...)
 - when CPU utilization drops below some threshold
 - just before granting a resource, check if granting it would lead to a cycle (it is deadlock avoidance rather than detection in this case)
- what do Linux and Windows do?
 - ignore the problem altogether; leave it to the programmer/application
 - why? **because it's cheaper**, especially if deadlock occurs infrequently

Deadlock recovery

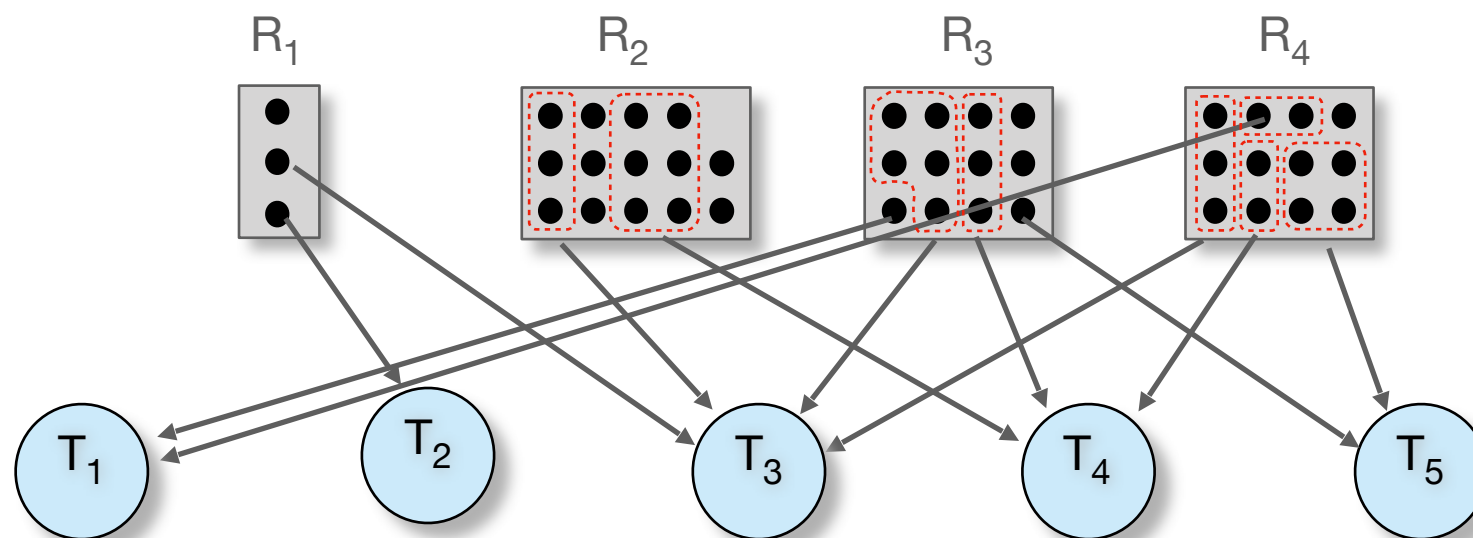
- inform the operator (manual recovery)

Deadlock recovery

- inform the operator (manual recovery)
- abort all deadlocked threads and reclaim their resources

Deadlock recovery

- inform the operator (manual recovery)
- abort all deadlocked threads and reclaim their resources
- abort one thread at a time and reclaim its resources; continue until all cycles in the resource allocation graph are eliminated
 - where to start? minimum cost is not well-defined? (a) low priority thread, (b) thread with most allocation of resources, (c) thread running for the longest time



Deadlock recovery

- inform the operator (manual recovery)
- abort all deadlocked threads and reclaim their resources
- abort one thread at a time and reclaim its resources; continue until all cycles in the resource allocation graph are eliminated
 - where to start? minimum cost is not well-defined? (a) low priority thread, (b) thread with most allocation of resources, (c) thread running for the longest time
- preempt some resources from one of multiple deadlocked threads?
 - which resources?

Deadlock recovery

- inform the operator (manual recovery)
- abort all deadlocked threads and reclaim their resources
- abort one thread at a time and reclaim its resources; continue until all cycles in the resource allocation graph are eliminated
 - where to start? minimum cost is not well-defined? (a) low priority thread, (b) thread with most allocation of resources, (c) thread running for the longest time
- preempt some resources from one of multiple deadlocked threads?
 - which resources?
- caveat: make sure that the system is in consistent state

Today's class

- Dealing with deadlock
 - Deadlock avoidance
 - Deadlock detection with several instances of a resource type

Reserve resources instead of obtaining them

- obtaining all resources at startup reduces resource utilization

Reserve resources instead of obtaining them

- obtaining all resources at startup reduces resource utilization
- a better approach is to state maximum resources that will be needed at startup and allocate resources dynamically when needed
 - assumptions: (a) the maximum resource need is known a priori; (b) when a thread gets all its resources it must return them in a finite amount of time

Reserve resources instead of obtaining them

- obtaining all resources at startup reduces resource utilization
- a better approach is to state maximum resources that will be needed at startup and allocate resources dynamically when needed
 - assumptions: (a) the maximum resource need is known a priori; (b) when a thread gets all its resources it must return them in a finite amount of time
 - when is a request granted?
 - the sum of the stated maximum resources exceeds the total available resources? **No, threads do not request resources at once**

Reserve resources instead of obtaining them

- obtaining all resources at startup reduces resource utilization
- a better approach is to state maximum resources that will be needed at startup and allocate resources dynamically when needed
 - assumptions: (a) the maximum resource need is known a priori; (b) when a thread gets all its resources it must return them in a finite amount of time
 - when is a request granted?
 - the sum of the stated maximum resources exceeds the total available resources? **No, threads do not request resources at once**
 - some sequential ordering of threads is deadlock free? **Yes!**

Reserve resources instead of obtaining them

- obtaining all resources at startup reduces resource utilization
- a better approach is to state maximum resources that will be needed at startup and allocate resources dynamically when needed
 - assumptions: (a) the maximum resource need is known a priori; (b) when a thread gets all its resources it must return them in a finite amount of time
 - when is a request granted?
 - the sum of the stated maximum resources exceeds the total available resources? **No, threads do not request resources at once**
 - some sequential ordering of threads is deadlock free? **Yes!**
 - there just needs to be some way for all threads to finish

Deadlock avoidance with resource reservation

- threads provide advance information about the maximum resources they may need during execution

Deadlock avoidance with resource reservation

- threads provide advance information about the maximum resources they may need during execution
- define a sequence of threads $\{t_1, \dots, t_n\}$ as **safe** if for each t_i , the resources that t_i will eventually request can be satisfied by the currently available resources plus the resources held by all t_j where $j < i$
 - hence a safe state is a state in which there is at least one safe sequence for the threads

Deadlock avoidance with resource reservation

- threads provide advance information about the maximum resources they may need during execution
- define a sequence of threads $\{t_1, \dots, t_n\}$ as **safe** if for each t_i , the resources that t_i will eventually request can be satisfied by the currently available resources plus the resources held by all t_j where $j < i$
 - hence a safe state is a state in which there is at least one safe sequence for the threads
- an unsafe state is not equivalent to deadlock, it **may** lead to deadlock
 - since some threads might not actually use the maximum resources they have declared

Deadlock avoidance with resource reservation


- threads provide advance information about the maximum resources they may need during execution
- define a sequence of threads $\{t_1, \dots, t_n\}$ as **safe** if for each t_i , the resources that t_i will eventually request can be satisfied by the currently available resources plus the resources held by all t_j where $j < i$
 - hence a safe state is a state in which there is at least one safe sequence for the threads
- an unsafe state is not equivalent to deadlock, it **may** lead to deadlock
 - since some threads might not actually use the maximum resources they have declared
- to summarize
 - a resource is given to a thread if the new state is safe
 - otherwise, the thread must wait **even if the resource is currently available**

Deadlock avoidance with resource reservation

- threads provide advance information about the maximum resources they may need during execution
- define a sequence of threads $\{t_1, \dots, t_n\}$ as **safe** if for each t_i , the resources that t_i will eventually request can be satisfied by the currently available resources plus the resources held by all t_j where $j < i$
 - hence a safe state is a state in which there is at least one safe sequence for the threads
- an unsafe state is not equivalent to deadlock, it **may** lead to deadlock
 - since some threads might not actually use the maximum resources they have declared
- to summarize
 - a resource is given to a thread if the new state is safe
 - otherwise, the thread must wait **even if the resource is currently available**
- this algorithm ensures no circular-wait condition exists

Definitions

- safe state: for any possible sequence of future resource requests, it is possible to eventually grant all requests
 - if a system is in safe state, then no deadlock
- unsafe state: some sequence of resource requests can result in deadlock
 - if a system is in unsafe state, possibility of deadlock
- doomed state: all possible computations lead to deadlock
- deadlock avoidance: ensure that the system never enters an unsafe state



Example

- threads t_1 , t_2 , and t_3 are competing for 12 tape drives; 11 drives are currently allocated to the threads, only 1 tap drive is available at this point
- the current state is **safe**
 - there exists a safe sequence, $\{t_1, t_2, t_3\}$, where all threads may obtain their maximum number of resources

Example

- threads t_1 , t_2 , and t_3 are competing for 12 tape drives; 11 drives are currently allocated to the threads, only 1 tap drive is available at this point
- the current state is **safe**
 - there exists a safe sequence, $\{t_1, t_2, t_3\}$, where all threads may obtain their maximum number of resources

	max need	in use	could want
t_1	4	3	1
t_2	8	4	4
t_3	12	4	8

Example

- if t_3 requests one more drive, then it must wait since allocating the drive would lead to an unsafe state
- granting this request will result in having no resource available, hence none of the threads can finish their tasks

Example

- if t_3 requests one more drive, then it must wait since allocating the drive would lead to an unsafe state
- granting this request will result in having no resource available, hence none of the threads can finish their tasks

	max need	in use	could want
t_1	4	3	1
t_2	8	4	4
t_3	12	5	7

Banker's algorithm

- each thread must a priori claim maximum use
 - grant a request if and only if it results in a safe state
 - otherwise the thread has to wait

Data structures for the Banker's algorithm

Let n be the number of threads and m be the number of resource types

- `available` is a vector of length m
 - if `available[j]=k` there are k instances of resource type R_j available
- `max` is an n by m matrix
 - if `max[i][j]=k` then thread T_i may request at most k instances of resource type R_j
- `allocation` is an n by m matrix
 - if `allocation[i][j]=k` then thread T_i is currently allocated k instances of R_j
- `need` is an n by m matrix
 - if `need[i][j]=k` then T_i may need k more instances of R_j to complete its task
 - thus `need[i][j] = max[i][j] - allocation[i][j]`

these data structures vary over time in size and value

Resource-request algorithm

let $\text{request}[i]$ be the request vector of thread T_i

if $\text{request}[i][j]=k$ then T_i wants k instances of resource type R_j

1. if $\text{request}[i] \leq \text{need}[i]$ go to step 2, otherwise, raise error condition since T_i has exceeded its maximum claim
2. if $\text{request}[i] \leq \text{available}$, go to step 3, otherwise T_i must wait since resources are not available
3. pretend to allocate requested resources to T_i by modifying the state as follows:
 $\text{available} -= \text{request}[i]$
 $\text{allocation}[i] += \text{request}[i]$
 $\text{need}[i] -= \text{request}[i]$

run the **safety check algorithm**. If safe, allocate resources to T_i

otherwise, T_i must wait and the old resource-allocation state is restored

Safety check algorithm

Let `work` and `finish` be vectors of length `m` and `n`, respectively.

1. Initialize:

```
work = available  
finish[i] = false for i={0, 1, ..., n-1}
```

2. Find an index `i` such that both:

(a) `finish[i] = false`

(b) `need[i] ≤ work`

if no such `i` exists, go to step 4

3. Set

```
work = work + allocation[i]
```

```
finish[i] = true
```

and go to step 2

4. If `finish[i]==true` for all `i`, then the system is in a safe state

Example

- 5 threads (T_0, T_1, T_2, T_3, T_4), 3 resource types (R_0, R_1, R_2),
 $W_0=10, W_1=5, W_2=7$

available: (3,3,2)

	allocation (R_0, R_1, R_2)	max (R_0, R_1, R_2)	need (R_0, R_1, R_2)
t_0	(0,1,0)	(7,5,3)	(7,4,3)
t_1	(2,0,0)	(3,2,2)	(1,2,2)
t_2	(3,0,2)	(9,0,2)	(6,0,0)
t_3	(2,1,1)	(2,2,2)	(0,1,1)
t_4	(0,0,2)	(4,3,3)	(4,3,1)

- the system is in a safe state since the sequence $\{T_1, T_3, T_4, T_2, T_0\}$ is a safe sequence

Example: can request for (3,3,0) by T_4 be granted?

- Pretend that the request is granted

available: (0,0,2)

	allocation (R_0, R_1, R_2)	max (R_0, R_1, R_2)	need (R_0, R_1, R_2)
t_0	(0,1,0)	(7,5,3)	(7,4,3)
t_1	(2,0,0)	(3,2,2)	(1,2,2)
t_2	(3,0,2)	(9,0,2)	(6,0,0)
t_3	(2,1,1)	(2,2,2)	(0,1,1)
t_4	(3,3,2)	(4,3,3)	(1,0,1)

Example: can request for (3,3,0) by T_4 be granted?

- Pretend that the request is granted

available: (0,0,2)

	allocation (R_0, R_1, R_2)	max (R_0, R_1, R_2)	need (R_0, R_1, R_2)
t_0	(0,1,0)	(7,5,3)	(7,4,3)
t_1	(2,0,0)	(3,2,2)	(1,2,2)
t_2	(3,0,2)	(9,0,2)	(6,0,0)
t_3	(2,1,1)	(2,2,2)	(0,1,1)
t_4	(3,3,2)	(4,3,3)	(1,0,1)

- the system is still in an unsafe state
- hence T_4 must wait and the saved state is restored

Example: can request for (1,0,2) by T_1 be granted?

- Pretend that the request is granted

available: (2,3,0)

	allocation (R_0, R_1, R_2)	max (R_0, R_1, R_2)	need (R_0, R_1, R_2)
t_0	(0,1,0)	(7,5,3)	(7,4,3)
t_1	(3,0,2)	(3,2,2)	(0,2,0)
t_2	(3,0,2)	(9,0,2)	(6,0,0)
t_3	(2,1,1)	(2,2,2)	(0,1,1)
t_4	(0,0,2)	(4,3,3)	(4,3,1)

Example: can request for (1,0,2) by T_1 be granted?

- Pretend that the request is granted

available: (2,3,0)

	allocation (R_0, R_1, R_2)	max (R_0, R_1, R_2)	need (R_0, R_1, R_2)
t_0	(0,1,0)	(7,5,3)	(7,4,3)
t_1	(3,0,2)	(3,2,2)	(0,2,0)
t_2	(3,0,2)	(9,0,2)	(6,0,0)
t_3	(2,1,1)	(2,2,2)	(0,1,1)
t_4	(0,0,2)	(4,3,3)	(4,3,1)

- the system is still in a safe state since the sequence $\{T_1, T_3, T_4, T_2, T_0\}$ is a safe sequence
- so the request will be granted

Problem with the banker algorithm

- the banker's algorithm is slow
 - it's too slow to run on every allocation

Problem with the banker algorithm

- the banker's algorithm is slow
 - it's too slow to run on every allocation
- what else can we do?
 - combine the three basic approaches:
 - prevention: ensure that the system will never enter a deadlock state
 - avoidance: ensure that the system will never enter a deadlock state
 - detection: let the system enter a deadlock state and then recover
 - partition resources into hierarchically ordered classes
 - use the most appropriate technique for handling deadlocks within each class

Homework

modify the safety check algorithm such that it can be used for detecting deadlock when there are multiple instances of each resource