

Operating System Concepts

Lecture 24: Segmentation and Paging

Omid Ardakanian
oardakan@ualberta.ca
University of Alberta

MWF 12:00-12:50 VVC 2 215

Today's class


- Memory-management strategies
 - contiguous memory allocation
 - segmentation
 - paging
- Hardware support
 - mapping tables in addition to registers

Segmentation

- process generates contiguous virtual addresses from 0 to MAX_{proc}
- OS lays the process down on **arbitrary size segments** and hardware translates virtual addresses to physical addresses in memory
 - using **a segment table** that keeps track of the segment locations in memory

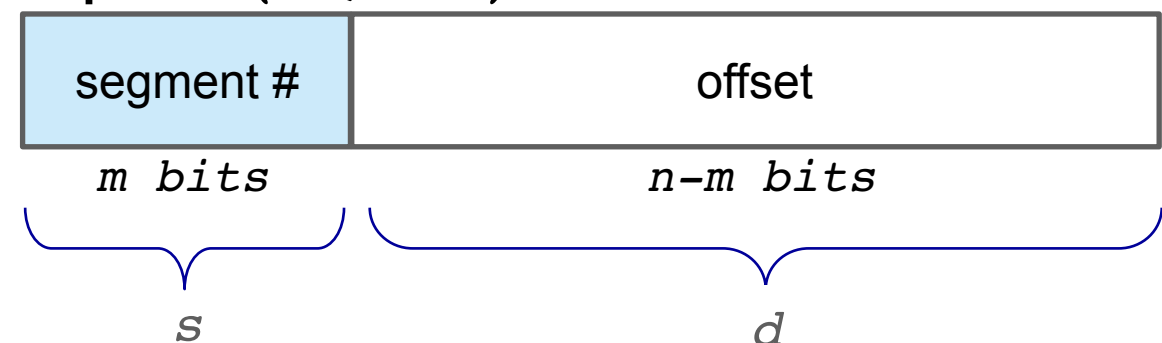
Segmentation

- process generates contiguous virtual addresses from 0 to MAX_{proc}
- OS lays the process down on **arbitrary size segments** and hardware translates virtual addresses to physical addresses in memory
 - using **a segment table** that keeps track of the segment locations in memory
- the virtual address is identified by a pair (s, d)
 - s is the segment number
 - d is the offset within the segment



Segmentation

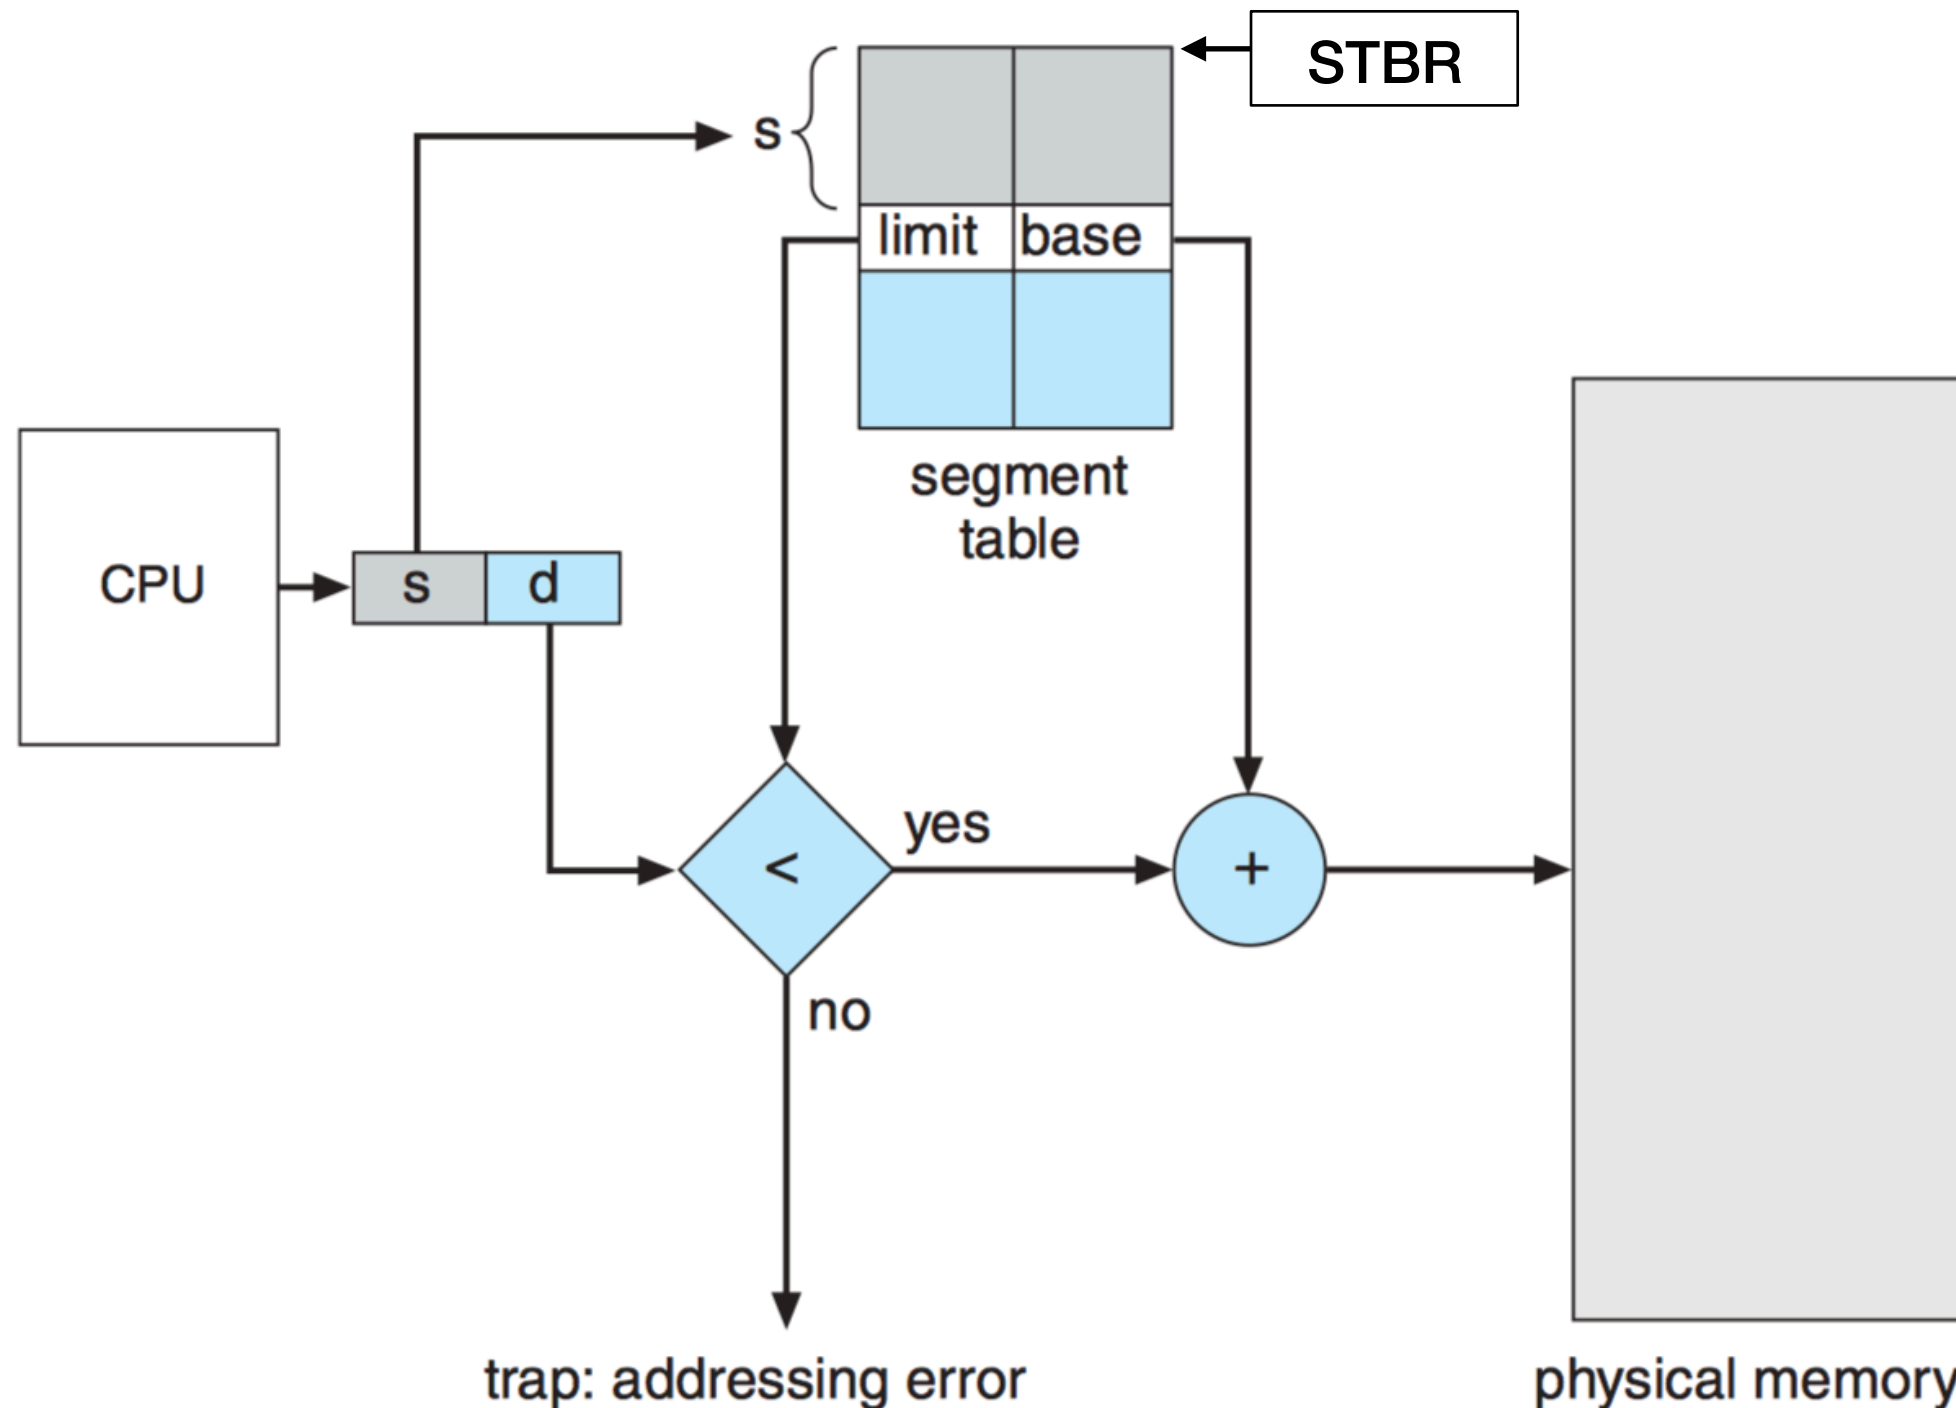
- process generates contiguous virtual addresses from 0 to MAX_{proc}
- OS lays the process down on **arbitrary size segments** and hardware translates virtual addresses to physical addresses in memory
 - using **a segment table** that keeps track of the segment locations in memory
- the virtual address is identified by a pair (s, d)
 - s is the segment number
 - d is the offset within the segment
- enable sharing code and data between processes (for example when linking libraries)



How to find addresses when segments are not allocated contiguously in memory?

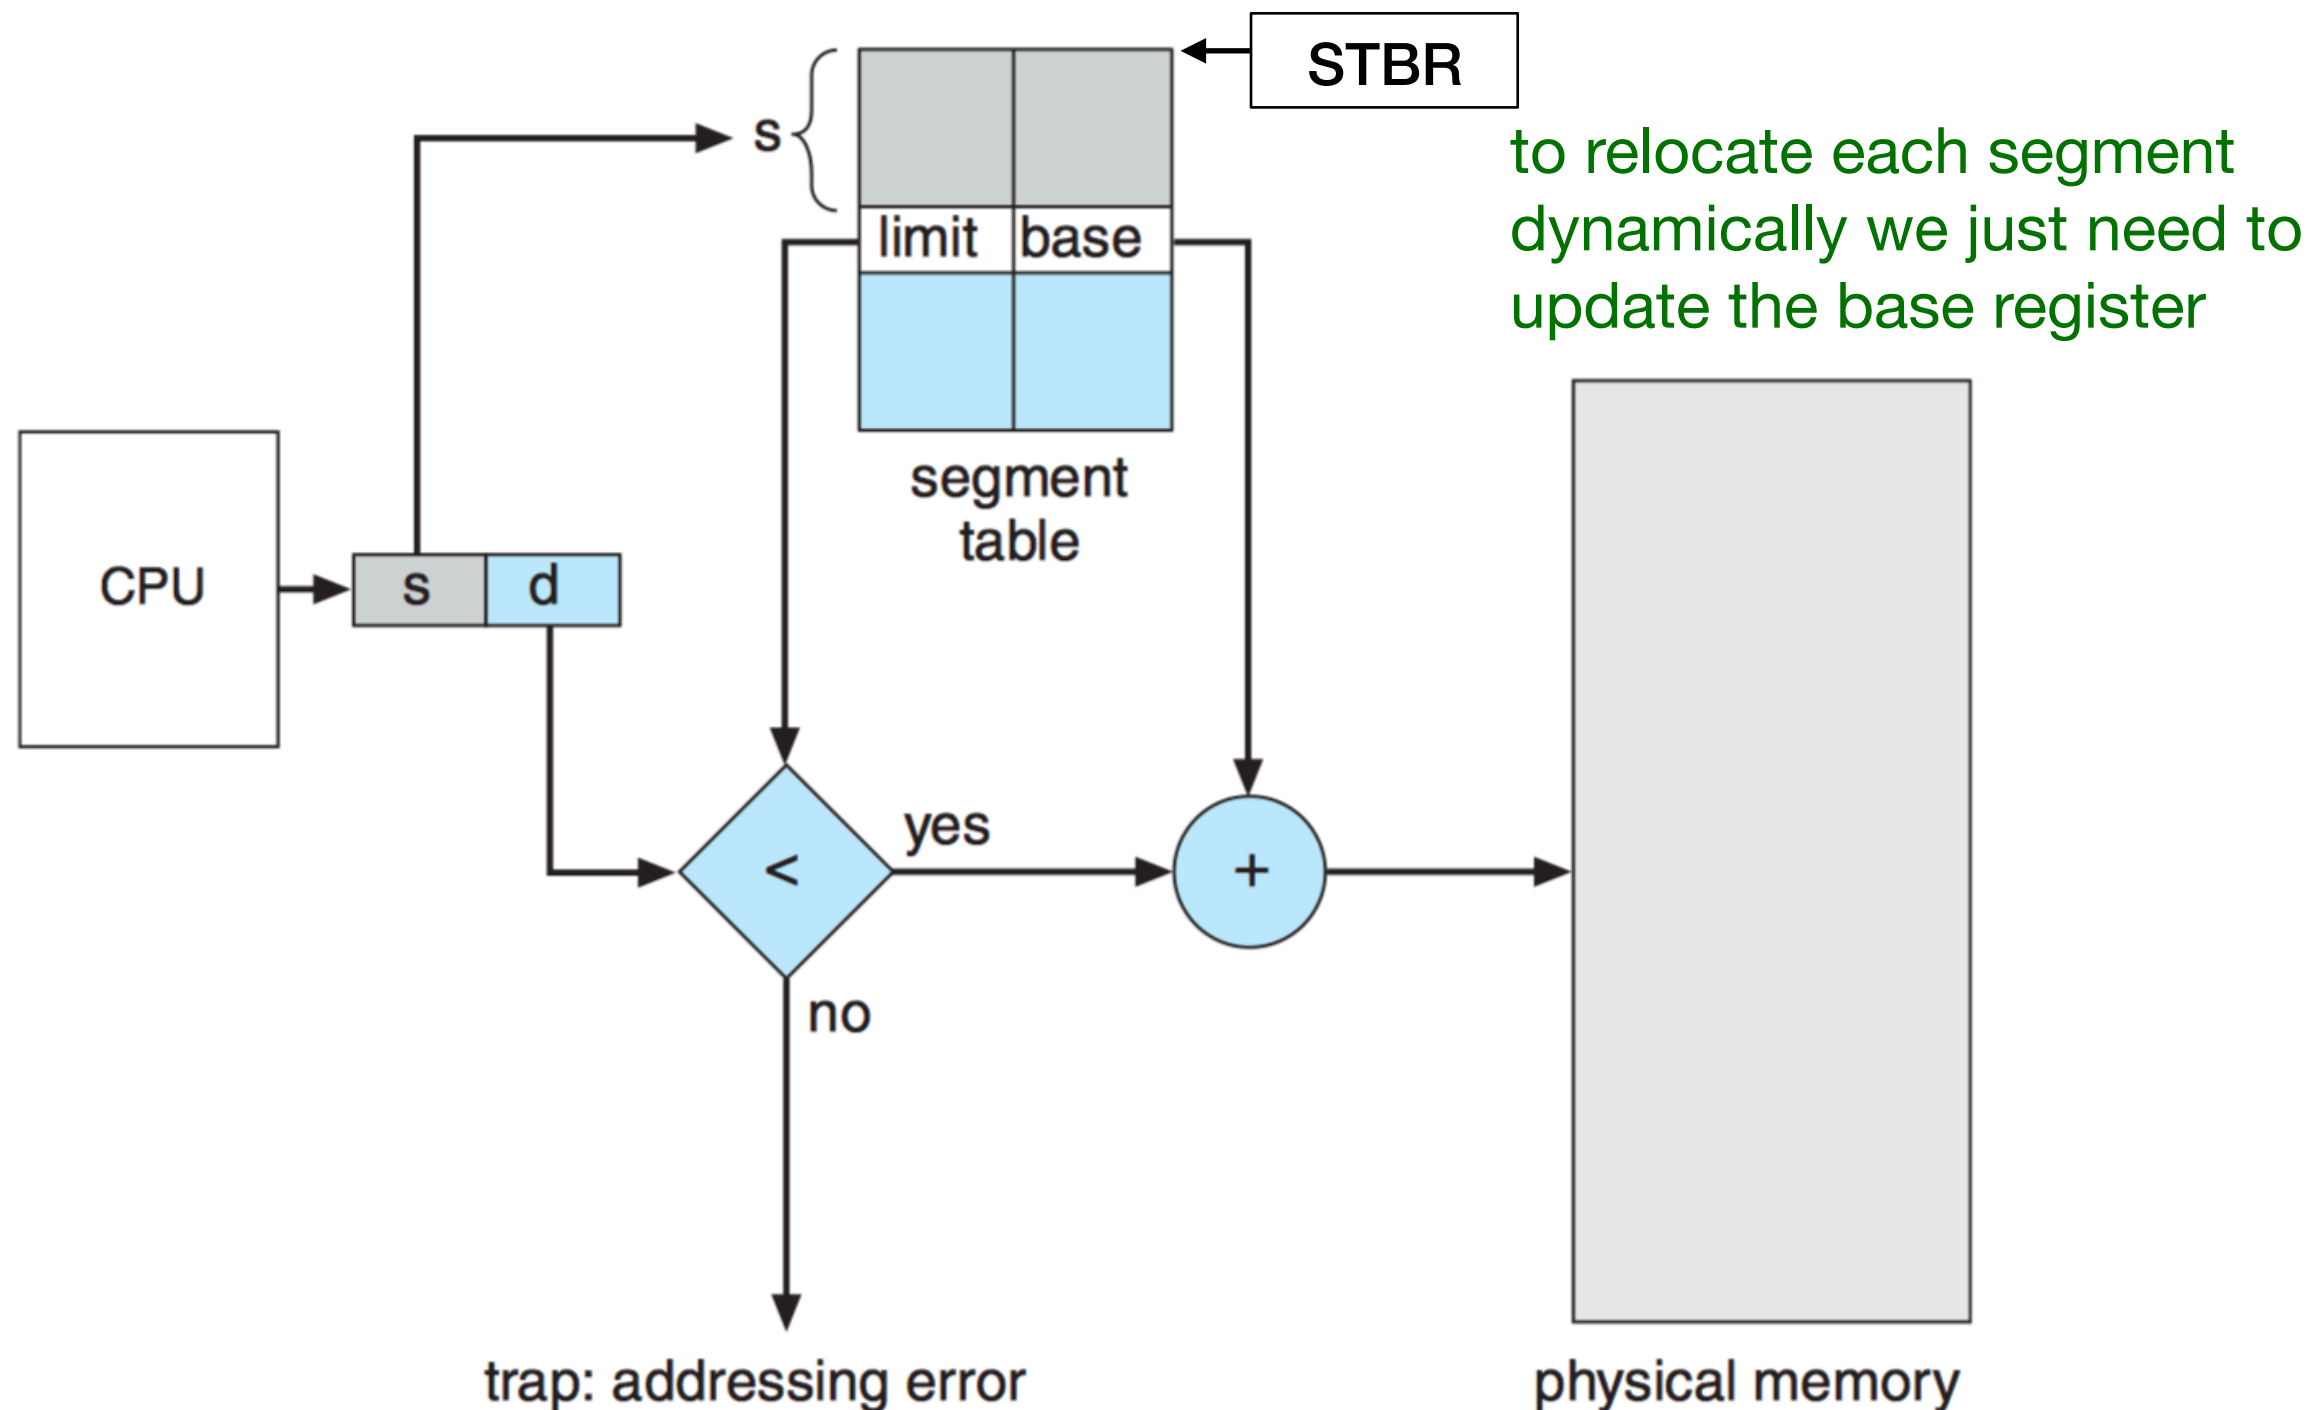
add a segment table containing base & limit register values

How to find addresses when segments are not allocated contiguously in memory?



add a segment table containing base & limit register values


How to find addresses when segments are not allocated contiguously in memory?



add a segment table containing base & limit register values


Hardware support for segmentation

programmer's view



	limit	base
0	1000	1400
1	400	6300
2	400	4300
3	1100	3200
4	1000	4700

segment table



the compiler generates references that identify the segment (e.g., code, global variables, stack, heap) and the offset in the segment, e.g., a code segment with offset = 399

Is segmentation enough?

- segmentation leads to poor memory utilization
 - we might not use much of a large segment, but we must keep the whole thing in memory
 - suffers from external fragmentation
 - allocation/deallocation of arbitrary size segments is complex
- how can we improve memory management? **using fixed-size allocation units**

Motivation for paging

- 90/10 rule
 - processes spend 90% of their time accessing 10% of their space in memory
 - so keep only those parts in memory that are being used

Motivation for paging

- 90/10 rule
 - processes spend 90% of their time accessing 10% of their space in memory
 - so keep only those parts in memory that are being used
- logical memory of a process is contiguous, but pages need not be allocated contiguously in memory

Motivation for paging

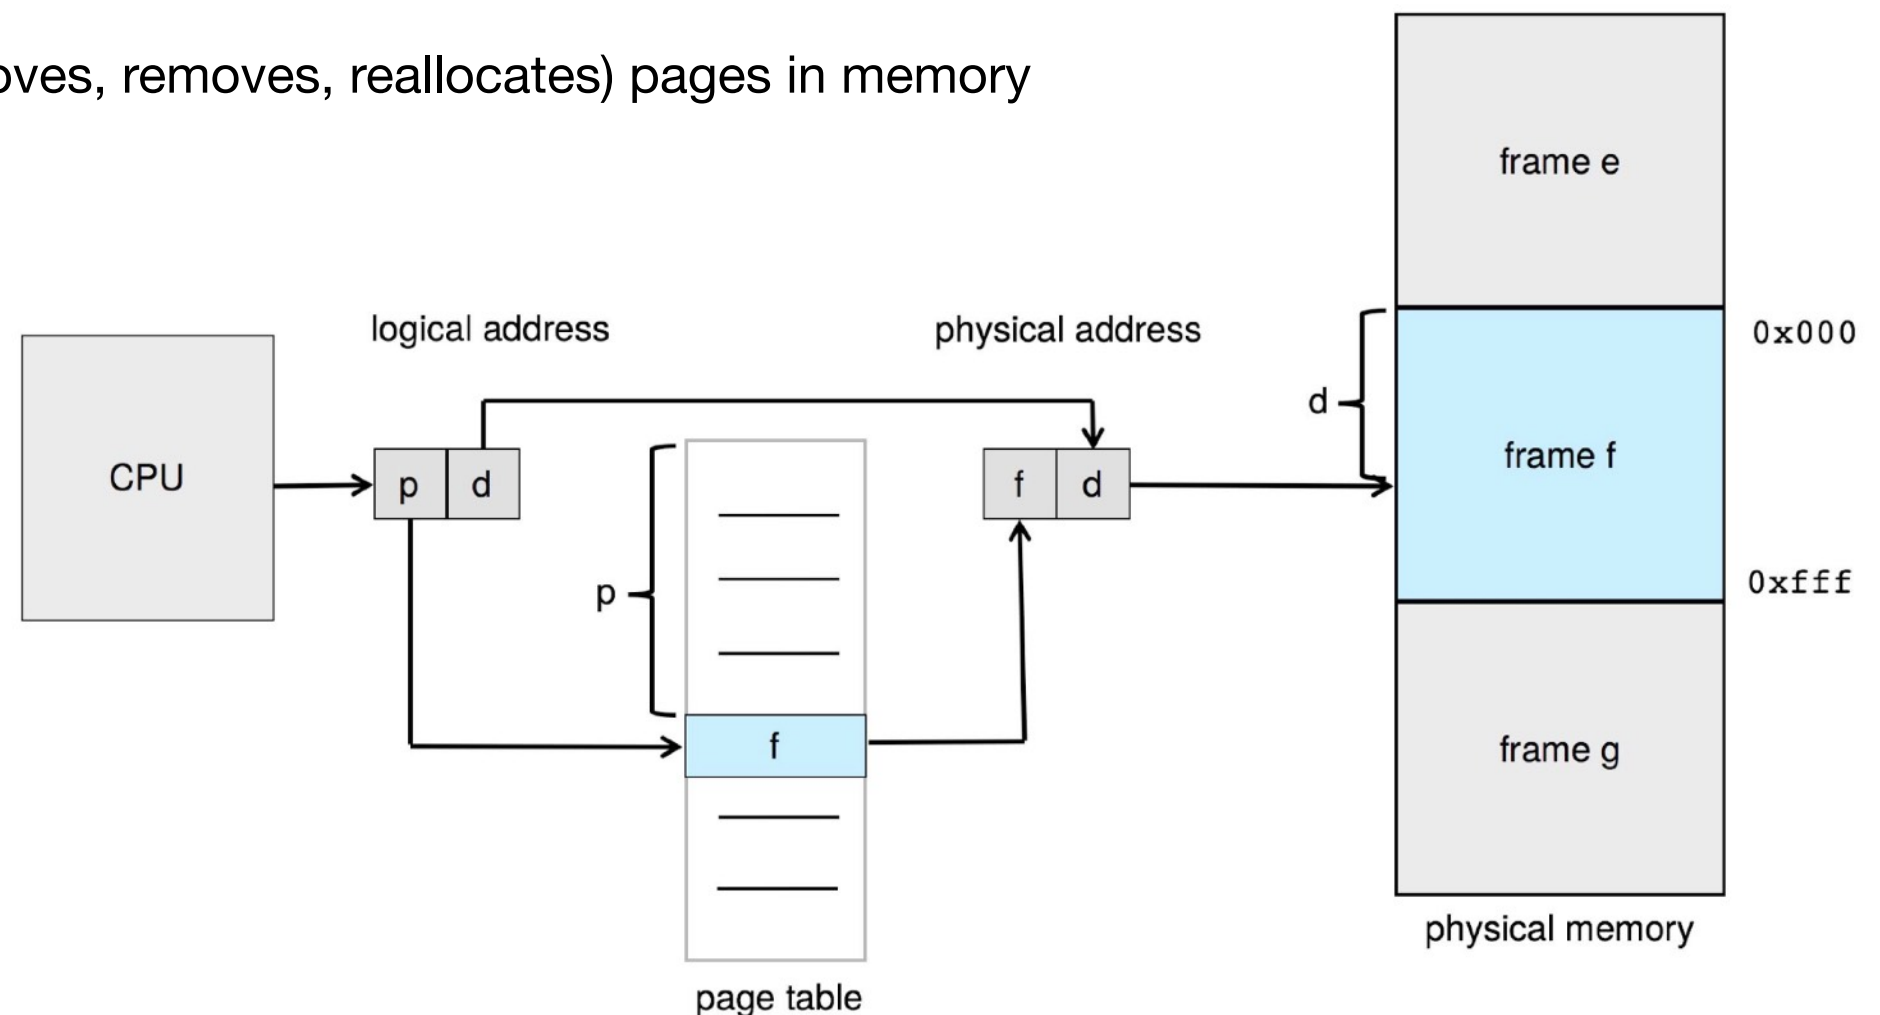
- 90/10 rule
 - processes spend 90% of their time accessing 10% of their space in memory
 - so keep only those parts in memory that are being used
- logical memory of a process is contiguous, but pages need not be allocated contiguously in memory
- paging eliminates external fragmentation by dividing memory into **fixed-size** pages
 - the `getpagesize()` system call returns the page size (e.g., 4096 bytes)
 - this command in linux returns the page size as well: `getconf PAGESIZE`

Motivation for paging

- 90/10 rule
 - processes spend 90% of their time accessing 10% of their space in memory
 - so keep only those parts in memory that are being used
- logical memory of a process is contiguous, but pages need not be allocated contiguously in memory
- paging eliminates external fragmentation by dividing memory into **fixed-size** pages
 - the `getpagesize()` system call returns the page size (e.g., 4096 bytes)
 - this command in linux returns the page size as well: `getconf PAGESIZE`
- paging causes internal fragmentation though
 - on average half a page is lost per process

How to find addresses when pages are not allocated contiguously in memory?

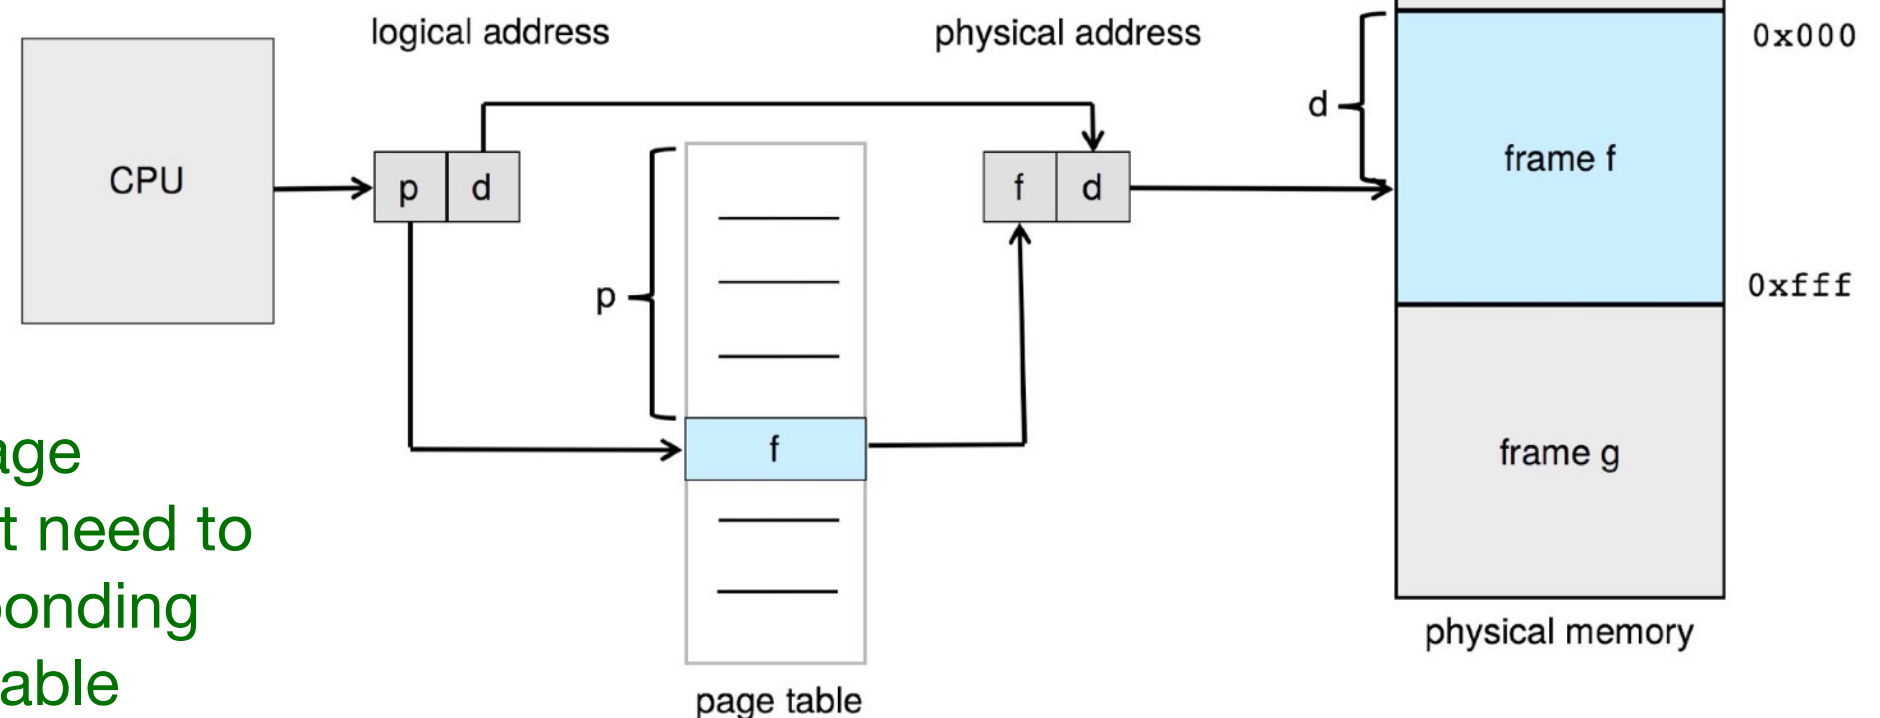
- OS divides processes into fixed sized pages
- selectively allocates pages to **page frames** in memory
 - the virtual address contains the page number (p) and the page offset (d)
 - p is used as an index into the **page table**, allowing to extract the frame number (f)
 - the page table keeps track of the page frame in memory in which the page is located
- keeps track of free frames in a **frame table** which has an entry for each physical page frame
- OS manages (moves, removes, reallocates) pages in memory



How to find addresses when pages are not allocated contiguously in memory?


- OS divides processes into fixed sized pages
- selectively allocates pages to **page frames** in memory
 - the virtual address contains the page number (p) and the page offset (d)
 - p is used as an index into the **page table**, allowing to extract the frame number (f)
 - the page table keeps track of the page frame in memory in which the page is located
- keeps track of free frames in a **frame table** which has an entry for each physical page frame
- OS manages (moves, removes, reallocates) pages in memory

to relocate each page
dynamically we just need to
update the corresponding
frame in the page table




Paging details

- page sizes (and frame sizes) are typically a power of 2 between 512 bytes and 8192 bytes per page
 - makes the translation of virtual addresses into physical addresses easier
- for example, given a memory of size 2^m bytes and a page of size 2^n then
 - the high order $m-n$ bits of a virtual address select the page
 - the low order n bits select the offset in the page



Example

mem. size: $65,536 = 2^{16}$
page size: $8,196 = 2^{13}$



logical
memory

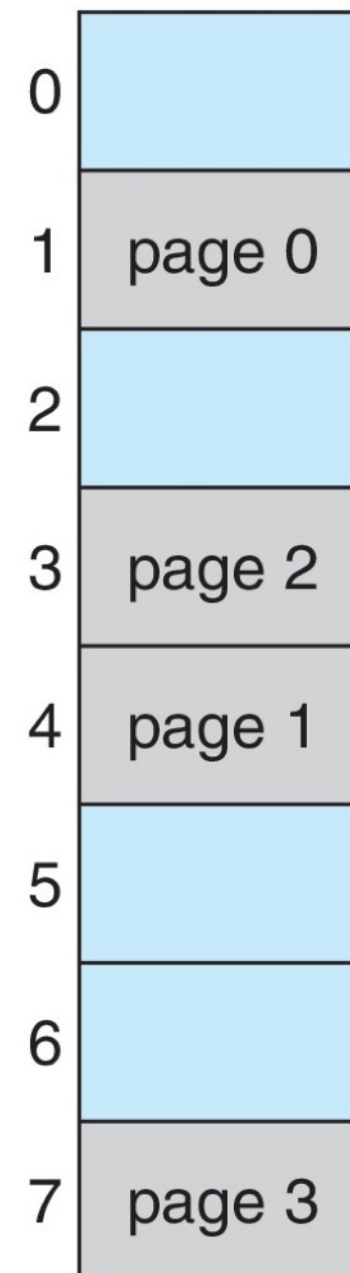
a process is
comprised of
fixed size pages

0	1
1	4
2	3
3	7

page table

(page #, frame #)


page table is like a set of
relocation registers, one
for each frame



physical
memory

Example

mem. size: $65,536 = 2^{16}$
page size: $8,196 = 2^{13}$



logical
memory

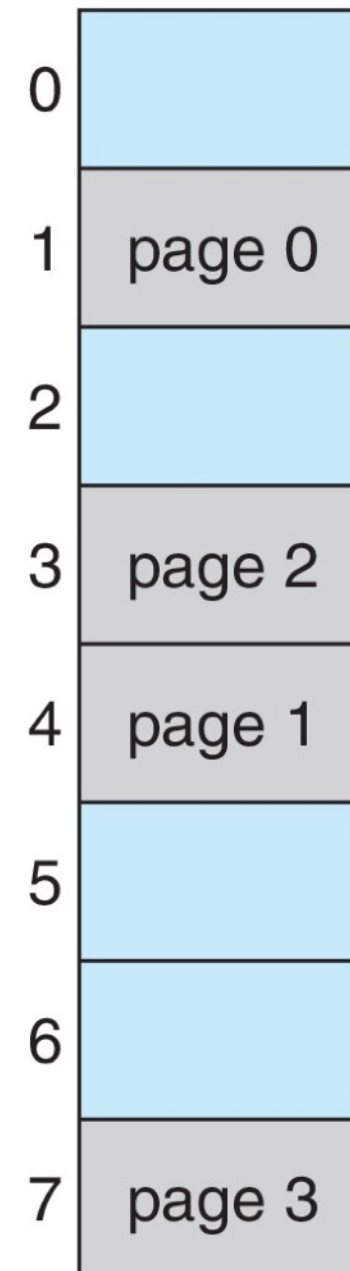
a process is
comprised of
fixed size pages

0	1
1	4
2	3
3	7

page table

(page #, frame #)

page table is like a set of
relocation registers, one
for each frame



physical
memory

How many bits for an address (assuming we can
address 1 byte increments)? **16 in this example**

Example of address translation

- how many frames do we have? 8 frames
- what is the page size? 2^{13} bytes
- assuming we can address 1 byte increments, how many bits for an address?
16bits (3 bits for page table, 13 bits for offset)
- assuming we can address 1 word (4 bytes) increments, how many bits for an address?
14bits (3 bits for page table, 11 bits for offset)

Calculating internal fragmentation

- recall that
 - worst-case internal fragmentation: 1 frame – 1 byte
 - average-case internal fragmentation: 1/2 frame size
- let the page size be 2,048 bytes and the process size be 72,766 bytes (i.e., 35 pages + 1,086 bytes)
 - internal fragmentation of $2,048 - 1,086 = 962$ bytes

Calculating internal fragmentation

- recall that
 - worst-case internal fragmentation: 1 frame – 1 byte
 - average-case internal fragmentation: 1/2 frame size
- let the page size be 2,048 bytes and the process size be 72,766 bytes (i.e., 35 pages + 1,086 bytes)
 - internal fragmentation of $2,048 - 1,086 = 962$ bytes
- can we say small frame sizes are desirable?
 - No, each page table entry takes memory to track

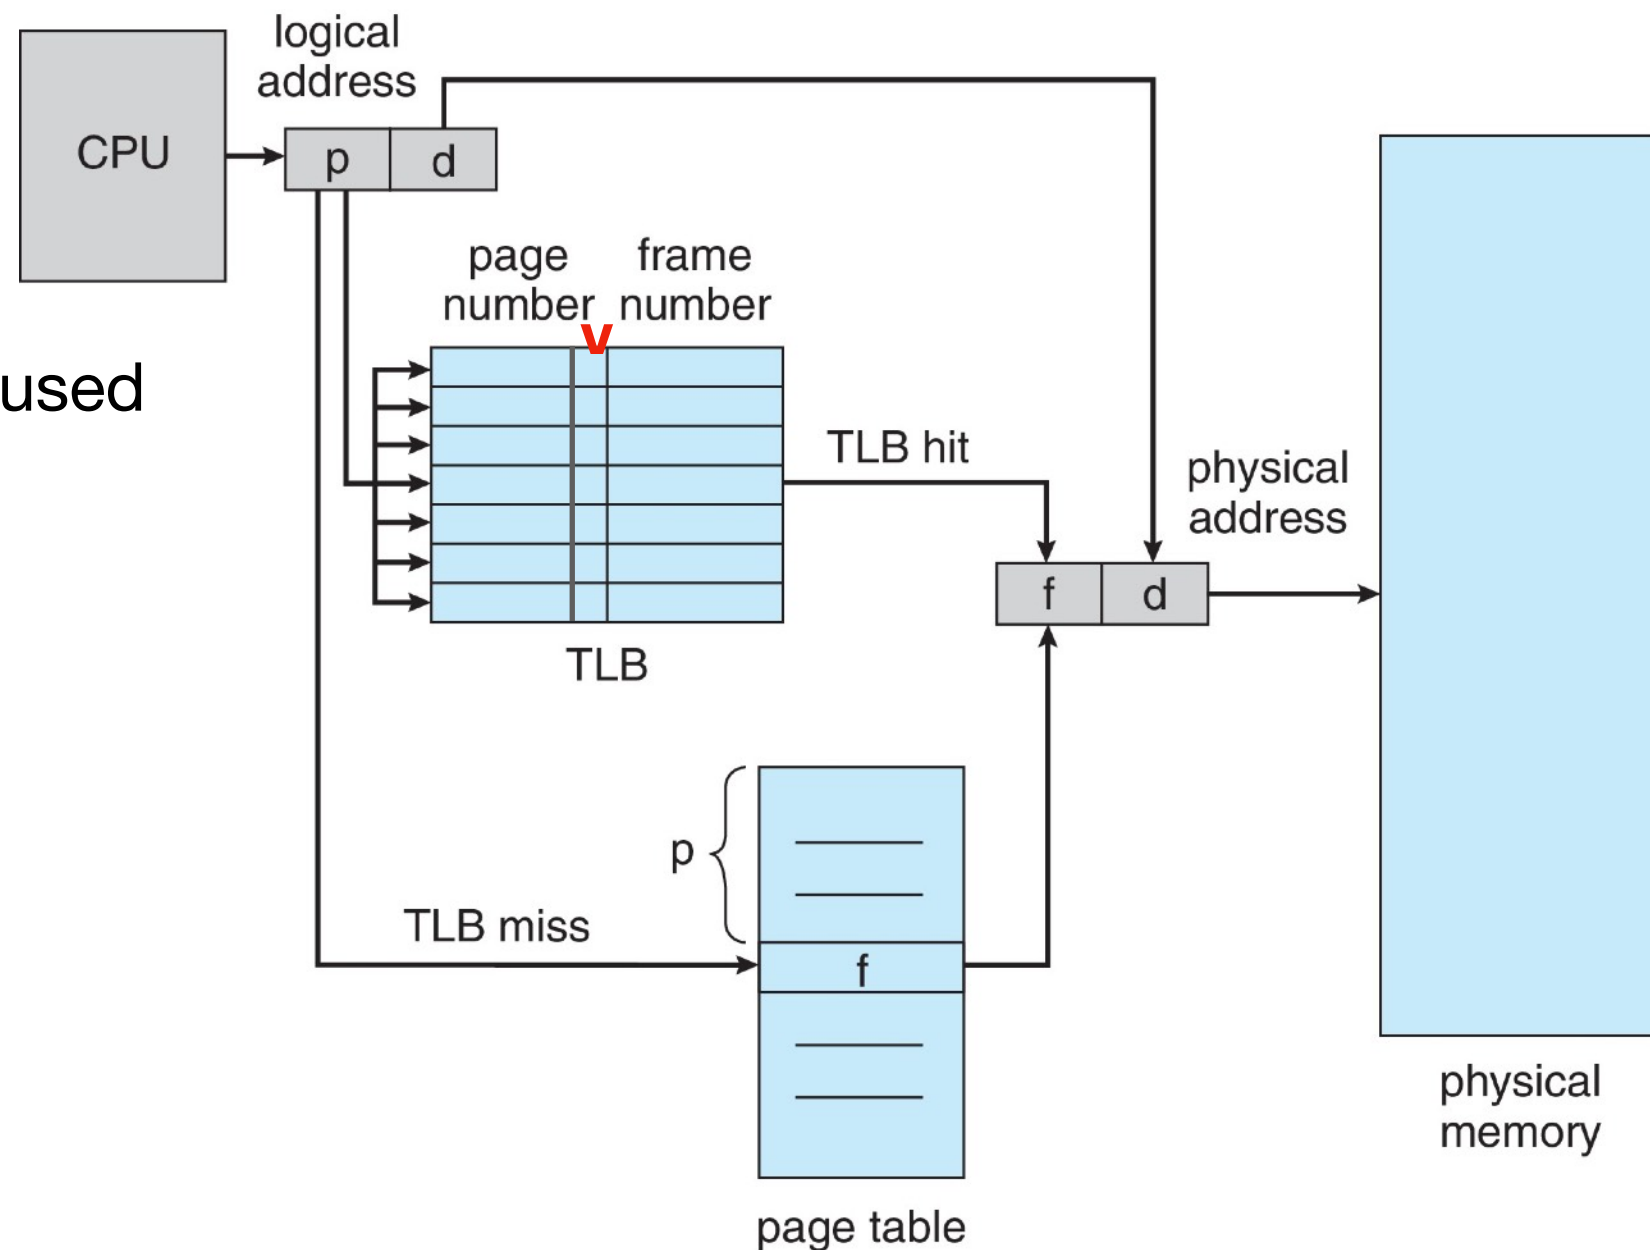
Where to store the page table?

- put it in memory?
 - page-table base register (PTBR) points to the page table
 - page-table length register (PTLR) indicates size of the page table
 - every data/instruction access requires two memory accesses: one for the page table and one for the data/instruction
- **translation look-aside buffer** (TLB) is a fast-lookup hardware cache in MMU (fully associative memory) that stores page numbers (key) and the frame (value) in which they are stored
 - if memory accesses have **locality**, address translation has locality too
 - typical TLB sizes range from 64 to 1,024 entries (TLB is small)
 - TLB is typically on chip, access time of a few nanoseconds instead of several hundred of nanoseconds for main memory

Paging with TLB

- on a TLB miss, value is loaded into the TLB for faster access next time
- replacement policies must be considered




v: valid/invalid bit used for protection



Valid/invalid bit in TLB

v: valid/invalid bit used for protection

- “valid” indicates that the associated page is in the process logical address space, and is thus a legal page
- “invalid” indicates that the page is not in the process logical address space




Effective memory access cost with TLB

- the effective access time (EAT) if the page table is in memory?
$$\text{EAT} = 2 \times C_{ma}$$
- the effective access time with a TLB?
$$\text{EAT} = \rho_{hit}(C_{ma} + C_{TLB}) + (1 - \rho_{hit})(2 \times C_{ma} + C_{TLB})$$
 - where ρ_{hit} is the hit ratio (i.e., percentage of times that a page number is found in the TLB)
- a large TLB improves hit ratio, decreases average memory cost
- example: let $C_{ma} = 10\text{ns}$, $C_{TLB} = 0\text{ns}$
 - if $\rho_{hit}=80\%$ then $\text{EAT}=12\text{ns}$ (20% slowdown in access time)
 - if $\rho_{hit}=99\%$ then $\text{EAT}=10.1\text{ns}$ (1% slowdown in access time)


How to allocate memory when starting a process?

suppose the process needs k pages

- if k page frames are free, then allocate these frames to pages; otherwise, free frames that are no longer needed and allocate them to the process
- the OS puts each page in a frame and then puts the frame number in the corresponding entry in the page table
- OS marks all TLB entries as invalid (flushes the TLB) and starts the process
- as process executes, OS loads TLB entries as each page is accessed, replacing an existing entry if the TLB is full



before allocation



after allocation

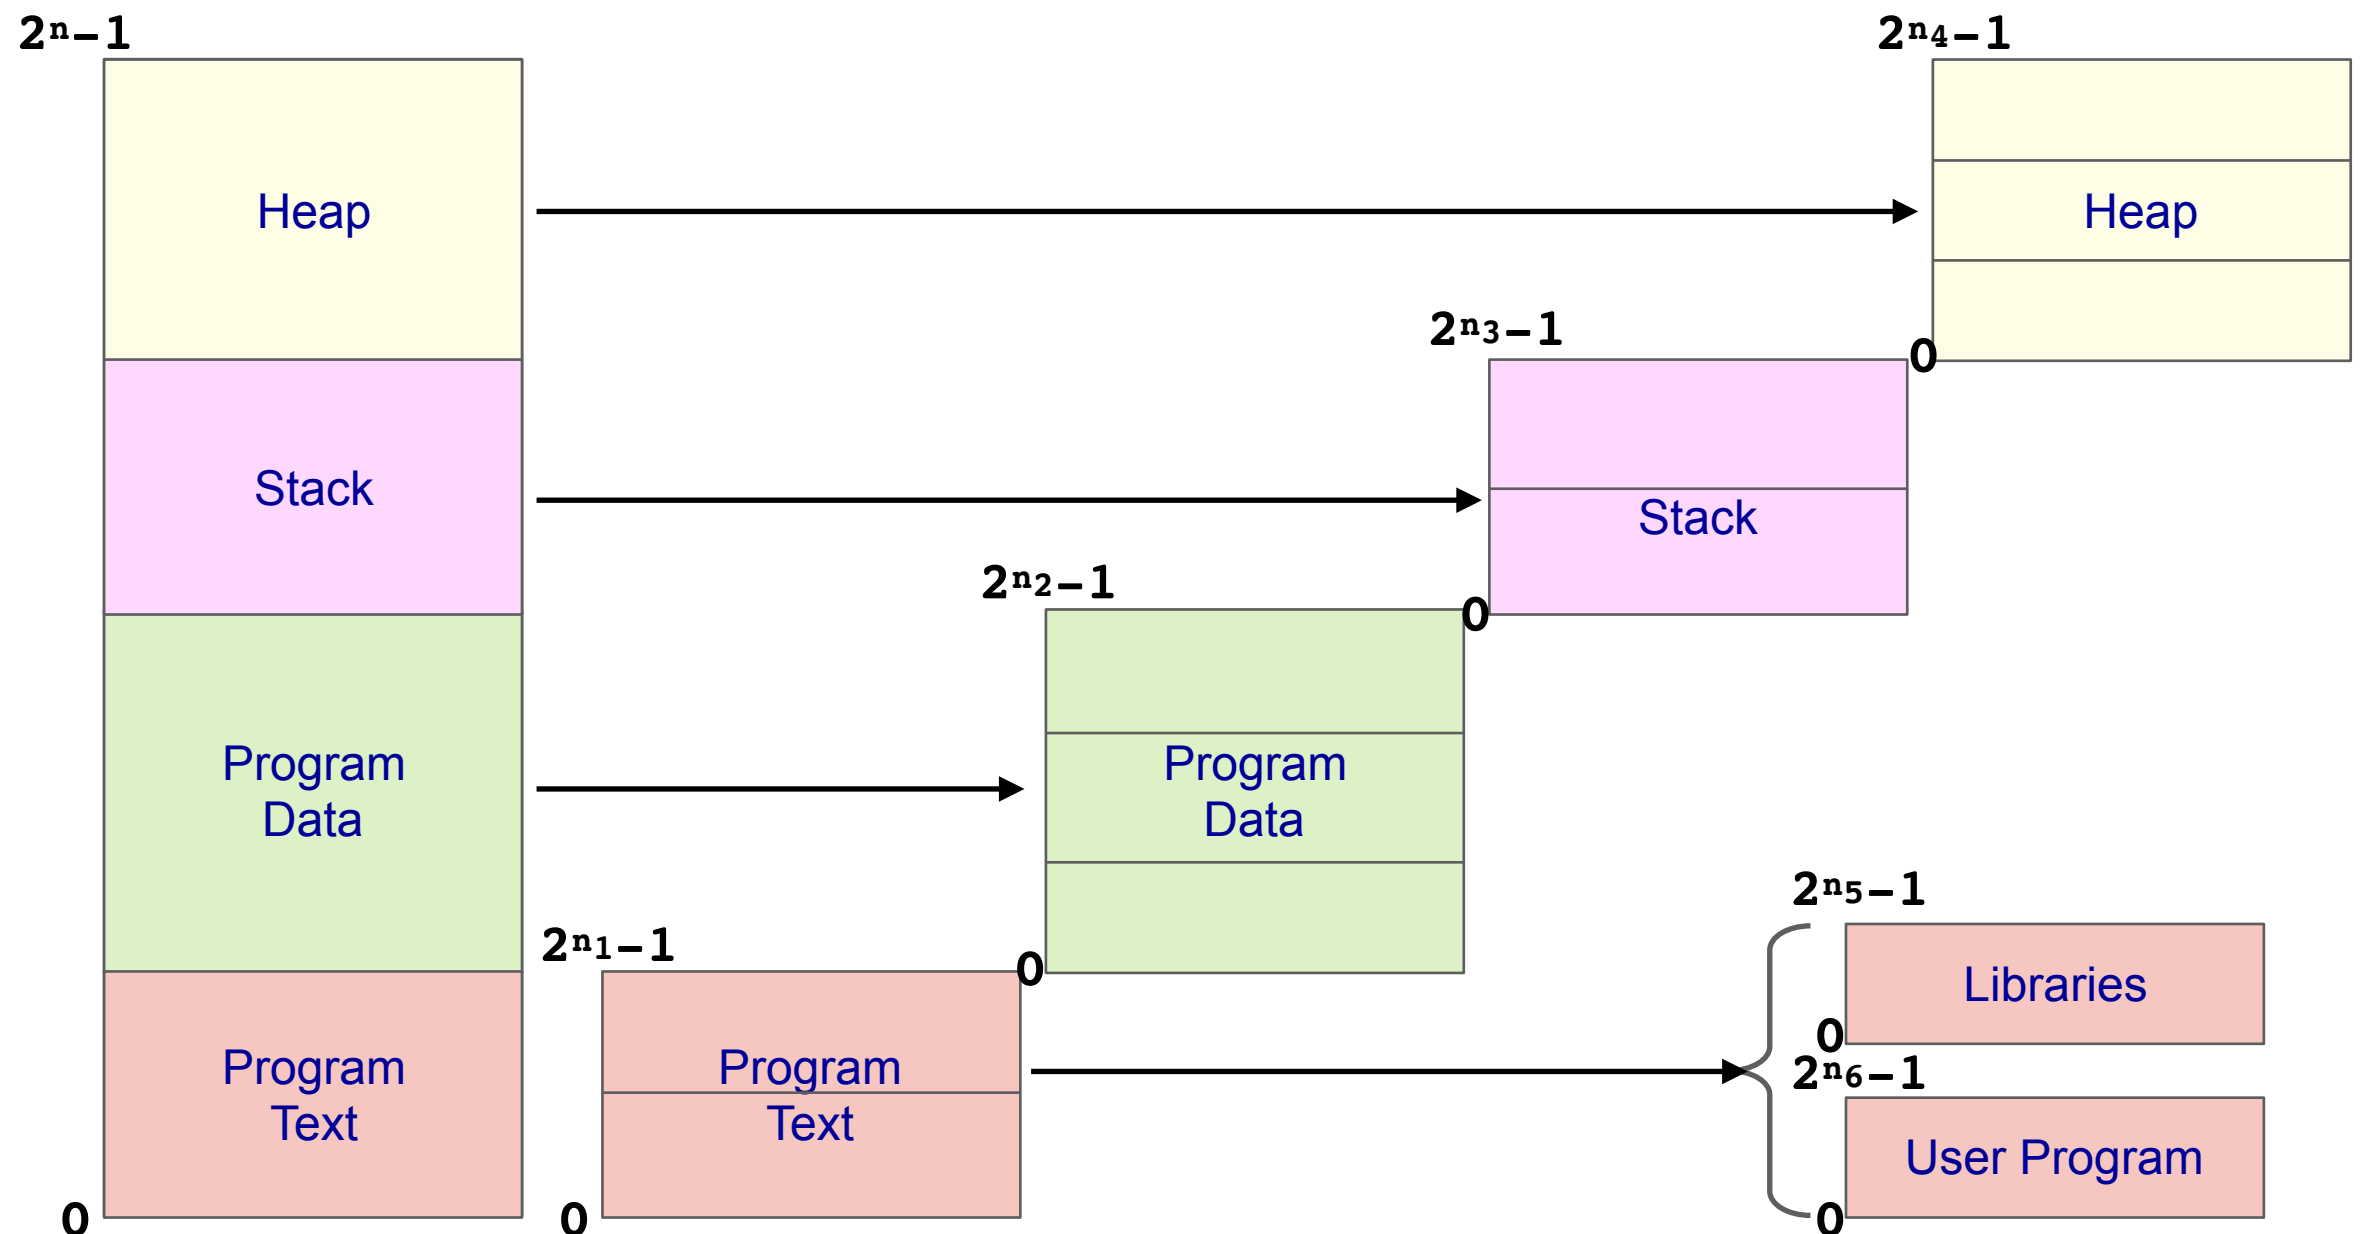
What happens on a context switch?

- OS saves the PTBR value to the PCB
- OS saves the process's page table in its PCB (i.e., copying the TLB to the PCB)
- OS marks all TLB entries as invalid (i.e., flushes the TLB)
- OS restores the PTBR value of the new process
- OS restores the TLB if it was saved

Sharing

- increases the level of multiprogramming by sharing code and data between different processes
 - shared code must be reentrant, that means the processes that are using it cannot change it
 - OS keeps track of available reentrant code in memory and reuses them if a new process requests the same program
- allows running many processes given a limited amount of memory
- sharable units: segments or pages
- paging allows sharing of memory across processes, since memory used by a process no longer needs to be contiguous

How to implement sharing?



Segmentation versus paging

- paging is a big improvement over segmentation
 - eliminates the problem of external fragmentation and therefore the need for compaction
 - enables processes to run when they are only partially loaded in main memory

Segmentation versus paging

- paging is a big improvement over segmentation
 - eliminates the problem of external fragmentation and therefore the need for compaction
 - enables processes to run when they are only partially loaded in main memory
- however, paging has its cost:
 - causes internal fragmentation
 - translating from a virtual address to a physical address is more time-consuming
 - segment table tends to be much smaller than page table
 - paging requires hardware support in the form of a TLB to be efficient enough
 - paging requires more complex OS to maintain the page table