

Operating System Concepts

Lecture 25: Segmented Paging

Omid Ardakanian
oardakan@ualberta.ca
University of Alberta

MWF 12:00-12:50 VVC 2 215

Today's class

- Memory-management strategies
 - contiguous memory allocation
 - non-contiguous memory allocation
 - segmentation
 - paging
 - segmented paging

Putting it all together

- base & bounds registers
 - OS loads the registers
 - simple and fast, but does not support sharing, incremental increase of heap/stack, etc.
 - complex memory allocation

Putting it all together

- base & bounds registers
 - OS loads the registers
 - simple and fast, but does not support sharing, incremental increase of heap/stack, etc.
 - complex memory allocation
- segment table
 - top few bits encode the segment number. Each segment has base and bounds
 - supports sharing
 - complex memory allocation

Putting it all together

- base & bounds registers
 - OS loads the registers
 - simple and fast, but does not support sharing, incremental increase of heap/stack, etc.
 - complex memory allocation
- segment table
 - top few bits encode the segment number. Each segment has base and bounds
 - supports sharing
 - complex memory allocation
- page table
 - **each process has its own page table**, but TLB contains entries for several different processes simultaneously (**address-space identifiers** are used to determine which entry belongs to which process)
 - memory allocation is done in small page sizes (4K – 16K)
 - pages are contiguous in virtual address space, but they are arbitrarily located in physical memory
 - avoids external fragmentation and the need for compaction
 - paging supports sharing


Efficient forking in UNIX

- UNIX fork makes a complete copy of a process
- segmentation and paging allow for a more efficient implementation of the `fork()` system call
 - copy segment/page table into the child
 - mark parent and child segments/pages read-only (they share the same pages)
 - start child process and return to the parent
- most of the time child calls `exec` right after fork
 - without writing to read-only segments/pages
 - so duplicating the pages belonging to the parent is unnecessary most of the time

Efficient forking in UNIX


- UNIX fork makes a complete copy of a process
- segmentation and paging allow for a more efficient implementation of the `fork()` system call
 - copy segment/page table into the child
 - mark parent and child segments/pages read-only (they share the same pages)
 - start child process and return to the parent
- most of the time child calls `exec` right after fork
 - without writing to read-only segments/pages
 - so duplicating the pages belonging to the parent is unnecessary most of the time
- but if the child or the parent writes to a segment/page (e.g., stack, heap)
 - trap into the kernel
 - make a copy of the segment/page (this technique is known as **copy-on-write**)
 - mark both segments/pages as writable
 - resume execution

Efficient forking in UNIX

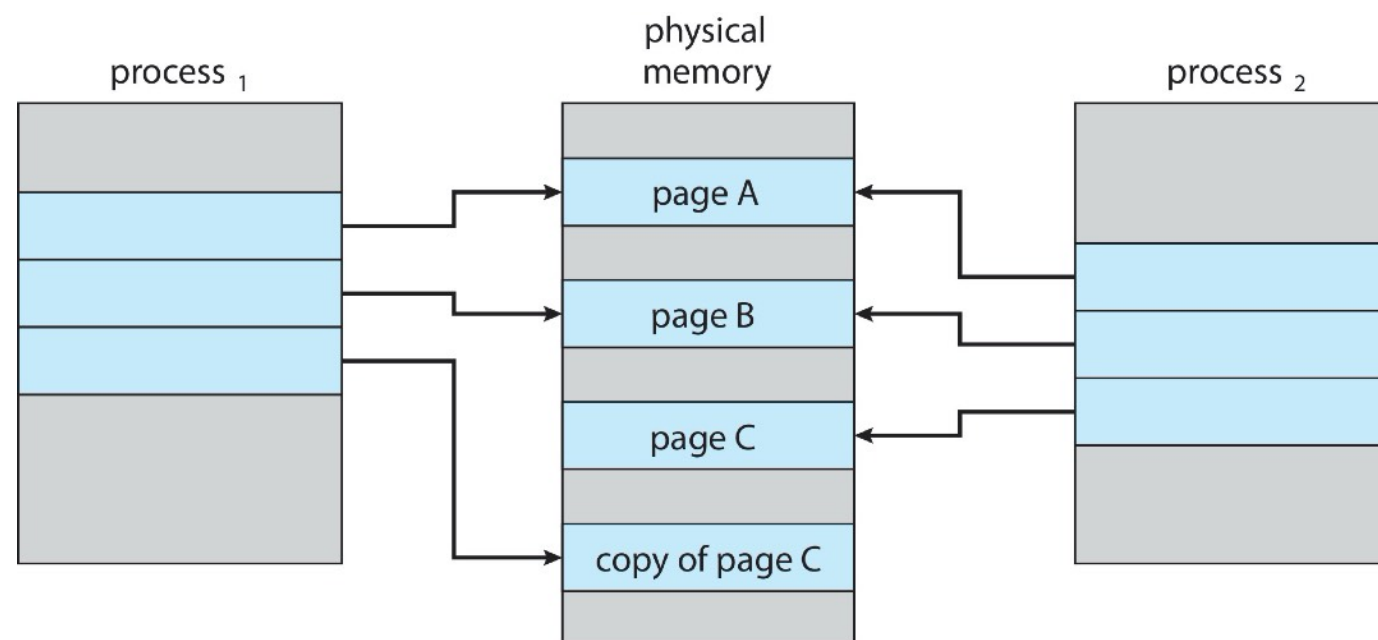


before process 1 modifies Page C

Efficient forking in UNIX



before process 1 modifies Page C



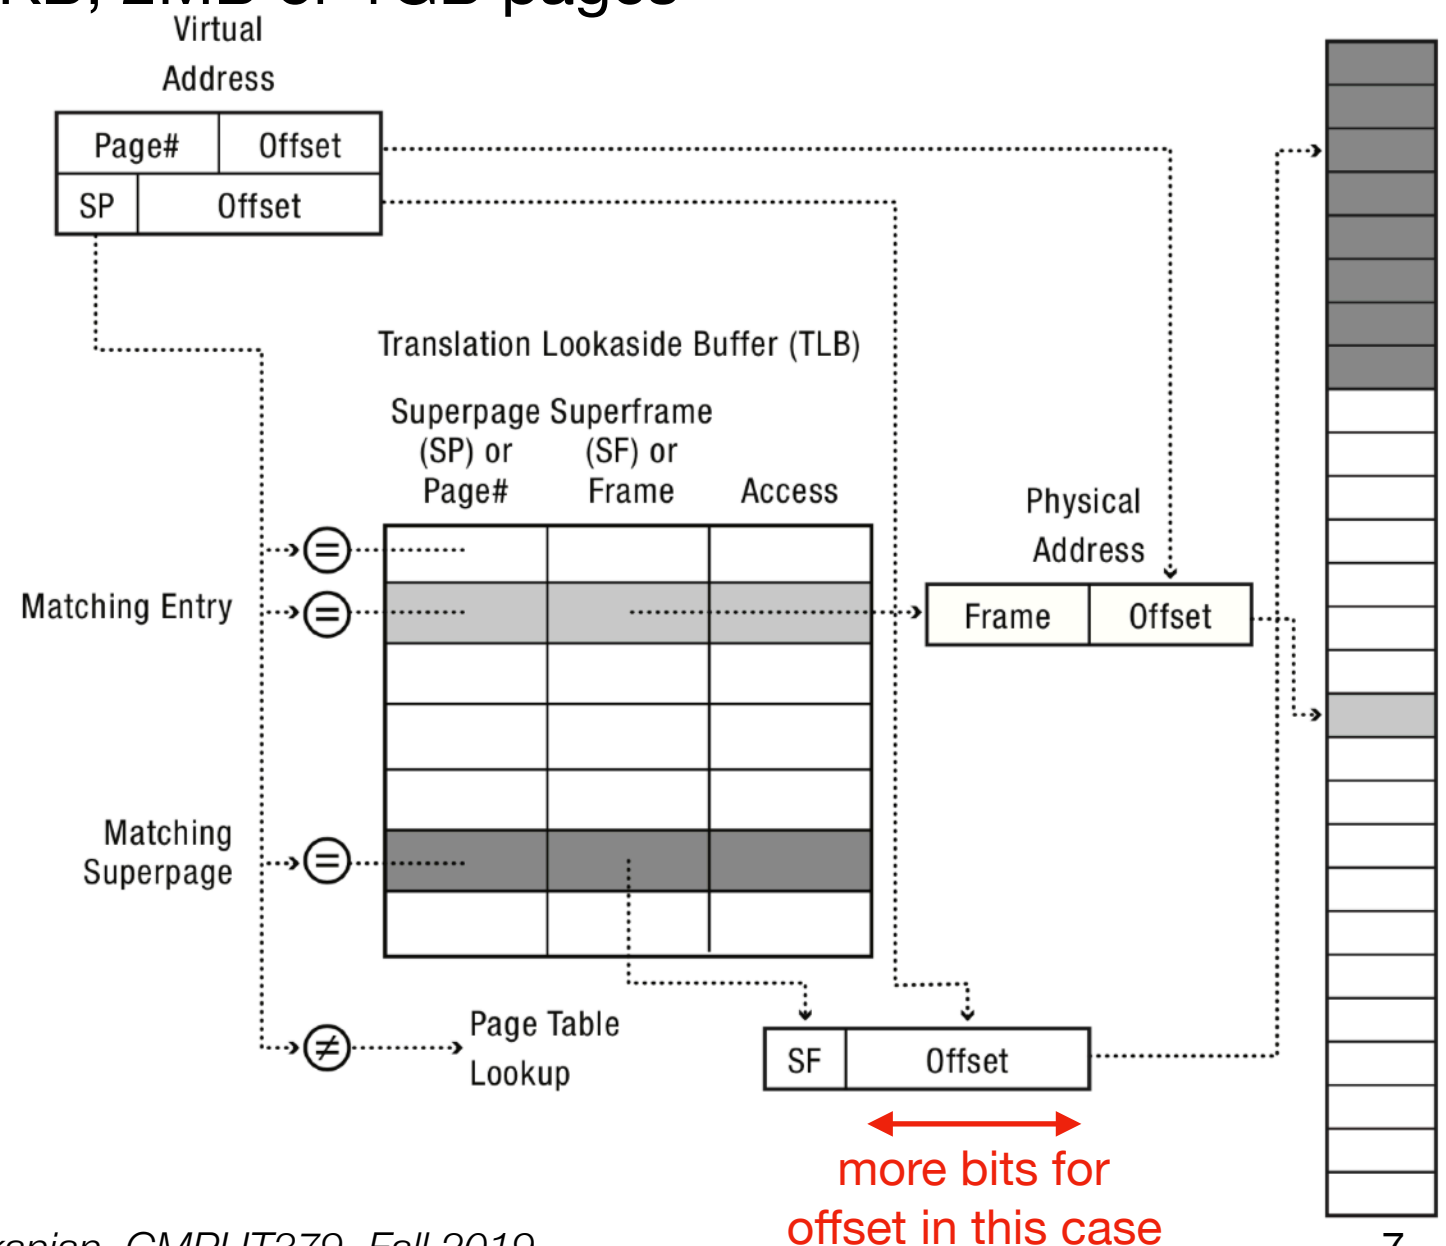
after process 1 modifies Page C
(a copy of the shared page is created)

Growing heap or stack

- when a process uses memory beyond the end of stack or dynamically allocates memory beyond the end of heap
 - segmentation fault occurs
 - OS allocates some memory
 - OS zeros the newly allocated memory (this is known as **zero-on-reference**) to avoid accidentally leaking information!
 - OS modifies the segment table
 - the process is resumed

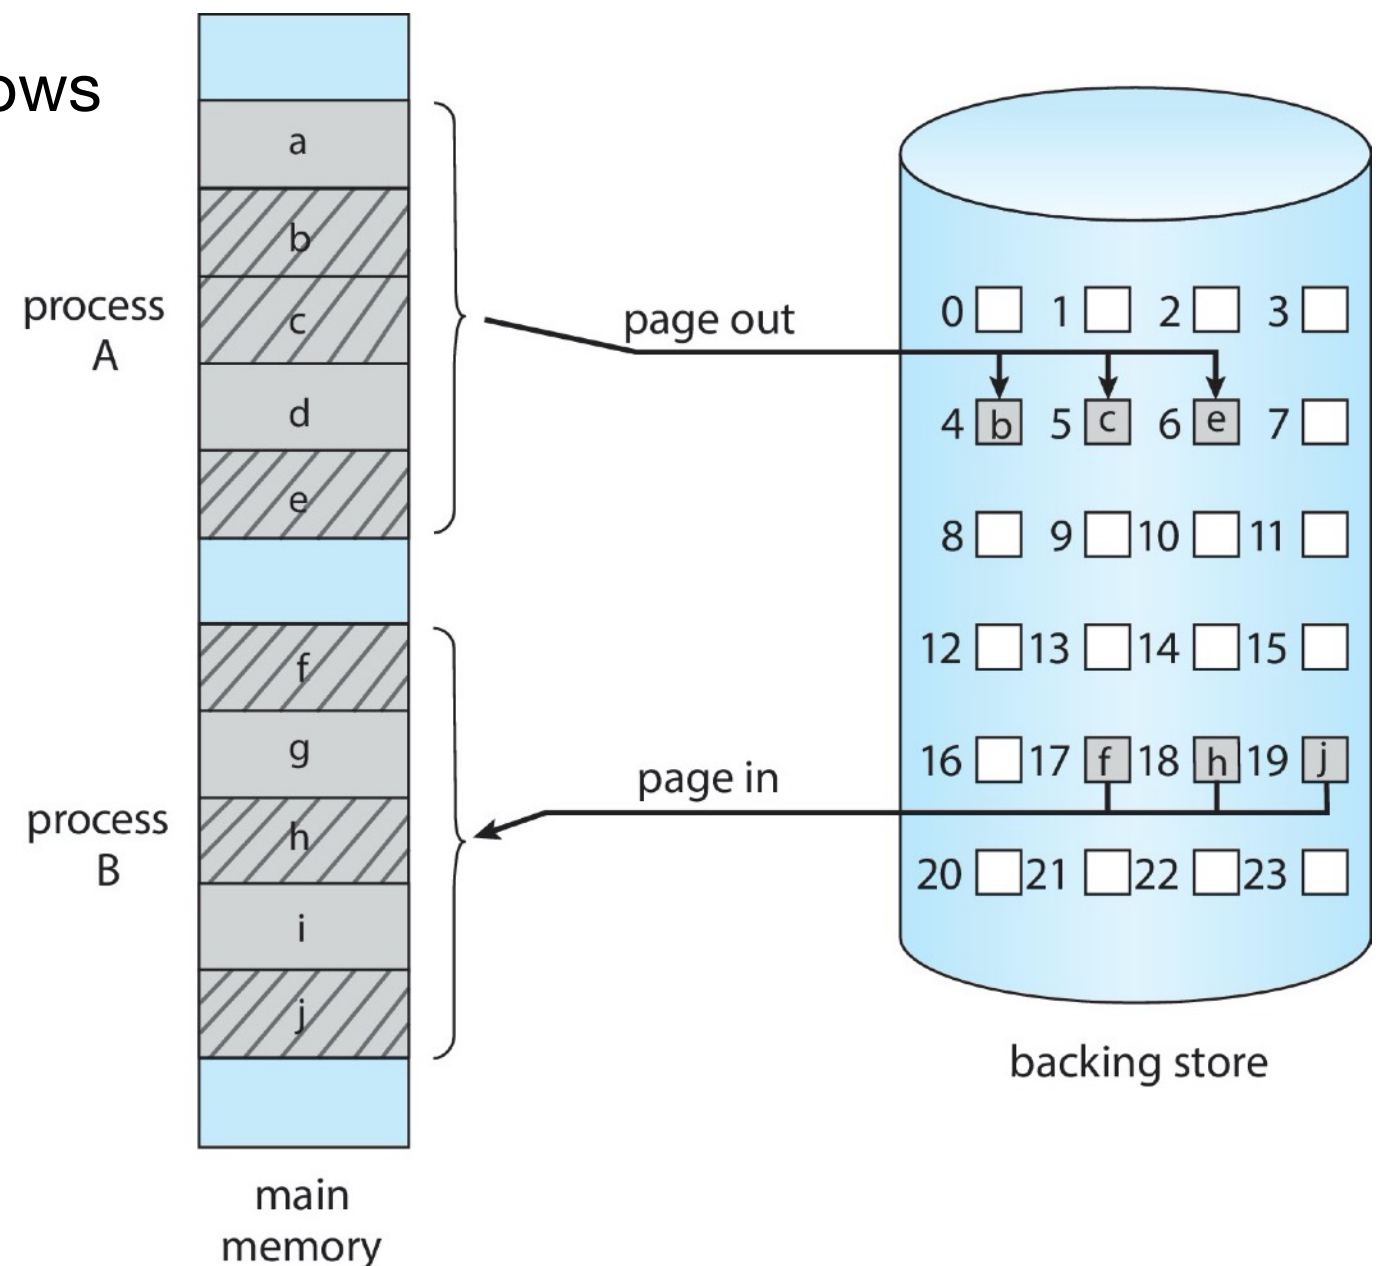
Superpages

- on many systems, the first part of a TLB entry can be
 - a page number or a superpage (a set of contiguous pages in one page table) number
- x86 TLB entries can point to 4KB, 2MB or 1GB pages



Swapping with paging

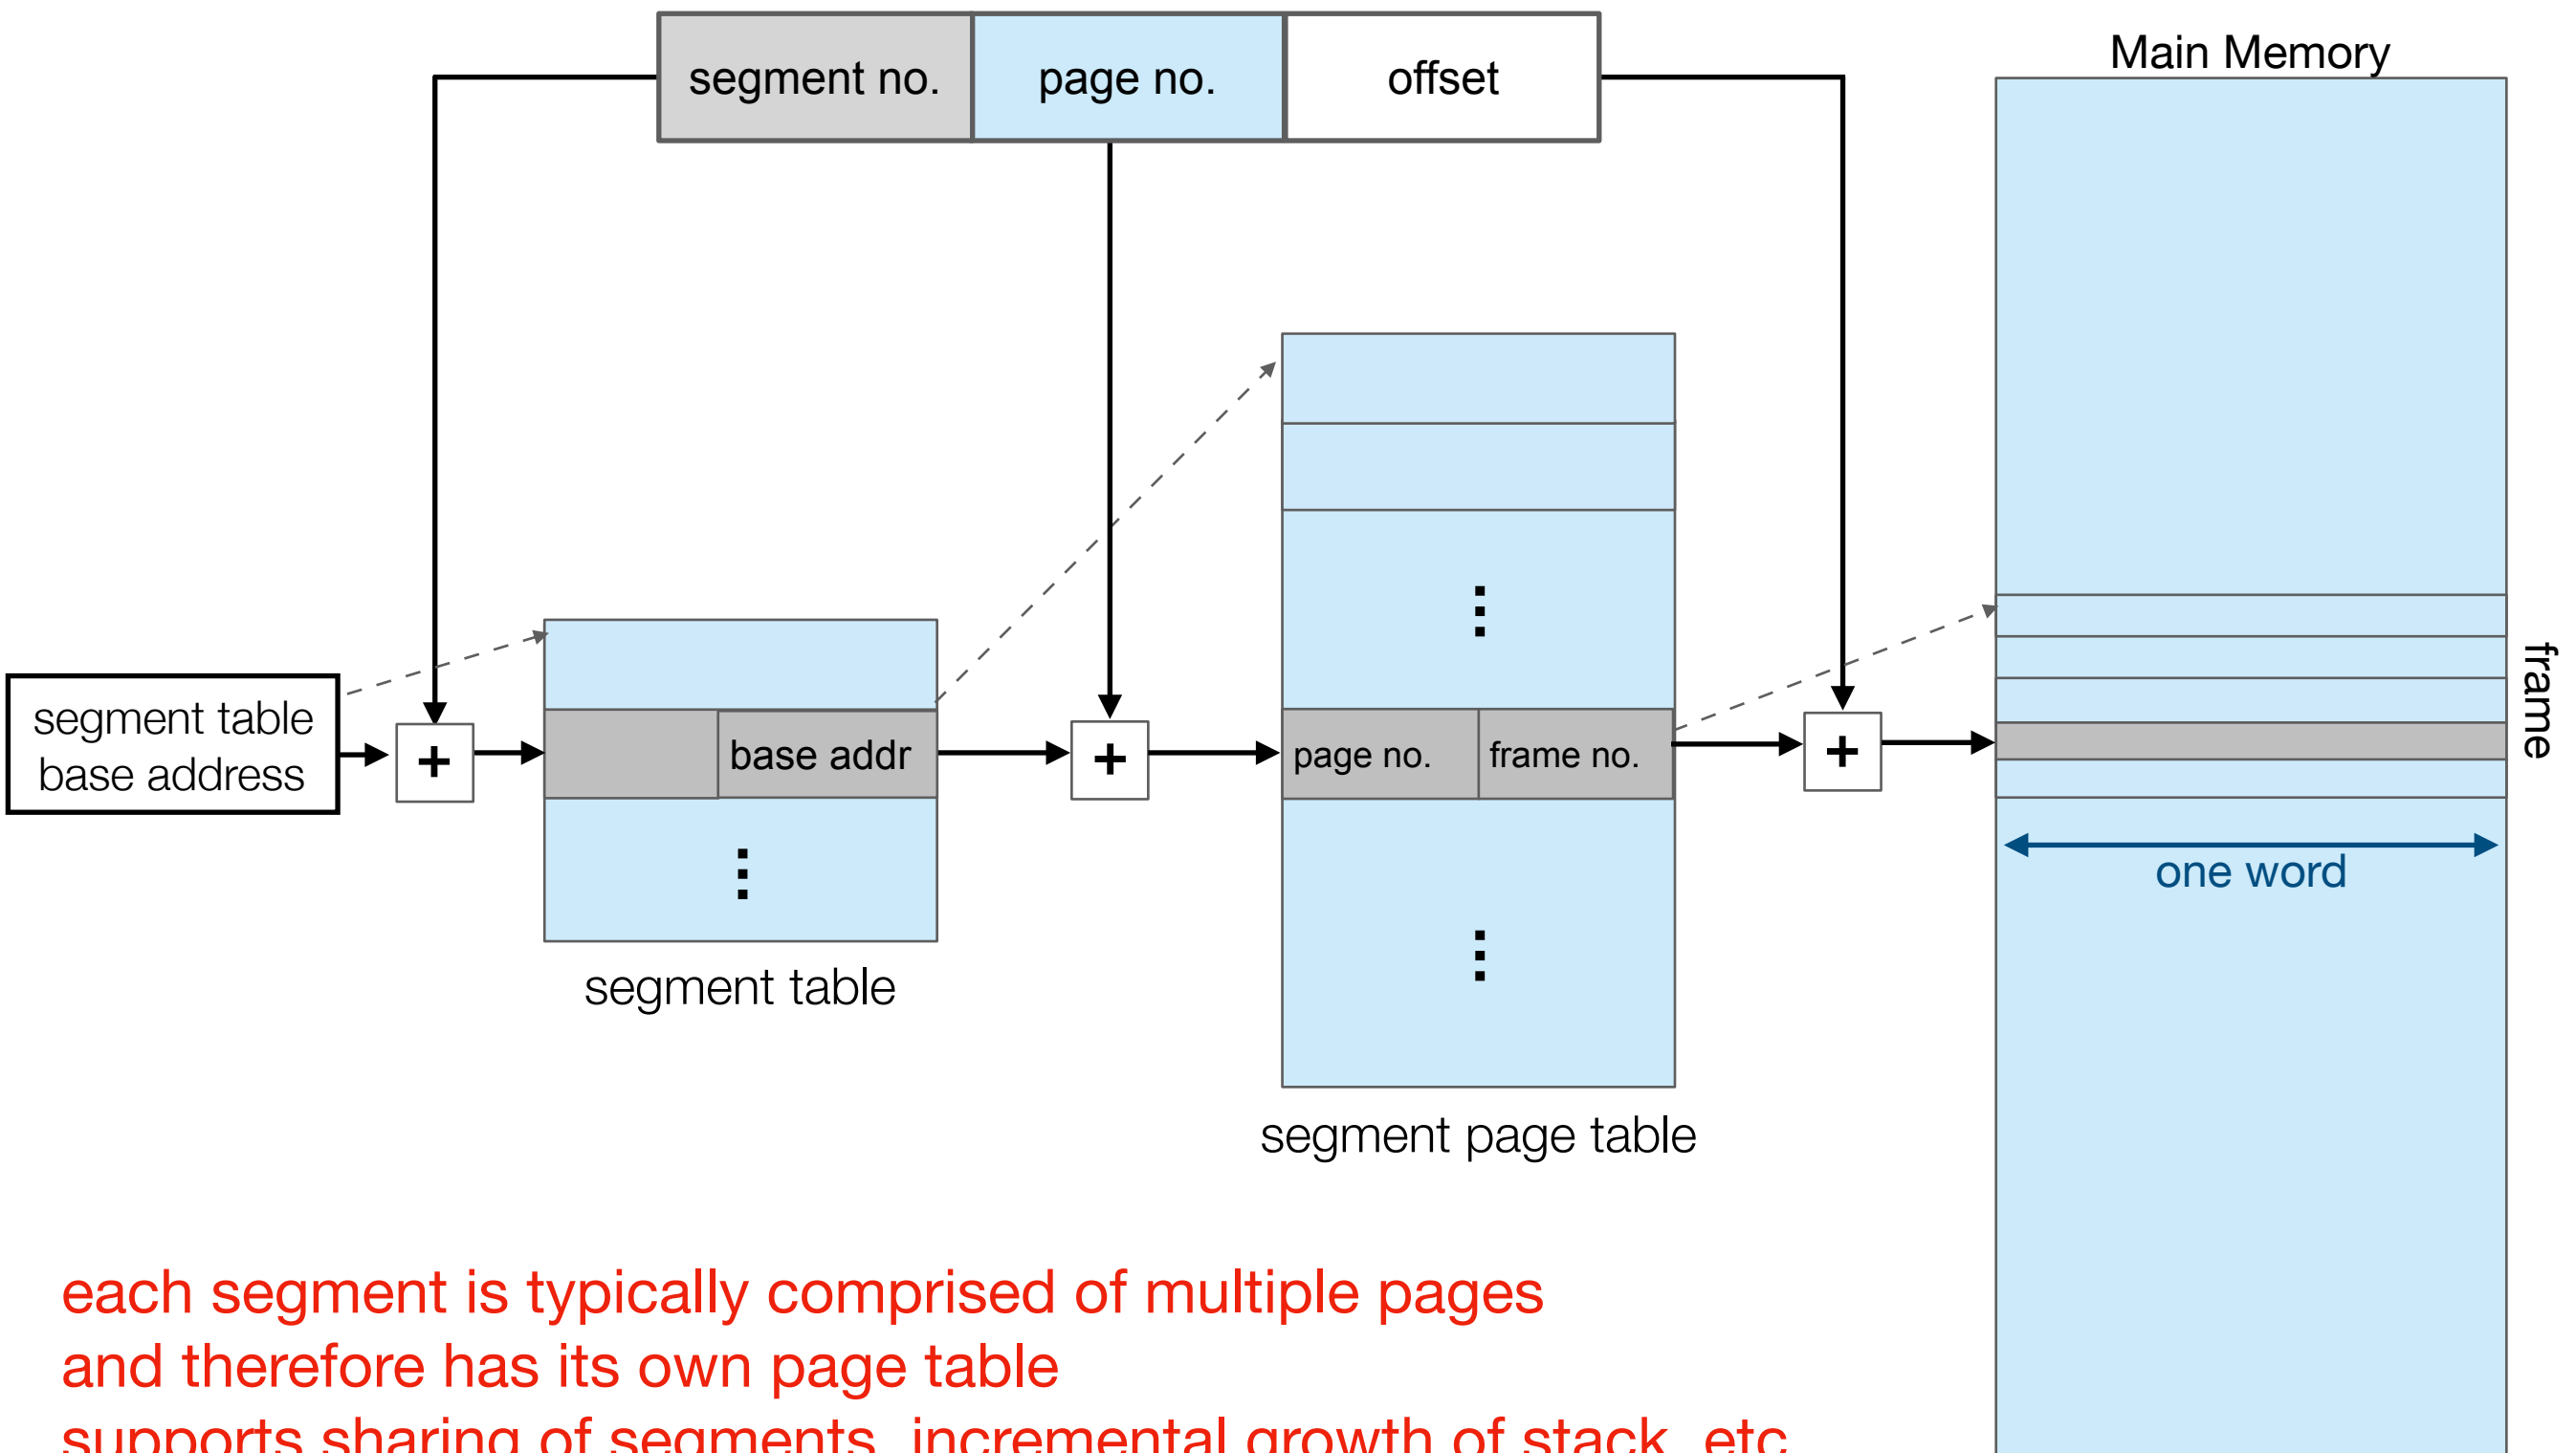
- specific pages of a process (rather than an entire process) are swapped
 - common in Linux and Windows



Combining segmentation and paging

- think of virtual address space as a collection of segments (logical units) of arbitrary sizes
- think of physical memory as a sequence of fixed size page frames
- segments are typically larger than page frames
- map a logical segment onto multiple page frames by paging the segments

Segmented paging



each segment is typically comprised of multiple pages
and therefore has its own page table
supports sharing of segments, incremental growth of stack, etc.

Addresses in segmented paging

- a virtual address becomes a segment number, a page within that segment, and an offset within the page
- the segment number indexes into the segment table which yields the base address of the page table for that segment
- check the remainder of the address (page number and offset) against the limit of the segment
- use the page number to index the page table
 - the entry is the frame (just like paging)
- add the frame and the offset to get the physical address

Example of segmented paging

given a memory size of 256 addressable words, a segment page table indexing 8 pages, a page size of 32 words, and 8 logical segments

- how many bits is a physical address?
- how many bits for the seg #, page #, offset?
- how many bits is a virtual address?

Example of segmented paging

given a memory size of 256 addressable words, a segment page table indexing 8 pages, a page size of 32 words, and 8 logical segments

- how many bits is a physical address? **8 bits**
- how many bits for the seg #, page #, offset? **3, 3, and 5 bits**
- how many bits is a virtual address? **$3 + 3 + 5 = 11$ bits**

What if the virtual address space is large?

- 32-bit logical address space, 4 KB page size (2^{12}) → over 1 million page table entries
 - each page table entry is typically 4 bytes => the size of the page table is over 4 MB!
 - hence the page table does not fit in one page

What if the virtual address space is large?


- 32-bit logical address space, 4 KB page size (2^{12}) → over 1 million page table entries
 - each page table entry is typically 4 bytes => the size of the page table is over 4 MB!
 - hence the page table does not fit in one page
- 64-bit logical address space, 4 KB page size → over 4 quadrillion page table entries

Paging in modern computer systems

memory structures for paging can get huge using straight-forward methods for modern systems that support a larger address space (2^{32} to 2^{64})

- consider a 32-bit logical address space

- page size of 4KB (2^{12})
- page table would have 1 million entries ($2^{32} / 2^{12}$)



- if each entry is 4 bytes then we need 4MB of physical address space for storing the page table alone
 - it is larger than the page size and we don't want to allocate that contiguously in main memory
 - in this case, we cannot use more than 10 bits for the page number if we want to fit the entire page table into one page

