

# Operating System Concepts

## Lecture 26: Page Table Structure & Demand Paging

Omid Ardakanian  
[oardakan@ualberta.ca](mailto:oardakan@ualberta.ca)  
University of Alberta

MWF 12:00-12:50 VVC 2 215

# Today's class

---

- Dealing with large page tables
  - hierarchical
  - hashed
  - inverted
- Demand paged virtual memory


# Paging in modern computer systems

---


memory structures for paging can get huge using straight-forward methods for modern systems that support a larger address space ( $2^{32}$  to  $2^{64}$ )

- consider a 32-bit logical address space

- page size of 4KB ( $2^{12}$ )
- page table would have 1 million entries ( $2^{32} / 2^{12}$ )



- if each entry is 4 bytes then we need 4MB of physical address space for storing the page table alone
  - it is larger than the page size and we don't want to allocate that contiguously in main memory
  - in this case, we cannot use more than 10 bits for the page number if we want to fit the entire page table into one page




# Structure of the page table

---

- a simple solution is to divide the page table into smaller units
  - hierarchical Paging
  - hashed Page Tables
  - inverted Page Tables

# Hierarchical page table


- break up the logical address space into multiple page tables
- a simple technique is a two-level page table
  - we page the page table
  - this requires two lookups per memory reference



# Hierarchical page table: example


---

- consider a 32-bit logical address space with the page size of 4KB ( $2^{12}$ )
  - since the page table is paged, the page number is further divided into:
    - a 10-bit page number
    - a 12-bit page offset

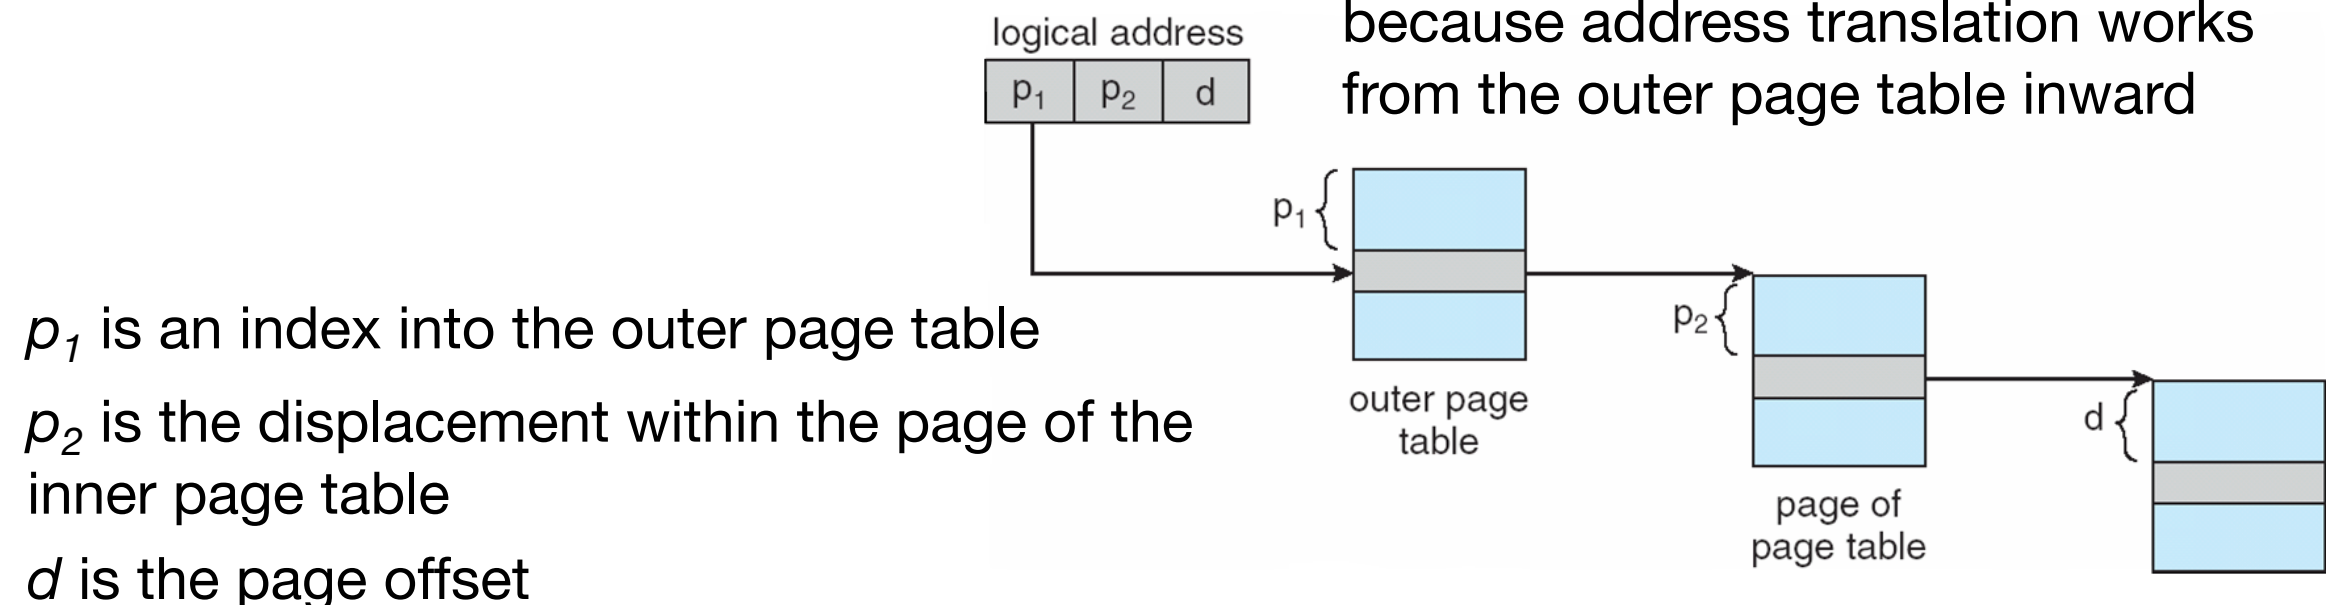


# Hierarchical page table: example

- consider a 32-bit logical address space with the page size of 4KB ( $2^{12}$ )
  - since the page table is paged, the page number is further divided into:
    - a 10-bit page number
    - a 12-bit page offset




known as **forward-mapped page table** because address translation works from the outer page table inward



# What about a 64-bit logical address space?


---






# What about a 64-bit logical address space?

---




should we page the outer page too (three-level paging scheme)? **but the 2nd outer page table will be  $2^{34}$  bits (16 GB) in size and can't fit in one page**




# What about a 64-bit logical address space?

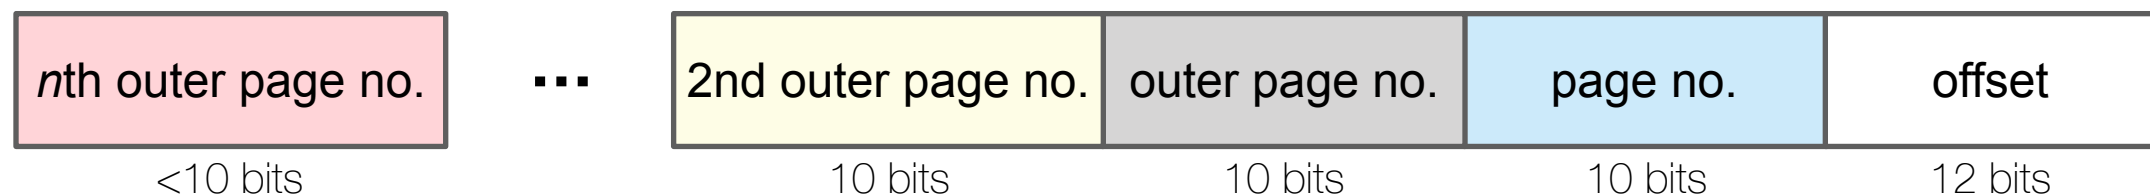
---




should we page the outer page too (three-level paging scheme)? **but the 2nd outer page table will be  $2^{34}$  bits (16 GB) in size and can't fit in one page**




this requires a prohibitive number of memory accesses! e.g., 5 levels of tables



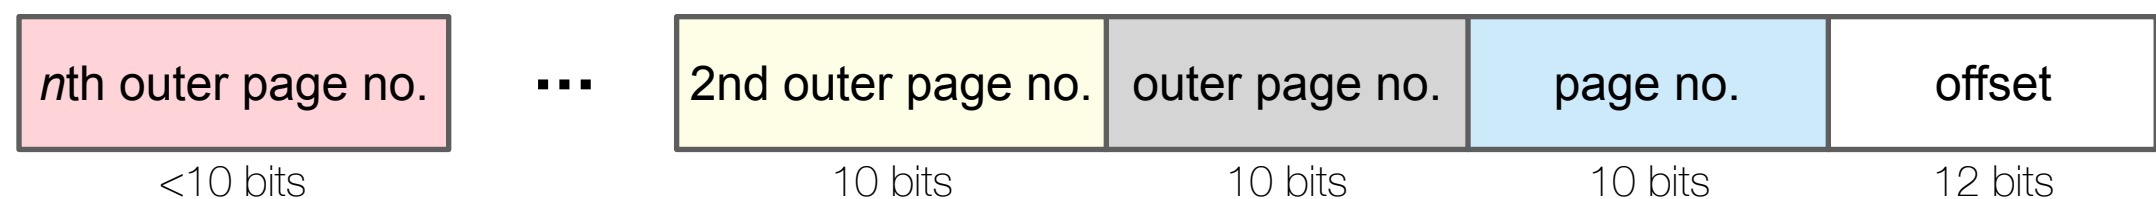
# What about a 64-bit logical address space?



should we page the outer page too (three-level paging scheme)? **but the 2nd outer page table will be  $2^{34}$  bits (16 GB) in size and can't fit in one page**

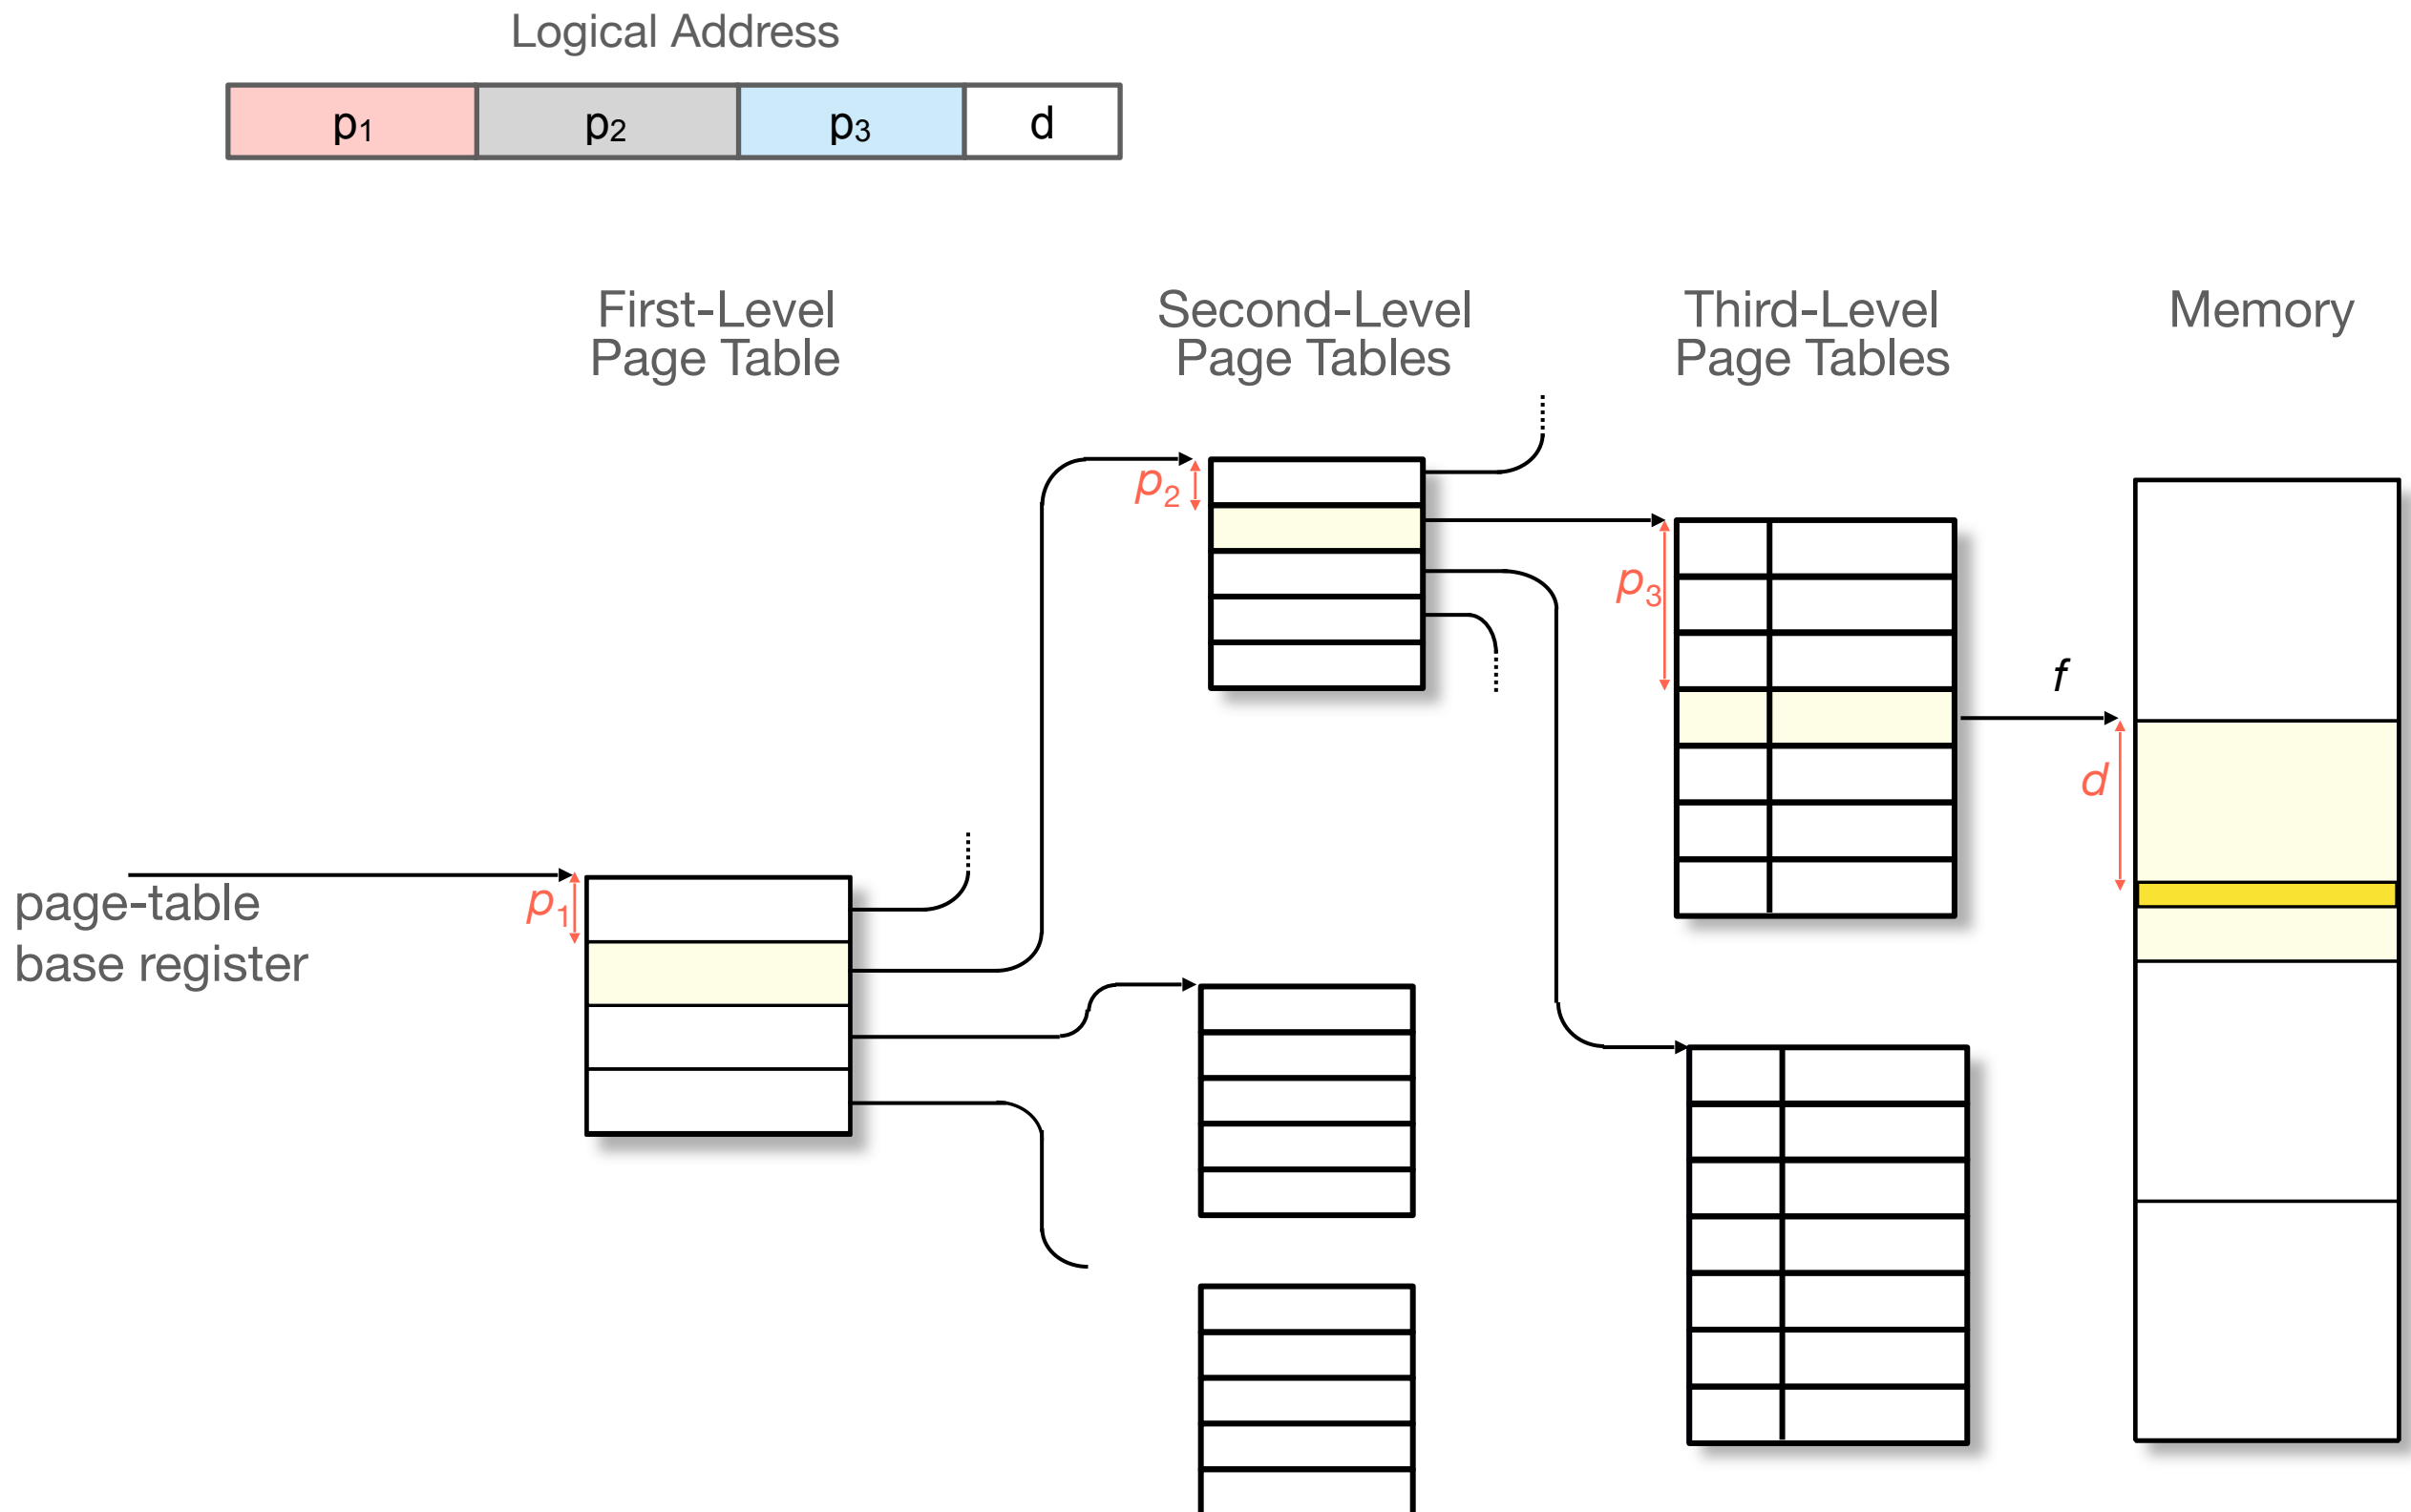


this requires a prohibitive number of memory accesses! e.g., 5 levels of tables



thus, for 64-bit architectures, hierarchical page tables are generally considered inappropriate

# Multilevel address translation




# Hashed page table

---

- common in address spaces  $> 32$  bits
  - the virtual page number is hashed into a page table, so page  $i$  is placed in slot  $H(i)$  where  $H$  is an agreed-upon hash function
- each slot in a chain of elements hashing to the same location
  - each element contains (1) the virtual page number (2) the value of the mapped page frame (3) a pointer to the next element
- virtual page numbers are compared in this chain searching for a match
  - if a match is found, the corresponding physical frame is extracted

# Hashed page table

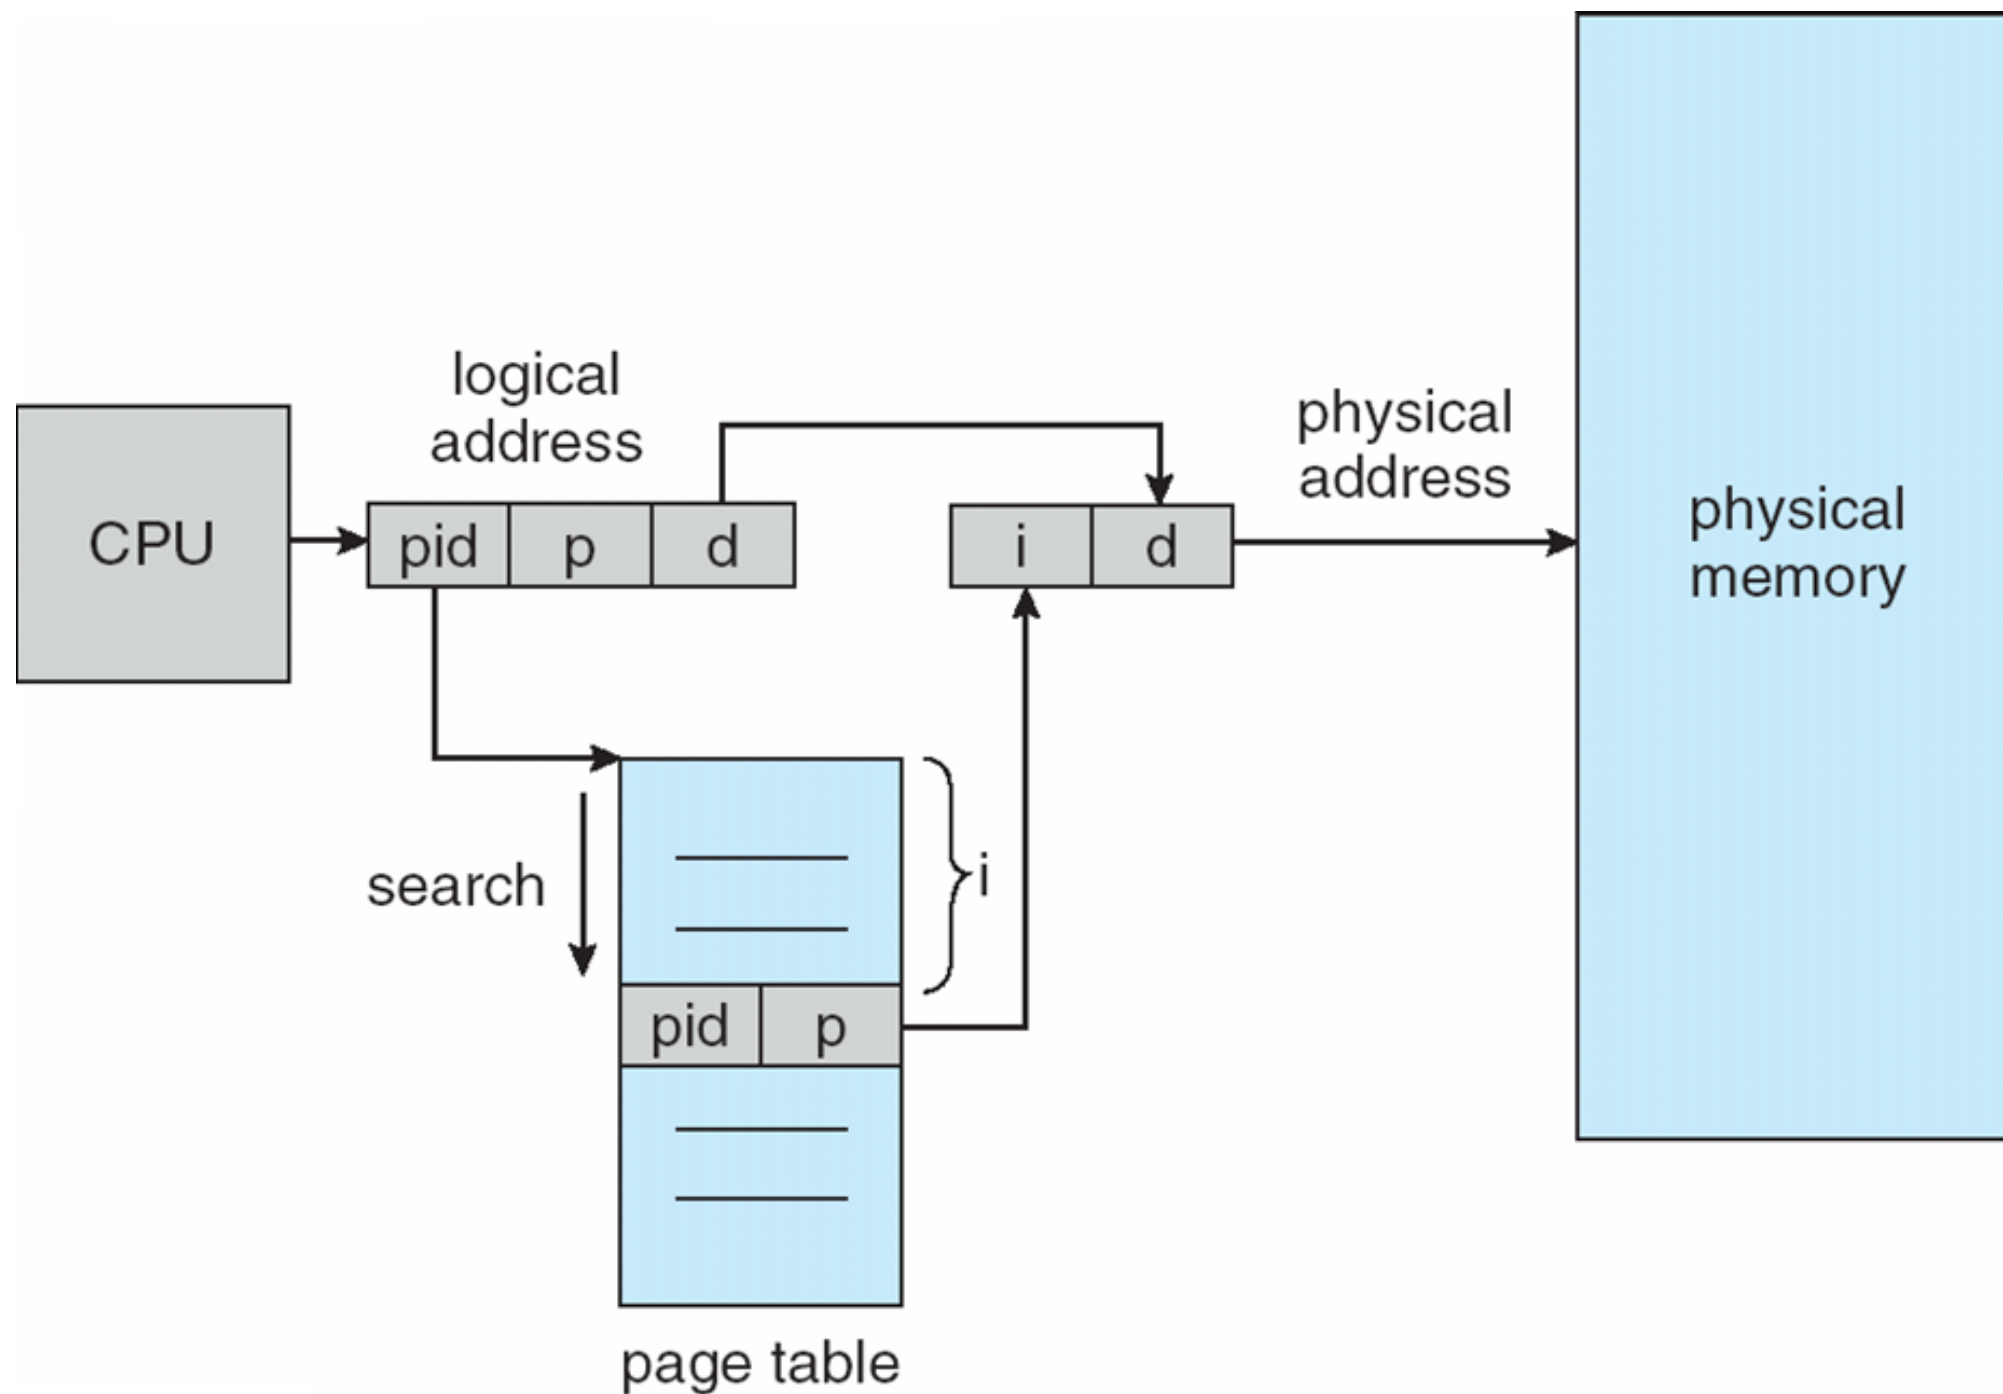


# Inverted page table

---

- rather than each process having a page table and keeping track of all possible logical pages, track all physical pages (frames) in one page table
  - the inverted page table has one entry for each frame of memory (sorted by physical address)
  - each page table entry consists of the virtual address of the page stored in that real memory location and an **address-space identifier** (e.g., which process owns that page)
- it **decreases** memory needed to store each page table, but **increases** time to search the table when a page reference occurs
- use hash table to limit the search to one — or at most a few — page table entries
  - TLB can accelerate access
- sharing cannot be used with inverted page table
  - because one physical page cannot have two or more shared virtual addresses

# Inverted page table





# Hierarchical versus hashed/inverted page table

---

- page table has one entry per virtual page
- hashed/inverted page table has one entry per page frame (i.e., physical frame)

# VIRTUAL MEMORY

the illusion of an infinite virtual memory!

# Virtual memory

---

- **fact:** the entire address space of a process is rarely used (e.g., unusual routines, large data structures) and we can run a process even if all its pages are not in memory
  - so why to load all pages into memory before running a process?

# Virtual memory

---

- **fact:** the entire address space of a process is rarely used (e.g., unusual routines, large data structures) and we can run a process even if all its pages are not in memory
  - so why to load all pages into memory before running a process?
- **basic idea:** only portions of a process's virtual address space (working set) are mapped at any point in time; the remainder is **in multiple files on disk** (in file system or in swap space)
  - code pages are stored in a file on disk (e.g., `a.out`)
  - data, stack, and heap pages are also stored in a special file that is not visible to users and exists only when the program is executing
  - some of these pages are loaded in the main memory; hence, the main memory acts as a cache for the disk
  - the page table indicates if the page is on disk or memory using the **valid/invalid** bit

# Virtual memory

---

- **fact:** the entire address space of a process is rarely used (e.g., unusual routines, large data structures) and we can run a process even if all its pages are not in memory
  - so why to load all pages into memory before running a process?
- **basic idea:** only portions of a process's virtual address space (working set) are mapped at any point in time; the remainder is **in multiple files on disk** (in file system or in swap space)
  - code pages are stored in a file on disk (e.g., `a.out`)
  - data, stack, and heap pages are also stored in a special file that is not visible to users and exists only when the program is executing
  - some of these pages are loaded in the main memory; hence, the main memory acts as a cache for the disk
  - the page table indicates if the page is on disk or memory using the **valid/invalid** bit
- **consequences:** (a) the virtual address space is no longer constrained by the size of the physical address space; (b) more processes can fit in memory to run concurrently

# Virtual memory

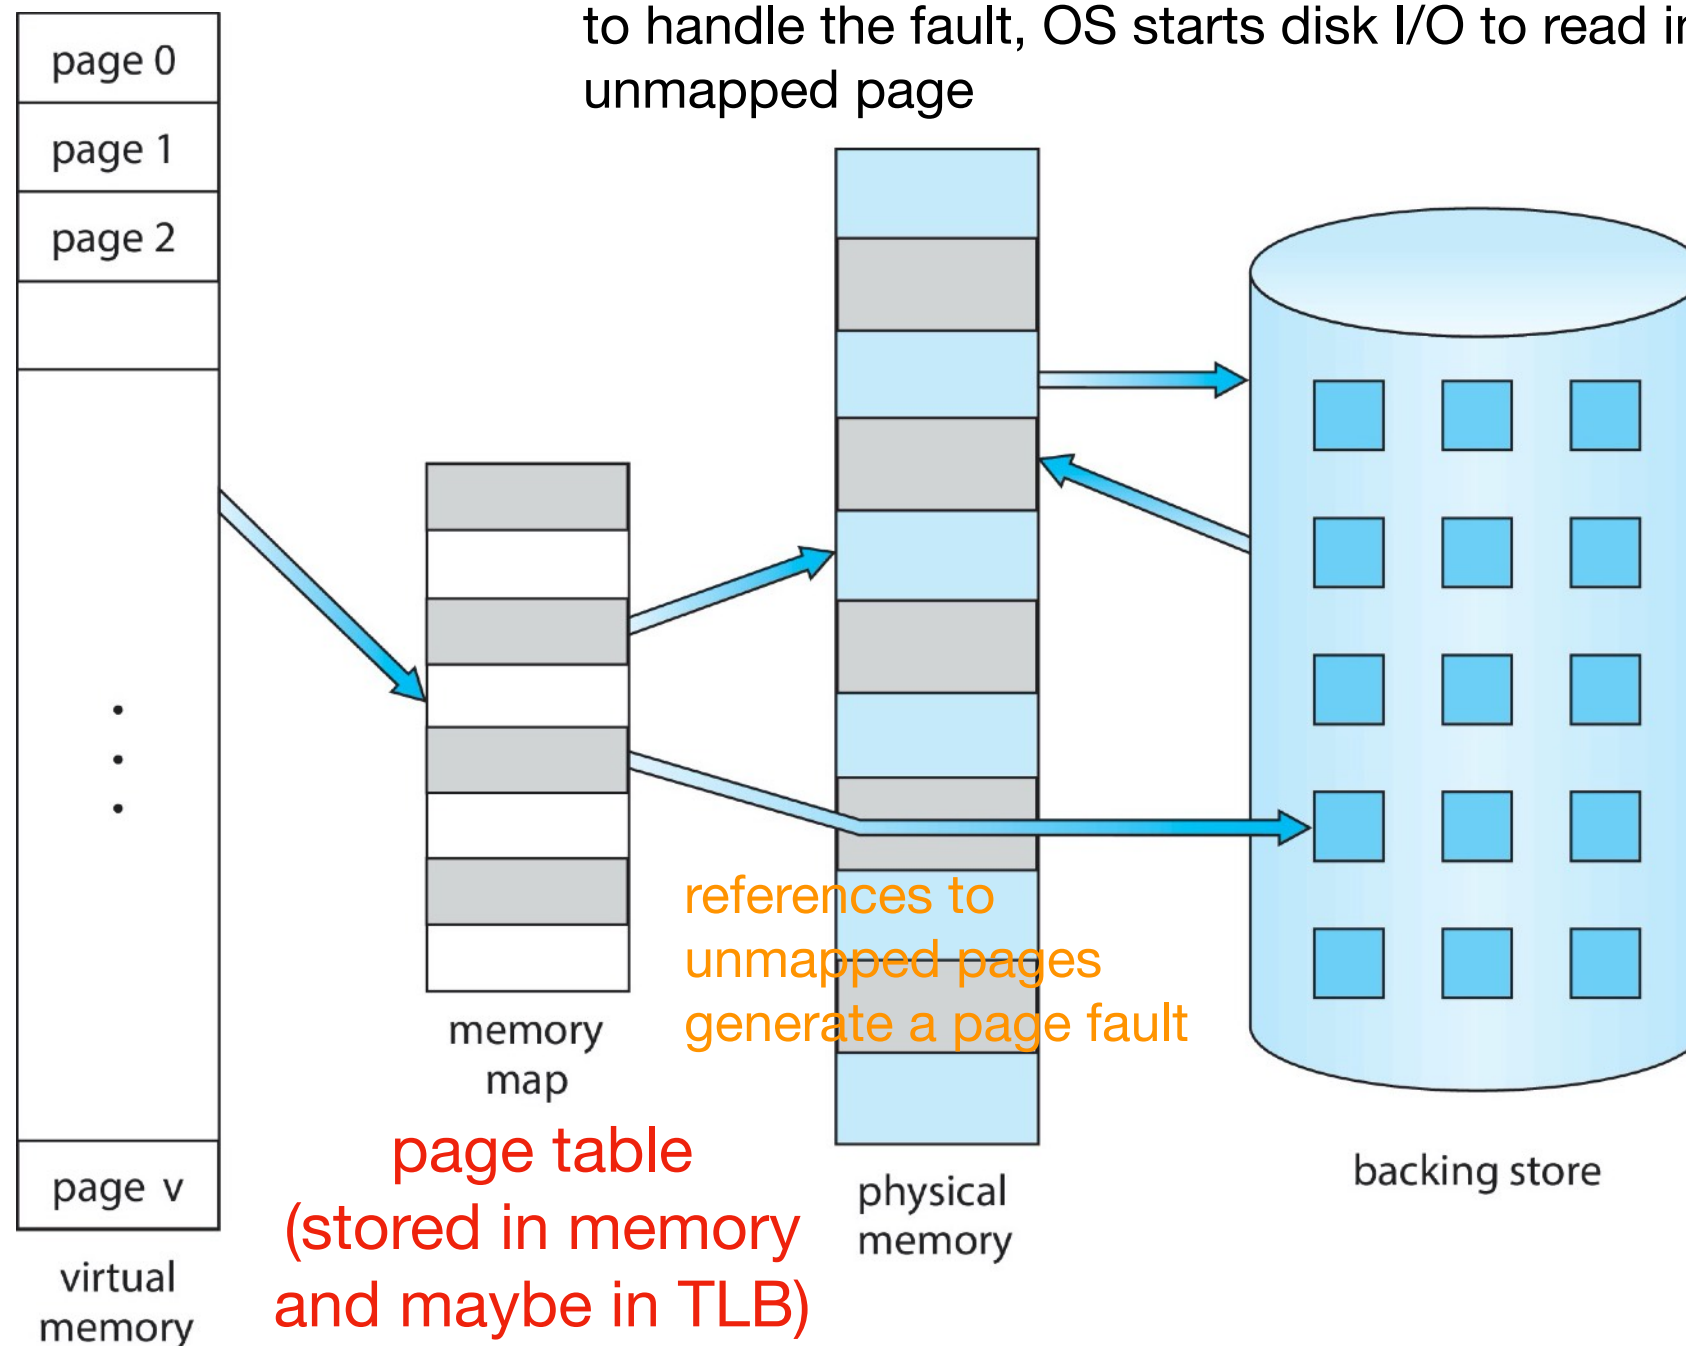
---

- **fact:** the entire address space of a process is rarely used (e.g., unusual routines, large data structures) and we can run a process even if all its pages are not in memory
  - so why to load all pages into memory before running a process?
- **basic idea:** only portions of a process's virtual address space (working set) are mapped at any point in time; the remainder is **in multiple files on disk** (in file system or in swap space)
  - code pages are stored in a file on disk (e.g., `a.out`)
  - data, stack, and heap pages are also stored in a special file that is not visible to users and exists only when the program is executing
  - some of these pages are loaded in the main memory; hence, the main memory acts as a cache for the disk
  - the page table indicates if the page is on disk or memory using the **valid/invalid** bit
- **consequences:** (a) the virtual address space is no longer constrained by the size of the physical address space; (b) more processes can fit in memory to run concurrently
- **caveat:** memory accesses must be for pages that are in memory for the vast majority of time, otherwise the effective memory access time will approach that of the disk

# What happens if a program references a page that is not memory resident?

a **page fault** occurs when the page table is accessed and the entry for the page does not have the valid bit set

to handle the fault, OS starts disk I/O to read in the unmapped page



# When to load a page into memory?

---

- load all pages of a process at its start time
  - disadvantages: unnecessary I/O for bringing all pages into memory, and wasting a lot of memory
  - the virtual address space must be no larger than the physical memory (it's not necessarily true)



# When to load a page into memory?

---

- load all pages of a process at its start time
  - disadvantages: unnecessary I/O for bringing all pages into memory, and wasting a lot of memory
  - the virtual address space must be no larger than the physical memory (it's not necessarily true)
- leave it to the programmer to decide when to load and evict pages
  - is difficult to do and error-prone

# When to load a page into memory?

---

- load all pages of a process at its start time
  - disadvantages: unnecessary I/O for bringing all pages into memory, and wasting a lot of memory
  - the virtual address space must be no larger than the physical memory (it's not necessarily true)
- leave it to the programmer to decide when to load and evict pages
  - is difficult to do and error-prone
- request paging: process tells OS before it needs a page, and when it is through with a page

# When to load a page into memory?

---

- load all pages of a process at its start time
  - disadvantages: unnecessary I/O for bringing all pages into memory, and wasting a lot of memory
  - the virtual address space must be no larger than the physical memory (it's not necessarily true)
- leave it to the programmer to decide when to load and evict pages
  - is difficult to do and error-prone
- request paging: process tells OS before it needs a page, and when it is through with a page
- **demand paging**: OS loads a page the first time it is referenced (transparent to the process)
  - may need to remove a page from memory to make room for the new page (**will discuss this later**)
  - process must give up the CPU while the page is being loaded (transitions to the waiting state)


# When to load a page into memory?

---

- load all pages of a process at its start time
  - disadvantages: unnecessary I/O for bringing all pages into memory, and wasting a lot of memory
  - the virtual address space must be no larger than the physical memory (it's not necessarily true)
- leave it to the programmer to decide when to load and evict pages
  - is difficult to do and error-prone
- request paging: process tells OS before it needs a page, and when it is through with a page
- **demand paging**: OS loads a page the first time it is referenced (transparent to the process)
  - may need to remove a page from memory to make room for the new page (**will discuss this later**)
  - process must give up the CPU while the page is being loaded (transitions to the waiting state)
- pre-paging: OS guesses in advance which pages the process will need and pre-loads them into memory
  - allows more overlap of CPU and I/O if the OS guesses correctly, otherwise it causes a page fault
  - errors may result in removing useful pages
  - difficult to get right due to branches in code


# Implementing demand paging

- a copy of the entire program must be stored on disk



# Implementing demand paging

- a copy of the entire program must be stored on disk
- the valid bit in the page table indicates that the page is in memory
  - ‘v’ indicates that the page is in-memory
  - ‘i’ indicates that it is not-in-memory (initiate disk I/O) or is an invalid memory reference (abort the process in this case)



# What happens on a page fault?

---


- instruction faults on a page whose valid bit is not set
- page fault causes a trap to kernel
  - saves the registers and state of the faulting process
  - checks if the memory reference was legal and determines the location of the page on the disk (requires a more complex page table)

# What happens on a page fault?

---

- instruction faults on a page whose valid bit is not set
- page fault causes a trap to kernel
  - saves the registers and state of the faulting process
  - checks if the memory reference was legal and determines the location of the page on the disk (requires a more complex page table)
- OS puts the faulting process on the wait queue of disk and starts reading the unmapped page from disk to a frame
  - requires selecting a free frame (or a page to replace using a page replacement algorithm) and zeroing it out

**free-frame list: a pool of free frames**





# What happens on a page fault?

---

- instruction faults on a page whose valid bit is not set
- page fault causes a trap to kernel
  - saves the registers and state of the faulting process
  - checks if the memory reference was legal and determines the location of the page on the disk (requires a more complex page table)
- OS puts the faulting process on the wait queue of disk and starts reading the unmapped page from disk to a frame
  - requires selecting a free frame (or a page to replace using a page replacement algorithm) and zeroing it out
- while I/O is being done, OS context switches to another process
- OS gets interrupt when I/O completes (i.e., the page is read from disk and loaded into memory)
  - saves the registers and state of the other process

# What happens on a page fault?

---

- instruction faults on a page whose valid bit is not set
- page fault causes a trap to kernel
  - saves the registers and state of the faulting process
  - checks if the memory reference was legal and determines the location of the page on the disk (requires a more complex page table)
- OS puts the faulting process on the wait queue of disk and starts reading the unmapped page from disk to a frame
  - requires selecting a free frame (or a page to replace using a page replacement algorithm) and zeroing it out
- while I/O is being done, OS context switches to another process
- OS gets interrupt when I/O completes (i.e., the page is read from disk and loaded into memory)
  - saves the registers and state of the other process
- OS updates the page table entry

# What happens on a page fault?

---

- instruction faults on a page whose valid bit is not set
- page fault causes a trap to kernel
  - saves the registers and state of the faulting process
  - checks if the memory reference was legal and determines the location of the page on the disk (requires a more complex page table)
- OS puts the faulting process on the wait queue of disk and starts reading the unmapped page from disk to a frame
  - requires selecting a free frame (or a page to replace using a page replacement algorithm) and zeroing it out
- while I/O is being done, OS context switches to another process
- OS gets interrupt when I/O completes (i.e., the page is read from disk and loaded into memory)
  - saves the registers and state of the other process
- OS updates the page table entry
- OS puts the faulting process back on the ready queue
  - resumes the interrupted instruction when CPU is allocated to this process again

# Handling a page fault

