

# Operating System Concepts

## Lecture 32: File System Reliability & Examples

Omid Ardakanian  
[oardakan@ualberta.ca](mailto:oardakan@ualberta.ca)  
University of Alberta

MWF 12:00-12:50 VVC 2 215

# Free space management

---

- disk space is limited so file system has to allocate the space released after deleting files to new files
  - must be able to find free space quickly and release disk space quickly

# Free space management

---

- disk space is limited so file system has to allocate the space released after deleting files to new files
  - must be able to find free space quickly and release disk space quickly
- to keep track of free disk blocks, file system maintains a free-space list
  - when blocks are allocated to a file, they are removed from this list
  - when a file is deleted, freed blocks are added to this list

# Free space management

---

- disk space is limited so file system has to allocate the space released after deleting files to new files
  - must be able to find free space quickly and release disk space quickly
- to keep track of free disk blocks, file system maintains a free-space list
  - when blocks are allocated to a file, they are removed from this list
  - when a file is deleted, freed blocks are added to this list
- the free-space list may not be implemented as a list!

# Free space management

---

- disk space is limited so file system has to allocate the space released after deleting files to new files
  - must be able to find free space quickly and release disk space quickly
- to keep track of free disk blocks, file system maintains a free-space list
  - when blocks are allocated to a file, they are removed from this list
  - when a file is deleted, freed blocks are added to this list
- the free-space list may not be implemented as a list!
- need to maintain a free-inode list too

# Bit vector

---

- each disk block is represented by one bit
  - the block is free if the corresponding bit is set

00001001011 ...

# Bit vector

---

- each disk block is represented by one bit
  - the block is free if the corresponding bit is set

00001001011 ...

- marking a block as free is simple because the block number is the index into the bitmap

# Bit vector

---

- each disk block is represented by one bit
  - the block is free if the corresponding bit is set

00001001011 ...

- marking a block as free is simple because the block number is the index into the bitmap
- to find a free block, the file system sequentially checks each word in the bitmap with zero
  - the first non-zero word is scanned for the first 1 bit then
  - expensive to find a free block if most blocks are used!



# Bit vector

---

- each disk block is represented by one bit

- the block is free if the corresponding bit is set

00001001011 ...

- marking a block as free is simple because the block number is the index into the bitmap
- to find a free block, the file system sequentially checks each word in the bitmap with zero
  - the first non-zero word is scanned for the first 1 bit then
  - expensive to find a free block if most blocks are used!
- problem? the bitmap might be too big to keep in memory


block size = 512 bytes

disk size = 2GB

$n = 2^{31}/2^9 = 2^{22}$  blocks  $\rightarrow 2^{22}$  bits (**512KB!**)


# Linked free-space list

- link together all free disk blocks; each block contains a pointer to the next free block (saving space)
- the head of the list is cached in memory
- no need to traverse the entire list (if the number of free blocks is recorded)



# Linked free-space list

- link together all free disk blocks; each block contains a pointer to the next free block (saving space)
- the head of the list is cached in memory
- no need to traverse the entire list (if the number of free blocks is recorded)
- problems?
  - linked list gets disorganized over time
  - cannot get contiguous space easily



# Extensions

---

- grouping
  - modify linked list to store address of next  $n-1$  free blocks in first free block, plus a pointer to next block that contains free-block-pointers (like this block)

# Extensions

---

- grouping
  - modify linked list to store address of next  $n-1$  free blocks in first free block, plus a pointer to next block that contains free-block-pointers (like this block)
- counting
  - space is contiguously used and freed frequently so
    - keep the address of the first free block and the number of following free blocks
    - free-space list has entries containing addresses and counts

# Summary

---

- many of the concerns, tradeoffs, and design decisions discussed for file system are similar to those of virtual memory (**and disk as we will see later**)
  - contiguous allocation is simple, but suffers from external fragmentation, the need for compaction, and the need to move files as they grow
  - indexed allocation is very similar to page tables
    - a mapping from logical file blocks to physical disk blocks
  - free space can be managed using a bitmap or a linked list

# Today's class

---

- How to make sure that data and metadata previously stored can be retrieved regardless of software crashes and hardware failures?
  - data consistency
  - metadata consistency
- How real file systems are implemented?
  - FFS
  - NTFS

# **File System Reliability**



# File system reliability problem

---

- a single file system operation can involve updates to multiple physical disk blocks (inode, data blocks, bitmap, etc.)
  - to move a file between directories, it has to be deleted from the old directory and added to the new directory
  - to create a new file, file system must allocate space on disk for file's data, write inode to disk, add the file number and name to the directory containing the file

# File system reliability problem

---

- a single file system operation can involve updates to multiple physical disk blocks (inode, data blocks, bitmap, etc.)
  - to move a file between directories, it has to be deleted from the old directory and added to the new directory
  - to create a new file, file system must allocate space on disk for file's data, write inode to disk, add the file number and name to the directory containing the file
- problem 1: physical disk blocks are updated one at a time
  - crash may occur between multiple updates causing inconsistencies among file-system data structures

# File system reliability problem

---

- a single file system operation can involve updates to multiple physical disk blocks (inode, data blocks, bitmap, etc.)
  - to move a file between directories, it has to be deleted from the old directory and added to the new directory
  - to create a new file, file system must allocate space on disk for file's data, write inode to disk, add the file number and name to the directory containing the file
- problem 1: physical disk blocks are updated one at a time
  - crash may occur between multiple updates causing inconsistencies among file-system data structures
- problem 2: **write-through** (i.e., to write modified data back to disk synchronously) is really slow, so changes are cached (written to disk when convenient)
  - crash may occur before disk I/O causing loss of data

# File system reliability problem

---

- a single file system operation can involve updates to multiple physical disk blocks (inode, data blocks, bitmap, etc.)
  - to move a file between directories, it has to be deleted from the old directory and added to the new directory
  - to create a new file, file system must allocate space on disk for file's data, write inode to disk, add the file number and name to the directory containing the file
- problem 1: physical disk blocks are updated one at a time
  - crash may occur between multiple updates causing inconsistencies among file-system data structures
- problem 2: **write-through** (i.e., to write modified data back to disk synchronously) is really slow, so changes are cached (written to disk when convenient)
  - crash may occur before disk I/O causing loss of data
- **Goal:** guarantee consistency regardless of when crash occurs

# Data consistency

---

- UNIX file system uses the **write-back strategy** (i.e., delaying writing the modified data back to disk) with periodic forced writes to disk (e.g., every 30 seconds)
  - so other processes read data from cache instead of disk
  - potential for loss of 30 seconds worth of cached changes
  - user can use the `sync` system call to flush buffer cache to disk immediately

# Data consistency

---

- UNIX file system uses the **write-back strategy** (i.e., delaying writing the modified data back to disk) with periodic forced writes to disk (e.g., every 30 seconds)
  - so other processes read data from cache instead of disk
  - potential for loss of 30 seconds worth of cached changes
  - user can use the `sync` system call to flush buffer cache to disk immediately
- how to write changes to a file onto disk?
  - naive approach:
    - delete old version; create new version
  - correct approach:
    - write new version in temp file; move old version to another temp file; move new version into real file; unlink old version
    - on a crash look at temp area; if there is any files out there notify user that there might be a problem

# Metadata consistency

---

- if multiple updates to metadata is needed, perform them in a specific order which allows for an operation to be safely interrupted

# Metadata consistency

---

- if multiple updates to metadata is needed, perform them in a specific order which allows for an operation to be safely interrupted
- recover after crash by checking if there were any in-progress operations and undo them
  - run a consistency checker (called `fsck` in UNIX) which scans the entire disk for internal consistency
  - this approach is used in FAT and UNIX file system



# Metadata consistency

---

- if multiple updates to metadata is needed, perform them in a specific order which allows for an operation to be safely interrupted
- recover after crash by checking if there were any in-progress operations and undo them
  - run a consistency checker (called `fsck` in UNIX) which scans the entire disk for internal consistency
  - this approach is used in FAT and UNIX file system
- disadvantages:
  - post-crash recovery is time consuming (can take minutes or hours) and may not be possible
  - difficult to reduce every operation to a safely interruptible sequence of writes
  - difficult to achieve consistency when multiple operations occur concurrently

# Example

---

- let's say you want to extend a given file by one block; the necessary operations are as follows
  - 1.find a free block
  - 2.write block bitmap
  - 3.write inode with pointer to free block and new file size
  - 4.write data

# Example

---

- let's say you want to extend a given file by one block; the necessary operations are as follows
  - 1.find a free block
  - 2.write block bitmap
  - 3.write inode with pointer to free block and new file size
  - 4.write data
- so in the case of a crash
  - if a bit is set in the bitmap but a pointer to this block is not added to any inode, writing inode must have been in progress when system crashed
  - if a bit is set in the bitmap and a pointer to this block is added to an inode but file data is not written to that block, writing data must have been in progress when system crashed

# Transaction concept

---

- **Definition:** a group of operations that are atomic and durable
  - atomic: happen as a group or not at all
  - durable: future failures do not affect/corrupt previously committed transactions

# Transaction concept

---

- **Definition:** a group of operations that are atomic and durable
  - atomic: happen as a group or not at all
  - durable: future failures do not affect/corrupt previously committed transactions
- **Basic idea:** do a set of metadata updates tentatively; if you don't get to **commit** (due to crash or failure), then **roll-back** the operations as if the transaction never happened
  - commit makes transaction durable by writing a single sector on disk (we assume this happens atomically)

# Implementation

---

- write operations to be performed **sequentially** and synchronously in an on-disk data structure (known as log, journal, or intention list)
  - ignore changes if crash occurs in the middle of these operations

# Implementation

---

- write operations to be performed **sequentially** and synchronously in an on-disk data structure (known as log, journal, or intention list)
  - ignore changes if crash occurs in the middle of these operations
- commit transaction when all changes are on log

# Implementation

---

- write operations to be performed **sequentially** and synchronously in an on-disk data structure (known as log, journal, or intention list)
  - ignore changes if crash occurs in the middle of these operations
- commit transaction when all changes are on log
- write the changes asynchronously to appropriate blocks of disk (random-access)
  - if crash occurs after commit, replay the log (from a pointer) to make sure updates get to disk



# Implementation

---

- write operations to be performed **sequentially** and synchronously in an on-disk data structure (known as log, journal, or intention list)
  - ignore changes if crash occurs in the middle of these operations
- commit transaction when all changes are on log
- write the changes asynchronously to appropriate blocks of disk (random-access)
  - if crash occurs after commit, replay the log (from a pointer) to make sure updates get to disk
- **Note:** sequential I/O is faster than random I/O, and therefore, can be done synchronously

# Log-based transaction-oriented file system

---

- almost all new file systems such as NTFS, ext3, and ext4 use **journaling** (or **write-ahead logging**)
  - write all changes in a transaction to log (update directory, allocate block, etc.) before sending any changes to disk
- journaling eliminates the need for doing post-crash file system consistency check (e.g., via `fsck`)

# Log-based transaction-oriented file system


---

- almost all new file systems such as NTFS, ext3, and ext4 use **journaling** (or **write-ahead logging**)
  - write all changes in a transaction to log (update directory, allocate block, etc.) before sending any changes to disk
- journaling eliminates the need for doing post-crash file system consistency check (e.g., via `fsck`)
- advantage: a general solution to reliability problem
- disadvantage: data is written twice!

# File System Examples

# Traditional UNIX file system

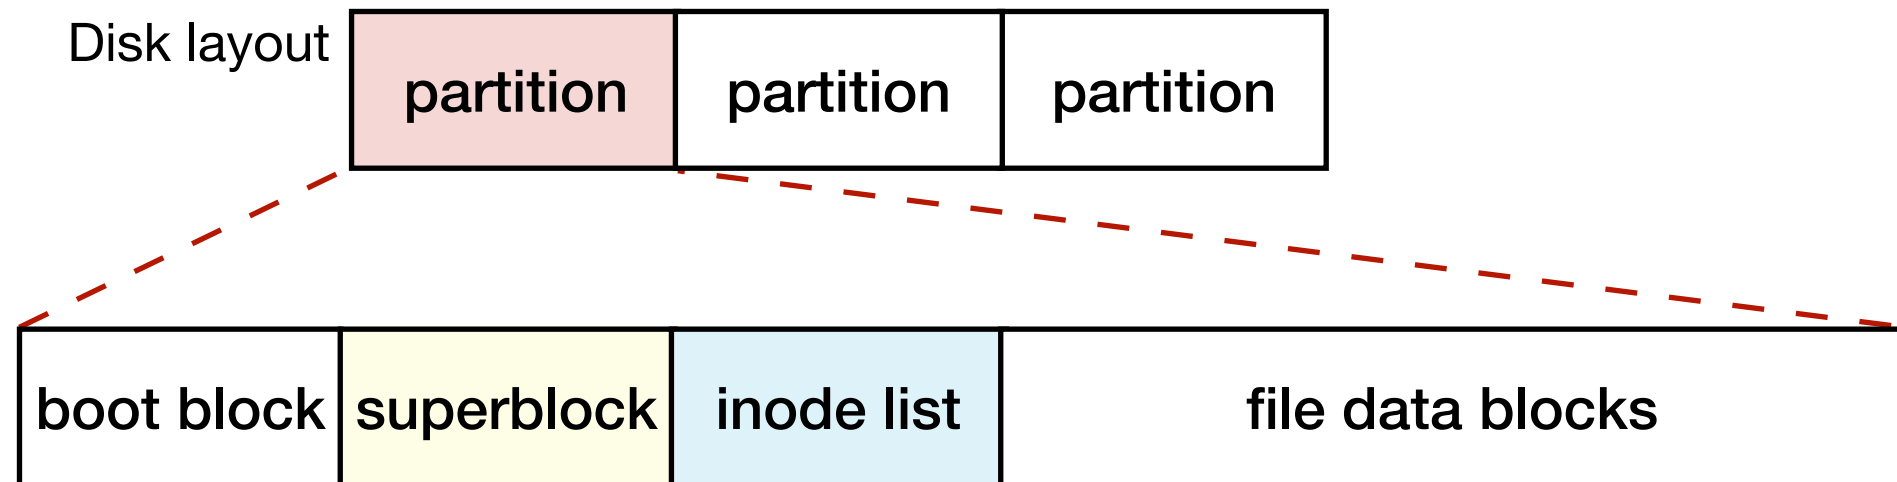
---



- boot block contains the bootstrap program that runs at boot time to load the OS from that partition

# Traditional UNIX file system

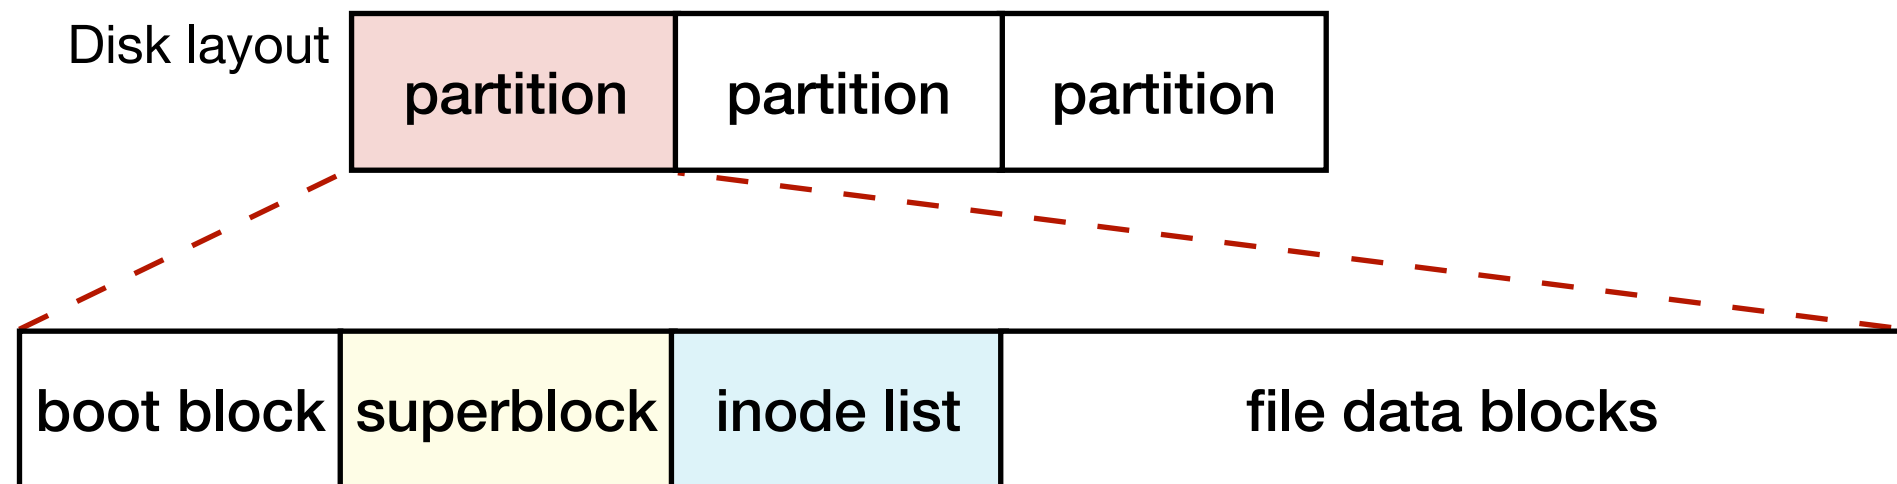
---



- boot block contains the bootstrap program that runs at boot time to load the OS from that partition
- superblock defines a file system (there is only one per file system)
  - size of the file system and size of the inode list
  - list of free disk blocks (and index of the next free block)
  - list of free inodes (and index of the next free inode)
  - location of the inode of the root directory (inode #2)

# Traditional UNIX file system

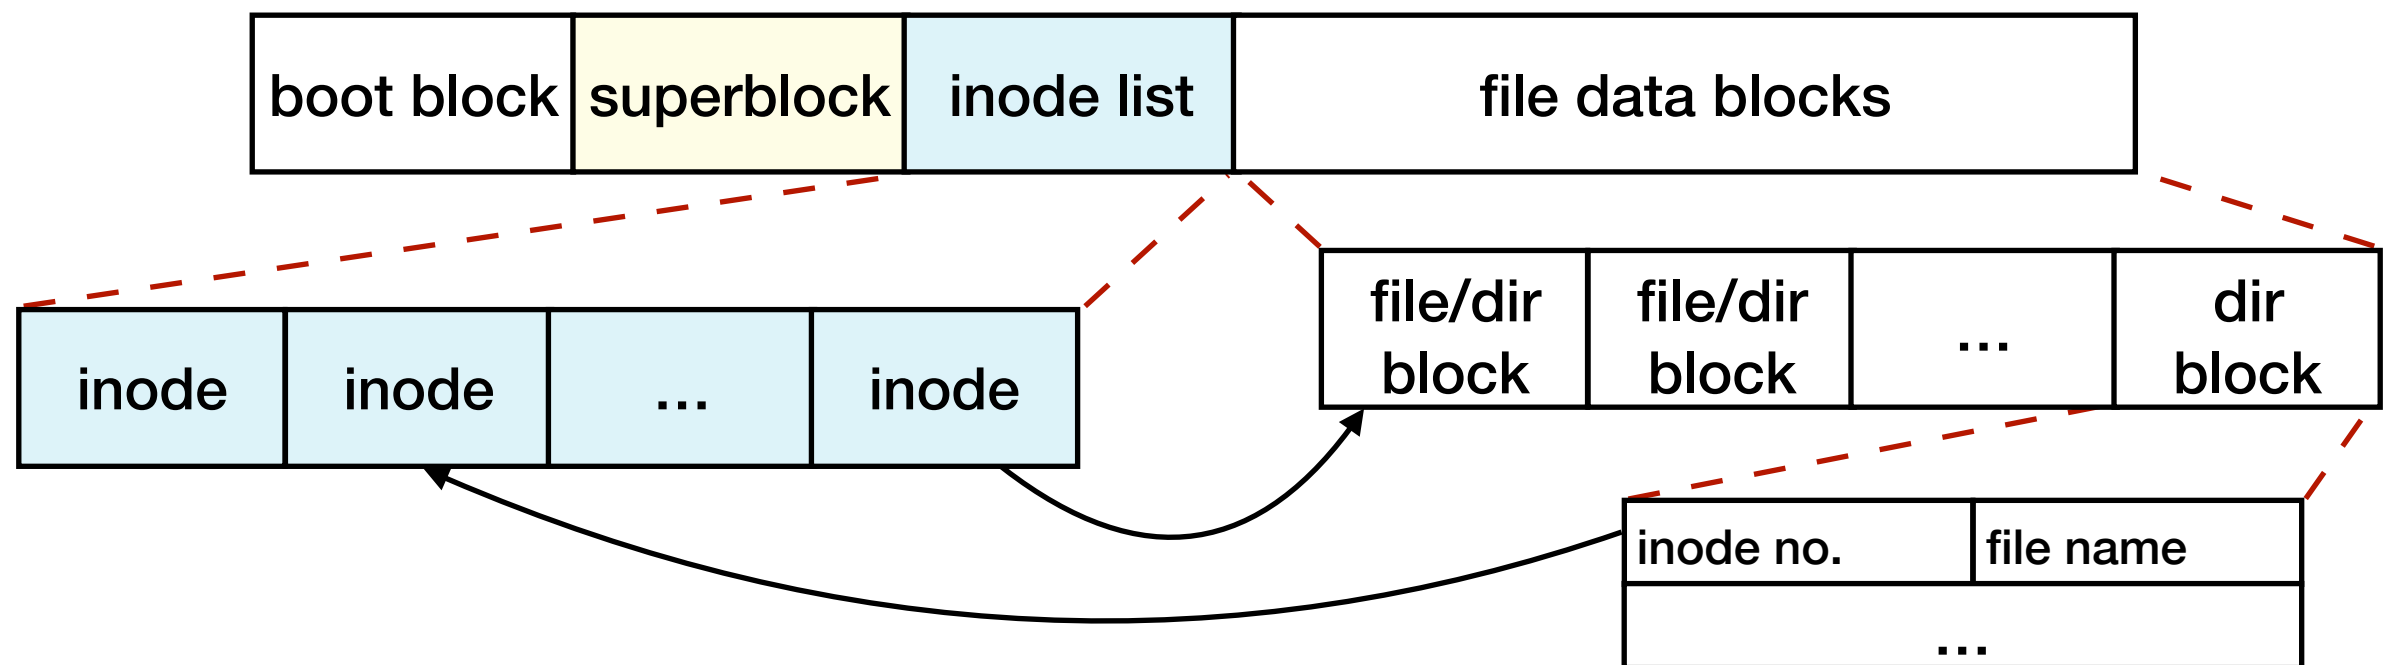
---



- boot block contains the bootstrap program that runs at boot time to load the OS from that partition
- superblock defines a file system (there is only one per file system)
  - size of the file system and size of the inode list
  - list of free disk blocks (and index of the next free block)
  - list of free inodes (and index of the next free inode)
  - location of the inode of the root directory (inode #2)
- inodes contain file metadata; each inode is identified by a nonnegative integer (i.e., the inode number)
  - can translate the inode number to a location on the disk (inode contains pointers to file data blocks)
  - inodes are either put together as one group (in the inode list) or spread across the disk

# Naming and directory structure


- a directory entry (**dentry**) is a collection of (name, inode number) pairs for files and directories therein
  - is stored as a regular file (its inode has a special flag bit set)
  - only OS can modify directory; users can just read them
  - . and .. are stored as ordinary file names with inode numbers pointing to the inodes of the same directory and the parent directory respectively





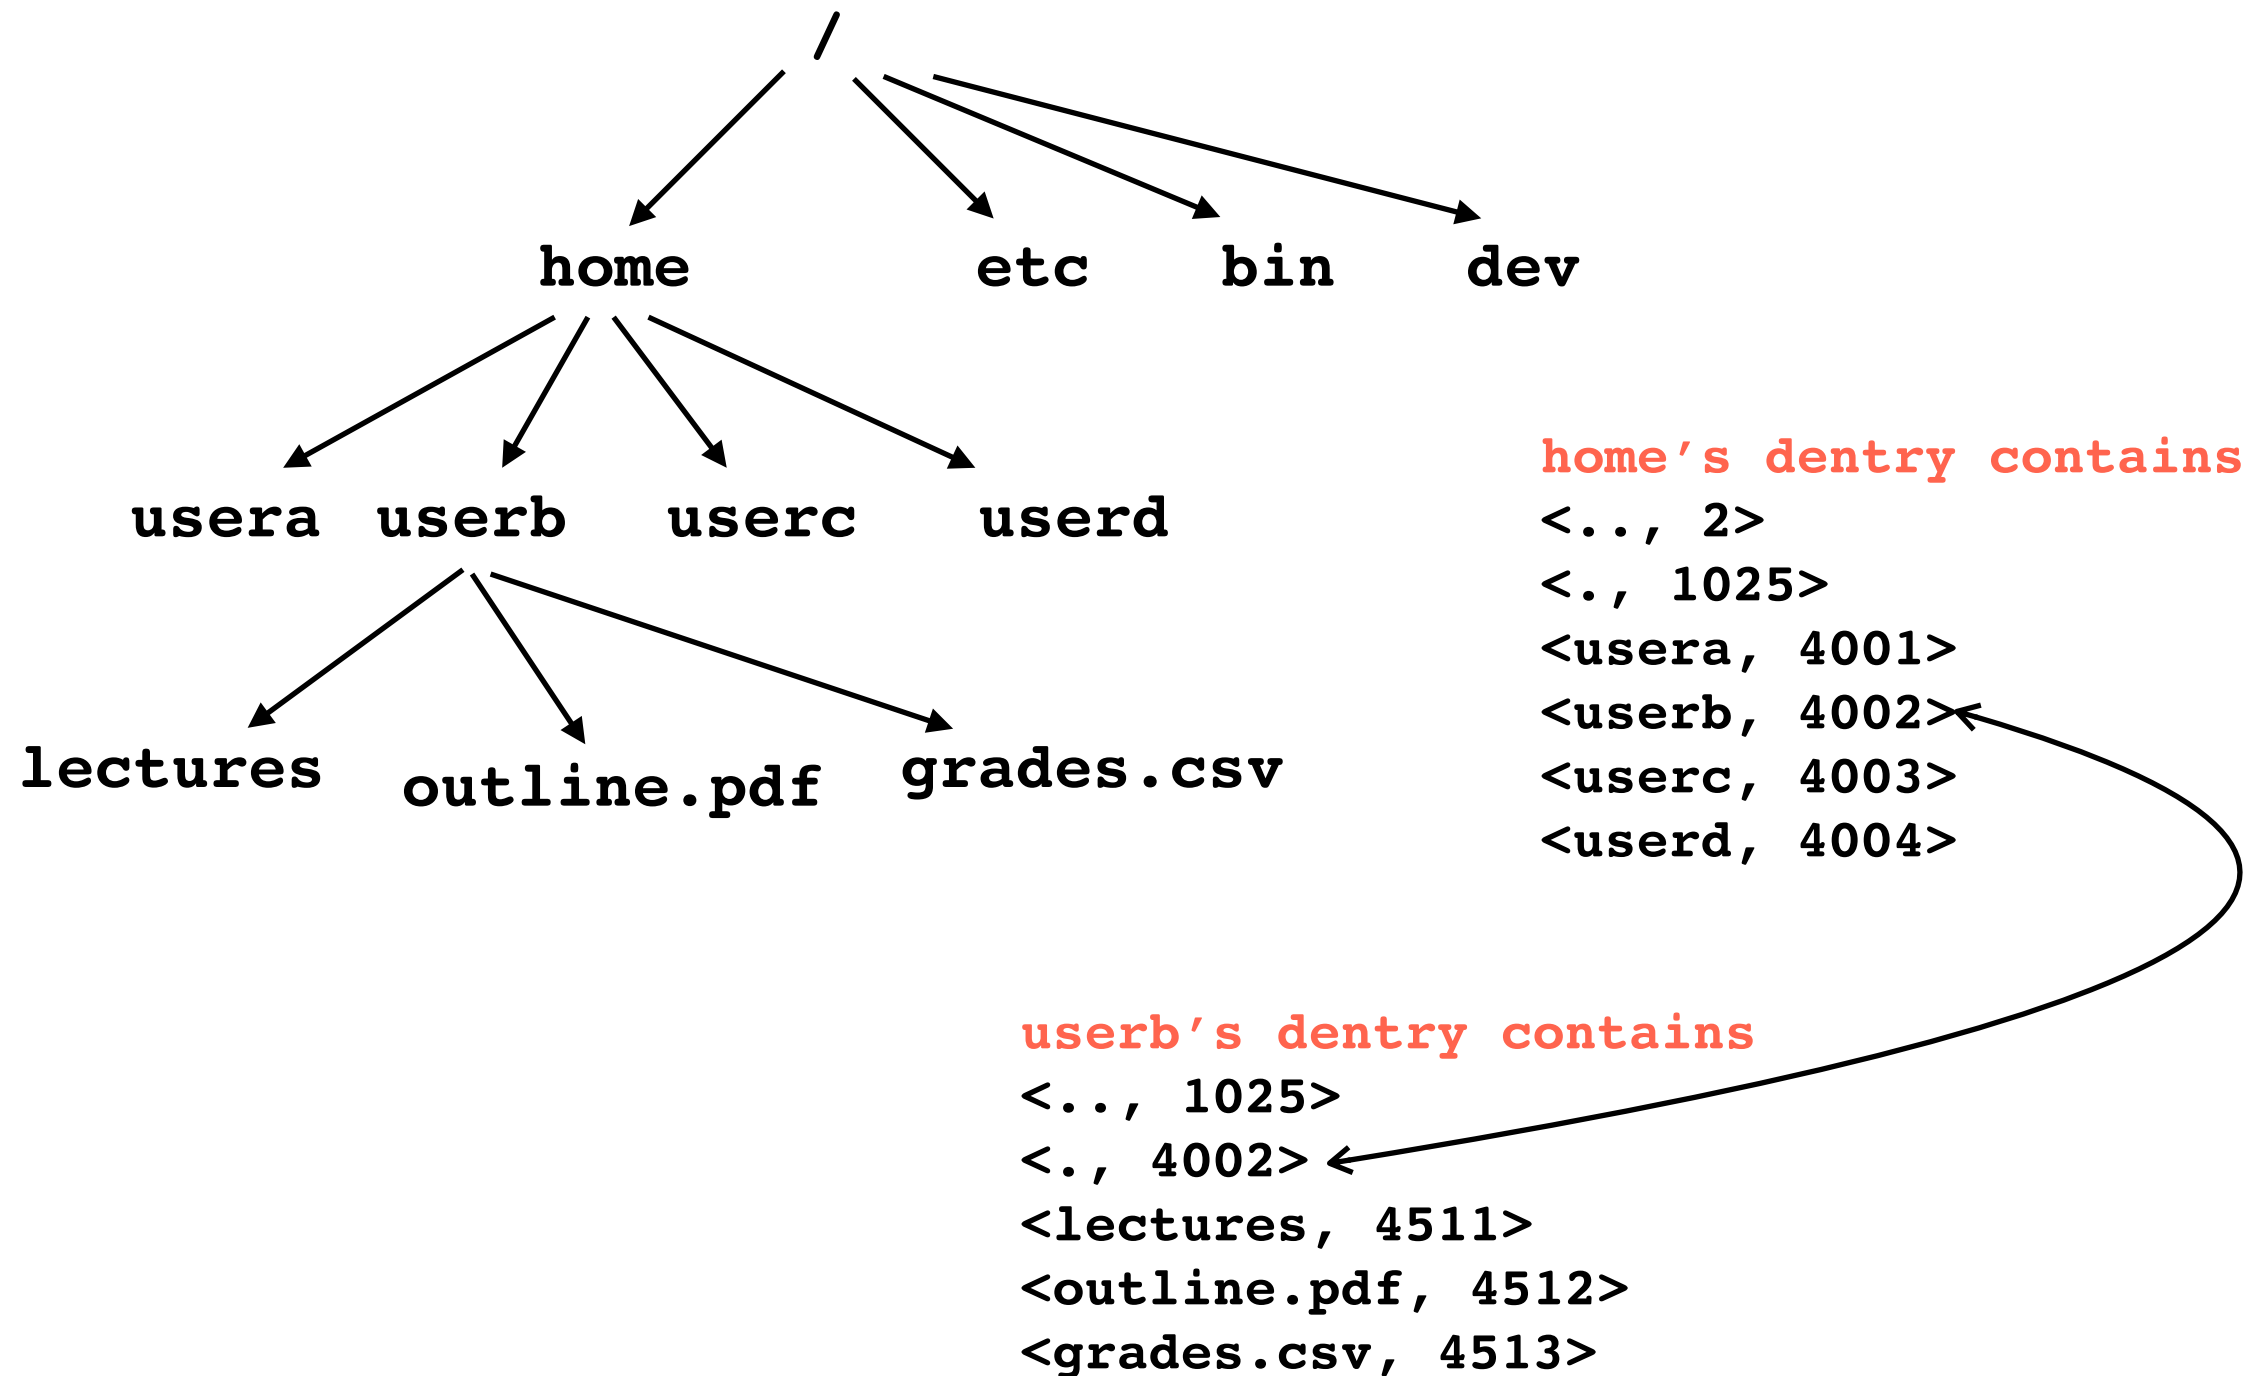
# A tree-structured hierarchy

---



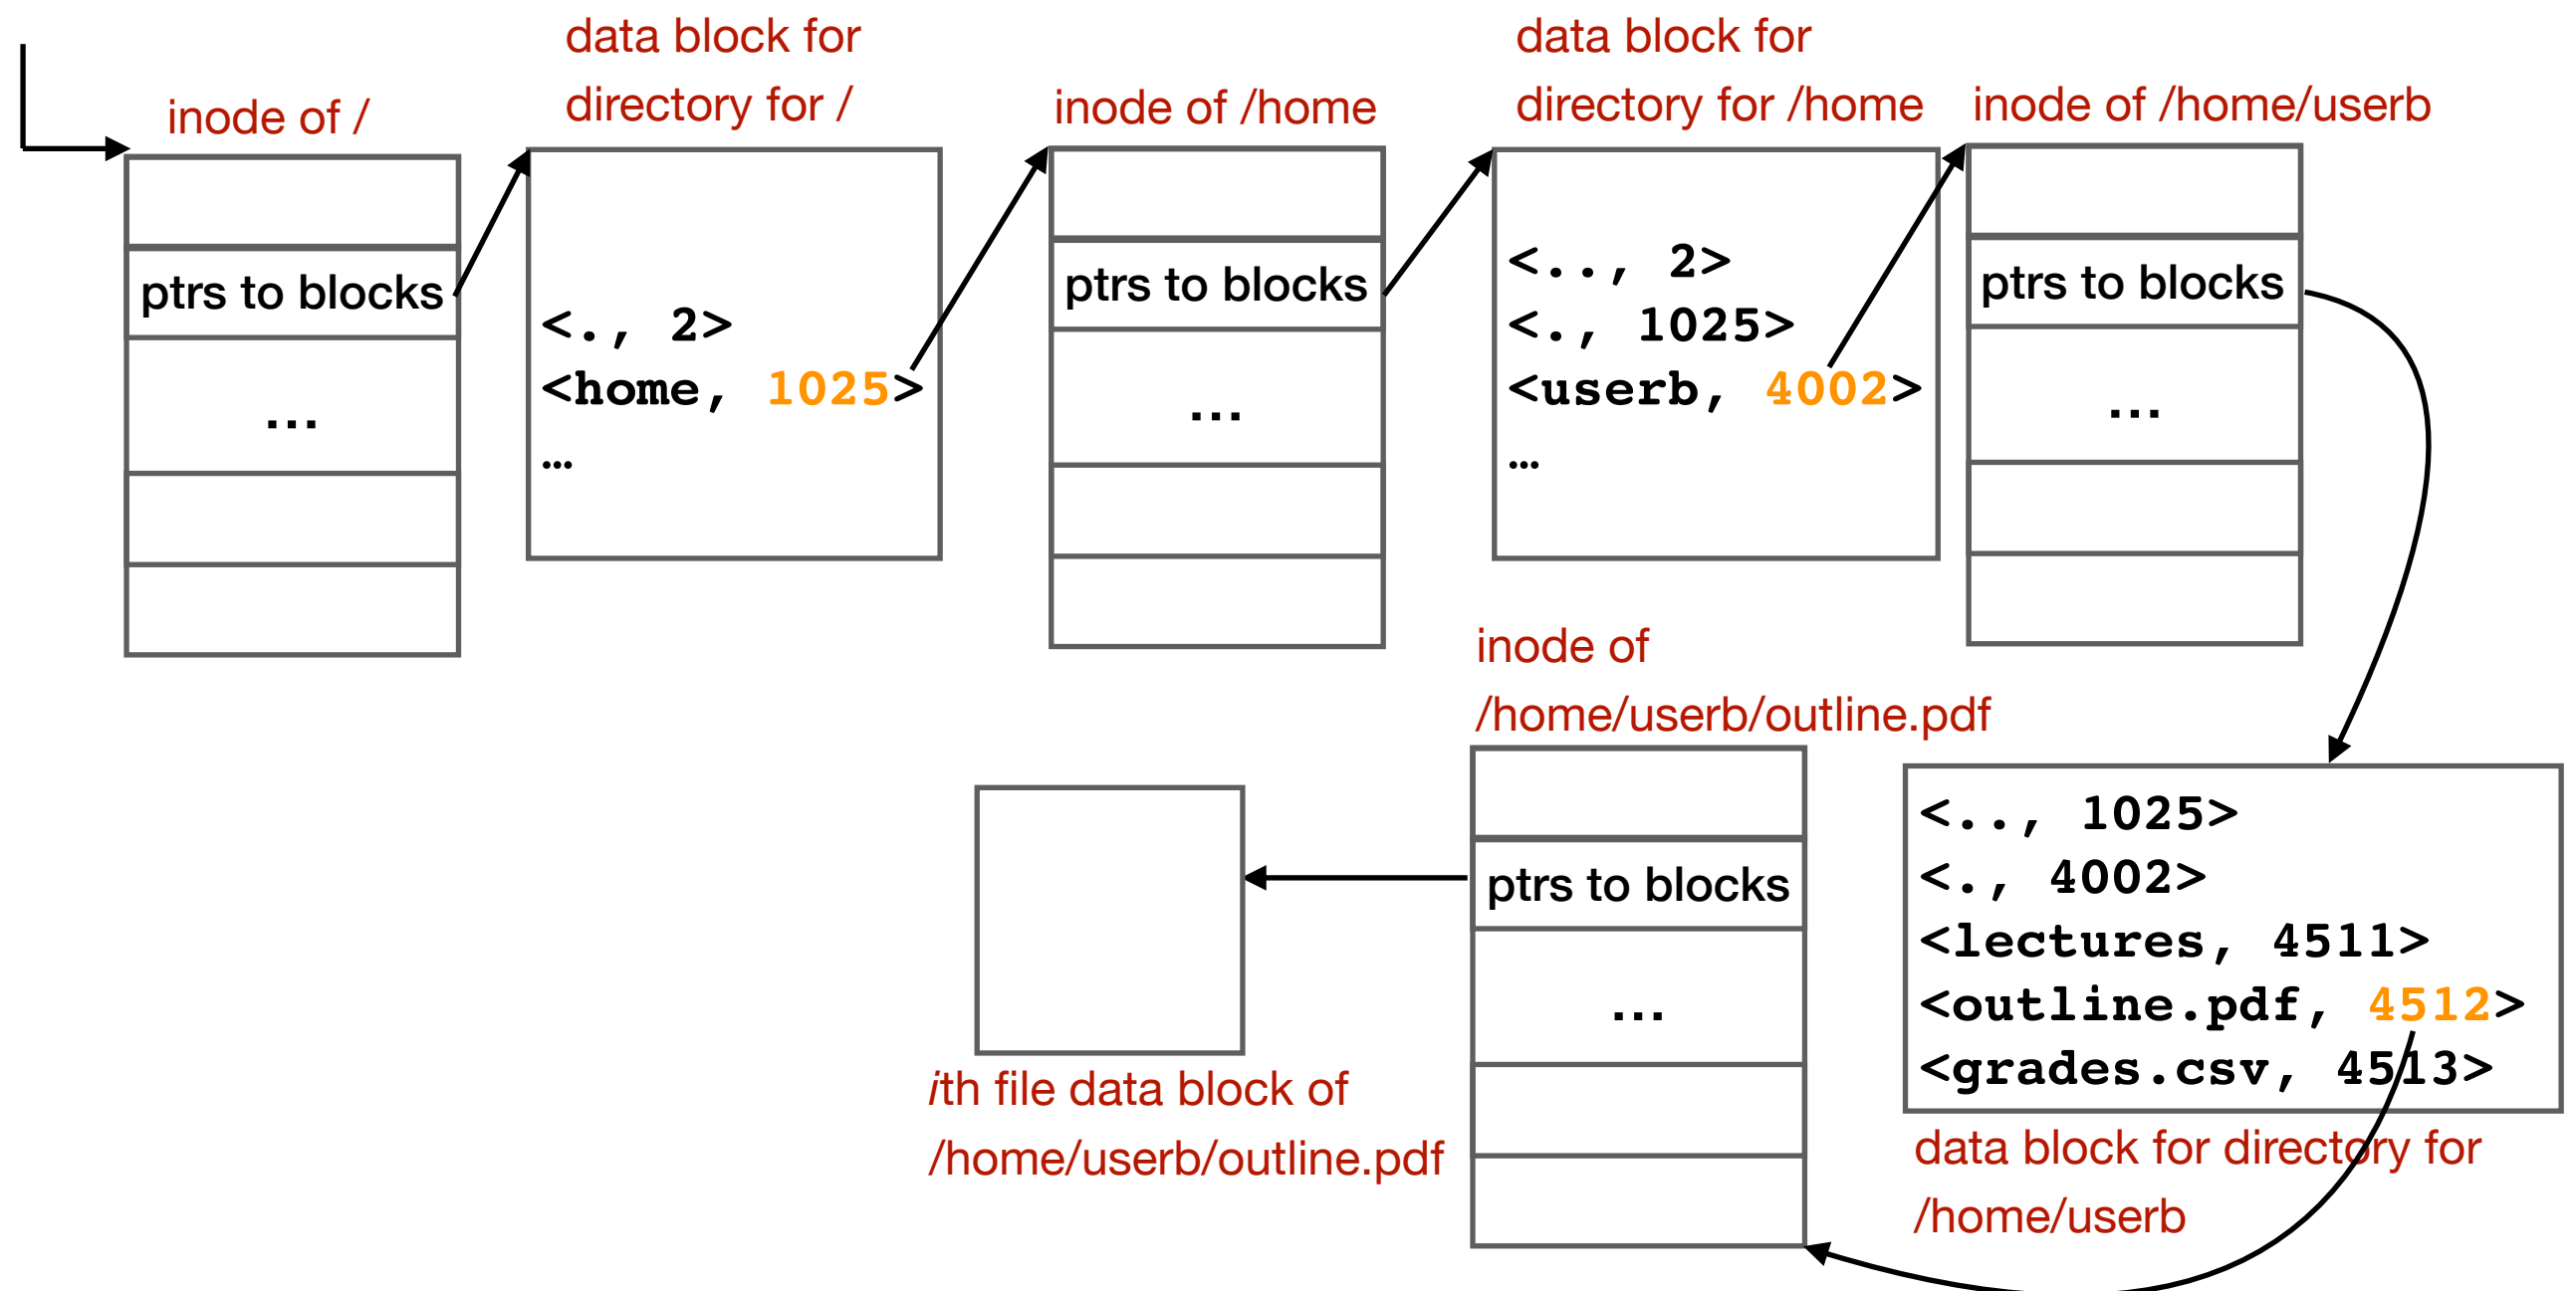
# A tree-structured hierarchy

---



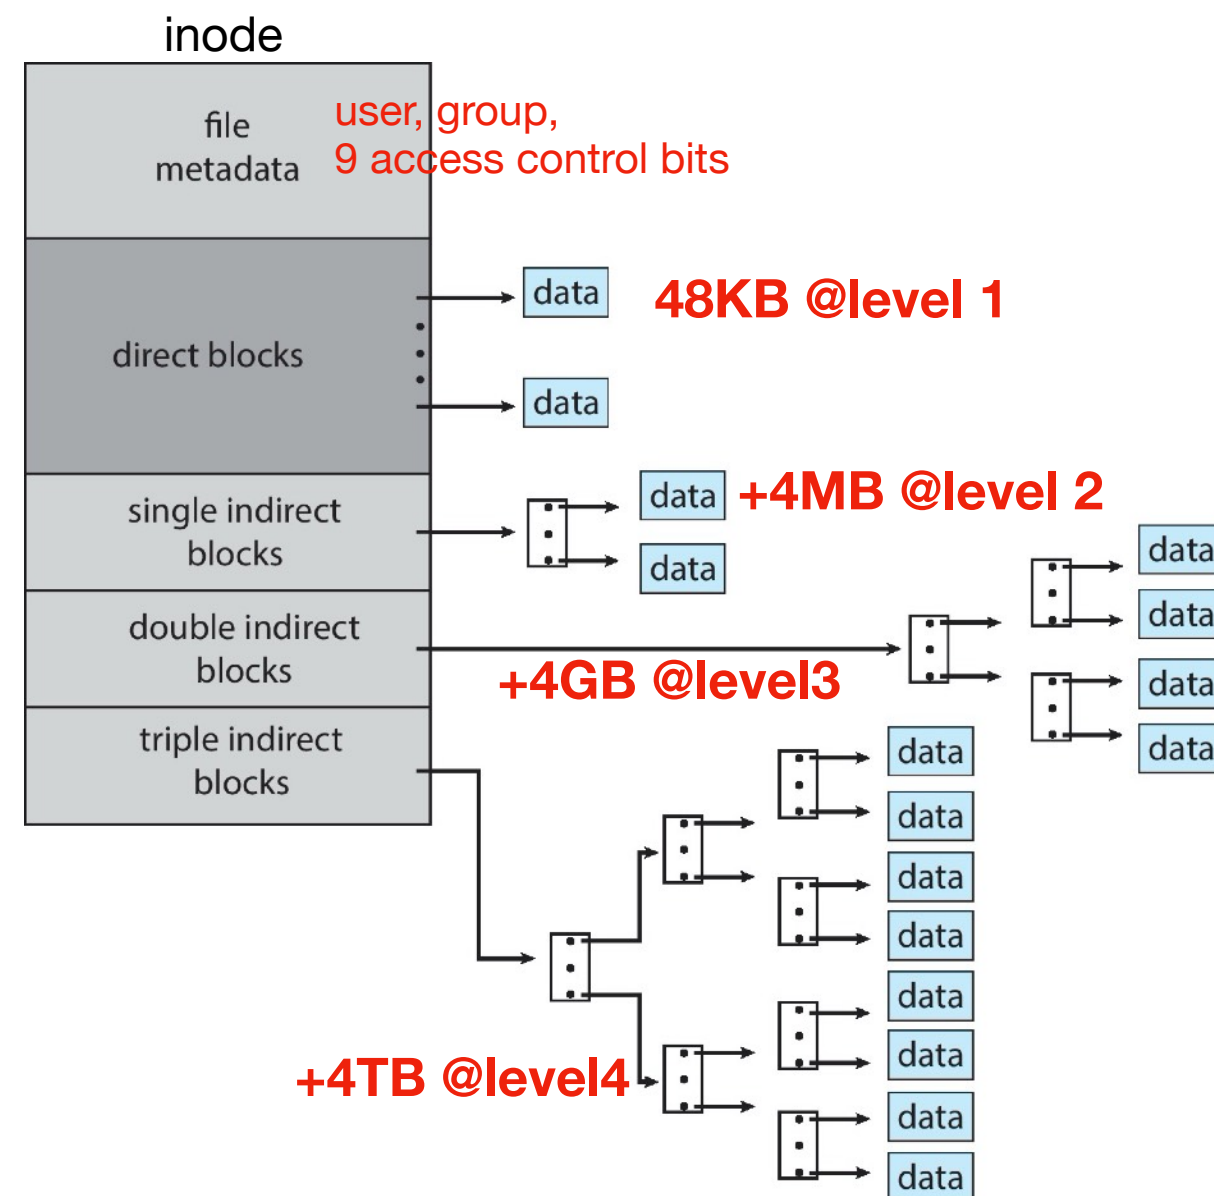
# Resolving file location

Where are data blocks of `/home/userb/outline.pdf`?



# Asymmetric tree

- original inode format appeared in BSD 4.1 (and later in BSD 4.2)
- each inode contains 15 block pointers (with 32-bit file pointer)
  - the first 12 pointers point to data blocks
    - small files (up to 48KB) use direct blocks only
  - the 13th pointer points to a block of 1024 pointers to 4KB data blocks (1-level indirection)
  - the 14th pointer points to a block of pointers to indirect blocks (2-level indirection)
  - the 15th pointer points to a block of pointers to indirect blocks (3-level indirection)
- shallow tree for small files
  - efficient storage for small files
- deep tree for large files
  - efficient lookup for random access in large files



# Problems

---

- but early versions of UNIX file system had a couple of problems
  - all inodes are put together as one group far from file data blocks;  
so resolving file address requires multiple disk seeks

# Problems

---

- but early versions of UNIX file system had a couple of problems
  - all inodes are put together as one group far from file data blocks;  
so resolving file address requires multiple disk seeks
    - put file data and inode in the same cylinder (hence no disk seeks)

# Problems

---

- but early versions of UNIX file system had a couple of problems
  - all inodes are put together as one group far from file data blocks;  
so resolving file address requires multiple disk seeks
    - put file data and inode in the same cylinder (hence no disk seeks)
  - a directory's contents are scattered over the whole disk;  
so listing a directory is slow

# Problems

---

- but early versions of UNIX file system had a couple of problems
  - all inodes are put together as one group far from file data blocks;  
so resolving file address requires multiple disk seeks
    - put file data and inode in the same cylinder (hence no disk seeks)
  - a directory's contents are scattered over the whole disk;  
so listing a directory is slow
    - store files from same directory near each other



# Problems

---

- but early versions of UNIX file system had a couple of problems
  - all inodes are put together as one group far from file data blocks;  
so resolving file address requires multiple disk seeks
    - put file data and inode in the same cylinder (hence no disk seeks)
  - a directory's contents are scattered over the whole disk;  
so listing a directory is slow
    - store files from same directory near each other
  - free-block list becomes disorganized over time (since blocks that are released are randomly placed in the free-block list) and blocks allocated to file get scattered;  
so sequential access becomes slow as it needs many seeks

# Problems

---

- but early versions of UNIX file system had a couple of problems
  - all inodes are put together as one group far from file data blocks;  
so resolving file address requires multiple disk seeks
    - put file data and inode in the same cylinder (hence no disk seeks)
  - a directory's contents are scattered over the whole disk;  
so listing a directory is slow
    - store files from same directory near each other
  - free-block list becomes disorganized over time (since blocks that are released are randomly placed in the free-block list) and blocks allocated to file get scattered;  
so sequential access becomes slow as it needs many seeks
    - use a bitmap and allocate files contiguously if possible

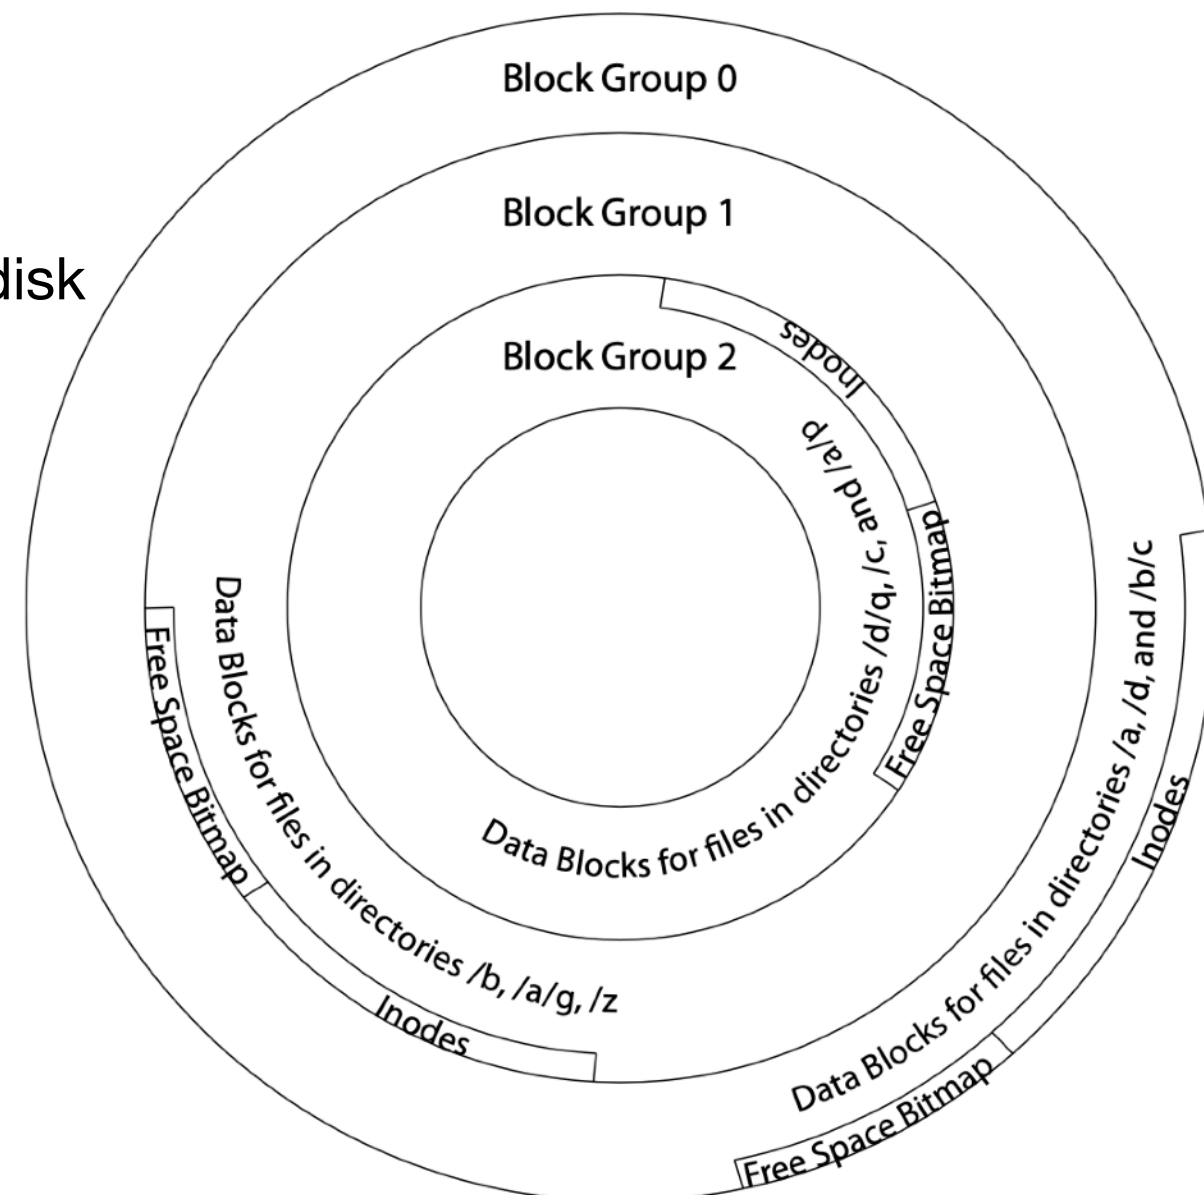
# Problems

---

- but early versions of UNIX file system had a couple of problems
  - all inodes are put together as one group far from file data blocks;  
so resolving file address requires multiple disk seeks
    - put file data and inode in the same cylinder (hence no disk seeks)
  - a directory's contents are scattered over the whole disk;  
so listing a directory is slow
    - store files from same directory near each other
  - free-block list becomes disorganized over time (since blocks that are released are randomly placed in the free-block list) and blocks allocated to file get scattered;  
so sequential access becomes slow as it needs many seeks
    - use a bitmap and allocate files contiguously if possible
- Fast File System (FFS) proposed allocation and placement policies for BSD 4.2 addressing these problems

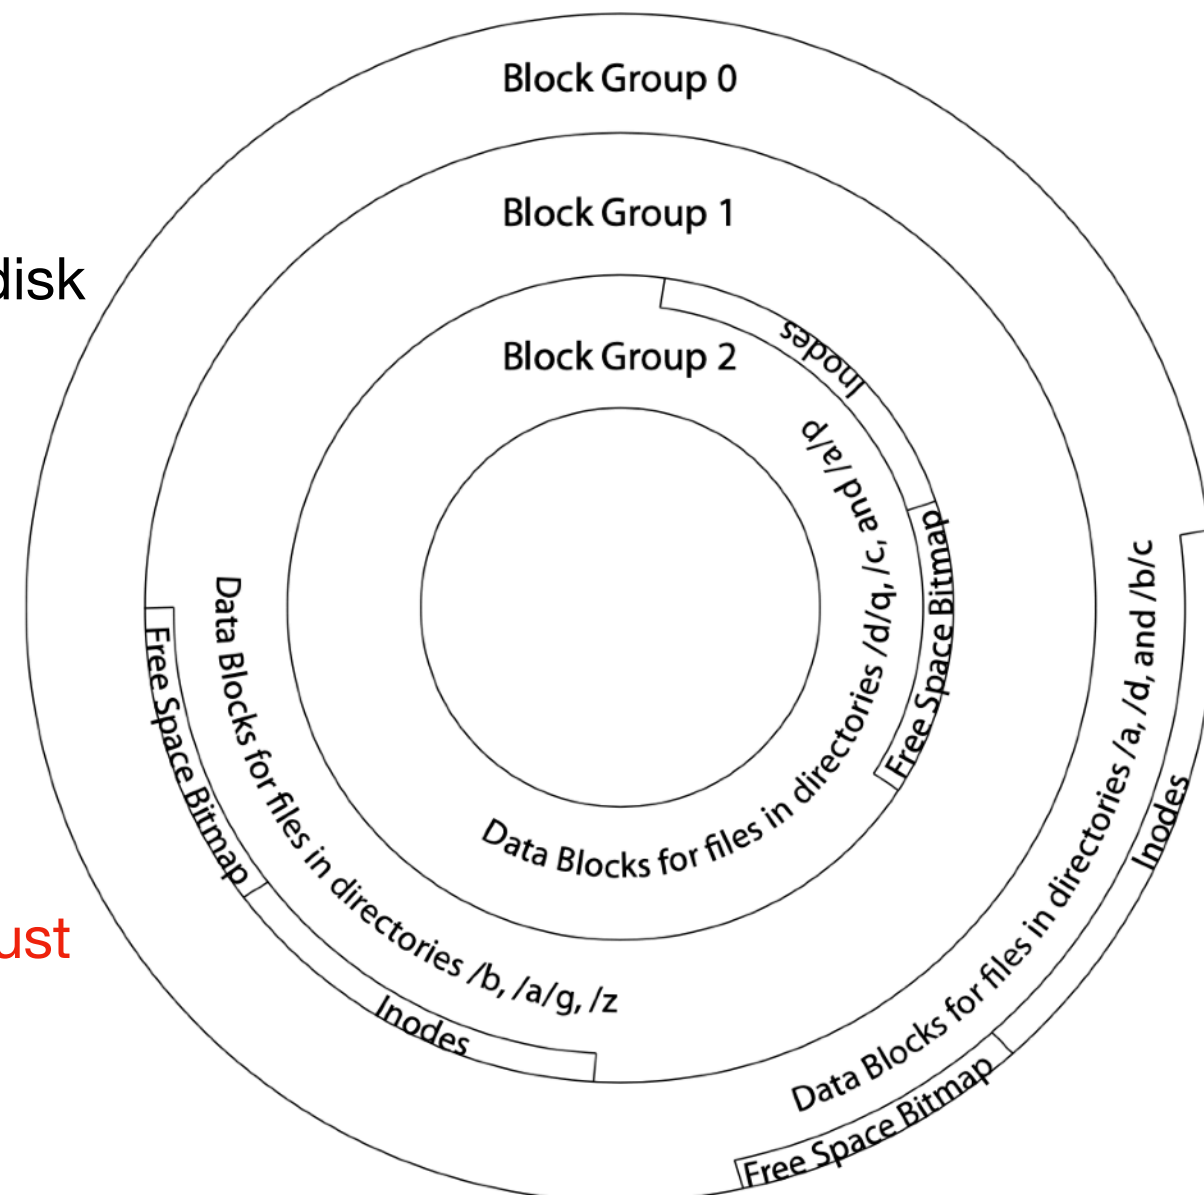
# Achieving data locality in FFS

- block group is a set of nearby cylinders
  - fast seek between cylinders in same block group
- block group allocation
  - files in same directory are located in same group
  - subdirectories are located in different block groups
- inode list and bitmap are spread throughout disk
  - near file blocks



# Achieving data locality in FFS

- block group is a set of nearby cylinders
  - fast seek between cylinders in same block group
- block group allocation
  - files in same directory are located in same group
  - subdirectories are located in different block groups
- inode list and bitmap are spread throughout disk
  - near file blocks



to allocate according to cylinder groups, disk must have free space scattered across cylinders  
(10% of disk space is reserved for this purpose)

# FFS first-fit block allocation

---

- when allocating space to a new file, search for a free block from the start of the block group
- when extending a file, search for a free block from the last block allocated to that file
  - if there is not enough free blocks in that range, allocate the remaining blocks from a new range
- there will be few little holes at start and a big hole at end of a block group
  - so small files will be fragmented, and large files will be mostly contiguous

start of block group




# FFS first-fit block allocation

---

- when allocating space to a new file, search for a free block from the start of the block group
- when extending a file, search for a free block from the last block allocated to that file
  - if there is not enough free blocks in that range, allocate the remaining blocks from a new range
- there will be few little holes at start and a big hole at end of a block group
  - so small files will be fragmented, and large files will be mostly contiguous

start of block group





# FFS first-fit block allocation

---

- when allocating space to a new file, search for a free block from the start of the block group
- when extending a file, search for a free block from the last block allocated to that file
  - if there is not enough free blocks in that range, allocate the remaining blocks from a new range
- there will be few little holes at start and a big hole at end of a block group
  - so small files will be fragmented, and large files will be mostly contiguous

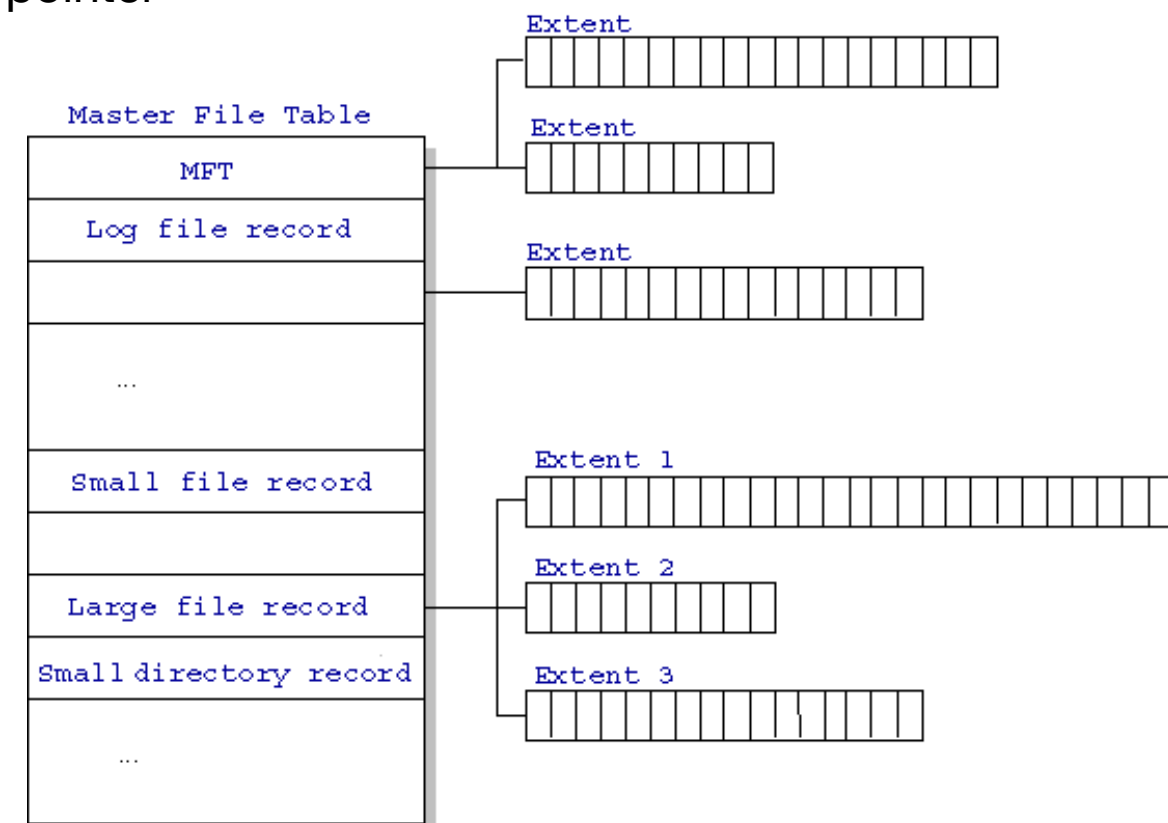
start of block group





# NTFS

- master file table is a flexible (1KB to 4KB) storage for metadata and data
- directories organized using B-trees
- index structure is a variable-depth tree
  - mixes direct and indirect block pointers
- variable-length extents (i.e., adjacent disk blocks)
  - have to store start and length of each extent in block pointer
  - user can provide hint as to size of file being created
- journalling for consistency checking
  - all metadata changes are written sequentially to a log



to learn more visit: [http://ntfs.com/ntfs\\_basics.htm](http://ntfs.com/ntfs_basics.htm)